

ECE 1508: Applied Deep Learning

Chapter 2: Feedforward Neural Networks

Ali Bereyhi

ali.bereyhi@utoronto.ca

Department of Electrical and Computer Engineering
University of Toronto

Fall 2025

Fully-Connected FNNs: Training

- + So, what do we do with forward pass?
- Say the FNN is fixed and all the weights and biases are given. Forward pass determines the label of a given data-point x .
- + Well! But we need to train the network! Right?!
- Yes! We define the loss and then train it via gradient descent
- + Then, we need to determine the gradient! It sounds complicated!
- Well! there is an efficient algorithm for that called backpropagation

Let's see what backpropagation is!

Fully-Connected FNNs: Training

Let's recall how we train the network: in our FNN, we considered M outputs; thus, we could assume that the dataset is of the form

$$\mathbb{D} = \{(\mathbf{x}_b, \mathbf{v}_b) \text{ for } b = 1, \dots, B\}$$

Here, we have denoted the true labels by $\mathbf{v}_b = [v_{b,1}, \dots, v_{b,M}]^\top$ to avoid confusion with the FNN's outputs. Now, let's denote the forward pass by $\mathbf{y}_b = \text{PassF}(\mathbf{x}_b | \mathbf{w})$ with \mathbf{w} is a vector collecting $\{\mathbf{W}_\ell\}$ for $\ell = 1, \dots, L + 1$

Given data-point \mathbf{x}_b , by forward pass we get \mathbf{y}_b at the output layer of the FNN with weights \mathbf{w} . This output is desired to be the true label \mathbf{v}_b

How do we do the training?

$$\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmin}} \hat{R}(\mathbf{w}) = \underset{\mathbf{w}}{\operatorname{argmin}} \frac{1}{B} \sum_{b=1}^B \mathcal{L}(\mathbf{y}_b, \mathbf{v}_b) \quad (\text{Training})$$

Fully-Connected FNNs: Training

Let's recall how we train the network

$$\begin{aligned}\mathbf{w}^* &= \underset{\mathbf{w}}{\operatorname{argmin}} \hat{R}(\mathbf{w}) = \underset{\mathbf{w}}{\operatorname{argmin}} \frac{1}{B} \sum_{b=1}^B \mathcal{L}(\mathbf{y}_b, \mathbf{v}_b) \\ &= \underset{\mathbf{w}}{\operatorname{argmin}} \frac{1}{B} \sum_{b=1}^B \mathcal{L}(\text{PassF}(\mathbf{x}_b | \mathbf{w}), \mathbf{v}_b)\end{aligned}$$

- 1: Initiate at some $\mathbf{w}^{(0)} \in \mathbb{R}^D$ and deviation $\Delta = +\infty$
- 2: Choose some small ϵ and η , and set $t = 1$
- 3: **while** $\Delta > \epsilon$ **do**
- 4: Update weights as $\mathbf{w}^{(t)} \leftarrow \mathbf{w}^{(t-1)} - \eta \nabla \hat{R}(\mathbf{w}^{(t-1)})$
- 5: Update the deviation $\Delta = |\hat{R}(\mathbf{w}^{(t)}) - \hat{R}(\mathbf{w}^{(t-1)})|$
- 6: **end while**

In this algorithm, the main challenge is to calculate $\nabla \hat{R}(\mathbf{w}^{(t-1)})$

Fully-Connected FNNs: Training

The main challenge is to calculate $\nabla \hat{R}(\mathbf{w})$

First, let's see what are the entries of \mathbf{w} : \mathbf{w} contains all weights and biases.
Following our notations, we can say

$$\mathbf{w} = \begin{bmatrix} \mathbf{w}_1 [1, :] \\ \vdots \\ \mathbf{w}_1 [\mathcal{W}_1, :] \\ \vdots \\ \mathbf{w}_{L+1} [1, :] \\ \vdots \\ \mathbf{w}_{L+1} [\mathcal{W}_{L+1}, :] \end{bmatrix} \xrightarrow{\text{layer } \ell = 1} \mathbf{w}_\ell [j, :] = \begin{bmatrix} w_\ell [j, 0] \\ w_\ell [j, 1] \\ \vdots \\ w_\ell [j, \mathcal{W}_{\ell-1}] \end{bmatrix}$$

$\xrightarrow{\text{layer } \ell = L + 1}$

So, the entries of \mathbf{w} are $w_\ell [j, i]$ for different choices of i , j and ℓ

Fully-Connected FNNs: Training

The main challenge is to calculate $\nabla \hat{R}(\mathbf{w})$

Let's try to open up the gradient: we need partial derivatives of $\hat{R}(\mathbf{w})$ with respect to $w_{\ell}[j, i]$ for $i = 0 : \mathcal{W}_{\ell-1}$, $j = 1 : \mathcal{W}_\ell$, and $\ell = 1 : L + 1$

$$\begin{aligned}\frac{\partial}{\partial w_{\ell}[j, i]} \hat{R}(\mathbf{w}) &= \frac{\partial}{\partial w_{\ell}[j, i]} \frac{1}{B} \sum_{b=1}^B \mathcal{L}(\mathbf{y}_b, \mathbf{v}_b) \\ &= \frac{1}{B} \sum_{b=1}^B \boxed{\frac{\partial}{\partial w_{i,j}[\ell]} \mathcal{L}(\mathbf{y}_b, \mathbf{v}_b)}\end{aligned}$$

So, it's enough to develop an algorithm that computes partial derivative for a single **data-point**. Derivative of the risk is then average of these point-wise derivatives.

Fully-Connected FNNs: Training

Let's make an agreement: we consider a *data-point* \mathbf{x} with *label* \mathbf{v} and write

$$\frac{\partial}{\partial w_{\ell}[j, i]} \mathcal{L}(\text{PassF}(\mathbf{x}|\mathbf{w}), \mathbf{v}) = \frac{\partial}{\partial w_{\ell}[j, i]} \mathcal{L}(\mathbf{y}, \mathbf{v})$$

while keeping in mind that \mathbf{y} is a function of \mathbf{w}

To determine the partial derivatives, we note that

\mathbf{y} is a nested function of \mathbf{w}

so, we can determine the derivative via *chain rule*. Let's recall the *chain rule* and see how we can *apply it on a graph*

Review: Chain Rule

Assume $z = g(x)$ and $y = f(z)$: y is a nested function of x , as we can write

$$y = f(g(x))$$

Intuitively, we can say: if at point x we move with tiny step $\text{d}x$, z varies as

$$\text{d}z = \dot{g}(x) \text{d}x$$

This variation also varies y : moving from $z = g(x)$ with tiny step $\text{d}z$ leads to

$$\text{d}y = \dot{f}(z) \text{d}z$$

So, we have

$$\text{d}y = \dot{f}(z) \dot{g}(x) \text{d}x$$

Review: Chain Rule

We have concluded that by moving x with $\text{d}x$, we get

$$\text{d}y = \dot{f}(z)\dot{g}(x)\text{d}x$$

On the other hand, we know that

$$\text{d}y = \frac{\text{d}}{\text{d}x}f(g(x))\text{d}x$$

This concludes the *chain rule*

Chain Rule: Scalar Form

The derivative of *nested function* $y = f(g(x))$ with respect to x is given by

$$\frac{\text{d}y}{\text{d}x} = \frac{\text{d}}{\text{d}x}f(g(x)) = \dot{f}(z)\dot{g}(x) = \frac{\text{d}y}{\text{d}z}\frac{\text{d}z}{\text{d}x}$$

Computation Graph

We can extend this idea to deeper nested functions:

Let $z_1 = g_1(x)$ and $z_{n+1} = g_{n+1}(z_n)$ for $n = 1, \dots, N-1$; then, derivative of $y = f(z_N)$ with respect to x is given by

$$\frac{dy}{dx} = \frac{dy}{dz_N} \left(\prod_{n=1}^{N-1} \frac{dz_{n+1}}{dz_n} \right) \frac{dz_1}{dx} = f'(z_N) \left(\prod_{n=1}^{N-1} g'_{n+1}(z_n) \right) g_1'(x)$$


We can represent the chain rule, using a *computation graph*: for the deep nested function given above, the computation graph is given by



In this graph, we start from x and pass forward to $z_1 \rightarrow z_2 \rightarrow \dots$ until we get to y . In each pass, we determine next variable via the function on the link

Computation Graph

The derivative of y with respect to any variable on this graph is determined by a *backward pass from y towards the variable*



$$\frac{dy}{dz_N} = \dot{f}(z_N)$$

$$\frac{dy}{dz_{N-1}} = \frac{dy}{dz_N} \frac{dz_N}{dz_{N-1}} = \dot{f}(z_N) \dot{g}_N(z_{N-1})$$

2

$$\frac{dy}{dz_1} = \frac{dy}{dz_N} \frac{dz_N}{dz_{N-1}} \dots \frac{dz_2}{dz_1} = \dot{f}(z_N) \dot{g}_N(z_{N-1}) \dots \dot{g}_2(z_2)$$

$$\frac{dy}{dx} = \frac{dy}{dz_N} \frac{dz_N}{dz_{N-1}} \cdots \frac{dz_2}{dz_1} \frac{dz_1}{dx} = \dot{f}(z_N) \dot{g}_N(z_{N-1}) \cdots \dot{g}_2(z_2) \dot{g}_1(x)$$


Computation Graph: Example

Example: y is a nested function of x through the following chain of functions:

$$z_1 = 2x \quad z_2 = z_1^2 \quad z_3 = e^{z_2} \quad y = \log z_3$$


Determine the derivative of y with respect to x at $x = 0.5$.

Let's first plot the computation graph




During the *forward pass* we get

$$z_1 = 2 \times 0.5 = 1 \rightarrow z_2 = 1^2 = 1 \rightarrow z_3 = e^1 = e \rightarrow y = \log e = 1$$




Computation Graph: Example

Now, we pass *backward* to determine the derivative




Let's first enter the values into the backward links



We now *navigate backward* to each variable that we want to determine the derivative y with respect to it

Computation Graph: Example



The derivatives are easily determined *recursively*


$$\frac{dy}{dz_3} = \frac{1}{e}$$

$$\frac{dy}{dz_2} = \frac{dy}{dz_3} \frac{dz_3}{dz_2} = \frac{e}{e} = 1$$

$$\frac{dy}{dz_1} = \frac{dy}{dz_2} \frac{dz_2}{dz_1} = 1 \times 2 = 2$$

$$\frac{dy}{dx} = \frac{dy}{dz_1} \frac{dz_1}{dx} = 2 \times 2 = 4$$

Computation Graph



In the example, we had to *first compute the value of each variable*, in order to be able to compute the *values of the backward links*

This is an important fact that we should remember

Backward pass is only possible if we have already taken the forward pass

Review: Chain Rule

The nested function can be a multivariate: *assume for $n = 1, \dots, N$*

$$z_n = g_n(x)$$

and let the nested function be

$$y = f(z_1, \dots, z_N)$$

Let's follow the same logic: starting from point x , we move with tiny step dx

$$\text{dz}_n = \dot{g}_n(x) \text{dx}$$

These variations lead to variation dy in the nested function

$$\text{dy} = \nabla f(\mathbf{z})^\top \text{dz} = \sum_{n=1}^N \frac{\partial y}{\partial z_n} \text{dz}_n = \sum_{n=1}^N \dot{f}_n(\mathbf{z}) \text{dz}_n$$

Review: Chain Rule

We can hence write

$$\text{d}y = \sum_{n=1}^N \dot{f}_n(\mathbf{z}) \text{d}z_n = \sum_{n=1}^N \dot{f}_n(\mathbf{z}) \dot{g}_n(x) \text{d}x$$

This concludes the vector form of the chain rule


Chain Rule: Vector Form

Let $\mathbf{z} = [z_1, \dots, z_N]^\top$ and $z_n = g_n(x)$. The derivative of nested function $y = f(\mathbf{z})$ with respect to x is given by

$$\frac{\text{d}y}{\text{d}x} = \sum_{n=1}^N \frac{\partial y}{\partial z_n} \frac{\text{d}z_n}{\text{d}x} = \sum_{n=1}^N \dot{f}_n(\mathbf{z}) \dot{g}_n(x)$$

Computation Graph


We can again represent the function via its vectorized *computation graph* and compute it by *passing forward* on this graph



We start from x and pass forward to $\mathbf{z} = [z_1, z_2, \dots, z_N]$. We then pass forward \mathbf{z} to y . In each pass, we compute next variable via function on the link.

Computation Graph

The derivative with respect to any node is then given by *backward pass towards the node on the computation graph*




We add all backward passes towards x to determine the derivative

$$\frac{dy}{dx} = \sum_{n=1}^N \frac{\partial y}{\partial z_n} \frac{dz_n}{dx} = \sum_{n=1}^N \dot{f}_n(\mathbf{z}) \dot{g}_n(x)$$

Computation Graph: Single Neuron


Let's now plot the *computation graph* of a single neuron and determine the gradient of the loss by *backward pass*



After passing the data-point \mathbf{x} through the neuron, we get \mathbf{y} and we calculate the loss for the *true label* v as $\mathcal{L}(\mathbf{y}, v)$

Computation Graph: Single Neuron


The *computation graph* is hence given by



Here, the computation nodes are the *weights and bias of the neuron*
once we fix them, we can pass forward and get to the loss \hat{R}

Computation Graph: Single Neuron

Once passed forward, we can move backward to determine the derivatives




For a particular weight w_n , we can write (we drop arguments whenever clear)

$$\frac{\partial \hat{R}}{\partial w_n} = \frac{d\hat{R}}{dy} \frac{dy}{dz} \frac{\partial z}{\partial w_n} = \dot{\mathcal{L}} \dot{f} x_n$$

Computation Graph: Multivariate Form

Let's get back to the following *computation graph*



We define **vector-valued functions**, and show the graph **compactly**: let's define


$$\mathbf{g}(x) = \begin{bmatrix} g_1(x) \\ \vdots \\ g_N(x) \end{bmatrix}$$

Computation Graph: Multivariate Form

Function $g(\cdot)$ gets x as the input and returns all z_n 's in a **vector** \mathbf{z} , i.e.,

$$g(x) = \begin{bmatrix} g_1(x) \\ \vdots \\ g_N(x) \end{bmatrix} = \begin{bmatrix} z_1 \\ \vdots \\ z_N \end{bmatrix} = \mathbf{z}$$

We now use this vectorized notation to simplify the *computation graph* as



The **forward pass** on this graph is exactly the same: we give x to the **vectorized function** $g(\cdot)$ to get \mathbf{z} which is then **passed forward** to $f(\cdot)$ to get y


*How does the **backward pass** look like then?*

Computation Graph: Multivariate Form

We could define the derivative $\dot{\mathbf{g}}(\cdot)$ as the vector of derivatives $\dot{g}(\cdot)$

$$\dot{\mathbf{g}}(x) = \begin{bmatrix} \dot{g}_1(x) \\ \vdots \\ \dot{g}_N(x) \end{bmatrix} = \begin{bmatrix} \frac{dz_1}{dx} \\ \vdots \\ \frac{dz_N}{dx} \end{bmatrix} = \frac{d\mathbf{z}}{dx}$$

Let's show this **vectorized derivative** and **gradient** of f on the backward links




Well, we can **pass backward** as follows

$$\frac{dy}{dx} = \sum_{n=1}^N \frac{\partial y}{\partial z_n} \frac{dz_n}{dx} = \nabla f(\mathbf{z})^\top \dot{\mathbf{g}}(x)$$

Computation Graph: Multivariate Form

- + What can we conclude then?
- We can sketch the computation graph very compactly using **vectorized derivatives** and **gradients**
- + Does it mean that we should then **pass backward** exactly the same as in a computation graph with **scalar variables and derivatives**?
- Pretty much Yes! Only one delicate detail: we should **know how to multiply those gradients and vectorized derivatives!**




In our example, we determined the *inner product*

$$\frac{dy}{dx} = \nabla f^T \dot{g}$$

Computation Graph: Multivariate Form

- + How do we know **which type of product** we should use?
- Well! If you were in doubt, we could always do it by **expanding in terms of entries**; however, we are going to practice **all key functions** that appear in NN computation graphs!

Before we start with all key functions, let's get back to a *single neuron*



Computation Graph: Multivariate Form


Let's define, as we did earlier, the following vectors

$$\mathbf{x} = \begin{bmatrix} x_0 = 1 \\ x_1 \\ \vdots \\ x_N \end{bmatrix} \quad \text{and} \quad \mathbf{w} = \begin{bmatrix} w_0 = b \\ w_1 \\ \vdots \\ w_N \end{bmatrix}$$

Recall that **output** of the neuron is determined as $y = f(z)$ for


$$z = \mathbf{x}^T \mathbf{w}$$

So, we can show the **computation graph** compactly as




Computation Graph: Multivariate Form

Let's look at each link carefully: we pass **backward**, so we start with **last link**




- \hat{R} is a *scalar* function of *scalar* y , i.e., $\hat{R} = \mathcal{L}(y, v)$
↳ the backward link contains the *scalar derivative* $\dot{\mathcal{L}}$
- y is a *scalar* function of *scalar* z , i.e., $y = f(z)$
↳ the backward link contains the *scalar derivative* \dot{f}
- z is a *scalar* function of *vector* w , i.e., $z = x^T w$
↳ the backward link contains the *gradient* ∇z

So, the graph with the **backward** links looks like




Computation Graph: Multivariate Form



We are almost complete; only we need to calculate ∇z


$$z = w_0 + w_1 x_1 + \dots + w_N x_N \rightsquigarrow \nabla z = \begin{bmatrix} \partial z / \partial w_0 \\ \partial z / \partial w_1 \\ \vdots \\ \partial z / \partial w_N \end{bmatrix} = \begin{bmatrix} 1 = x_0 \\ x_1 \\ \vdots \\ x_N \end{bmatrix} = \mathbf{x}$$

So, we are complete! Here is the *vectorized computation graph* of the neuron



Computation Graph: Multivariate Form

Now, how do we pass **backward** on this graph?



We arrived at y at **the end of forward pass**: at this point, we can determine

$$\frac{d\hat{R}}{dy} = \dot{\mathcal{L}}(y, v) = \dot{\mathcal{L}}$$


and we are at the computing node y . We then pass **backward** $\dot{\mathcal{L}}$. At node z , we can compute $\dot{f}(z)$, and use what we received from y to compute

$$\frac{d\hat{R}}{dz} = \frac{d\mathcal{L}}{dy} \frac{dy}{dz} = \dot{f}\dot{\mathcal{L}}$$

and pass it **backward**

Computation Graph: Multivariate Form

Now, how do we pass **backward** on this graph?



Arriving at w , we can determine $\nabla z = x$ and use what we received from z to compute **what we want**

$$\nabla \hat{R} \text{ w.r.t. } w \equiv \nabla_w \hat{R} = \frac{d\hat{R}}{dz} \nabla z = \dot{f} \dot{\mathcal{L}} x$$


- + Well! That seems easier!
- Right! Let's now try some important cases

Backpropagation: Local Operations

Let's consider a general problem: an *objective scalar \hat{R}* is a function of K -dimensional vector $\mathbf{y} \in \mathbb{R}^K$. Clearly in this case, we have a *gradient*

$$\nabla_{\mathbf{y}} \hat{R} = \begin{bmatrix} \partial \hat{R} / \partial y_1 \\ \vdots \\ \partial \hat{R} / \partial y_K \end{bmatrix}$$

Assume that we know this gradient. The *vector \mathbf{y}* is also function of *another variable*. We want to compute *gradient of \hat{R}* with respect to this *other variable*




We now consider different cases for the *other variable* and its link to \mathbf{y}

Backpropagation: Local Operation 1

Entry-wise Functional Operation

$\mathbf{y} \in \mathbb{R}^K$ is a function of $\mathbf{z} \in \mathbb{R}^K$ as $\mathbf{y} = f(\mathbf{z})$ with $f(\cdot)$ operating entry-wise



For this case, we note that y_k is only a function of z_k ; thus we have

$$\frac{\partial \hat{R}}{\partial z_k} = \frac{\partial \hat{R}}{\partial y_k} \frac{\partial y_k}{\partial z_k} = \frac{\partial \hat{R}}{\partial y_k} \dot{f}(z_k)$$

So, we can use entry-wise product \odot to get from $\nabla_{\mathbf{y}} \hat{R}$ to $\nabla_{\mathbf{z}} \hat{R}$


$$\nabla_{\mathbf{z}} \hat{R} = \nabla_{\mathbf{y}} \hat{R} \odot \dot{f}(\mathbf{z})$$

Backpropagation: Local Operation 1

Reminder: Entry-wise product of two vectors of the same size is

$$\mathbf{z} \odot \mathbf{y} \begin{bmatrix} z_1 \\ \vdots \\ z_K \end{bmatrix} \odot \begin{bmatrix} y_1 \\ \vdots \\ y_K \end{bmatrix} = \begin{bmatrix} y_1 z_1 \\ \vdots \\ y_K z_K \end{bmatrix}$$

So, we can compactly perform this local operation as follows




with the backward step

$$\nabla_{\mathbf{z}} \hat{R} = \nabla_{\mathbf{y}} \hat{R} \odot f'(\mathbf{z})$$

Backpropagation: Local Operation 2

Linear Vector-to-Vector Operation

$\mathbf{y} \in \mathbb{R}^K$ is a function of $\mathbf{z} \in \mathbb{R}^N$ as $\mathbf{y} = \mathbf{A}\mathbf{z}$ with $\mathbf{A} \in \mathbb{R}^{K \times N}$



Here, y_k is a linear function of z_1, \dots, z_N

$$y_k = \sum_{n=1}^N \mathbf{A}[k, n] z_n$$

where $\mathbf{A}[k, n]$ is entry of \mathbf{A} at row k and column n . We thus can write

$$\frac{\partial \hat{R}}{\partial z_n} = \sum_{k=1}^K \frac{\partial \hat{R}}{\partial y_k} \frac{\partial y_k}{\partial z_n} = \sum_{k=1}^K \frac{\partial \hat{R}}{\partial y_k} \mathbf{A}[k, n]$$

Backpropagation: Local Operation 2

Let's denote column n of \mathbf{A} by notation $\mathbf{A}[:, n]$; so, we can write

$$\frac{\partial \hat{R}}{\partial z_n} = \sum_{k=1}^K \frac{\partial \hat{R}}{\partial y_k} \mathbf{A}[k, n] = \mathbf{A}[:, n]^T \nabla_y \hat{R}$$


Now, if we collect them in a vector form we get

$$\nabla_z \hat{R} = \begin{bmatrix} \partial \hat{R} / \partial z_1 \\ \vdots \\ \partial \hat{R} / \partial z_N \end{bmatrix} = \begin{bmatrix} \mathbf{A}[:, 1]^T \\ \vdots \\ \mathbf{A}[:, N]^T \end{bmatrix} \nabla_y \hat{R} = \mathbf{A}^T \nabla_y \hat{R}$$

This makes sense! Since we are **changing dimensions from K to N** , we need a product that **does such dimensionality change** for us

Backpropagation: Local Operation 2

Long story short . . .




with backward step

$$\nabla_z \hat{R} = A^T \nabla_y \hat{R}$$

Backpropagation: Local Operation 3

Linear Matrix-to-Vector Operation

$\mathbf{y} \in \mathbb{R}^K$ is a function of $\mathbf{A} \in \mathbb{R}^{K \times N}$ as $\mathbf{y} = \mathbf{Ax}$ with $\mathbf{x} \in \mathbb{R}^N$



- + Wait a moment! The **other variable** is a matrix! How do we define $\nabla_{\mathbf{A}} \hat{R}$?
- Right! Let's first extend the definition

Backpropagation: Local Operation 3

Assume scalar \hat{R} is a function of matrix $\mathbf{A} \in \mathbb{R}^{K \times N}$, we define

$$\nabla_{\mathbf{A}} \hat{R} = \begin{bmatrix} \partial \hat{R} / \partial \mathbf{A}[1, 1] & \dots & \partial \hat{R} / \partial \mathbf{A}[1, N] \\ \vdots & & \vdots \\ \partial \hat{R} / \partial \mathbf{A}[K, 1] & \dots & \partial \hat{R} / \partial \mathbf{A}[K, N] \end{bmatrix}$$

with $\mathbf{A}[k, n]$ being the entry of \mathbf{A} at row k and column n


It is worth to also think of **gradient descent** in this case: assume we are minimizing \hat{R} over \mathbf{A} using **gradient descent** with **learning rate** η . At iteration t we got point $\mathbf{A}^{(t)}$; now, in the next iteration we can readily write

$$\mathbf{A}^{(t+1)} = \mathbf{A}^{(t)} - \nabla_{\mathbf{A}} \hat{R}|_{\mathbf{A}=\mathbf{A}^{(t)}}$$

so apparently everything is as before!

Backpropagation: Local Operation 3

Back to our problem, we can write



Entry k of \mathbf{y} is a linear function of the k -th row of \mathbf{A} , i.e.,

$$y_k = \sum_{n=1}^N x_n \mathbf{A}[k, n]$$

So, we can write

$$\frac{\partial \hat{R}}{\partial \mathbf{A}[j, n]} = \sum_{k=1}^K \frac{\partial \hat{R}}{\partial y_k} \frac{\partial y_k}{\partial \mathbf{A}[j, n]} = \sum_{k=1}^K \frac{\partial \hat{R}}{\partial y_k} \mathbf{1}\{k = j\} x_n = \frac{\partial \hat{R}}{\partial y_j} x_n$$

Backpropagation: Local Operation 3

Let's now put them in a matrix


$$\nabla_{\mathbf{A}} \hat{R} = \begin{bmatrix} \frac{\partial \hat{R}}{\partial y_1} x_1 & \dots & \frac{\partial \hat{R}}{\partial y_1} x_N \\ \vdots & & \vdots \\ \frac{\partial \hat{R}}{\partial y_K} x_1 & \dots & \frac{\partial \hat{R}}{\partial y_K} x_N \end{bmatrix} = \nabla_{\mathbf{y}} \hat{R} x^T$$

So, we should now apply outer product! This again makes sense!

We have a K -dimensional gradient $\nabla_{\mathbf{y}} \hat{R}$ and an N -dimensional vector x , we need an outer product to convert it into the $K \times N$ matrix $\nabla_{\mathbf{A}} \hat{R}$

Backpropagation: Local Operation 3

So, we could conclude




with the backward step

$$\nabla_A \hat{R} = \nabla_y \hat{R} x^T$$


Now, we are ready to “**backpropagate**” over an FNN

Backpropagation: Algorithm


Let's recall the compact diagram of an FNN with L hidden layers



We can easily expand it into a *computation graph*




Our objective is *the empirical risk*; so let's include it also in the graph



Backpropagation: Algorithm

Given data-point \mathbf{x} and its true label v , we once complete a forward pass




At the end of forward pass,

we know the value of all variables, i.e., \mathbf{z}_ℓ and \mathbf{y}_ℓ for all ℓ

Now, let's assume we want to find $\nabla_{\mathbf{W}_{L+1}} \hat{R}$

Backpropagation: Algorithm

We now cut the graph at the link \mathbf{W}_{L+1}



Let's recall ...

- + what is the variable here?
- It's \mathbf{W}_{L+1}
- + Can we modify the graph such that it becomes a node?
- Sure! We note that $\mathbf{z}_{L+1} = \mathbf{W}_{L+1}\mathbf{y}_L$. We can look at it as a linear **matrix**-to-vector operation; so, we could modify the graph as

Backpropagation: Algorithm


Let's now move backward to \mathbf{W}_{L+1}

- ① We have \mathbf{y}_{L+1} , so we compute $\nabla_{\mathbf{y}_{L+1}} \hat{R}$, and pass it
- ② We have \mathbf{z}_{L+1} , so we compute $\dot{f}_{L+1}(\mathbf{z}_{L+1})$, and then we get

$$\nabla_{\mathbf{z}_{L+1}} \hat{R} = \nabla_{\mathbf{y}_{L+1}} \hat{R} \odot \dot{f}_{L+1}(\mathbf{z}_{L+1})$$

- ③ We have \mathbf{y}_L , so we compute $\nabla_{\mathbf{W}_{L+1}} \hat{R}$ from the last pass as

$$\nabla_{\mathbf{W}_{L+1}} \hat{R} = \nabla_{\mathbf{z}_{L+1}} \hat{R} \mathbf{y}_L^\top$$




Backpropagation: Algorithm

We can *propagate* backward deeper and deeper

- We cut at the link that *we want to compute the gradient with respect to*
- We exchange the linear vector-to-vector function *at that particular link* to a linear *matrix*-to-vector function
- We move backwards till we get to the source of this graph

Let's see the example for \mathbf{W}_L



Backpropagation: Algorithm


We already have passed backward messages till \mathbf{z}_{L+1} : we have $\nabla_{\mathbf{z}_{L+1}} \hat{R}$

① We now pass $\nabla_{\mathbf{z}_{L+1}} \hat{R}$ to \mathbf{y}_L : $\nabla_{\mathbf{y}_L} \hat{R} = \mathbf{W}_{L+1}^\top \nabla_{\mathbf{z}_{L+1}} \hat{R}$

② We then pass $\nabla_{\mathbf{y}_L} \hat{R}$ to \mathbf{z}_L : $\nabla_{\mathbf{z}_L} \hat{R} = \nabla_{\mathbf{y}_L} \hat{R} \odot \dot{f}_L$

↳ We should remove entry of $\nabla_{\mathbf{y}_L} \hat{R}$ at index 0: we don't want $\partial \hat{R} / \partial \mathbf{y}_L[0]$

③ We finally pass $\nabla_{\mathbf{z}_L} \hat{R}$ to \mathbf{W}_L : $\nabla_{\mathbf{W}_L} \hat{R} = \nabla_{\mathbf{z}_L} \hat{R} \mathbf{y}_{L-1}^\top$



Backpropagation: Algorithm


As we arrive backward at layer ℓ , we already have messages till $\mathbf{z}_{\ell+1}$

① We pass $\nabla_{\mathbf{z}_{\ell+1}} \hat{R}$ to \mathbf{y}_ℓ : $\nabla_{\mathbf{y}_\ell} \hat{R} = \mathbf{W}_{\ell+1}^\top \nabla_{\mathbf{z}_{\ell+1}} \hat{R}$

② We then pass $\nabla_{\mathbf{y}_\ell} \hat{R}$ to \mathbf{z}_ℓ : $\nabla_{\mathbf{z}_\ell} \hat{R} = \nabla_{\mathbf{y}_\ell} \hat{R} \odot \dot{f}_\ell$

↳ We should remove entry of $\nabla_{\mathbf{y}_\ell} \hat{R}$ at index 0: [we don't want $\partial \hat{R} / \partial \mathbf{y}_\ell[0]$]

③ We finally pass $\nabla_{\mathbf{z}_\ell} \hat{R}$ to \mathbf{W}_ℓ : $\nabla_{\mathbf{W}_\ell} \hat{R} = \nabla_{\mathbf{z}_\ell} \hat{R} \mathbf{y}_{\ell-1}^\top$



Once we propagate back to the input, i.e., $\ell = 1$, then we have all gradients!

Backpropagation: Few Notations

To formally present backpropagation, let us define a few notations

For $\ell = 1, \dots, L + 1$, we define

$$\overleftarrow{\mathbf{y}}_\ell = \nabla_{\mathbf{y}_\ell} \hat{R}$$

$$\overleftarrow{\mathbf{z}}_\ell = \nabla_{\mathbf{z}_\ell} \hat{R}$$

and keep in mind that


- \mathbf{y}_ℓ and $\overleftarrow{\mathbf{y}}_\ell$ are totally different things
- \mathbf{z}_ℓ and $\overleftarrow{\mathbf{z}}_\ell$ are totally different things

Backpropagation: Pseudo Code

```

1: Initiate with  $\overleftarrow{\mathbf{y}}_{L+1} = \nabla \mathcal{L}(\mathbf{y}_{L+1}, \mathbf{v})$  and  $\overleftarrow{\mathbf{z}}_{L+1} = \overleftarrow{\mathbf{y}}_{L+1} \odot \dot{f}_{L+1}(\mathbf{z}_{L+1})$ 
2: for  $\ell = L, \dots, 1$  do
3:   Determine  $\overleftarrow{\mathbf{y}}_\ell = \mathbf{W}_{\ell+1}^\top \overleftarrow{\mathbf{z}}_{\ell+1}$  and drop  $\overleftarrow{\mathbf{y}}_\ell[0]$  # backward affine
4:   Determine  $\overleftarrow{\mathbf{z}}_\ell = \dot{f}_\ell(\mathbf{z}_\ell) \odot \overleftarrow{\mathbf{y}}_\ell$  # backward activation
5: end for
6: for  $\ell = 1, \dots, L+1$  do
7:   Return  $\nabla_{\mathbf{w}_\ell} \hat{R} = \overleftarrow{\mathbf{z}}_\ell \mathbf{y}_{\ell-1}^\top$ 
8: end for

```



Backpropagation: Pseudo Code

```

1: Initiate with  $\overleftarrow{\mathbf{y}}_{L+1} = \nabla \mathcal{L}(\mathbf{y}_{L+1}, \mathbf{v})$  and  $\overleftarrow{\mathbf{z}}_{L+1} = \overleftarrow{\mathbf{y}}_{L+1} \odot \dot{f}_{L+1}(\mathbf{z}_{L+1})$ 
2: for  $\ell = L, \dots, 1$  do
3:   Determine  $\overleftarrow{\mathbf{y}}_\ell = \mathbf{W}_{\ell+1}^\top \overleftarrow{\mathbf{z}}_{\ell+1}$  and drop  $\overleftarrow{\mathbf{y}}_\ell[0]$            # backward affine
4:   Determine  $\overleftarrow{\mathbf{z}}_\ell = \dot{f}_\ell(\mathbf{z}_\ell) \odot \overleftarrow{\mathbf{y}}_\ell$                   # backward activation
5: end for
6: for  $\ell = 1, \dots, L+1$  do
7:   Return  $\nabla_{\mathbf{w}_\ell} \hat{R} = \overleftarrow{\mathbf{z}}_\ell \mathbf{y}_{\ell-1}^\top$ 
8: end for

```

- + This looks very similar to forward propagation! Right?!
- Yeah! Just we go **backward**! That's the whole point of backpropagation

You need to go once **forth** and then **back** to determine all gradients

Let's put them next to each other

Backpropagation: *Pseudo Code*

`ForwardProp():`

- ```

1: Initiate with $\mathbf{y}_0 = \mathbf{x}$
2: for $\ell = 0, \dots, L$ do
3: Add $y_\ell[0] = 1$ and determine $\mathbf{z}_{\ell+1} = \mathbf{W}_{\ell+1}\mathbf{y}_\ell$ # forward affine
4: Determine $\mathbf{y}_{\ell+1} = f_{\ell+1}(\mathbf{z}_{\ell+1})$ # forward activation
5: end for
6: for $\ell = 1, \dots, L + 1$ do
7: Return \mathbf{y}_ℓ and \mathbf{z}_ℓ
8: end for

```

BackProp() :

- ```

1: Initiate with  $\widehat{\mathbf{y}}_{L+1} = \nabla \mathcal{L}(\mathbf{y}_{L+1}, \mathbf{v})$  and  $\widehat{\mathbf{z}}_{L+1} = \widehat{\mathbf{y}}_{L+1} \odot \dot{f}_{L+1}(\mathbf{z}_{L+1})$ 
2: for  $\ell = L, \dots, 1$  do
3:   Determine  $\widehat{\mathbf{y}}_\ell = \mathbf{W}_{\ell+1}^\top \widehat{\mathbf{z}}_{\ell+1}$  and drop  $\widehat{\mathbf{y}}_\ell[0]$            # backward affine
4:   Determine  $\widehat{\mathbf{z}}_\ell = \dot{f}_\ell(\mathbf{z}_\ell) \odot \widehat{\mathbf{y}}_\ell$                       # backward activation
5: end for
6: for  $\ell = 1, \dots, L+1$  do
7:   Return  $\nabla \mathbf{w}_\ell \hat{R} = \widehat{\mathbf{z}}_\ell \mathbf{y}_{\ell-1}^\top$ 
8: end for

```

Complete Training Loop via Gradient Descent

- + Say we use backpropagation; then, how does **gradient descent** look?
 - Well! We should go back and forth with all data-points

Say we have dataset

$$\mathbb{D} = \{(\textcolor{blue}{x}_b, \textcolor{red}{v}_b) \text{ for } b = 1, \dots, B\}$$

`GradientDescent():`

- ```

1: Initiate with some initial values $\{\mathbf{W}_\ell^{(0)}\}$ and set a learning rate η
2: while weights not converged do
3: for $b = 1, \dots, B$ do
4: NN.values \leftarrow ForwardProp ($\mathbf{x}_b, \{\mathbf{W}_\ell^{(t)}\}$) # forward
5: $\{\nabla_{\mathbf{W}_\ell^{(t)}} \hat{R}_b\} \leftarrow$ BackProp ($\mathbf{x}_b, \mathbf{v}_b, \{\mathbf{W}_\ell^{(t)}\}, \text{NN.values}$) # backward
6: end for
7: for $\ell = 1, \dots, L + 1$ do
8: $\mathbf{W}_\ell^{(t+1)} \leftarrow \mathbf{W}_\ell^{(t)} - \eta \text{ mean}(\nabla_{\mathbf{W}_\ell^{(t)}} \hat{R}_1, \dots, \nabla_{\mathbf{W}_\ell^{(t)}} \hat{R}_B)$
9: end for
10: end while

```