

# ECE 1508: Applied Deep Learning

## Chapter 2: Feedforward Neural Networks

Ali Bereyhi

[ali.bereyhi@utoronto.ca](mailto:ali.bereyhi@utoronto.ca)

Department of Electrical and Computer Engineering  
University of Toronto

Fall 2025

# Binary Classification via FNN

Let's now design a deep FNN for *binary classification*

We have a set of *images of hand-written numbers*, something like this<sup>a</sup>

A 10x10 grid of handwritten digits from 0 to 9. Each digit is represented by a different stroke pattern. The digits are arranged in a 10x10 grid, with each digit appearing once in every row and column.

0	0	0	0	0	0	0	0	0	0
1	1	1	1	1	1	1	1	1	1
2	2	2	2	2	2	2	2	2	2
3	3	3	3	3	3	3	3	3	3
4	4	4	4	4	4	4	4	4	4
5	5	5	5	5	5	5	5	5	5
6	6	6	6	6	6	6	6	6	6
7	7	7	7	7	7	7	7	7	7
8	8	8	8	8	8	8	8	8	8
9	9	9	9	9	9	9	9	9	9

We intend to train a fully-connected FNN that given a new hand-written image, it finds out whether it is "2" or not

---

<sup>a</sup> Source: Wikipedia

This is a *binary classification!*

## Binary Classification via FNN: Data

Let's get clear about the data: *our dataset looks like*

$$\mathbb{D} = \{(\mathbf{x}_b, v_b) \text{ for } b = 1, \dots, B\}$$

where in this set each component is defined as follows:

- $B$  is the number of images we have
  - ↳ we also call the *set* of images: a *batch* of images
- $\mathbf{x}_b \in \mathbb{R}^N$  is the pixel vector of image  $b$ 
  - ↳  $N$  is the number of pixels in the image
- $v_b \in \{0, 1\}$  is a binary label indicating whether it is "2" or not
  - ↳ if the image is a hand-written "2" we set  $v_b = 1$
  - ↳ if the image is not a hand-written "2" we set  $v_b = 0$

# Binary Classification via FNN: Model

Let's now set the model: we use a fully-connected FNN

What are the **hyperparameters**?

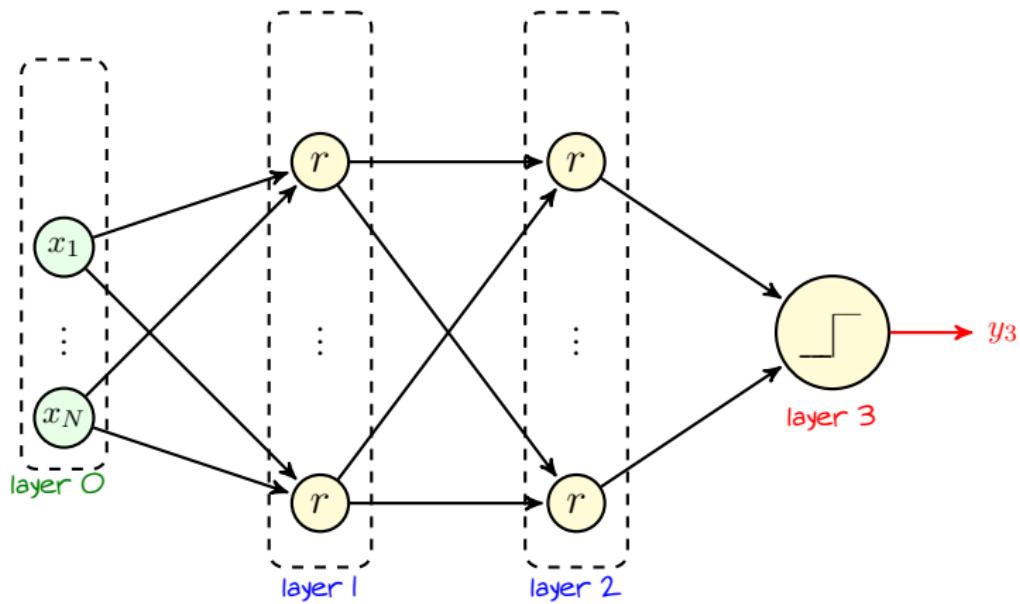
- We want it to be **deep**; so, we consider **2 hidden layers**
  - ↳ the **depth** is hence **3**
- We specify the width of each hidden layer
  - ↳ first hidden layer has **width  $K$**
  - ↳ second hidden layer has **width  $J$**
- All hidden neurons use **ReLU activation**
  - ↳  $f_1(\cdot) = f_2(\cdot) = \text{ReLU}(\cdot)$ : let's **show ReLU by  $r$** , i.e.,

$$r(x) = \text{ReLU}(x)$$

- **Output layer** has a single **perceptron**

We can now write down the model!

# Binary Classification via FNN: Model



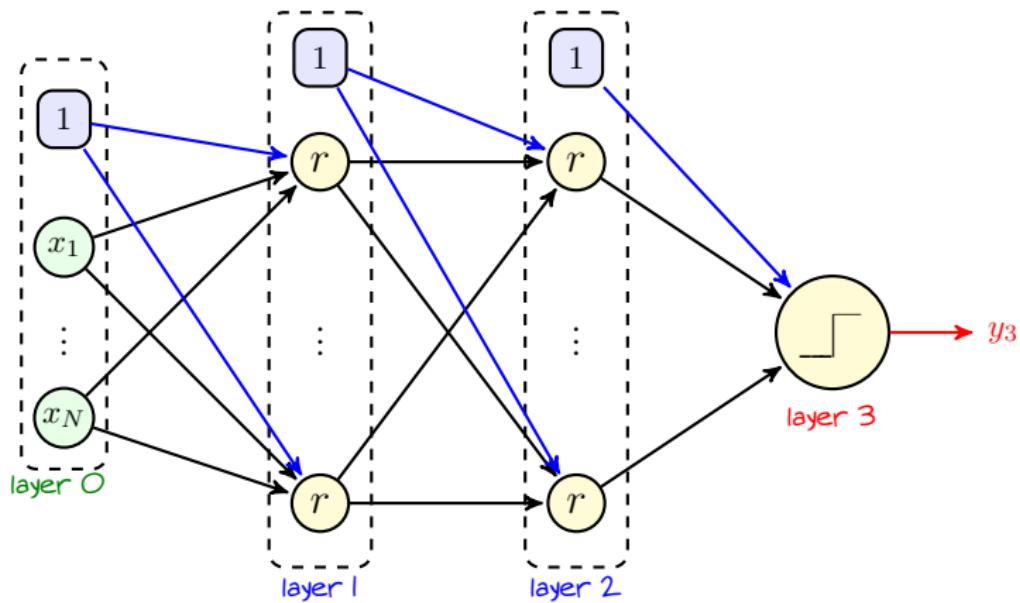
# Binary Classification via FNN: Model

Let's now set the model: we use a fully-connected FNN

What are the learnable parameters?

- Layer 1 has  $(N + 1)K$  links
  - ↳  $NK$  of them are weights
  - ↳  $K$  of them are biases  $\equiv$  weights of dummy node  $x_0 = 1$
- Layer 2 has  $(K + 1)J$  links
  - ↳  $KJ$  of them are weights
  - ↳  $J$  of them are biases  $\equiv$  weights of dummy node  $y_1[0] = 1$
- Output layer has  $J + 1$  links
  - ↳  $J$  of them are weights
  - ↳ one is bias  $\equiv$  weight of dummy node  $y_2[0] = 1$

# Binary Classification via FNN: Model



## Binary Classification via FNN: Loss

How to calculate the loss? Let's do what we did before

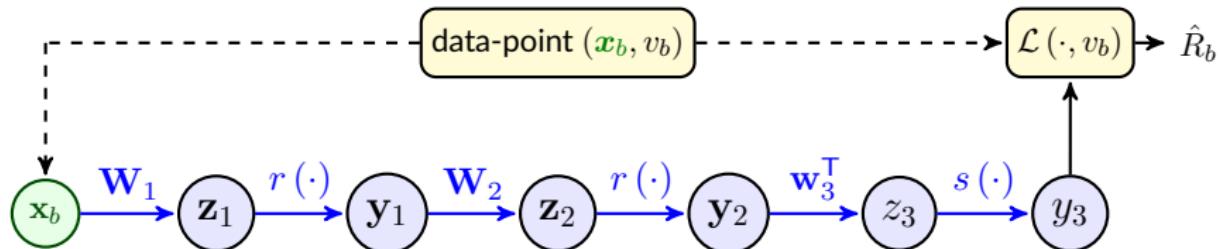
We use the error indicator as the loss function

$$\mathcal{L}(y_b, v_b) = \mathbb{1}\{y_b \neq v_b\} = \begin{cases} 1 & y_b \neq v_b \\ 0 & y_b = v_b \end{cases}$$

- + Wait a moment! Didn't you say that this was a **bad choice?**
- Yeah! So said I also for the **perceptron's activation!** Let's try it out to find out really **why they are bad!** We should be able to understand it now

# Binary Classification via FNN: Training

Let's look at the computation graph: for a given data-point  $(\mathbf{x}_b, v_b)$ , we have

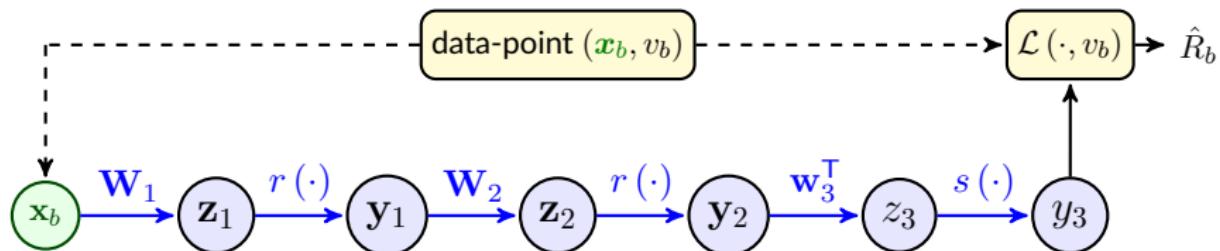


Here, we have 3 *linear operations*

- First operation is  $\mathbf{z}_1 = \mathbf{W}_1 \mathbf{x}_b$  with  $\mathbf{W}_1 \in \mathbb{R}^{K \times (N+1)}$   
 ↳ first column of  $\mathbf{W}_1$  is *bias* and the remaining columns are *weights*
- Second operation is  $\mathbf{z}_2 = \mathbf{W}_2 \mathbf{y}_1$  with  $\mathbf{W}_2 \in \mathbb{R}^{J \times (K+1)}$   
 ↳ first column of  $\mathbf{W}_2$  is *bias* and the remaining columns are *weights*
- Last operation is  $\mathbf{z}_3 = \mathbf{w}_3^T \mathbf{x}$  with  $\mathbf{w}_3 \in \mathbb{R}^{J+1}$   
 ↳ first entry of  $\mathbf{w}_3$  is *bias* and the remaining entries are *weights*

## Binary Classification via FNN: Training

Let's look at the computation graph: *for a given data-point  $(\mathbf{x}_b, v_b)$ , we have*



We have 3 *functional* operations

- The first two are  $\mathbf{y}_1 = r(\mathbf{z}_1)$  and  $\mathbf{y}_2 = r(\mathbf{z}_2)$
- The last one is  $y_3 = s(z_3)$ , and recall that  $s(\cdot)$  is the **step function**

$$y_3 = s(z_3) = \begin{cases} 1 & z_3 \geq 0 \\ 0 & z_3 < 0 \end{cases}$$

# Binary Classification via FNN: *Training*

Let's write **gradient descent** for training of our model

**GradientDescent()**:

- 1: Initiate with some initial values  $\{\mathbf{W}_1^{(0)}, \mathbf{W}_2^{(0)}, \mathbf{w}_3^{(0)}\}$  and set a learning rate  $\eta$
- 2: **while** weights not converged **do**
- 3:   **for**  $b = 1, \dots, B$  **do**
- 4:     NN.values  $\leftarrow$  ForwardProp ( $\mathbf{x}_b$ ,  $\{\mathbf{W}_1^{(t)}, \mathbf{W}_2^{(t)}, \mathbf{w}_3^{(t)}\}$ )
- 5:      $\{\mathbf{G}_{1,b}, \mathbf{G}_{2,b}, \mathbf{g}_{3,b}\} \leftarrow$  BackProp ( $\mathbf{x}_b$ ,  $\mathbf{v}_b$ ,  $\{\mathbf{W}_1^{(t)}, \mathbf{W}_2^{(t)}, \mathbf{w}_3^{(t)}\}$ , NN.values)
- 6:   **end for**
- 7:   Update

$$\mathbf{W}_1^{(t+1)} \leftarrow \mathbf{W}_1^{(t)} - \eta \text{ mean}(\mathbf{G}_{1,1}, \dots, \mathbf{G}_{1,B})$$

$$\mathbf{W}_2^{(t+1)} \leftarrow \mathbf{W}_2^{(t)} - \eta \text{ mean}(\mathbf{G}_{2,1}, \dots, \mathbf{G}_{2,B})$$

$$\mathbf{w}_3^{(t+1)} \leftarrow \mathbf{w}_3^{(t)} - \eta \text{ mean}(\mathbf{g}_{3,1}, \dots, \mathbf{g}_{3,B})$$

- 8: **end while**

Let's look at **forward** and **backward** propagation!

## Binary Classification via FNN: Forward Pass

Forward pass is very straightforward: say we are at iteration  $t$

- ① For each pixel vector  $x_b$ , we determine  $z_1$  as

$$\mathbf{x} \leftarrow \begin{bmatrix} 1 \\ \mathbf{x}_b \end{bmatrix} \rightsquigarrow \mathbf{z}_1 = \mathbf{W}_1^{(t)} \mathbf{x}$$

The output of first layer is then given by  $y_1 = r(z_1)$ :  $r(\cdot)$  is ReLU, so

we keep positive entries of  $z_1$  and replace negative ones with zero

- ② We add 1 at index 0 of  $y_1$  and determine  $z_2$  as

$$\mathbf{y}_1 \leftarrow \begin{bmatrix} 1 \\ \mathbf{y}_1 \end{bmatrix} \rightsquigarrow \mathbf{z}_2 = \mathbf{W}_2^{(t)} \mathbf{y}_1$$

The output of second layer is given by  $y_2 = r(z_2)$

## Binary Classification via FNN: Forward Pass

- ③ We add 1 at index 0 of  $\mathbf{y}_2$  and determine  $z_3$  as

$$\mathbf{y}_2 \leftarrow \begin{bmatrix} 1 \\ \mathbf{y}_2 \end{bmatrix} \rightsquigarrow z_3 = \mathbf{w}_3^{(t)T} \mathbf{y}_2$$

The network output is given by  $y_3 = s(z_3)$ : *s(·) is step function, so*

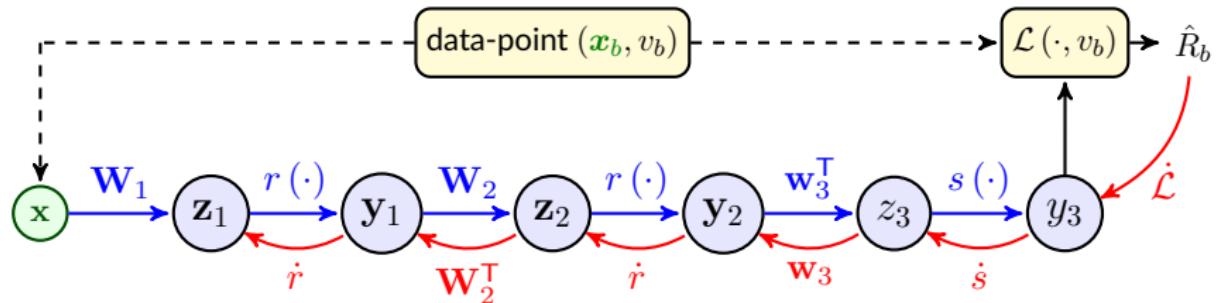
*it's 0 if  $z_3$  is negative, and 1 if it is not negative*

At this point, we have all that we need, i.e.,

$\mathbf{x}, \mathbf{z}_1, \mathbf{y}_1, \mathbf{z}_2, \mathbf{y}_2, z_3$  and  $y_3$

# Binary Classification via FNN: Training

How does the graph look like on the backward pass?



Let's move backward!

# Binary Classification via FNN: Backward Pass

We know all the derivatives, i.e.,

$$\dot{\mathcal{L}}(\textcolor{red}{y}, v_b) = \frac{d}{dy} \mathbb{1}\{\textcolor{red}{y} \neq v_b\} \quad \dot{s}(z) = \frac{d}{dz} s(z) \quad \dot{r}(z) = \frac{d}{dz} r(z)$$

For backward pass we start at node  $\textcolor{red}{y}_3$ :

- ① We find derivative w.r.t. output  $\overleftarrow{y}_3 = \dot{\mathcal{L}}(\textcolor{red}{y}_3, v_b)$  and set

$$\overleftarrow{z}_3 = \overleftarrow{y}_3 \dot{s}(z_3)$$

- ② We compute  $\overleftarrow{\mathbf{y}}_2 = \mathbf{w}_3 \overleftarrow{\mathbf{z}}_3$  and drop its first entry; then, compute

$$\overleftarrow{\mathbf{y}}_2 \leftarrow \begin{bmatrix} \overleftarrow{\mathbf{y}}_2[0] \\ \overleftarrow{\mathbf{y}}_2[1:] \end{bmatrix} \rightsquigarrow \overleftarrow{\mathbf{z}}_2 = \dot{r}(\mathbf{z}_2) \odot \overleftarrow{\mathbf{y}}_2$$

## Binary Classification via FNN: Backward Pass

- ③ We compute  $\widehat{\mathbf{y}}_1 = \mathbf{W}_2^T \widehat{\mathbf{z}}_2$  and drop its first entry; then, compute

$$\widehat{\mathbf{y}}_1 \leftarrow \begin{bmatrix} \widehat{\mathbf{y}}_1[0] \\ \widehat{\mathbf{y}}_1[1:] \end{bmatrix} \rightsquigarrow \widehat{\mathbf{z}}_1 = \dot{r}(\mathbf{z}_1) \odot \widehat{\mathbf{y}}_1$$

At this point, we can calculate all gradients

$$\mathbf{G}_{1,b} = \nabla_{\mathbf{W}_1} \hat{R}_b = \widehat{\mathbf{z}}_1 \mathbf{y}_0^T = \widehat{\mathbf{z}}_1 \mathbf{x}^T$$

$$\mathbf{G}_{2,b} = \nabla_{\mathbf{W}_2} \hat{R}_b = \widehat{\mathbf{z}}_2 \mathbf{y}_1^T$$

$$\mathbf{g}_{3,b}^T = \nabla_{\mathbf{w}_3^T} \hat{R}_b = \widehat{\mathbf{z}}_3 \mathbf{y}_2^T$$

All done! We repeat it for every image in the batch and then average gradients. Finally, we move one step in gradient descent and find the weights of the next iteration

## Binary Classification via FNN: Differentiability Issue

- + Where is then the issue with *perceptron* and *indicator error*?
- $\dot{\mathcal{L}}(\mathbf{y}, v_b)$  and  $\dot{s}(z)$  are not well-defined!
  - ↳ Recall that they are *discontinuous*

---

In fact, the empirical risk is not a *smooth function* of the weights and biases; therefore, using *gradient descent* we do not end up with a well-trained network

- + How can we get over it?
- Well! There is a very well-established trick!

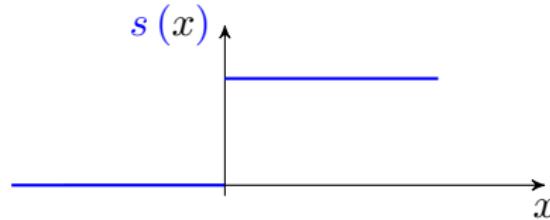
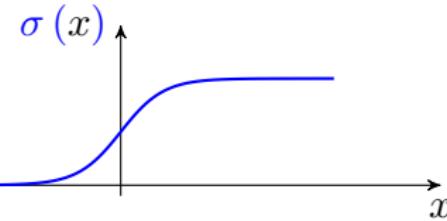
## Binary Classification via FNN: Differentiability Issue

We first replace the **perceptron** with a neuron whose **activation** is a good approximation of *step function* and **differentiable**<sup>1</sup>

*We already have seen the **sigmoid** function*

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

*which looks pretty close to step function*



Using sigmoid instead of step function **resolves** the **differentiability** issue

<sup>1</sup>Or at least, we can easily calculate a sub-gradient for it

## Binary Classification via FNN: Differentiability Issue

But, replacing *perceptron* by *sigmoid-activated neuron* makes a new problem

The *output* of the network is now *not binary*!

How can we address this problem?

We now interpret the *output* as *probability*, i.e.,

$y_3$  is the *probability* of the *label being 1*

- + OK! But how can we define the *loss* now?
- Well! We could look at the *true label* from the same point of view

Say  $v \in \{0, 1\}$  is *true label*: if  $v = 1$  then the *true label is 1* with probability 1; if  $v = 0$  then the *true label is 1* with probability 0. So, we could say

*the true label is 1 with probability  $v$*

## Binary Classification via FNN: Differentiability Issue

true label is 1 with probability  $v \leftrightarrow y_3$  is probability of the label being 1

Apparently,  $v$  and  $y_3$  are of the same nature: we can still define a loss that evaluates the difference between  $y_3$  and  $v$

- + What should be the loss then?
- Definitely not the indicator error!

Indicator error is not suitable because

- ① we already know that it is not differentiable
- ② more importantly, with sigmoid activation becomes useless

$$\mathbb{1}\{\sigma(z_3) \neq 1\} = \mathbb{1}\{\sigma(z_3) \neq 0\} = 1$$

## Binary Classification via FNN: MSE

*One may suggest that we use the squared error, i.e.,*

$$\mathcal{L}(y_3, v) = (y_3 - v)^2$$

*in this case the empirical risk is called*

*Mean Squared Error (MSE)*

This loss is differentiable

$$\dot{\mathcal{L}}(y_3, v) = 2(y_3 - v)$$

and proportional to the distance between  $y_3$  and  $v$

*It's a good choice but not best*

## Binary Classification via FNN: Cross-Entropy

A better choice is to determine the *cross-entropy loss*

$$\begin{aligned}\mathcal{L}(y_3, v) &= \text{CE}(y_3, v) = -v \log y_3 - (1-v) \log(1-y_3) \\ &= \begin{cases} \log \frac{1}{y_3} & v = 1 \\ \log \frac{1}{1-y_3} & v = 0 \end{cases}\end{aligned}$$

This loss function is sometimes *wrongly* called *KL-divergence*: it is proportional to the *Kullback-Leibler divergence* but it's *different*

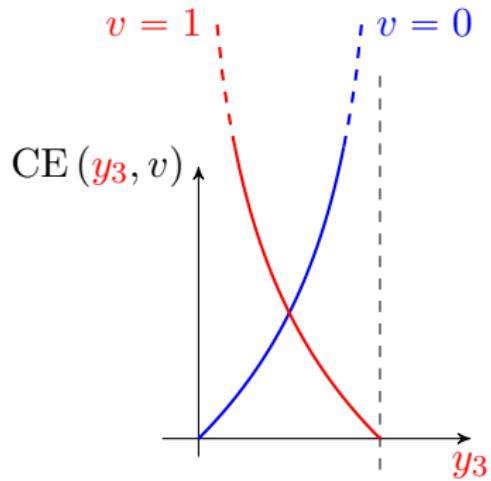
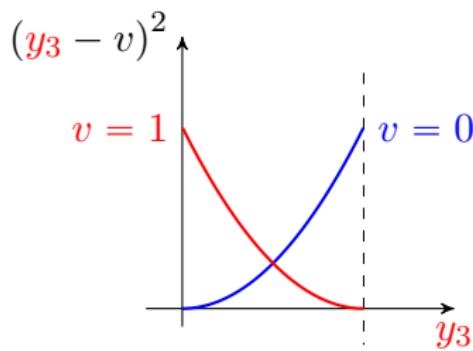
This loss is again *differentiable*

$$\dot{\mathcal{L}}(y_3, v) = \dot{\text{CE}}(y_3, v) = -\frac{v}{y_3} + \frac{1-v}{1-y_3}$$

**Note:** The logarithm is usually in natural base, i.e.,  $\log x = \ln x$

## Binary Classification via FNN: Cross-Entropy

- + But why cross-entropy is a better loss?
- It pushes  $y_3$  more towards the edges of interval  $[0, 1]$

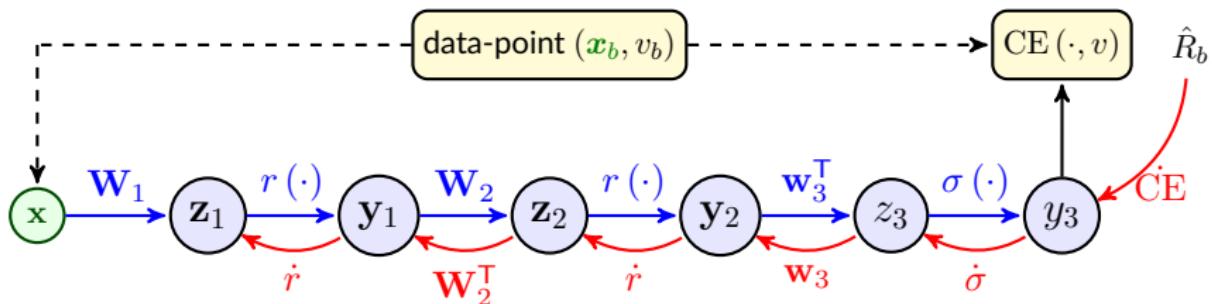


Cross entropy returns much higher loss when  $y_3$  is different from  $v$

# Binary Classification via FNN: Training with Cross-Entropy

- + What changes in the training loop in this case?
- Pretty much *nothing!* Just replace

- $\mathcal{L}(y_3, v)$  with CE( $y_3, v$ )
- $\dot{\mathcal{L}}(y_3, v)$  with  $\dot{\text{CE}}(y_3, v)$
- $s(z_3)$  with  $\sigma(z_3)$
- $\dot{s}(z_3)$  with  $\dot{\sigma}(z_3)$



## Binary Classification via FNN: Training with Cross-Entropy

- + How do we use the output of network then, when we give a new **image** to it for classification? **It's not binary!**
- Just follow the **interpretation**

$y_3$  gives the **probability** of the **image** being **hand-written "2"**; therefore,  $(1 - y_3)$  gives the **probability** of image being **any other hand-written number**. So, we select the outcome with **higher chance**, i.e.,

- if  $y_3 \geq 0.5$ , we label the new **image** as a hand-written "2"
- if  $y_3 < 0.5$ , we label the new **image** as **not being** a hand-written "2"

- + Can't we classify more classes? Like hand-written "0", "1", ..., "9"?
- Now that we have this nice interpretation: Yes! We can!

# Multiclass Classification

We initially saw that any multiclass classification can be seen as a **sequence of binary classifications**; however, for that, we need **multiple NNs!**

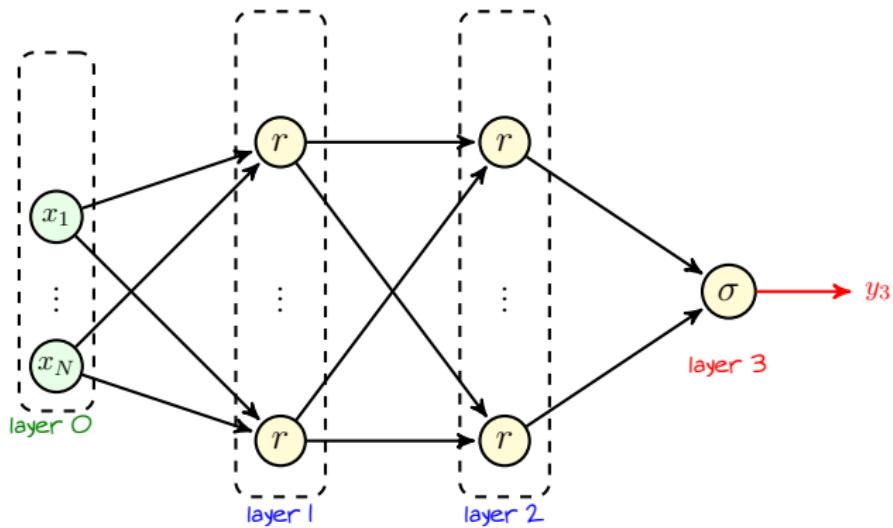
- + Why not follow the same idea and determine the probability of input belonging to each class?
- Yes! That's actually the effective way!

Let's get back to our **image recognition**, but now with **multiple** classes!

We have **images** of **hand-written numbers from "0" to "9"** and want to train a NN that **recognizes any hand-written number**

We first draw our earlier FNN

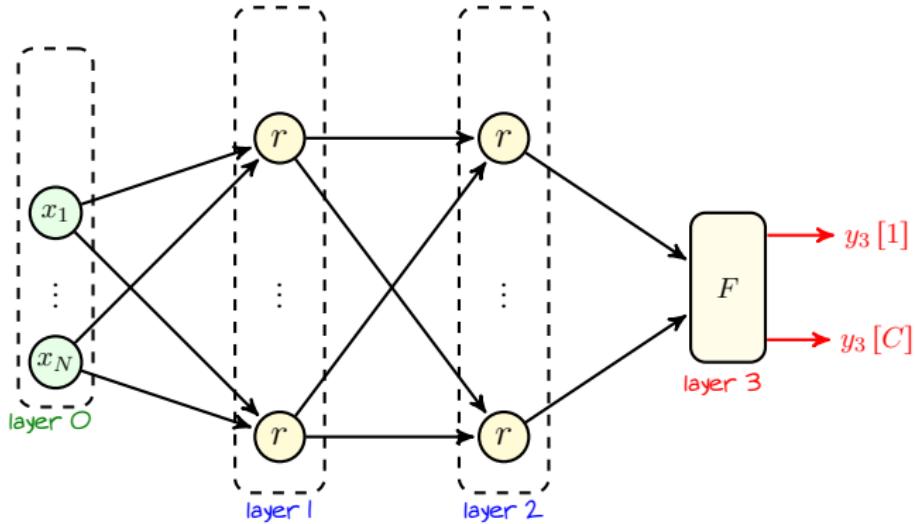
# Multiclass Classification via FNN



In this FNN,  $y_3$  is interpreted as a probability of label being 1  
probability of label being 0 is hence  $1 - y_3$

This was done by a standard **single-output neuron**, since we had only **2 classes**

# Multiclass Classification via FNN



With  $C$  classes, we need a module that computes probabilities of all  $C$  classes

this module can be seen as a **neuron** with **vector output**

# Multiclass Classification: Vector-Activated Neuron

## Vector-Activated Neuron

A **vector-activated neuron** is an artificial neuron with multivariate **activation function**: let  $x \in \mathbb{R}^N$  be the input to this neuron and  $C$  be its output dimension; then, the output vector  $y \in \mathbb{R}^C$  is given by

$$y = F(\tilde{\mathbf{W}}x + \mathbf{b})$$

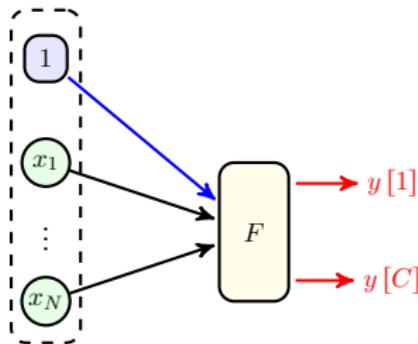
for weight matrix  $\tilde{\mathbf{W}} \in \mathbb{R}^{C \times N}$ , bias  $\mathbf{b} \in \mathbb{R}^C$  and activation  $F(\cdot) : \mathbb{R}^C \mapsto \mathbb{R}^C$

First thing first: let's get rid of the bias before we go on

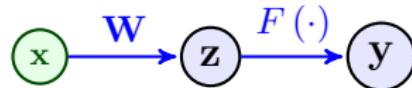
$$y = F\left(\begin{bmatrix} \mathbf{b} & \tilde{\mathbf{W}} \end{bmatrix} \begin{bmatrix} 1 \\ x \end{bmatrix}\right) = F(\mathbf{W}x)$$

So, we keep on with our dummy 1 input here as well

# Multiclass Classification: Vector-Activated Neuron



Next, let's see how its computation graph looks

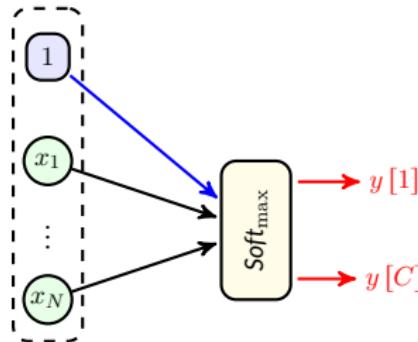


This looks exactly like a standard layer with a **minor difference**

$F(\cdot)$  does **not necessarily** perform **entry-wise**

# Multiclass Classification: Softmax

A very well-known example of **vector activation** is *softmax*



## Softmax Function

For  $\mathbf{z} \in \mathbb{R}^C$ , softmax function returns  $\text{Soft}_{\max}(\mathbf{z}) = \mathbf{y} \in \mathbb{R}^C$  whose entry  $c$  is

$$y[c] = \frac{e^{z[c]}}{\sum_{j=1}^C e^{z[j]}}$$

## Multiclass Classification: Softmax

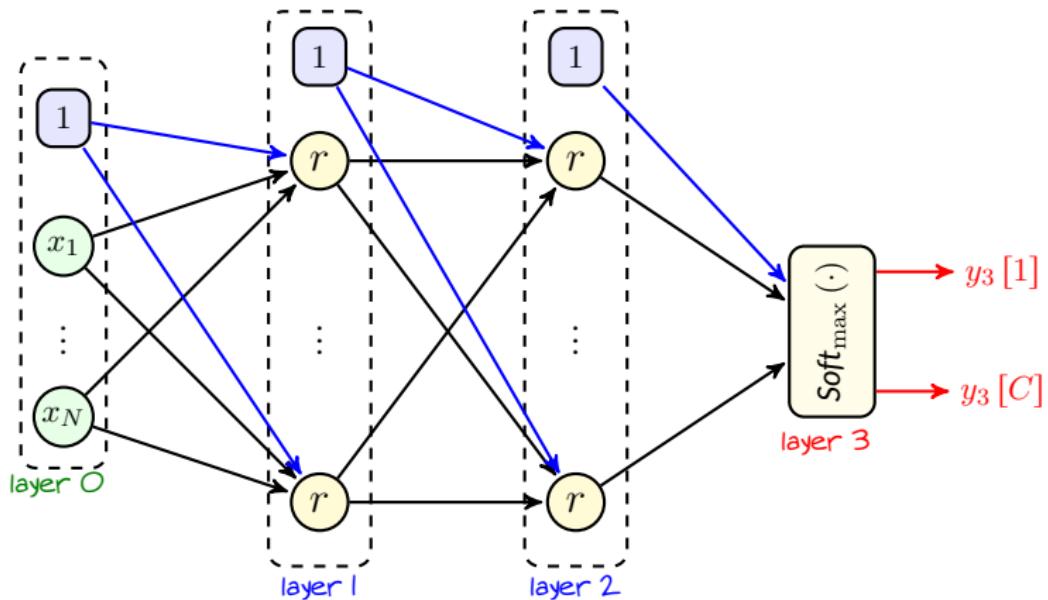
Softmax always returns a probability distribution on the set of classes

$$\sum_{c=1}^C y[c] = \sum_{c=1}^C \frac{e^{z[c]}}{\sum_{j=1}^C e^{z[j]}} = \frac{\sum_{c=1}^C e^{z[c]}}{\sum_{j=1}^C e^{z[j]}} = 1$$

We can hence use it to extend our FNN to a **multiclass classifier**

We replace **layer 3** with a softmax-activated multivariate neuron and treat its outcome as the **chance of input belonging to each class**; then, we select the class with highest chance

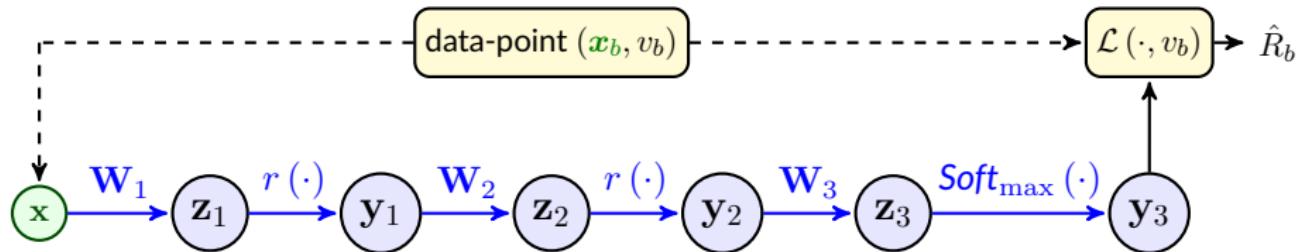
# Multiclass Classification via FNN: Softmax Activation



Let's try again the forward pass

# Multiclass Classification via FNN: Softmax Activation

Let's look at the computation graph: *for a given data-point  $(\mathbf{x}_b, v_b)$ , we have*



Note that **the output layer** has been changed

- We now have a **vector**  $\mathbf{z}_3 \in \mathbb{R}^C$ 
  - ↳ So we have a **matrix of weights**  $\mathbf{W}_3 \in \mathbb{R}^{C \times (J+1)}$
- We now have a **vector**  $\mathbf{y}_3 \in \mathbb{R}^C$
- We get from  $\mathbf{z}_3$  to  $\mathbf{y}_3$  via **softmax**
  - ↳ This is **not** an entry-wise activation anymore!

# Multiclass Classification via FNN: Softmax Activation

- ① For each pixel vector  $x_b$ , we determine  $z_1$  as

$$\mathbf{x} \leftarrow \begin{bmatrix} 1 \\ \mathbf{x}_b \end{bmatrix} \rightsquigarrow \mathbf{z}_1 = \mathbf{W}_1^{(t)} \mathbf{x}$$

The **output** of first layer is then given by  $\mathbf{y}_1 = r(\mathbf{z}_1)$

- ② We add 1 at index 0 of  $\mathbf{y}_1$  and determine  $z_2$  as

$$\mathbf{y}_1 \leftarrow \begin{bmatrix} 1 \\ \mathbf{y}_1 \end{bmatrix} \rightsquigarrow \mathbf{z}_2 = \mathbf{W}_2^{(t)} \mathbf{y}_1$$

The output of second layer is given by  $\mathbf{y}_2 = r(\mathbf{z}_2)$

- ③ We add 1 at index 0 of  $\mathbf{y}_2$  and determine  $z_3$  as

$$\mathbf{y}_2 \leftarrow \begin{bmatrix} 1 \\ \mathbf{y}_2 \end{bmatrix} \rightsquigarrow \mathbf{z}_3 = \mathbf{W}_3^{(t)T} \mathbf{y}_2$$

The network **output** is given by  $\mathbf{y}_3 = \text{Soft}_{\max}(\mathbf{z}_3)$

## Multiclass Classification: Loss

- + How can we define the loss now? On one side we have a vector of probabilities; one the other side an integer label!
- Again we need to convert **true labels** to **probabilities**

Let's say we have  $C$  classes: the vector of probabilities contains  $C$  entries

$$\mathbf{p} = \begin{bmatrix} p_1 \\ \vdots \\ p_C \end{bmatrix} = \begin{bmatrix} \Pr \{ \text{image belongs to class 1} \} \\ \vdots \\ \Pr \{ \text{image belongs to class } C \} \end{bmatrix}$$

If we know that the image  $b$  belongs to class  $v_b$ , we could say that

$$\mathbf{p} \text{ of image } b = \begin{bmatrix} p_1 \\ \vdots \\ p_{v_b} \\ \vdots \\ p_C \end{bmatrix} = \begin{bmatrix} \Pr \{ \text{image } b \text{ belongs to class 1} \} = 0 \\ \vdots \\ \Pr \{ \text{image } b \text{ belongs to class } v_b \} = 1 \\ \vdots \\ \Pr \{ \text{image } b \text{ belongs to class } C \} = 0 \end{bmatrix} = \begin{bmatrix} 0 \\ \vdots \\ 1 \\ \vdots \\ 0 \end{bmatrix}$$

## Multiclass Classification: Loss

So, we could say that *label  $v$*  is corresponding to a vector of size  $C$  whose entry  $v$  is 1 and the remaining entries are all 0: this vector is called a one-hot vector

### One-hot Vector

The one-hot vector  $\mathbf{1}_v \in \{0, 1\}^C$  is a  $C$ -dimensional vector whose entry  $v$  is 1 and all remaining entries are 0

For instance: say  $C = 3$ ; then, we have

$$\mathbf{1}_1 = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} \quad \mathbf{1}_2 = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix} \quad \mathbf{1}_3 = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$$

### Moral of Story

We can interpret true label  $v$  as a probability vector  $\mathbf{1}_v$

## Multiclass Classification: Loss

We can interpret **true label  $v$**  as a probability vector  $\mathbf{1}_v$

Now for **image  $b$**  with **label  $v_b$**  we compare network's output  $\mathbf{y}_3$  to  $\mathbf{1}_{v_b}$

$$\hat{R}_b = \mathcal{L}(\mathbf{y}_3, \mathbf{1}_{v_b})$$

for **loss  $\mathcal{L}(\cdot)$**  that determines distance between two **probability vectors**

- + What kind of **loss functions** do we use usually?
  - Like binary case: squared error is **good**, cross-entropy is the **best**
- + How do we define them in this case?
  - Just extend them to multi-dimensional vectors

## Multiclass Classification: Loss

We can extend squared error to vector form as

$$\begin{aligned}\mathcal{L}(\mathbf{y}, \mathbf{1}_v) &= \|\mathbf{y} - \mathbf{1}_v\|^2 = \sum_{c=1}^C (y[c] - \mathbb{1}_{\{c=v\}})^2 \\ &= \sum_{c=1, c \neq v}^C y[c]^2 + (y[v] - 1)^2\end{aligned}$$

This gradient of this loss is

$$\nabla \mathcal{L}(\mathbf{y}, \mathbf{1}_v) = 2 \begin{bmatrix} y[1] \\ \vdots \\ y[v] - 1 \\ \vdots \\ y[C] \end{bmatrix} = 2(\mathbf{y} - \mathbf{1}_v)$$

## Multiclass Classification: Loss

Cross entropy can also be extended as follows

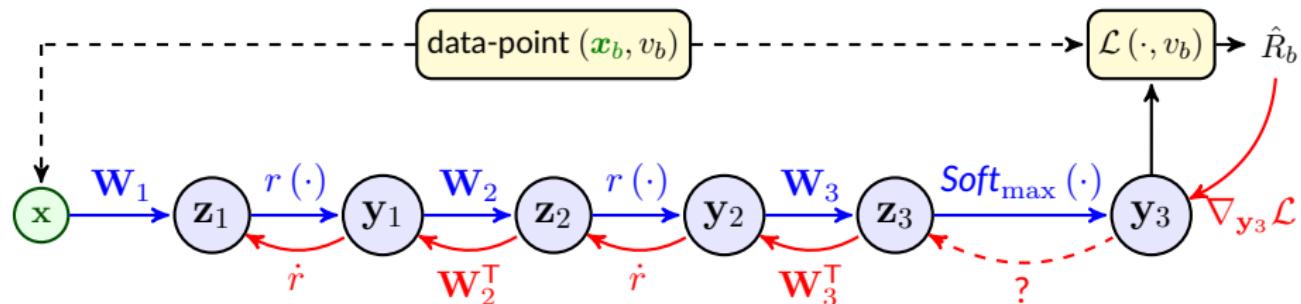
$$\begin{aligned}\mathcal{L}(\mathbf{y}, \mathbf{1}_v) = \text{CE}(\mathbf{y}, \mathbf{1}_v) &= - \sum_{c=1}^C \mathbb{1}\{c = v\} \log \mathbf{y}(c) \\ &= - \log \mathbf{y}[v]\end{aligned}$$

The gradient of this loss is

$$\nabla \mathcal{L}(\mathbf{y}, \mathbf{1}_v) = \nabla \text{CE}(\mathbf{y}, \mathbf{1}_v) = \begin{bmatrix} 0 \\ \vdots \\ -1/\mathbf{y}[v] \\ \vdots \\ 0 \end{bmatrix} = -\frac{1}{\mathbf{y}[v]} \mathbf{1}_v$$

# Backpropagation Through Vector Activation

How can we backpropagate through this neural network?

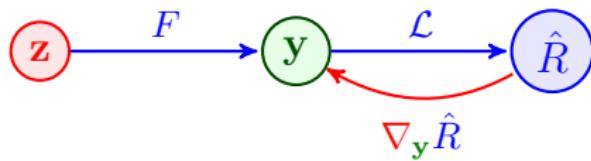


- 1 Compute  $\nabla_{y_3} \hat{R}_b = \nabla_{y_3} \mathcal{L}(y_3, \mathbf{1}_{v_b})$
- 2 Compute  $\nabla_{z_3} \hat{R}_b$  from  $\nabla_{y_3} \hat{R}_b$  by its *backward* link that [we don't know yet]
- 3 Compute  $\nabla_{y_2} \hat{R}_b = \mathbf{W}_3^T \nabla_{z_3} \hat{R}_b$
- 4 Remove entry at index 0 of  $\nabla_{y_2} \hat{R}_b$  and compute  $\nabla_{z_2} \hat{R}_b = \dot{r}(z_2) \odot \nabla_{y_2} \hat{R}_b$
- 5 Compute  $\nabla_{y_1} \hat{R}_b = \mathbf{W}_2^T \nabla_{z_2} \hat{R}_b$
- 6 Remove entry at index 0 of  $\nabla_{y_1} \hat{R}_b$  and compute  $\nabla_{z_1} \hat{R}_b = \dot{r}(z_1) \odot \nabla_{y_1} \hat{R}_b$

## Backpropagation Through Vector Activation

How is  $\nabla_{\mathbf{z}_3} \hat{R}_b$  related to  $\nabla_{\mathbf{y}_3} \hat{R}_b$ ? Let's do what we did before

In this graph,  $F(\cdot)$  is a **vector activation**. We know  $\nabla_{\mathbf{y}} \hat{R}$



We want to find  $\nabla_{\mathbf{z}} \hat{R}$

As mentioned before: we can always extend things entry-wise

With **vector activation**, we need to use the notion of **Jacobian**

## Recap: Jacobian Matrix

Consider vector activation  $F(\cdot)$  that maps  $C$ -dimensional  $\mapsto C$ -dimensional

$$\begin{bmatrix} y_1 \\ \vdots \\ y_C \end{bmatrix} = F\left(\begin{bmatrix} z_1 \\ \vdots \\ z_C \end{bmatrix}\right)$$

When we use this function, we can say

Any entry  $y_j$  is function of all<sup>a</sup>  $z_1, \dots, z_C$ , so we have

$$\nabla_{\mathbf{z}} y_j = \begin{bmatrix} \partial y_j / \partial z_1 \\ \vdots \\ \partial y_j / \partial z_C \end{bmatrix}$$

---

<sup>a</sup>It is not any more an entry-wise functional operation

## Recap: Jacobian Matrix

Consider vector activation  $F(\cdot)$  that maps  $C$ -dimensional  $\mapsto C$ -dimensional

$$\begin{bmatrix} y_1 \\ \vdots \\ y_C \end{bmatrix} = F\left(\begin{bmatrix} z_1 \\ \vdots \\ z_C \end{bmatrix}\right)$$

When we use this function, we can say

We can collect all these gradients into a matrix

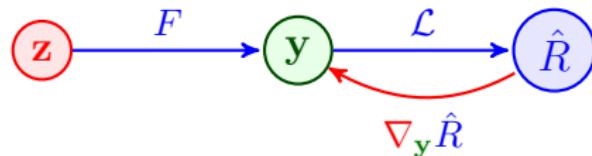
$$\mathbf{J}_{\mathbf{z}} \mathbf{y} = \mathbf{J}_{\mathbf{z}} F = \begin{bmatrix} \nabla_{\mathbf{z}} y_1^T \\ \vdots \\ \nabla_{\mathbf{z}} y_C^T \end{bmatrix} = \begin{bmatrix} \partial y_1 / \partial z_1 & \dots & \partial y_1 / \partial z_C \\ \vdots & & \vdots \\ \partial y_C / \partial z_1 & \dots & \partial y_C / \partial z_C \end{bmatrix}$$

and we call it the *Jacobian matrix*

# Backpropagation Through Vector Activation

Now, let's get back to our problem

In this graph,  $F(\cdot)$  is a **vector activation**. We know  $\nabla_{\mathbf{y}} \hat{R}$



We want to find  $\nabla_{\mathbf{z}} \hat{R}$ : let's write down a partial derivative  $\hat{R}$  w.r.t.  $z_c$

$$\frac{\partial \hat{R}}{\partial z_c} = \sum_{j=1}^C \frac{\partial \hat{R}}{\partial y_j} \frac{\partial y_j}{\partial z_c} = \underbrace{\left[ \frac{\partial y_1}{\partial z_c} \quad \dots \quad \frac{\partial y_C}{\partial z_c} \right]}_{\text{transpose of column } c \text{ of } J_{\mathbf{z}} \mathbf{y} \equiv \text{row } c \text{ of } J_{\mathbf{z}}^T \mathbf{y}} \nabla_{\mathbf{y}} \hat{R}$$

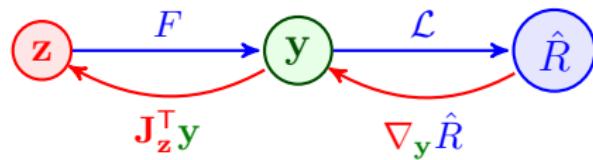
transpose of column  $c$  of  $J_{\mathbf{z}} \mathbf{y} \equiv$  row  $c$  of  $J_{\mathbf{z}}^T \mathbf{y}$

# Backpropagation Through Vector Activation

So, the gradient of  $\hat{R}$  w.r.t.  $\mathbf{z}$  is given by

$$\nabla_{\mathbf{z}} \hat{R} \begin{bmatrix} \partial \hat{R} / \partial z_1 \\ \vdots \\ \partial \hat{R} / \partial z_C \end{bmatrix} = \begin{bmatrix} \text{row 1 of } \mathbf{J}_{\mathbf{z}}^T \mathbf{y} \nabla_{\mathbf{y}} \hat{R} \\ \vdots \\ \text{row } C \text{ of } \mathbf{J}_{\mathbf{z}}^T \mathbf{y} \nabla_{\mathbf{y}} \hat{R} \end{bmatrix} = (\mathbf{J}_{\mathbf{z}}^T \mathbf{y}) \nabla_{\mathbf{y}} \hat{R}$$

So, we can complete the computation graph as follows

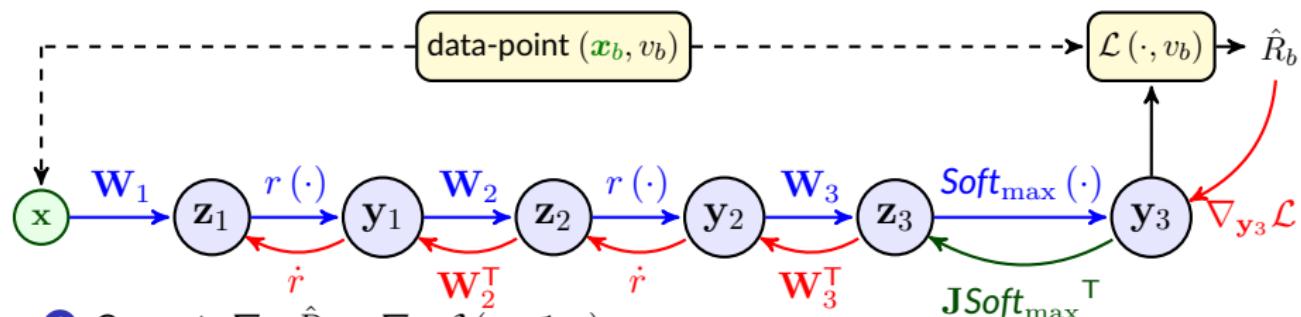


## Backward Pass of Vector Activation

To pass backward on a vector activation, we use the **transpose** of its **Jacobian**

# Backpropagation Through Vector Activation

How can we backpropagate through this neural network? Let's complete



- ① Compute  $\nabla_{y_3} \hat{R}_b = \nabla_{y_3} \mathcal{L}(y_3, \mathbf{1}_{v_b})$
- ② Compute  $\nabla_{z_3} \hat{R}_b = (\mathbf{J} \text{Soft}_{\max})^T \nabla_{y_3} \hat{R}_b$
- ③ Compute  $\nabla_{y_2} \hat{R}_b = \mathbf{W}_3^T \nabla_{z_3} \hat{R}_b$
- ④ Remove entry at index 0 of  $\nabla_{y_2} \hat{R}_b$  and compute  $\nabla_{z_2} \hat{R}_b = \dot{r}(z_2) \odot \nabla_{y_2} \hat{R}_b$
- ⑤ Compute  $\nabla_{y_1} \hat{R}_b = \mathbf{W}_2^T \nabla_{z_2} \hat{R}_b$
- ⑥ Remove entry at index 0 of  $\nabla_{y_1} \hat{R}_b$  and compute  $\nabla_{z_1} \hat{R}_b = \dot{r}(z_1) \odot \nabla_{y_1} \hat{R}_b$

# Multiclass Classification via FNN: Training

Let's now recall **gradient descent** for training of multiclass classifier

**GradientDescent()**:

- 1: Initiate with some initial values  $\{\mathbf{W}_1^{(0)}, \mathbf{W}_2^{(0)}, \mathbf{W}_3^{(0)}\}$  and set a learning rate  $\eta$
- 2: **while** weights not converged **do**
- 3:   **for**  $b = 1, \dots, B$  **do**
- 4:     NN.values  $\leftarrow$  ForwardProp ( $\mathbf{x}_b$ ,  $\{\mathbf{W}_1^{(t)}, \mathbf{W}_2^{(t)}, \mathbf{W}_3^{(t)}\}$ )
- 5:      $\{\mathbf{G}_{1,b}, \mathbf{G}_{2,b}, \mathbf{G}_{3,b}\} \leftarrow$  BackProp ( $\mathbf{x}_b$ ,  $v_b$ ,  $\{\mathbf{W}_1^{(t)}, \mathbf{W}_2^{(t)}, \mathbf{W}_3^{(t)}\}$ , NN.values)
- 6:   **end for**
- 7:   Update

$$\mathbf{W}_1^{(t+1)} \leftarrow \mathbf{W}_1^{(t)} - \eta \text{ mean}(\mathbf{G}_{1,1}, \dots, \mathbf{G}_{1,B})$$

$$\mathbf{W}_2^{(t+1)} \leftarrow \mathbf{W}_2^{(t)} - \eta \text{ mean}(\mathbf{G}_{2,1}, \dots, \mathbf{G}_{2,B})$$

$$\mathbf{W}_3^{(t+1)} \leftarrow \mathbf{W}_3^{(t)} - \eta \text{ mean}(\mathbf{G}_{3,1}, \dots, \mathbf{G}_{3,B})$$

- 8: **end while**

We call this form of gradient descent **full-batch**