

ECE 1508: Applied Deep Learning

Chapter 2: Feedforward Neural Networks

Ali Bereyhi

ali.bereyhi@utoronto.ca

Department of Electrical and Computer Engineering
University of Toronto

Fall 2025

Full-Batch Training

Batch \equiv the dataset reserved for training

In full-batch training, we compute gradients at all data-points in the **batch**: so, we need to wait till forward and backward pass are over for all B data-points

This can be a **huge burdensome!**

- + Wait a moment! Don't we use all the dataset for training?
- No! As you may have noticed in the assignments, we reserve a part of it for **testing**
- + And, why should it be a burdensome?
- OK! Let's get more into datasets!

Public Datasets

Let's consider our example of **image recognition**: we want to recognize the **hand-written number** in an **image**. For this, we need to have access to **images** of **hand-written numbers**. This has been done before by people at **National Institute of Standards and Technology** and collected in a **database** called

*Modified National Institute of Standards and Technology (**MNIST**)*

that is available for **public** on internet

There are several of such **public databases**; some well-known examples are

- *CIFAR-10 and CIFAR-100 by Canadian Institute For Advanced Research*
- *ImageNet initiated by Fei-Fei Li at Princeton University*
- *Caltech-101 and Caltech-256 compiled at Caltech*
- *Fashion MNIST that collects fashion images and labels them*

You can find out more about public datasets [online](#)

Public Datasets: Accessing via PyTorch

PyTorch provides us a simple tool to access these public datasets, e.g.,

```
>> import torchvision.datasets as DataSets  
  
>> dataset = DataSets.MNIST( ... )  
>> dataset = DataSets.CIFAR10( ... )
```

In the example of MNIST, we load the dataset which contains the **pixel vectors of the images of size 28×28** . This means that we load a *list of pairs* where each pair contains

a 784-dimensional vector of pixel values and a label that is in $\{0, 1, \dots, 9\}$

Public Datasets: How Do They Look?

Public datasets include a large amount of data-points with their labels

MNIST includes 70,000 images of hand-written numbers with their true labels: from these 70,000 we use 60,000 for training and 10,000 for test

This means that once we load the MNIST dataset, we make a batch of 60,000 images to train our FNN. Once the training is over, we test the performance of the trained FNN on the remaining 10,000 images

Back to our problem, this means that our full-batch training performs each iteration of the gradient descent after

60,000 forward and backward passes over the FNN

which sounds a lot!

Full-Batch Training: Complexity

Given the example of MNIST, let's see roughly how long it takes to do a full-batch training: if we need 100 iterations of gradient descent, we need to pass back and forth for 6×10^6 times!

- + But do we really need to do this much? This sounds impossible in large NNs!
- No! We really don't need! We can do the training much faster

The full-batch training is really not practical: in practice, we use stochastic (mini-batch) gradient descent to train our NN with feasible complexity

Let's take a look at these approaches!

Sample-Level Training

The most primary idea is to apply one step of **gradient descent** after **each forward and backward pass**: *in our FNN this means that we do the following*

SampLevel_GradientDescent():

```
1: Initiate with some initial values { $\mathbf{W}_1^{(0)}, \mathbf{W}_2^{(0)}, \mathbf{W}_3^{(0)}$ } and set a learning rate  $\eta$ 
2: Start at  $b = 1$ 
3: while weights not converged do
4:   if  $b > B$  then
5:     Update  $b \leftarrow 1$                                 # start over with the dataset
6:   end if
7:   NN.values  $\leftarrow$  ForwardProp ( $x_b, \{\mathbf{W}_1^{(t)}, \mathbf{W}_2^{(t)}, \mathbf{W}_3^{(t)}\}$ )
8:    $\{\mathbf{G}_{1,b}, \mathbf{G}_{2,b}, \mathbf{G}_{3,b}\} \leftarrow$  BackProp ( $x_b, v_b, \{\mathbf{W}_1^{(t)}, \mathbf{W}_2^{(t)}, \mathbf{W}_3^{(t)}\}, \text{NN.values}$ )
9:   Update  $\mathbf{W}_\ell^{(t+1)} \leftarrow \mathbf{W}_\ell^{(t)} - \eta \mathbf{G}_{\ell,b}$           # sample_level update
10:  Update  $b \leftarrow b + 1$                                 # go for next data-point
11: end while
```

We call this approach **sample-level training**

Sample-Level Training: Meaning

- + But what does it mean in the sense of empirical risk minimization? Aren't we now doing something different from the standard **gradient descent**!?
- Yes! We are in fact performing an approximative **gradient descent**

Consider the an **ideal scenario** in which

$$\mathbf{G}_{\ell,1} = \mathbf{G}_{\ell,2} = \dots = \mathbf{G}_{\ell,B}$$

In this case, we do not need to wait for the **batch** to be fully over, since

$$\mathbf{G}_{\ell,1} = \text{mean } (\mathbf{G}_{\ell,1}, \mathbf{G}_{\ell,2}, \dots, \mathbf{G}_{\ell,B})$$

In other words, in this case

sample-level training \equiv **full-batch training**

Sample-Level Training: Meaning

In practice, at each data-point we calculate a **noisy-version** of a **ground truth gradient** $\bar{\mathbf{G}}_\ell$. In other words, we can think of $\mathbf{G}_{\ell,b}$ for each b as

$$\mathbf{G}_{\ell,b} = \bar{\mathbf{G}}_\ell + \text{Noise}$$

If this noise is small enough, we can say that

$$\mathbf{G}_{\ell,b} \approx \bar{\mathbf{G}}_\ell \approx \text{mean}(\mathbf{G}_{\ell,1}, \mathbf{G}_{\ell,2}, \dots, \mathbf{G}_{\ell,B})$$

and therefore, we can conclude that

sample-level (worse approx. $\bar{\mathbf{G}}_\ell$) \approx **full-batch (better approx. $\bar{\mathbf{G}}_\ell$)**

In this case, we say that $\mathbf{G}_{\ell,b}$ is an **estimator** of the **ground truth gradient**

Sample-Level Training: Repetitive Cycle Issue

Naive sample-level update can trap us into a repetitive cycle: in simple words, we can end up with our initial point at the end of the batch. For instance, consider the following dummy (but possible) scenario in our three-layer FNN

We start with $\mathbf{W}_\ell^{(0)}$ and get into the batch for the first time

- We update $\mathbf{W}_\ell^{(0)}$ after the first data-point to $\mathbf{W}_\ell^{(1)}$
- We update $\mathbf{W}_\ell^{(1)}$ after the second data-point to $\mathbf{W}_\ell^{(2)}$
- ...
- We update $\mathbf{W}_\ell^{(B-1)}$ after the last data-point to $\mathbf{W}_\ell^{(B)}$

Now, assume that $\mathbf{W}_\ell^{(B)} = \mathbf{W}_\ell^{(0)}$ for all layers again!

In the above dummy example, further looping over the batch is useless, since we always get back to the initial point: this is the most basic example of the repetitive cycle issue

Stochastic Sample-Level Training: SGD

- + How can we avoid such **cyclic behaviors**?
- We can use **Stochastic Gradient Descent (SGD)**

Each time we are to loop over our **training batch**, we **shuffle** the data-points **randomly**: this way we avoid next loop behave like the previous one

This idea is called Stochastic Gradient Descent (SGD)

SGD is the most common algorithm for training of NNs!

What does **random shuffling** mean?

It means **randomly permuting** the data-points

Stochastic Gradient Descent

SGD():

```
1: Initiate with some initial values  $\{\mathbf{W}_1^{(0)}, \mathbf{W}_2^{(0)}, \mathbf{W}_3^{(0)}\}$  and set a learning rate  $\eta$ 
2: Randomly shuffle the batch and start at  $b = 1$ 
3: while weights not converged do
4:   if  $b > B$  then
5:     Randomly shuffle the batch and set  $b \leftarrow 1$            # random shuffling
6:   end if
7:   NN.values  $\leftarrow$  ForwardProp ( $x_b, \{\mathbf{W}_1^{(t)}, \mathbf{W}_2^{(t)}, \mathbf{W}_3^{(t)}\}$ )
8:    $\{\mathbf{G}_{1,b}, \mathbf{G}_{2,b}, \mathbf{G}_{3,b}\} \leftarrow$  BackProp ( $x_b, v_b, \{\mathbf{W}_1^{(t)}, \mathbf{W}_2^{(t)}, \mathbf{W}_3^{(t)}\}, \text{NN.values}$ )
9:   Update  $\mathbf{W}_\ell^{(t+1)} \leftarrow \mathbf{W}_\ell^{(t)} - \eta \mathbf{G}_{\ell,b}$           # sample_level update
10:  Update  $b \leftarrow b + 1$                                      # go for next data-point
11: end while
```

- + But, doesn't sample-level training lead to any drawback?
- Sure! But we accept this drawback as a cost we pay for less complexity

Let's see how this trade-off looks like

Recap: Variance

For random variable x with mean μ , the variance is defined as

$$\text{Var}\{x\} = \mathbb{E}\{(x - \mu)^2\} = \mathbb{E}\{x^2\} - \mu^2$$

Clearly, when x is zero-mean, we can say $\text{Var}\{x\} = \mathbb{E}\{x^2\}$

Properties of Variance

For any random variable x and constant c , we have

$$\text{Var}\{cx\} = c^2 \text{Var}\{x\}$$

Let x_1, \dots, x_N be N independent random variables; then, we have

$$\text{Var}\left\{\sum_{n=1}^N x_n\right\} = \sum_{n=1}^N \text{Var}\{x_n\}$$

Recap: Variance

Now, assume x_1, \dots, x_N are N **independent zero-mean** random variables all with variance σ^2 : let \bar{x} be the *arithmetic average* of x_1, \dots, x_N , i.e.,

$$\bar{x} = \text{mean}(x_1, \dots, x_N) = \frac{1}{N} \sum_{n=1}^N x_n$$

We could then say

$$\begin{aligned}\text{Var}\{\bar{x}\} &= \text{Var}\left\{\frac{1}{N} \sum_{n=1}^N x_n\right\} = \frac{1}{N^2} \text{Var}\left\{\sum_{n=1}^N x_n\right\} \\ &= \frac{1}{N^2} \sum_{n=1}^N \underbrace{\text{Var}\{x_n\}}_{\sigma^2} = \frac{1}{N^2} (N\sigma^2) \\ &= \frac{\sigma^2}{N} \quad \text{variance of average drops by } 1/\#\text{ samples}\end{aligned}$$

Complexity-Accuracy Trade-off of SGD

Now, let's get back to our problem: when we talked about the meaning of symbol level update, we said

In practice, at each data-point we calculate a **noisy-version** of a **ground truth gradient** $\bar{\mathbf{G}}_\ell$. In other words, we can think of $\mathbf{G}_{\ell,b}$ for each b as

$$\mathbf{G}_{\ell,b} = \bar{\mathbf{G}}_\ell + \text{Noise}$$

and called $\mathbf{G}_{\ell,b}$ an **estimator** of the **ground truth**

Let's make the above statement a bit more formal: we assume that Noise for each b is a matrix with **independent** zero-mean entries all with **variance** σ^2 , i.e.,

$$\mathbf{G}_{\ell,b} = \bar{\mathbf{G}}_\ell + \mathbf{N}_{\ell,b}$$

where we define $\bar{\mathbf{G}}_\ell$ to be the gradient of the **true risk**

Complexity-Accuracy Trade-off of SGD

What is the *true risk*? If you remember, when we started with training

Our goal was to *minimize* the *risk* $R(\mathbf{w})$

However, we could not do this: since we *did not know* (1) the *true function*, and (2) the *data distribution*. Thus,

we approximated the *true risk* $R(\mathbf{w})$ with *empirical risk* $\hat{R}(\mathbf{w})$

We assume that $\bar{\mathbf{G}}_\ell$ is the gradient of *true risk* with respect to \mathbf{W}_ℓ , i.e.,

$$\bar{\mathbf{G}}_\ell = \nabla_{\mathbf{W}_\ell} R(\mathbf{w})$$

- + Can we determine this gradient?
- Of course not! We can only approximate it with $\nabla_{\mathbf{W}_\ell} \hat{R}(\mathbf{w})$

Complexity-Accuracy Trade-off of SGD

Let's see accurate the gradient is approximated, when we do **full-batch** training

In **full-batch** training, we determine the gradient as

$$\begin{aligned}\hat{\mathbf{G}}_{\ell}^{\text{batch}} &= \text{mean}(\mathbf{G}_{\ell,1}, \mathbf{G}_{\ell,2}, \dots, \mathbf{G}_{\ell,B}) = \frac{1}{B} \sum_{b=1}^B (\bar{\mathbf{G}}_{\ell} + \mathbf{N}_{\ell,b}) \\ &= \bar{\mathbf{G}}_{\ell} + \underbrace{\frac{1}{B} \sum_{b=1}^B \mathbf{N}_{\ell,b}}_{\hat{\mathbf{N}}_{\ell}^{\text{batch}}} = \bar{\mathbf{G}}_{\ell} + \hat{\mathbf{N}}_{\ell}^{\text{batch}}\end{aligned}$$

Recall that by **arithmetic averaging variance drops by $1/\# \text{ of samples}$**

In **full-batch** training the **approximated gradient** is different from the **true gradient** by an error whose variance **drops as σ^2/B**

Complexity-Accuracy Trade-off of SGD

In **full-batch** training the **approximated gradient** is different from the **true gradient** by an error whose variance **drops as σ^2/B**

Now, let's compare it to SGD

In SGD, we **approximate** the gradient with a sample gradient, i.e.,

$$\hat{\mathbf{G}}_{\ell}^{\text{SGD}} = \mathbf{G}_{\ell,b} = \bar{\mathbf{G}}_{\ell} + \mathbf{N}_{\ell,b}$$

SGD is still **approximating** the true gradient but with much **larger variance**: entries of $\mathbf{N}_{\ell,b}$ have all **variance σ^2**

For instance, consider MNIST with 60,000 samples: by **full-batch** training we get gradient values whose difference from the entries of the **true gradient** is approximately 1.67×10^{-5} times smaller than those gradient entries calculated by **SGD!**

Complexity-Accuracy Trade-off of SGD

In the context of ML, we often say: *in the analyses of last slides,*

SGD and **full-batch** training are both **unbiased estimators of $\bar{\mathbf{G}}_\ell$**

We call them **unbiased**, since $\mathbb{E} \left\{ \hat{\mathbf{G}}_\ell^{\text{SGD}} \right\} = \mathbb{E} \left\{ \hat{\mathbf{G}}_\ell^{\text{batch}} \right\} = \bar{\mathbf{G}}_\ell$

Complexity-Accuracy Trade-off

Assume that a forward and backward pass takes time T and that gradient of the risk at each sample be an **unbiased estimators of the true gradient**; then,

- ① each step of **SGD** takes time T while each step of **full-batch** training takes BT with B being the **batch size**
- ② if we denote the variance of estimation given by **SGD** by σ^2 , the variance of **full-batch** estimator is σ^2/B

Training via Mini-Batches

- + But, can't we play with this trade-off? For instance, increase a bit the complexity to improve the accuracy!
- Yes! This is the idea of mini-batch training

In mini-batch training, we divide the whole batch of data into mini-batches:

- after each mini-batch is over, we average the gradients over the mini-batch
- we apply one step of gradient descent using this averaged gradient

To avoid cyclic behavior, we still shuffle the dataset randomly each time we start a new loop over it. This training approach is hence often called

Mini-Batch Stochastic Gradient Descent = Mini-Batch SGD

Mini-Batch SGD

`mBatchSGD() :`

```

1: Initiate with some initial values  $\{\mathbf{W}_\ell^{(0)}\}$  and set a learning rate  $\eta$ 
2: Randomly shuffle the batch and divide it into mini-batches of size  $\Omega$ 
3: Denote the number of mini-batches by  $\Xi = \lceil B/\Omega \rceil$  and start at  $\xi = 1$ 
4: while weights not converged do
5:   if  $\xi > \Xi$  then
6:     Randomly shuffle the batch and divide it into mini-batches of size  $\Omega$ 
7:     Set  $\xi \leftarrow 1$                                 # start over with the dataset
8:   end if
9:   for  $\omega = 1, \dots, \Omega$  do
10:    NN.values  $\leftarrow$  ForwardProp ( $x_\omega, \{\mathbf{W}_\ell^{(t)}\}$ )
11:     $\{\mathbf{G}_{\ell,\omega}\} \leftarrow$  BackProp ( $x_\omega, v_\omega, \{\mathbf{W}_\ell^{(t)}\}, \text{NN.values}$ )
12:   end for
13:   Update  $\mathbf{W}_\ell^{(t+1)} \leftarrow \mathbf{W}_\ell^{(t)} - \eta \text{ mean}(\mathbf{G}_{\ell,1}, \dots, \mathbf{G}_{\ell,\Omega})$ 
14:   Update  $\xi \leftarrow \xi + 1$                       # go for next mini-batch
15: end while

```

Complexity-Accuracy Trade-off

It is easy to see that

- mini-batch training reduces to full-batch training when we set the size of mini-batches to B , i.e., $\Omega = B$
- mini-batch training reduces to SGD when we set the size of mini-batches to 1, i.e., $\Omega = 1$

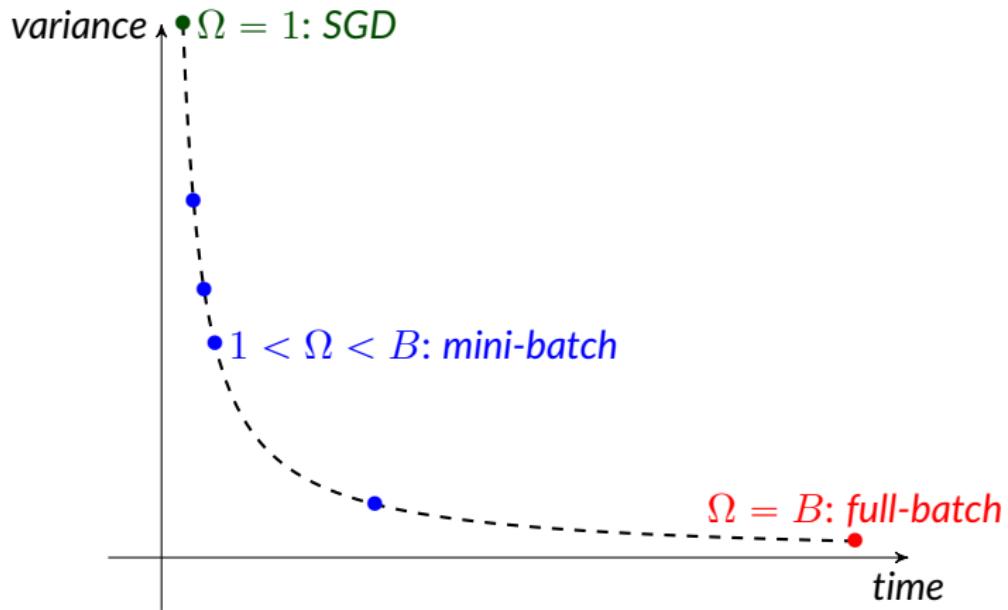
Complexity-Accuracy Trade-off

Assume that a forward and backward pass takes time T and that gradient of the risk at each sample be an unbiased estimators of the true gradient; then,

- ① each step of SGD takes time T while each step of mini-batch training takes ΩT with Ω being the mini-batch size
- ② if we denote the variance of estimation given by SGD by σ^2 , the variance of mini-batch estimator is σ^2/Ω

Complexity-Accuracy Trade-off

The complete trade-off can be visualized as



Mini-batch size is what specifies the trade-off point

Few Definitions: Epoch and Iteration

In the language of deep learning there are few terms that we must know

Batch Size

Through time, the term **mini-batch** has been transformed to **batch**, and the **complete batch** is referred to as **training dataset** or the **full batch**. People hence call the size of each **mini-batch**, i.e., Ω , the **batch size**

Iteration

When we take one step of gradient descent, we take one **iteration**. So, one **iteration** is over when we finish with a **mini-batch**

Epoch

An **epoch** is over when we finish **once** with the **whole training dataset**

Few Definitions: Epoch and Iteration

We can annotate these definitions in our algorithm

mBatchSGD():

- 1: Initiate with some initial values $\{\mathbf{W}_\ell^{(0)}\}$ and set a learning rate η
- 2: Randomly shuffle **training dataset** and make mini-batches of size $\Omega \equiv \text{batch-size}$
- 3: Denote the number of mini-batches by $\Xi = \lceil B/\Omega \rceil$ and start at $\xi = 1$
- 4: **while** weights not converged **do**
- 5: **if** $\xi > \Xi$ **then**
- 6: Randomly shuffle the batch and divide it into mini-batches of size Ω
- 7: Set $\xi \leftarrow 1$ **← one epoch is over, we start another epoch**
- 8: **end if**
- 9: **for** $\omega = 1, \dots, \Omega$ **do**
- 10: NN.values \leftarrow ForwardProp $(\mathbf{x}_\omega, \{\mathbf{W}_\ell^{(t)}\})$ **going through a min-batch**
- 11: $\{\mathbf{G}_{\ell,\omega}\} \leftarrow$ BackProp $(\mathbf{x}_\omega, v_\omega, \{\mathbf{W}_\ell^{(t)}\}, \text{NN.values})$
- 12: **end for**
- 13: Update $\mathbf{W}_\ell^{(t+1)} \leftarrow \mathbf{W}_\ell^{(t)} - \eta \text{ mean}(\mathbf{G}_{\ell,1}, \dots, \mathbf{G}_{\ell,\Omega})$ **← one iteration**
- 14: Update $\xi \leftarrow \xi + 1$ **# go for next mini-batch**
- 15: **end while**

Few Definitions: Epoch and Iteration

We can consider a simple example: say we train our FNN over MNIST using **mini-batch SGD** with **batch size $\Omega = 100$** . Our **training dataset** has 60,000 **data-points**; thus, we have

$$\Xi = \frac{60,000}{100} = 600$$

mini-batches. Each time we finish with a **mini-batch**, we do **one iteration** of gradient descent. After **600 iterations**, we finish with a single **epoch**

So, if we have trained the FNN for 10 **epochs**, it means that

we have done **$600 \times 10 = 6000$ iterations of gradient descent**

Testing NNs with New Data-Point

- + Say we are over with the **training**; then, what should we do?
- We need to test it with the data we reserved for **testing**

After **training**, we need to test our **trained** NN: say we get a new data-point x_{new} with label v_{new} . We can test our NN for this new **test data-point** by evaluating **classical metrics**

- ① **Test Risk** also called **Test Loss**: we pass x_{new} forward through our **trained NN** and get y_{new} . We then calculate the **test loss** as $\mathcal{L}(y_{\text{new}}, v_{\text{new}})$ using the same loss function \mathcal{L} we used for training
- ② **Test Accuracy**: we use y_{new} to classify x_{new} . We then compare it to the **true class** of x_{new} . If they are the same; then, the test accuracy is 1, if not, it is 0

Testing NNs over Test Dataset

Testing for a single new point is not reliable: *this is why we had reserved the test dataset.*

Given the **test dataset**, we go through every single test **data-point**

- we pass the data-point forward through the **trained NN**
- we compute the **test loss** and **test accuracy**
- we **average them** over the whole **test dataset**

Therefore, we get

- an **average loss that approximates the risk**
- a **test accuracy** between 0 and 1 that says how **accurate** our trained NN is

Learning Curves

- + What you said gives us **two numbers!** But, I have seen **curves!**
- Yes! They are **learning curves**

In practice, the SGD can take **very long** to converge, i.e., to stop iterating

*it needs **too many iterations** to get **too close** to the minimum*

But, it might be **not really needed** to get **that close!** So,

*we test our NN once every **epoch***

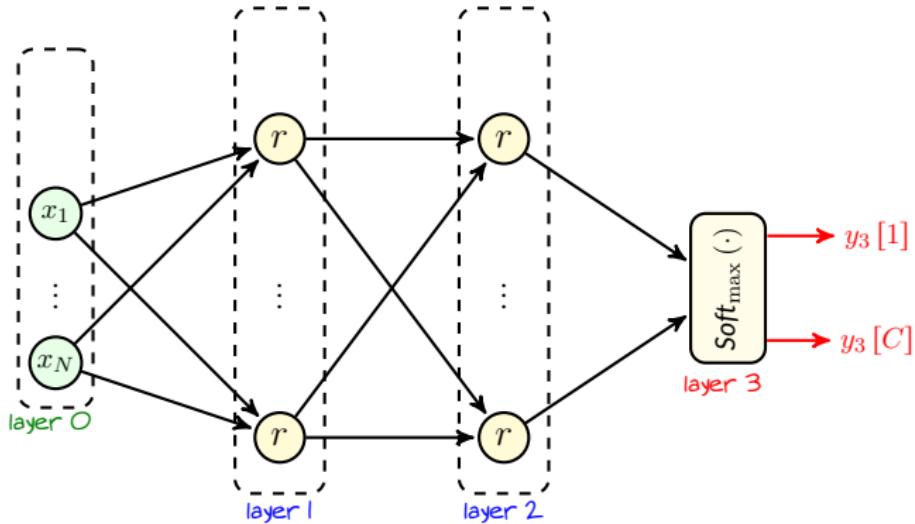
We then plot the **test risk** and **test accuracy** against **number of epochs** in a curve: *these curves are often called **learning curves***

*if we see that learning curves are **saturating**, we can **stop** the training*

In practice: **we always perform the **training** for a **fixed number of epochs****

Learning Curves: Example

Let's see an example: *recall our three-layer FNN. Say, we train it for image classification over MNIST which has 60,000 data-points for and 10,000 for test*



In MNIST, we have 10 classes, so $C = 10$. We use cross-entropy as loss function

Learning Curves: Example

We agree to do the following: we use *mini-batch SGD* with *batch size $\Omega = 100$* and train the FNN for *100 epochs*.

In epoch $\xi = 1, \dots, 100$

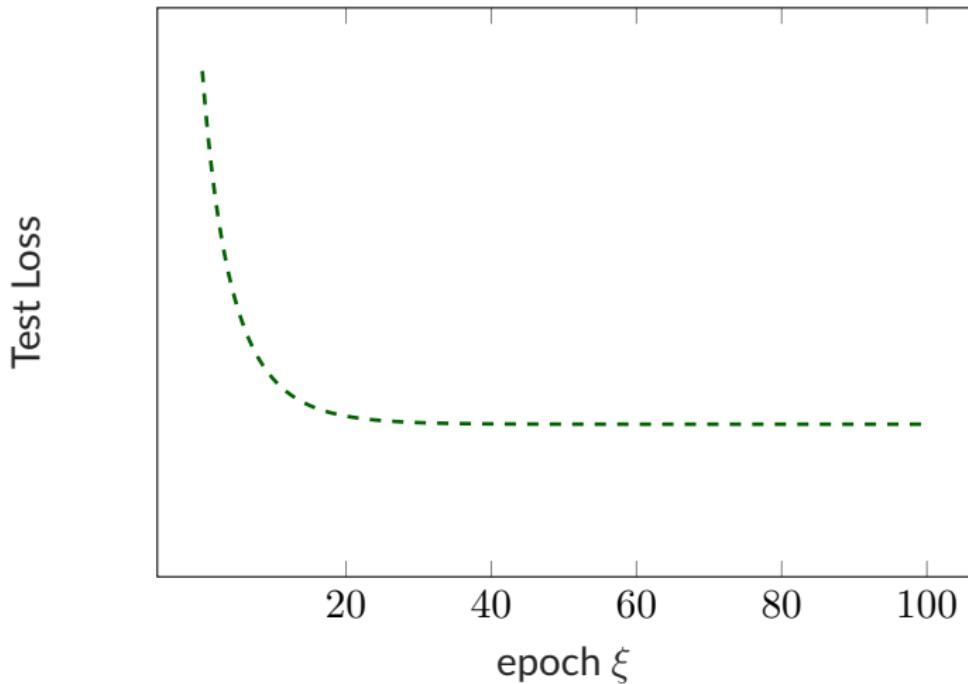
- ① we perform 600 iterations of *gradient descent*
- ② we fix the weights to what we computed at *last iteration of the epoch*
- ③ for each test data-point: we *pass it forward* and determine y_3
 - ① we compute $\text{CE}(y_3, 1_v)$, where v is the *true class of test data-point*
 - ② we find the index of *maximum term in y_3* and compare it to v
 - ↳ if they are the same, we set *accuracy to 1*; otherwise, we *set it 0*
- ④ we *average test loss and accuracy*

Now, for each epoch

we have a *test loss* and *test accuracy*: we plot them against ξ

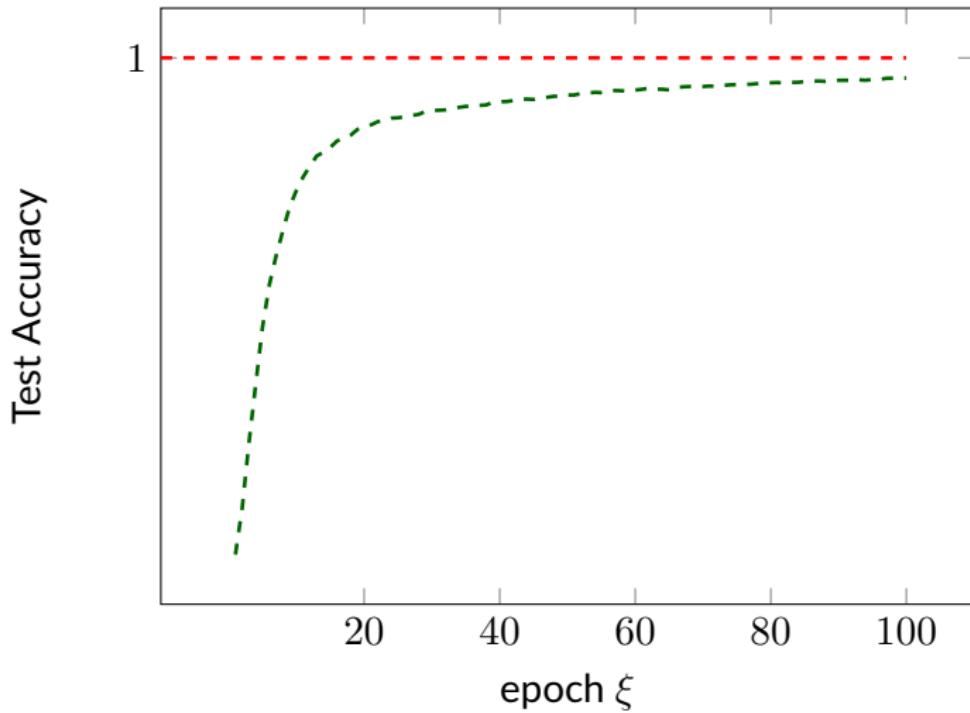
Learning Curves: Example

How should the learning curves look? A typical curve for test loss is



Learning Curves: Example

How should the learning curves look? A typical curve for test accuracy is



Summary of This Chapter

- To train a NN we need *gradients*
 - ↳ We can *calculate gradient* by *forward* and *backpropagation* over the NN
 - ↳ In FNNs, *forward propagation* uses simple linear and nonlinear operations
 - ↳ *Backpropagation* is readily derived using *computation graph*
- We tried Classification via FNNs
 - ↳ Better to work with *probabilities* instead of exact *labels*
 - ↳ For *multiclass* classification, we should use *vector-activated neurons*
- To minimize the *exact* empirical risk, we have to do *full-batch* training
 - ↳ This requires *huge computation complexity*
 - ↳ We can hugely reduce this cost by *SGD* which does *sample-level training*
 - ↳ *SGD* versus *full-batch* describes a *complexity-accuracy trade-off*
 - ↳ We can tune this *trade-off* by *mini-batch SGD*