

Applied Deep Learning

Chapter 3: Advancing Our Toolbox

Ali Bereyhi

ali.bereyhi@utoronto.ca

Department of Electrical and Computer Engineering
University of Toronto

Fall 2025

What's Next

Right now, we are familiar with FNN and how to train them

↳ We potentially can also handle *some other architectures*

But, if we try to **implement them from scratch**, we could get into trouble

We still need to learn **more tricks** for having a working implementation

In this chapter, we advance **our bag of tools** in four respects

① We learn more about **optimizers**

↳ more **advanced** tricks to make gradient descent work

② We learn about **hyperparameter tuning**

③ We learn about **data preprocessing**

↳ how to **handle data in practice**

④ Tricks to make training **faster and more robust**

Back to Gradient Descent

Let's take a look at *training in abstract form* once again

$$\min_{\mathbf{w}} \hat{R}(\mathbf{w}) \quad (\text{Training})$$

Recall that \mathbf{w} includes *all weights and biases*; for instance,

In FNN of Assignment 2, \mathbf{w} has 3,457 entries! In practice much higher!

Also recall the gradient descent

- 1: Initiate at some $\mathbf{w}^{(0)} \in \mathbb{R}^D$ and deviation $\Delta = +\infty$
- 2: Choose some small ϵ and η , and set $t = 1$
- 3: **while** $\Delta > \epsilon$ **do**
- 4: Update weights as $\mathbf{w}^{(t)} \leftarrow \mathbf{w}^{(t-1)} - \eta \nabla \hat{R}(\mathbf{w}^{(t-1)})$
- 5: Update the deviation $\Delta = |\hat{R}(\mathbf{w}^{(t)}) - \hat{R}(\mathbf{w}^{(t-1)})|$
- 6: **end while**

Keep in mind: we almost always use (mini-batch) SGD ↵ let's call it SGD

Convergence Rate of Optimizers

Recall: we can make gradient descent converging to **local minimum** if
we set the learning rate η small enough

So is it also with **SGD**. But, **how fast** does the algorithm converge?

Speed of an optimization algorithm is evaluated by **convergence rate**


We can understand its meaning from the eyes of **optimality gap**

Optimality Gap in Iteration t

Optimality gap is the gap between $\mathbf{w}^{(t)}$ and local minimizer \mathbf{w}^* , i.e.,

$$\text{optimality gap in iteration } t = |\hat{R}(\mathbf{w}^{(t)}) - \hat{R}(\mathbf{w}^*)|$$

Convergence Rate of Optimizers



An optimizer *shrinks this gap gradually*, i.e., we can *approximately* say¹

$$|\hat{R}(\mathbf{w}^{(t+1)}) - \hat{R}(\mathbf{w}^*)| \leq |\hat{R}(\mathbf{w}^{(t)}) - \hat{R}(\mathbf{w}^*)| \leq \dots \leq |\hat{R}(\mathbf{w}^{(0)}) - \hat{R}(\mathbf{w}^*)|$$

In practice, this drop can occur *with different speeds in terms of t*

¹It's just *approximately* correct: *gap may increase in one iteration only!* That's no problem

Convergence Rate of Optimizers

An *ideal scenario* is that the gap *drops by a constant factor each iteration*

This *intuitively* means that *for each iteration t, we see*

$$|\hat{R}(\mathbf{w}^{(t)}) - \hat{R}(\mathbf{w}^*)| \leq \alpha |\hat{R}(\mathbf{w}^{(t-1)}) - \hat{R}(\mathbf{w}^*)|$$

for *some* $\alpha < 1$: in this case, we can say

$$\begin{aligned} |\hat{R}(\mathbf{w}^{(t)}) - \hat{R}(\mathbf{w}^*)| &\leq \alpha |\hat{R}(\mathbf{w}^{(t-1)}) - \hat{R}(\mathbf{w}^*)| \\ &\leq \alpha (\alpha |\hat{R}(\mathbf{w}^{(t-2)}) - \hat{R}(\mathbf{w}^*)|) = \alpha^2 |\hat{R}(\mathbf{w}^{(t-2)}) - \hat{R}(\mathbf{w}^*)| \\ &\leq \dots \leq \alpha^t |\hat{R}(\mathbf{w}^{(0)}) - \hat{R}(\mathbf{w}^*)| \end{aligned}$$

Now say that $|\hat{R}(\mathbf{w}^{(0)}) - \hat{R}(\mathbf{w}^*)| = C$.² So, we can say

$$|\hat{R}(\mathbf{w}^{(t)}) - \hat{R}(\mathbf{w}^*)| \leq C\alpha^t \rightsquigarrow \mathcal{O}(\alpha^t)$$

²We don't really care how large C is. It's going to shrink at the end

Convergence Rate of Optimizers

An ideal scenario is that the gap drops by a constant factor each iteration

If we wish to end up somewhere in ϵ -neighborhood of \mathbf{w}^* ; then, we need

$$C\alpha^t \leq \epsilon$$

For this to happen, we need to have at least

$$t \geq \frac{\log 1/\epsilon + \log C}{\log 1/\alpha} \rightsquigarrow \mathcal{O}(\log 1/\epsilon)$$

iterations: the closer we need to get, the more we should iterate

For this required time, we say that the optimizer converges linearly

It's a fast rate, since number of iterations is proportional to logarithm of $1/\epsilon$

Convergence Rate of Optimizers

- + You said **ideal!** Isn't gradient descent **always** converging at **this rate?**
- Well! Only when empirical risk is **strongly convex** and we do **full-batch training**! You can guess it happens **almost never** for us!
- + But how it works with **realistic NNs** and **SGD**?

In general, it's hard to characterize **exact convergence**; however, we know that when *empirical risks are rather smooth functions*³, we have

$$|\hat{R}(\mathbf{w}^{(t)}) - \hat{R}(\mathbf{w}^*)| \leq \frac{L}{t\eta}$$

for some L that **gets larger as w becomes larger**. So, in practice, we have

$$|\hat{R}(\mathbf{w}^{(t)}) - \hat{R}(\mathbf{w}^*)| = \mathcal{O}(1/t)$$

³To be rigorous: when it's **Lipschitz** continuous, but we don't really need details on that

Convergence Rate of Optimizers

In practice, we have

$$|\hat{R}(\mathbf{w}^{(t)}) - \hat{R}(\mathbf{w}^*)| = \mathcal{O}(1/t)$$

So, if we want to end up somewhere in ϵ -neighborhood of \mathbf{w}^* , we need

$$t = \mathcal{O}(1/\epsilon)$$

This can potentially take long! We say in this case that

the optimizer converges sub-linearly

- + How bad can it be?
- Just set $\epsilon = 10^{-3}$: with linear convergence we need time in order of 3; with sub-linear one, we need in order of 1000!

Alternative Optimizers

Moral of Story

Vanilla gradient descent is not what we can use in practice!

In practice, we employ **improved** versions of **gradient descent**: there is a **long list** of them, but we check a few **important** ones that are typically used

- Gradient descent with *learning rate scheduling*
 - Gradient descent with *momentum*
 - Rprop: *Resilient backpropagation*
 - RMSprop: *Root mean square propagation*
 - Adam: *Adaptive moment estimation*
- + If gradient descent is not used in *practice*, why we did *backpropagation*?
- No worries! They all use *gradient!* This is why these algorithms are commonly referred to as *gradient-based training algorithms*

SGD with Learning Rate Scheduling

We had it in simple words in Chapter 1: we *vary learning rate* through *time*

$$\mathbf{w}^{(t)} \leftarrow \mathbf{w}^{(t-1)} - \eta^{(t-1)} \nabla \hat{R}(\mathbf{w}^{(t-1)})$$

It's called *learning rate scheduling*: start at large $\eta^{(0)} = \eta$ and reduce it with t

We may schedule learning rate with various approaches

- We could have *linear* decay

$$\eta^{(t)} = \frac{\eta}{t + 1}$$

- We could have *polynomial* decay with power P

$$\eta^{(t)} = \frac{\eta}{(t + 1)^P}$$

- We could have *exponential* decay with some exponent rate $\kappa > 0$


$$\eta^{(t)} = \eta e^{-\kappa t}$$

SGD with Learning Rate Scheduling

There is also a *more practical trick* for learning rate scheduling:

- ① Use a **fixed learning rate** η until the **test risk** saturates
- ② Reduce the learning rate as $\eta \leftarrow \alpha\eta$ for $\alpha < 1$
- ③ Repeat the above steps until test risk arrives at a **desired level**

A *good choice* of α is around $\alpha = 0.1$



SGD with Learning Rate Scheduling

Learning rate scheduling can lead us to a better local minima; however,
it does not change the convergence rate

It is nevertheless a good approach for easy problems

We can access pre-implemented scheduling techniques in PyTorch

```
>> import torch  
>> torch.optim.lr_scheduler
```

Momentum: Moving Average

Momentum is one of the key approaches to robust SGD

The idea is simple: we replace the gradient⁴ with its moving average

Say we are in iteration t and let the computed gradient to be $\mathbf{g}^{(t)}$, i.e.,

$$\mathbf{g}^{(t)} = \text{estimator} \left\{ \nabla \hat{R}(\mathbf{w}^{(t)}) \right\} \quad \text{e.g., we computed by SGD}$$

With standard gradient descent we update as

$$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \eta^{(t)} \mathbf{g}^{(t)}$$

In momentum approach, we replace $\mathbf{g}^{(t)}$ with its moving average

⁴Which can largely variate, especially with small mini-batches

Momentum: Moving Average

Moving Average \sim Momentum

Moving average with factor β in iteration t is

$$\mathbf{m}^{(t)} = \beta \mathbf{m}^{(t-1)} + (1 - \beta) \mathbf{g}^{(t)}$$

Moving average has less fluctuations; hence, it's an estimator of true gradient with less variance

SGD with momentum does the following update

```
Initiate  $\mathbf{m}^{(0)} = \mathbf{0}$ 
for  $t = 1, \dots$  do
    ...
     $\mathbf{m}^{(t)} = \beta \mathbf{m}^{(t-1)} + (1 - \beta) \mathbf{g}^{(t)}$ 
     $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \eta^{(t)} \mathbf{m}^{(t)}$ 
    ...
end for
```

Nesterov Momentum: Accelerated Gradient Computation

Nesterov approach adds an *intermediate* step: when we compute *gradient*, we use the *already-calculated momentum* to *estimate future weights*

```
Initiate  $\mathbf{m}^{(0)}$ 
for  $t = 1, \dots$  do
    ...
     $\hat{\mathbf{w}} = \mathbf{w}^{(t)} - \beta \mathbf{m}^{(t-1)}$  # approximate next point
    Compute gradient for weights  $\hat{\mathbf{w}}$ : call it  $\hat{\mathbf{g}}^{(t)}$ 
     $\mathbf{m}^{(t)} = \beta \mathbf{m}^{(t-1)} + (1 - \beta) \hat{\mathbf{g}}^{(t)}$ 
     $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \eta^{(t)} \mathbf{m}^{(t)}$ 
    ...
end for
```

You can combine these lines into a *single line of update*

that may looks a bit more complicated, but it's the same thing!

SGD with Momentum: *Implementation*

In PyTorch, SGD is already implemented with momentum

```
>> import torch  
>> torch.optim.SGD()
```

When using this implementation, we can specify momentum factor, i.e., β , and choose whether Nesterov being applied or not

Typical choice of momentum factor is $\beta = 0.9$, and remember

with $\beta = 0$, we return to the standard SGD

This is the default value in PyTorch

Rprop: Resilient Backpropagation

Rprop was introduced by Riedmiller and Braun in 1992; [check the paper here](#)

Riedmiller and Braun noticed that **gradient descent** can be **improved** if we could have **individual learning rate** in **each dimension of w**: you may recall the first question in Assignment 1!

They hence came up with **Rprop**: let's see first **one-dimensional** case

Rprop():

Initiate $\eta^{(0)}$ and choose $\mu^+ > 1$, $\mu^- < 1$, η_{\max} and η_{\min}

for $t = 1 : T$ do

...

Compute gradient at $w^{(t)}$ and call it $g^{(t)}$

Update learning rate as $\eta^{(t)} \leftarrow \text{Rprop_Scheduler}(\eta^{(t-1)}, g^{(t)}, g^{(t-1)})$

Update weight $w^{(t+1)} = w^{(t)} - \eta^{(t)} \text{sign}(g^{(t)})$ # only sign of gradient

...

end for

Rprop: Learning Rate Scheduler

The key point of **Rprop** is its **scheduler**

Rprop_Scheduler():

Use $\mu^+ > 1$ and η_{\max} as well as $\mu^- < 1$ and η_{\min}

if $\text{sign}(g^{(t)}) = \text{sign}(g^{(t-1)})$ **then**

Update learning rate $\eta^{(t)} = \min \left\{ \mu^+ \eta^{(t-1)}, \eta_{\max} \right\}$ # go faster

else

Update learning rate $\eta^{(t)} = \max \left\{ \mu^- \eta^{(t-1)}, \eta_{\min} \right\}$ # slow down


end if

It follows an **intuitive** strategy

- ↳ Keep **going faster** as the **sign of gradient is not changing**
- ↳ **Slow down** only when the **sign changes**: you have **passed the minimum!**

Rprop: Resilient Backpropagation

Let's look at it geometrically



In simple words, with **Rprop**

we keep on increasing until we pass the minimum

Rprop: Resilient Backpropagation

- + How does it extend to *multi-dimensional case that we have in training?*
- We apply this idea on every entry individually

```
Initiate  $\eta$  and choose  $\mu^+$ ,  $\mu^-$ ,  $\eta_{\max}$  and  $\eta_{\min}$ 
for  $t = 1 : T$  do
    ...
    Compute gradient  $\mathbf{g}^{(t)}$ 
    for every entry  $i$  of  $\mathbf{g}^{(t)}$  do
        Use the sign of entry and update  $\eta_i^{(t)}$  via Rprop_Scheduler()
        Apply one-dimensional Rprop() to update entry  $i$  of  $\mathbf{w}^{(t)}$ 
    end for
    ...
end for
```

Rprop: Resilient Backpropagation

Typical choices of parameters for this algorithm are

- initial learning rate $\eta = 0.01$
- factors $\mu^+ = 1.2$ and $\mu^- = 0.5$
- η_{\max} and η_{\min} are less important: they get automatically regulated

It's better to avoid choices that $1/\mu^+ = \mu^-$, because if we pass the minimum, we don't like to move exactly the previous point

We can again access Rprop through module optim in PyTorch

```
>> import torch  
>> torch.optim.Rprop()
```

RMSprop: Root Mean Square Propagation

It turns out the Rprop only works fine with full-batch training

↳ it's because, it ignores the magnitude of gradient

To understand why this happens, consider the following dummy example

Consider a one-dimensional case, i.e., $w = w$: we break the full batch into 4 mini-batches and come up with the following derivatives calculated in each step of mini-batch SGD

$$(1) \rightsquigarrow 0.1 \quad (2) \rightsquigarrow 0.1 \quad (3) \rightsquigarrow 0.1 \quad (4) \rightsquigarrow -0.5$$

We may approximately say that the derivative of full-batch risk, at the very first choice of w , was close to zero or negative; however, Rprop

- ① takes first three steps with larger and larger learning rates
- ② only comes back with smaller step at last iteration

It could be hence already lost at the last iteration!

RMSprop: Root Mean Square Propagation

Geoffrey Hinton in his [lecture notes](#)⁵ came up with a **solution**: we can use the idea of **moving average** to further normalize the **learning rate** according to **average gradient magnitude**

this way we do **not** completely **ignore** the **magnitude of gradient**

Hinton looks differently at **update rule** of **Rprop**: recall the **update rule**

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \eta^{(t)} \odot \text{sign}(\mathbf{g}^{(t)})$$

We can write **alternatively** as

learning rates are updated entry-wise

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \eta^{(t)} \odot \frac{\mathbf{g}^{(t)}}{|\mathbf{g}^{(t)}|}$$

where $|\cdot|$ operates **entry-wise**

⁵Click to check out the lecture notes!

RMSprop: Root Mean Square Propagation

Hinton suggests that we *replace the denominator* with

moving average of root mean square of the gradients

This means: *starting with some $\mathbf{v}^{(0)} = \mathbf{0}$ we determine in iteration t*

$$\mathbf{v}^{(t)} = \beta \mathbf{v}^{(t-1)} + (1 - \beta) |\mathbf{g}^{(t)}|^2$$

for some $\beta < 1$ and normalize $\mathbf{g}^{(t)}$ with $\sqrt{\mathbf{v}^{(t)}}$

Initiate $\mathbf{v}^{(0)} = \mathbf{0}$

for $t = 1 : T$ do

...

$\mathbf{v}^{(t)} = \beta \mathbf{v}^{(t-1)} + (1 - \beta) |\mathbf{g}^{(t)}|^2$ # compute moving average

$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \eta \frac{\mathbf{g}^{(t)}}{\sqrt{\mathbf{v}^{(t)}}}$ # normalize by RMS

...

end for

RMSprop: Root Mean Square Propagation

We may note that RMSprop can be observed as

SGD with Rprop-inspired learning rate scheduling

Note that the exact form of RMSprop has more details!

In typical implementations of RMSprop, we set

- learning rate to a some constant: typical choice $\eta = 0.01$
- β to be close to 1: typical choice $\beta > 0.9$

The complete form of RMSprop is also available in module optim of PyTorch

```
>> import torch  
>> torch.optim.RMSprop()
```

Adaptive Momentum Estimation

Most recent implementations use the **optimizer**

Adaptive Momentum Estimation: Adam

that was proposed by **Kingma and Ba in 2015**⁶

The idea of **Adam** is straightforward: it combines **RMSprop** with **momentum**

*it combines the **strength** of both approaches*

In simple words: **Adam** suggests that we use

- **momentum** for *updating the weights*
- **Rprop-inspired** approach for *normalization* \equiv **scheduling**

⁶Click to check out the original paper!

Adaptive Momentum Estimation

We can think of Adam as below

```
Initiate  $\mathbf{m}^{(0)}$  and  $\mathbf{v}^{(0)}$ 
for  $t = 1 : T$  do
    ...
     $\mathbf{m}^{(t)} = \beta_1 \mathbf{m}^{(t-1)} + (1 - \beta_1) \mathbf{g}^{(t)}$           # compute momentum
     $\mathbf{v}^{(t)} = \beta_2 \mathbf{v}^{(t-1)} + (1 - \beta_2) |\mathbf{g}^{(t)}|^2$       # compute RMS
     $\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \eta \frac{\mathbf{m}^{(t)}}{\sqrt{\mathbf{v}^{(t)}}}$           # move with normalized momentum
    ...
end for
```

Attention

Pseudo codes in this lecture give **the main ideas**: there are **further details** and **numerical tricks** to make these algorithms **robust and stable** in practice

Adaptive Momentum Estimation

In **typical** implementations of **Adam**, we set

- learning rate to a some **constant**: **typical choice** $0.001 < \eta < 0.01$
- β_1 to be **close** to 1: **typical choice** $\beta_1 = 0.9$
- β_2 to be **closer** to 1: **typical choice** $\beta_2 = 0.99$

Just check out the **PyTorch** implementation in the **module optim**

```
>> import torch  
>> torch.optim.Adam()
```

Other Optimizers: First Order vs Second Order

There is a long list of modified gradient descents

Press Tab after typing `torch.optim`. to see how long it is!

In some particular applications, we may need to learn a new one: it is hence good to know these two terms

- First-order optimizers that use only gradient, i.e., first-order derivatives
 - ↳ What we had in this section were all first-order
- Second-order optimizers also use Hessian, i.e., second-order derivatives
 - ↳ These approaches are inspired by Newton's method that shows convergence of gradient descent is boosted if we multiply gradient with inverse of Hessian
 - ↳ They have typically better convergence behavior
 - ↳ Finding Hessian is a huge computation: practical algorithms usually approximate Hessian; but, they still need high computation