

Applied Deep Learning

Chapter 3: Advancing Our Toolbox

Ali Bereyhi

ali.bereyhi@utoronto.ca

Department of Electrical and Computer Engineering
University of Toronto

Fall 2025

Tuning Hyperparameters

A problem that we intuitively discussed but *left open* in Chapter 1 was

How can we tune the hyperparameters of a model?

In this section, we are going to *get the genie out of the bottle!*

- + *Shouldn't we set hyperparameters as much as we could?! Make NNs as deep and wide as our computer let?*
- *Not really! We don't need very deep and wide NNs always*
- + *But, what if we are computationally strong?! Then, we are fine! Right?!*
- *No! In fact with large NNs we can do the so-called overfitting!*

Let's see a *very simple example!*

Example: Fitting Polynomial from Noisy Data

We start by a *classical example* which is *not that of NN* we expect

We have a *machine* which gets *real-valued x* and *returns*

$$y = x^2 + 3x + 3$$

We however *don't know this relation*: the only thing that we know is that the *inputs* and *labels* are related *via a polynomial*

We invoke ML to *learn this machine*

Let's start with making the *ML components*, i.e.,

- ① Dataset
- ② Model
- ③ Loss

Example: Polynomial Fitting - Dataset

We start by collecting data: we give input x_b to this machine and measure its output for a batch of inputs. Our measurements are however noisy, i.e.,

$$v_b = x_b^2 + 3x_b + 3 + \varepsilon_b$$

where ε_b is noise with bounded magnitude, i.e., $|\varepsilon_b| \leq \alpha$ for some constant α

We make our dataset as

$$\mathbb{D} = \{(x_b, v_b) : i = 1, \dots, B\}$$

Example: Polynomial Fitting - Model

We know that machine is polynomial: we assume a polynomial model

$$y = w_0 + w_1 x + w_2 x^2 + \dots + w_P x^P$$

for some integer order P


We can write it down as

$$\begin{aligned} y &= w_0 + w_1 x + w_2 x^2 + \dots + w_P x^P \\ &= \underbrace{\begin{bmatrix} w_0 & w_1 & w_2 & \dots & w_P \end{bmatrix}}_{\mathbf{w}^\top} \underbrace{\begin{bmatrix} x^0 \\ x^1 \\ x^2 \\ \vdots \\ x^P \end{bmatrix}}_{\mathbf{h}} = \mathbf{w}^\top \mathbf{h} \end{aligned}$$

Example: Polynomial Fitting - Model

We can look at our **model** as an **NN** with **dummy neurons**

h is what we get from **hidden layer** and **w** includes **weights of output layer**



The key hyperparameter in this network is P

Example: Polynomial Fitting - Loss

We have a **regression** problem: *recall that*

in **regression**, the **labels** are **real-valued**

We use **squared error** as the **loss** function, i.e.,

$$\mathcal{L}(y, v) = (y - v)^2$$

for **measurement** v and NN's **output** y

Now, the components are **ready**

*let's start **training***

Example: Polynomial Fitting - Training

For training, we follow what we *already learned* in previous lectures

- ① We split \mathbb{D} into a *training dataset* and *test dataset*
- ② We start use *gradient descent* to *train* over the *training dataset*
- ③ We *test* our *trained model* over the *test dataset*

Before we go on with *training*, let's take a look back

Naive conclusion was that *making the NN large* is *always good*; if so


when we *increase P*, we should always see *lower test risk*

However, it's *not* the case!

Let's see how the *test risk* changes against *hyperparameter P*

Example: Polynomial Fitting

Test risk against the hyperparameter P looks like the curve below!



Over and Underfitted Model

- + What is happening here?
- As we pass $P = 2$, we are overfitting!

Overfitting

Overfitting occurs when training fits the model, i.e., NN, to the training dataset, so that it does not generalize to new data-points

We may also pay attention to the term **generalize** in this definition

Generalization


We say a **model**, i.e., **NN**, **generalizes** well if not only its **training risk**, but also its **test risk** is **small**

In simple words:

trained NN generalizes \equiv it does what we want on new data


Overfitting: Polynomial Fitting

Let's take a look back on polynomial fitting example: for large P the NN fits very well to the training data, but it deviates greatly from the true function



Overfitting: Polynomial Fitting

We can see the importance of hyperparameter tuning: if we set P to a right choice; then, our NN generalizes well, i.e., it closely track the true function



Underfitting

The other side of the coin is **underfitting**: it happens when our NN does not have enough parameters to **train**

Underfitting

Underfitting occurs when the **model**, i.e., NN, **neither** fits to the **training** dataset, nor **generalizes** to new data


We would also need to **prevent underfitting**; however,

with **current NNs**, **underfitting** can **hardly** occur

This is why it's **less discussed** in the literature

Underfitting: Polynomial Fitting

A linear model **underfits** in our example: *setting $P = 1$ will lead to a line that can never fit our training dataset*



Validation: First Step Against Overfitting

- + What is connection to our main task, i.e., hyperparameter tuning?
- Well! Before everything, we need to tune the hyperparameters right to avoid over or underfitting

Hyperparameter tuning is done by validation: we change hyperparameters among possible choices and for each choice, we validate our model


In validation we train the NN with the specified hyperparameters and then test it on a validation set separated from training and test sets

We set hyperparameters to the choice that

gives minimal validation risk \equiv generalizes the best

Validation: Polynomial Fitting

In fact **what we did** in our **dummy example** was **validation**



Why Overfitting Happens

- + But can we really do *hyperparameter tuning* in a very deep NN?
- Not really! We may tune some *general hyperparameters* like *number of layers*, but *cannot really do a complete validation*

In NNs, we invoke *other approaches* as well to combat *overfitting*

- *Regularization*
- *Dropout*
- *Data Augmentation*
- ...

To understand *these approaches*, we should first answer *the following question*

When does *overfitting* happen in a NN?

Let's take a look!

Why Overfitting Happens: Model Capacity

We know the initial answer: in our *dummy example*, it happened because
we assumed *large polynomial order*

In other words

our model was too complex for our learning task

We can extend this to NNs: *overfitting* happens when *the model is too complex*, i.e., it's suited for learning *complicated functions*

When does *overfitting* happen in a NN? It happens when

- ① *the model has a large capacity*

Though *model capacity* has a concrete definition, for our purpose

model capacity \equiv ability of *model* to learn *different functions*

Why Overfitting Happens: Dataset Size

- + But how can we **find this** out? It does **not** seem to be **easy!**
- That's right! This is why we look into **other reasons** as well

Let's get back to our dummy example: *this time we check it a bit differently*


In our polynomial example, we consider an **overfitted model** with $P = 5$ and train it on two randomly generated datasets

- ① a dataset with 20 data-points
- ② a dataset with only 4 data-points

After training: we compare trained models with the true function

Why Overfitting Happens: Dataset Size

As we can see: **overfitting** is **reduced** as we **increase the number of data-points**



Why Overfitting Happens: Dataset Size

This is a general behavior: if we have a large enough dataset the model cannot really overfits too much

- + How large it should be?
- It depends on the NN

A general rule is that the more learnable parameters the model has, the larger the training dataset should be

So, we can add to our answers

When does overfitting happen in a NN? It happens when

- ① the model has a large capacity
- ② our training dataset is small

Why Overfitting Happens: Co-Adaptation

Another way to see **overfitting** is to look at how **model parameters** change as **optimizer** iterates. To see it, let's get back to our **dummy polynomial-fitting NN**

Consider the following setting: we have a **high-capacity NN with $P = 5$** and **dataset with 8 noisy samples**. We train this NN using **full-batch SGD**

We now take a look at the trained NN at different iterations: recall that the vector of **model parameters** is

$$\mathbf{w} = \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_5 \end{bmatrix}$$

We start with **vector of all zeros** and keep on going

Why Overfitting Happens: Co-Adaptation

Recall that the *ground truth* \mathbf{w}^* for our polynomial machine

Now, let's look at few iterations

$$\mathbf{w}^* = \begin{bmatrix} 3 \\ 3 \\ 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

$$\mathbf{w}^{(0)} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix} \rightsquigarrow \mathbf{w}^{(t_1)} = \begin{bmatrix} 2.61 \\ 2.36 \\ 0.71 \\ 0.01 \\ 0.02 \\ 0.01 \end{bmatrix} \rightsquigarrow \mathbf{w}^{(t_2)} = \begin{bmatrix} 3.03 \\ 2.97 \\ 0.98 \\ 0.21 \\ 0.12 \\ 0.09 \end{bmatrix}$$

$\mathbf{w}^{(t_2)}$ looks good! But, what if we keep on training


$$\mathbf{w}^{(t_2)} = \begin{bmatrix} 3.03 \\ 2.97 \\ 0.98 \\ 0.21 \\ 0.12 \\ 0.09 \end{bmatrix} \rightsquigarrow \dots \rightsquigarrow \mathbf{w}^{(t_n)} = \begin{bmatrix} 2.36 \\ 4.43 \\ 3.13 \\ -2.1 \\ 1.98 \\ -1.2 \end{bmatrix}$$

Why Overfitting Happens: Co-Adaptation

Let's formulate what we observed

Weights start to get *close to what we want* up to *intermediate number of iterations t_2* . But, by *further training* they start to *deviate* \equiv *overfit*

We can also see this behavior in the figure below



Why Overfitting Happens: Co-Adaptation

This behavior is *co-adaptation of the parameters*

In initial iterations, NN fits to *the true model*: since data comes from a quadratic function, the *first iterations* of SGD

update majorly w_0 , w_1 and w_2


After NN has *gone* close to *ground truth*, it starts to *overfit*: due to *noise*, *quadratic model can't perfectly fit*; thus,

w_3 , w_4 and w_5 try to *co-adapt*, i.e., *compensate the gap caused by noise*

In simple words: in *first iterations* NN learns *true function*; however, at some point it starts to learn *noise*!

Co-adaptation is the most observable implication of overfitting with NNs

Co-Adaptation: Typical Learning Curve



Why Overfitting Happens: Final List

So, let's complete the answer list

When does overfitting happen in NNs? *It happens when*

- ① the model has a **large capacity**
- ② our training dataset is **small**
- ③ due to large number of training iterations **co-adaptation** occurs

Attention: sources are mutually related

If we have a very **large model capacity**, i.e., **very deep** with **too many neurons**; then, training it **by a small dataset** leads to **overfitting**, especially if we **keep on training for too many epochs!**

- + Now, can we do anything to avoid **overfitting**?
- Yes! Depending on what we see as **source**, we use **different tricks**

Classical Solutions to Overfitting

Overfitting happens when

- ① the model has a **large capacity**
- ② our training dataset is **small**
- ③ due to large number of training iterations **co-adaptation** occurs

The key tricks to address overfitting in each of these cases are

Done

- ① We can tune the **hyperparameters** to **restrict** the **NN's capacity**
 - ↳ For instance, we can **validate** our FNN with **2, 3 and 4 hidden layers** and choose the model with **minimal validation risk**
- ② We can **increase** our dataset by the so-called **data augmentation**
 - ↳ For instance, we can **add rotated** and **shifted** versions of **images** inside the dataset with **the same label**: a **rotated** image of a **dog** is still a **dog**!
- ③ We can **regularize** our empirical risk to **penalize co-adapted solutions**
 - ↳ For instance, we can **drop out** randomly some **neurons** in each **mini-batch**

Wait

Next

Training by Penalized Risk: *Regularization*

Regularization aims to resolve **overfitting** by treating **co-adaptation**

Let's recall **co-adaptation** in our **dummy polynomial fitting NN**: we set $P = 5$ and train our NN via the **noisy samples** inside **training dataset**; clearly,

as **training** proceeds, **empirical risk drops**

In our particular example with

$$\mathbf{w}^{(0)} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix} \rightsquigarrow \mathbf{w}^{(t_2)} = \begin{bmatrix} 3.03 \\ 2.97 \\ 0.98 \\ 0.21 \\ 0.12 \\ 0.09 \end{bmatrix} \rightsquigarrow \mathbf{w}^{(t_n)} = \begin{bmatrix} 2.36 \\ 4.43 \\ 3.13 \\ -2.1 \\ 1.98 \\ -1.2 \end{bmatrix}$$

This means that $\hat{R}(\mathbf{w}^{(t_n)}) \leq \hat{R}(\mathbf{w}^{(t_2)}) \leq \hat{R}(\mathbf{w}^{(0)})$

Training by Penalized Risk: *Regularization*

Regularization follows this idea: can we modify empirical risk, such that it stops dropping after t_2 ?

Let's continue with our example: assume risk's value at each \mathbf{w} is

$$\hat{R}(\mathbf{w}^{(t_3)}) = 0.001 \quad \hat{R}(\mathbf{w}^{(t_2)}) = 0.01 \quad \hat{R}(\mathbf{w}^{(0)}) = 100$$

We may note that as training progresses, vector $\mathbf{w}^{(t)}$ becomes larger. So, what if we add a penalty to risk that is proportional to $\|\mathbf{w}^{(t)}\|^2$: this way when risk becomes too small, this penalty becomes large and thus the sum increases. Let's look at this sum at different iterations

$$\tilde{R}(\mathbf{w}^{(0)}) = \hat{R}(\mathbf{w}^{(0)}) + \|\mathbf{w}^{(0)}\|^2 = 100$$

$$\tilde{R}(\mathbf{w}^{(t_2)}) = \hat{R}(\mathbf{w}^{(t_2)}) + \|\mathbf{w}^{(t_2)}\|^2 = 19.04$$

$$\tilde{R}(\mathbf{w}^{(t_n)}) = \hat{R}(\mathbf{w}^{(t_n)}) + \|\mathbf{w}^{(t_n)}\|^2 = 44.761$$

Training by Penalized Risk: Regularization

Penalized risk shows a different behavior

$$\tilde{R}(\mathbf{w}^{(0)}) = 100 \quad \tilde{R}(\mathbf{w}^{(t_2)}) = 19.04 \quad \tilde{R}(\mathbf{w}^{(t_n)}) = 44.761$$

From above values, we can say: if we apply **SGD** to minimize **penalized risk** we **may get** from $\mathbf{w}^{(0)}$ to $\mathbf{w}^{(t_2)}$; however, we will **not** get from $\mathbf{w}^{(t_2)}$ to $\mathbf{w}^{(t_n)}$

This idea is called **regularization** which can **prevent** NNs from **overfitting**

Regularization

In **training** with **regularization**, we minimize a **penalized (regularized)** form of the **empirical risk**, i.e.,

$$\min_{\mathbf{w}} \hat{R}(\mathbf{w}) + \Pi(\mathbf{w}) \quad (\text{Regularized Training})$$

$\Pi(\mathbf{w})$ is a **penalty** that describes the behavior of \mathbf{w} in the case of **overfitting**

Classical Regularization Approaches

There are various regularization penalties: some **important** ones are

- ℓ_2 or **Tikhonov regularization** in which we add a term proportional to $\|\mathbf{w}\|^2$

$$\Pi(\mathbf{w}) = \lambda \|\mathbf{w}\|^2$$

- This way we **avoid** very **large weights**
- This **prevents** perfect fit to **training** dataset **reducing** chance of **overfitting**
- ℓ_1 or **Lasso regularization** in which we add a term proportional to $\|\mathbf{w}\|_1$

$$\Pi(\mathbf{w}) = \lambda \|\mathbf{w}\|_1 = \lambda \sum_{i=1}^D |w_i|$$

- This way enforce **w** to be **sparse**, i.e., to have too many zeros
- This way we **reduce** the **capacity** of **NN** and thus prevent **overfitting**

Regularizing by Dropout

A less conventional regularization approach is **dropout** that was proposed by Hinton et al. first in [their 2012 paper](#) and then in [their 2014 paper](#):¹ the idea is at the same time **easy** and **effective**


for each **training** iteration, we **deactivate** some **nodes** of NN **at random**

or in other words we **drop them out**

¹Click to check out the papers!


Dropout: Schematic

Let's say this is the **dense NN**




Dropout: Schematic

For *first forward-backward* we select few **nodes** in *each layer*




Dropout: Schematic

For second forward-backward we select new nodes in each layer at random




Dropout: Schematic

For *next forward-backward* we select again few *nodes* in *each layer* at random



Dropout: Schematic

At the end, we **average** the gradients determined over these **reduced NNs**



Dropout: Intuition

- + But, why does *dropout* work?
- We can explain it *heuristically*

Recall that observing *overfitting* means that *NN* is *larger* than *required*. With *dropout*, in each iteration we train a *smaller version of NN*

- We *randomly* switch among these *smaller versions*
- Many of these *smaller versions* do not *overfit*

We can look at the training loop with *dropout* as an averaged training of these *smaller NNs*; hence, training loop gets *less chances* to *overfit*

- + Why do we do it *randomly*? Why not *sticking* to one *smaller version*?
- Not all *smaller NNs* are *good*, and we *cannot* check all of them: it's *exponentially hard* to *check all smaller NNs*

Dropout: Training Loop

How does training change with **dropout**? Training with **dropout** is exactly as before. Say we use **mini-batches**: for each **mini-batch**

- we compute the **gradient** by **forward** and **backpropagation**
- we give the **gradient** to the **optimizer** to apply the next **iteration**

The **only thing** that is **different** now is that we set the **output** of some nodes to **zero** in the **forward pass** \equiv only **forward propagation changes**

Let's make it concrete: when we pass **forward**, we generate **random masks** for each **layer** $\ell = 0, \dots, L$. Mask of **layer** ℓ is a vector whose length is **layer's width** and **entries are 0 or 1**, i.e., $s_\ell \in \{0, 1\}^{W_\ell}$. Entries of s_ℓ are generated **randomly**

$$\text{each entry of } s_\ell = \begin{cases} 1 & \text{with probability } p_\ell \\ 0 & \text{with probability } 1 - p_\ell \equiv \text{dropout probability} \end{cases}$$

Dropout: Forward Propagation

Let's show generation of **random mask** s_ℓ by following notation

$$s_\ell = \text{mask}(\mathcal{W}_\ell | p_\ell)$$

We are going to do **forward** propagation for each **data-point** as

DropoutForwardProp():

- 1: Initiate with $y_0 = x$
- 2: **for** $\ell = 0, \dots, L$ **do**
- 3: Generate $s_\ell = \text{mask}(\mathcal{W}_\ell | p_\ell)$ # random mask
- 4: Set $y_\ell = y_\ell \odot s_\ell$ # dropout nodes
- 5: Add $y_\ell[0] = 1$ and determine $z_{\ell+1} = \mathbf{W}_{\ell+1}y_\ell$ # forward affine
- 6: Determine $y_{\ell+1} = f_{\ell+1}(z_{\ell+1})$ # forward activation
- 7: **end for**
- 8: **for** $\ell = 1, \dots, L + 1$ **do**
- 9: Return y_ℓ and z_ℓ
- 10: **end for**

Dropout: Backpropagation

The backpropagation goes exactly as before: of course those *outputs* that were *dropped out* participate with value *zero* in gradient computation

One final *piece of trick*

After *training* is over, we *scale weights* of each layer with its *retain probability* p_ℓ : say T is the *last iteration* of training loop; then, we finally do

$$\mathbf{W}_\ell^{(T)} \leftarrow p_\ell \mathbf{W}_\ell^{(T)}$$

- + Why do we do that?
 - Well! It's *practically* understood; however, we can justify it as follows: each weight could be *what has been computed* with probability p_ℓ and *zero* with probability $1 - p_\ell$. We hence compute the *average*

Dropout: Implementation

Dropout is implemented in almost all **deep learning libraries**

```
>> import torch  
>> torch.nn.Dropout()
```

Typical choices of retain probability p_ℓ are

- for **input layer**, i.e., **layer 0**, $p_\ell = 0.8$
- for **hidden layers** $p_\ell = 0.5$

It's generally suggested to **drop out** more at **hidden layers**

With dropout forward pass changes in **training** and **evaluation**

- In **training** we use random mask $\rightsquigarrow \text{model.train()}$
- In **evaluation** we **don't** use random mask $\rightsquigarrow \text{model.eval()}$