

Applied Deep Learning

Chapter 3: Advancing Our Toolbox

Ali Bereyhi

ali.bereyhi@utoronto.ca

Department of Electrical and Computer Engineering
University of Toronto

Fall 2025

Data Preparation

Frankly speaking, preparing data for **training** and **testing** is

most time-consuming part of a **practical project**

- + How **hard** it could be? Really **harder** than finding **right hyperparameters**, **regularizing** or **adjusting the optimizer**?
- Sure! We get used to such **design tasks** and start to **have feeling** about **NNs**, as they **repeat** so much. The **main thing** that is **new** is **data**

How to **prepare data** that **works well** for our **purpose** is an **individual topic** discussed in courses on **data science**: we only **briefly** touch it in this section

Data Preparation

Procedure of **processing raw data** into a form **suitable** for **underlying model**

Data Preprocessing Procedures

There is a long list of **techniques** for **data preparation**

- **Data augmentation** which we do when *training dataset is too small*
- **Data cleaning** that we do to either remove or modify unwanted samples in *training dataset*
 - ↳ Samples like outliers, duplicates and nulls
- **Data transform** which aims to transform data into a form that lead to more robust training
- **Dimensionality reduction** which reduces the redundancy by extracting lower-dimensional features with minimal confusion from samples
- ...

We discuss the first two items in this section: *let's start with the first one that we already have some idea about*

Expanding Dataset via Augmentation

Recall that one conclusion from overfitting was that dataset is too small

In practice, we may expanding our dataset by augmenting it

- + Say we have a set of cat and dog images! How can we expand it?
- Well! Let's take a look at this simple example!

Say we have a dataset of cat and dog N -pixel images. We write it as

$$\mathbb{D} = \{(\mathbf{x}_b, v_b) : b = 1, \dots, B\} \quad v_b = 0 : \text{cat} \quad v_b = 1 : \text{dog}$$

Say \mathbf{x}_1 is pixel-vector of a cat image. We are given function $\mathcal{A} : \mathbb{R}^N \mapsto \mathbb{R}^N$: it gets \mathbf{x}_1 and returns $\hat{\mathbf{x}} = \mathcal{A}(\mathbf{x}_1)$ which satisfies two conditions

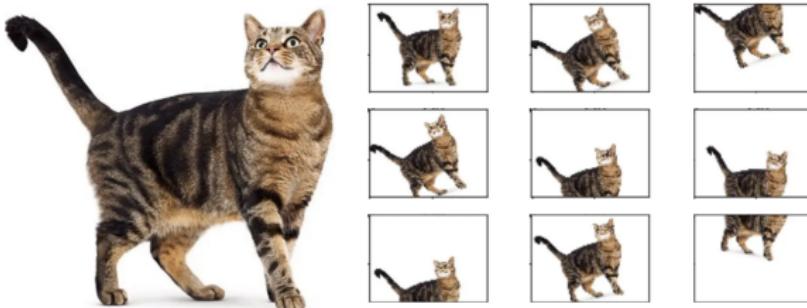
- ① After plotting $\hat{\mathbf{x}}$ we still see a cat
- ② This new cat image does not belong to the dataset, i.e., $(\hat{\mathbf{x}}, 0) \notin \mathbb{D}$

We can then expand our dataset as $\mathbb{D} \leftarrow \{(\hat{\mathbf{x}}, 0)\} \cup \mathbb{D}$!

Data Augmentation: Example

- + But, how could we **know** such a function $\mathcal{A}(\cdot)$?
- For **images** we **know** some!

We can simply **rotate**, **shift**, **zoom-in**, **zoom-out**, **change intensity** and so on!



How can we **rotate** an **image**? Multiply it by a rotation matrix $\mathbf{A} \in \mathbb{R}^{N \times N}$

These are examples of **data augmentation** by **engineering**

we can find $\mathcal{A}(\cdot)$ **analytically** using properties of our data

Data Augmentation by Engineering

Data augmentation by engineering depends on data and learning task

- If we are classifying a set of images
 - ↳ We can apply geometric transformations, e.g., random rotating, flipping, stretching, zooming, and cropping
 - ↳ We may use kernel filters to make random filtering, e.g., changing sharpness or blurring
 - ↳ We can apply random color-space transformations, e.g., changing intensity or brightness, and exchanging RGB channels
 - ↳ We can randomly remove pixels, i.e., set their value to some reference value
 - ↳ We can randomly mix images, e.g., get multiple cat images and make a new one out of them

Note that we may train a separate NN to do any of those transforms for us: just think about the last one!

Data Augmentation by Engineering

Data augmentation by engineering depends on data and learning task

- If we are working with audio signals
 - ↳ We can apply **noise injection**, e.g., add background noise
 - ↳ We may change sampling rate to change the speed of audio data
 - ↳ We can apply **random shifts**, e.g., sample signals at a bit deviated points
- If we are dealing with text data
 - ↳ We may apply **random replacements**, e.g., replace a word with its synonyms
 - ↳ We may manipulate syntax-tree, e.g., rephrase a sentence
 - ↳ We can do **random shuffling** in some applications
 - ↳ We may apply **removal and insertion** of redundant words, e.g., so

Synthetic Data Generation

An **alternative** approach to **expand** a small dataset is to

generate synthetic data

We did this in Assignment 1 for the **dummy projectile example**

Recall we had a **projectile** with **velocity v** and **height h** : we knew by **Newton's laws** that the **hitting distance d** is given by

$$d = 0.45v\sqrt{h}$$

To make **dataset**, we **generated** lots of **velocities v_i** and **heights h_i** **at random** and for each pair we determined **d_i** by above equation

Synthetic Data vs Augmented Data

When we **generate** synthetic data

- we need to **know the process** of **data being generated** from a **seed**
- we make **new data** by **simulating the process** with a **random seed**

When we **augment** data

- we need to **know transforms** that **keep data-points inside dataset**
- we **apply those transforms** on the **existing data**

- + It sounds that **synthetic data** is only feasible in scientific problems, where we **know physics!** **Right?!**
- Until few years ago the answer was **Yes!** But, currently **No!** Nowadays, we can **generate images of what we want from noise** using **generative adversarial networks (GANs)** or **diffusion models!**

Data Augmentation: Formulation

$\mathcal{A}(\cdot)$ gets *data-point x* and *returns new data-point $\hat{x} = \mathcal{A}(x)$*

In general, we do **not** need $\mathcal{A}(\cdot)$ operate on **single** data-point

$\mathcal{A}(\cdot)$ can get *multiple samples* and generate a *new one*

For instance, *it combines multiple cat images and makes a new one*

Also, $\mathcal{A}(\cdot)$ is **not** enforced to return *data-points with same labels*

label of what $\mathcal{A}(\cdot)$ returns is only required to be a valid label

For instance, *it gets multiple cat images and makes a new dog image*; however, if our *dataset includes only cats and dogs* it should **not** return a *horse image*

We can now *formulate data augmentation more precisely*

Augmentation and Synthetic Generation: Formulation

Data Augmentation

A data augmentation technique $\mathcal{A}(\cdot)$ takes the **training dataset** as the input and returns **new samples** from **data distribution**

We can think of it as the following block

$$\text{training dataset } \mathbb{D} \rightsquigarrow \boxed{\mathcal{A}(\cdot)} \rightsquigarrow (\mathbf{x}_{\text{new},1}, v_{\text{new},1}), (\mathbf{x}_{\text{new},2}, v_{\text{new},2}), \dots$$

Synthetic Data Generation

A synthetic data generator $\mathcal{S}(\cdot)$ takes a **random seed** as the input and returns **samples** from **data distribution**

We can think of it as the following block

$$\text{random seed } \mathbf{s} \rightsquigarrow \boxed{\mathcal{S}(\cdot)} \rightsquigarrow (\mathbf{x}_1, v_1), (\mathbf{x}_2, v_2), \dots$$

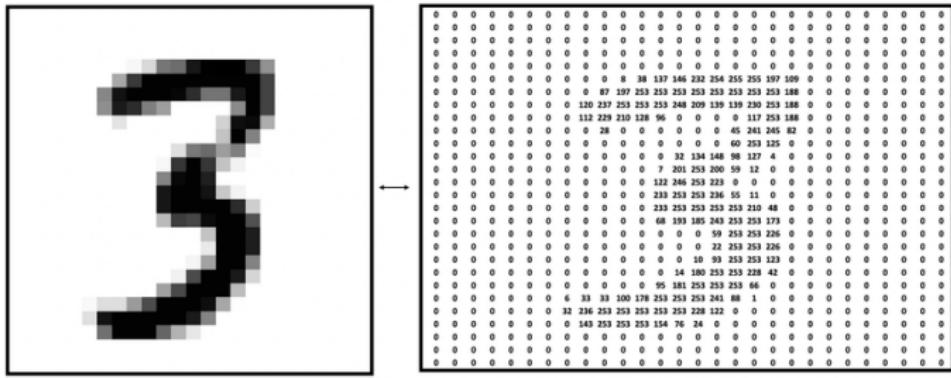
Augmentation and Synthetic Generation: Formulation

- + There is **one point** left **bothering** in **definitions!**

What does **data distribution** mean?

- In simple words it means an **abstract machine** that generates **only data-points** that **we need**
- + But, why we call it **distribution**
- Well! Let's get it **clear!**

Possible Samples of Data



Let's start with **MNIST** dataset: in **MNIST** we have 28×28 pixel images

- They are **8-bit images**: each **pixel** value is **an integer between 0 and 255**
 - These **images** are all **hand-written numbers**

But, we note that they are

- neither all 28×28 pixel 8-bit images
 - nor all possible images of hand-written numbers

Possible Samples of Data

MNIST images are *not all* 28×28 pixel 8-bit images

A 28×28 pixel image has in total 784 pixels with each pixel being one of 256 different possible values values: in total we have

$$\text{total number of images} = 256^{784} = 2^{6272} > 10^{1881}$$

MNIST has only $70,000 < 10^5$ images!

We also note that *not all* of those 2^{6272} can get into MNIST! For instance,



Possible Samples of Data

MNIST images are **not all possible** images of **hand-written numbers**

We can imagine that *among those 2^{6272} images there are **much more** than **only 70,000 images** of hand-written numbers*: just take an image of my handwriting and convert it into a 28×28 pixel image!

Space of Possible Data (Data Space)

Space of possible data is the set of all **labeled data-points** whose **labels** are valid

In our example, the *space of possible data* is

$$\mathbb{X} = \left\{ \mathbf{x} \in \{0, \dots, 255\}^{784} : \text{image of } \mathbf{x} \text{ is classified as hand-written number} \right\}$$

Of course *space of possible data* is **only a definition**: most of the time, it is **impossible** to specify it **explicitly** like above example

Data-Point as Sample of Random Object

We obviously see that a dataset is a subset of data space: MNIST is a subset of \mathbb{X} defined in the last slide and it is much much smaller

In machine learning, we have a specific way to look at datasets

dataset is collection of samples drawn randomly from the data space

This can be easily understood as the example below

Recall the data space \mathbb{X} defined in last slide

- We assume that there exist a machine with a button
 - ↳ each time we push this button the machine randomly generates data-point x from \mathbb{X}

MNIST is then generated by pushing this button for 70,000 times

Data Distribution

Data Distribution

Data distribution is the probability **distribution** by which the **dataset** has been generated from the **data space**

- + Do we know this **distribution**?!
- No! We can **neither fully** specify the **space of possible data** **nor** the **data distribution**! They are **mainly** abstract definitions!

But, we can have a **partial** understanding: assume I say such a sentence

“MNIST contains 70,000 samples drawn from **data distribution** $p(\mathbf{x})$ ”

We cannot find out what $p(\mathbf{x})$ is, but we know for sure



$$\equiv \mathbf{x}_1 \rightsquigarrow p(\mathbf{x}_1) \neq 0$$



$$\equiv \mathbf{x}_2 \rightsquigarrow p(\mathbf{x}_2) = 0$$

Data Distribution

From now on, if we get into **such a sentence** in a paper

the learner has access to a *small reference dataset* $S_T := \{(x_{T,1}, y_{T,1}), \dots, (x_{T,m_T}, y_{T,m_T})\}$ of m_T samples drawn i.i.d. from a target distribution \mathcal{D}_T

we know **what it means!**

Last note: although we do **not** have access to **data distribution**

we can in practice **approximate** it from **samples that we have collected**

For instance, if **data-points** are **heights** of different people: we can plot the **histogram** to **approximate data distribution**

Data Distribution: Practical Aspects

In practice, this looking at **data-points** as samples of a **random process** lets us use **statistical methods** to **preprocess data**

We can use **these methods** to

- realize if our **dataset** is a **good representative** of **data space**
 - ↳ Maybe we have too many non-typical data-points
- transform our **dataset** into a **better representative** of **data space**
 - ↳ Maybe we should add, remove or change some data-points

This is what we call **data cleaning**: this can be a **separate course!** So, we make it **very short** by discussing **only few practical techniques**

Data Cleaning: Duplicates

Duplication impacts model training: assume we have *training dataset*

$$\mathbb{D} = \{(\mathbf{x}_b, \mathbf{v}_b) : b = 1, \dots, B\}$$

Without any *duplication*, our *training loop* solves

$$\min_{\mathbf{w}} \frac{1}{B} \sum_{b=1}^B \mathcal{L}(y_b, \mathbf{v}_b) : y_b \text{ output of NN with weights } \mathbf{w} \text{ and input } \mathbf{x}_b$$

Now assume that we *copy* $(\mathbf{x}_1, \mathbf{v}_1)$ *by mistake* M times: the *training* on this *duplicated dataset* is

$$\min_{\mathbf{w}} \frac{1}{B + M} \sum_{b=1}^B \mathcal{L}(y_b, \mathbf{v}_b) + \frac{M}{B + M} \mathcal{L}(y_1, \mathbf{v}_1)$$

which is *not* the same thing!

Data Cleaning: Duplicates

Having the **same data-points** in dataset does **not necessarily** mean **duplication**: consider the following two simple examples

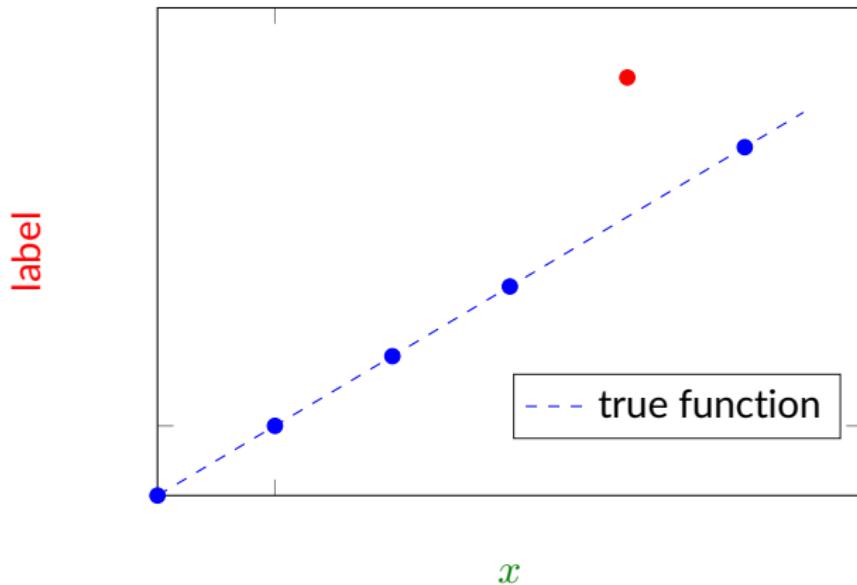
- ① We are training an NN that takes *age, education and place of birth* as input and returns *number of children* as output
 - ↳ Dataset has too many $x_b = [22, \text{Bachelor}, \text{Toronto}]$ and $v_b = 0$
 - ↳ These are **not duplicates** since they come from **independent samples**
- ② We are training an NN that takes *age and height* as input and returns *weight* as output
 - ↳ It is not **likely** to have multiple $x_b = [48, 176.42]$ and $v_b = 73.31$
 - ↳ They most probably come from the **same person** and are **duplicates**

Bingo! Terms like “come from **independent samples**” and “**not likely to have**” are used since we look at **data-points** as **samples** of a **random process**

Data Cleaning: Outliers

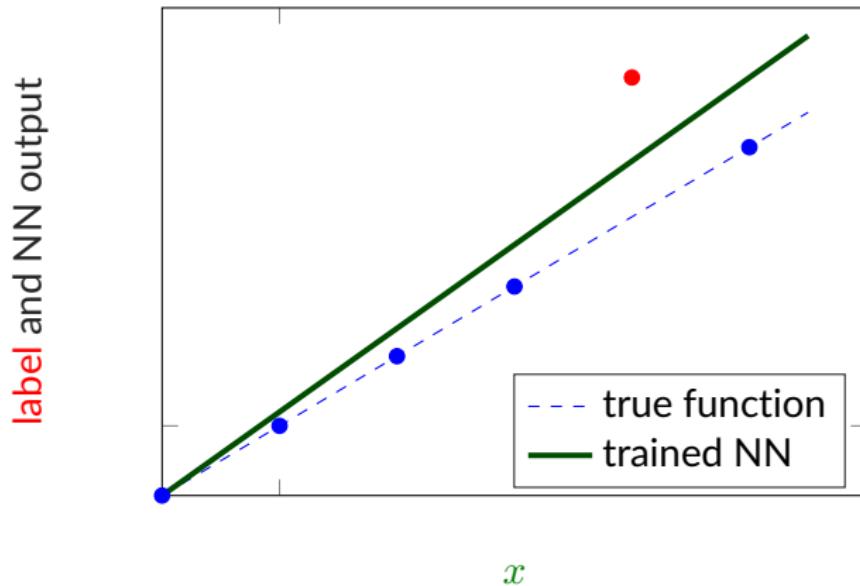
Outliers

Outliers are *data-points* that lie in an *abnormal* distance from other *samples*



Data Cleaning: Outliers

Outliers can hinder training even with good tuning and regularization



Data Cleaning: Outliers

Finding **outliers** is required for **training** of a model that **generalizes** well

There are two types of **outliers** in a **dataset**

- ① **Univariate outliers** which are detected from their **marginal distributions**
 - ↳ These **data-points** are understood to be **outliers** from an **individual variable** in them **without comparing to other variables**

We collect **heights** and **weights**: Sultan Kösen^a is **among our samples** with **height 2.51 m!** **Without checking weights**, we can say this is an **outlier**

^aTallest alive person in the world

Data Cleaning: Outliers

Finding **outliers** is required for **training** of a model that **generalizes** well

There are two types of **outliers** in a **dataset**

- ② **Multivariate outliers** which are detected from their **joint distributions**
 - ↳ These are understood to be **outliers** by comparing different variables

We collect **heights** and **weights**: a sample with **height 1.82 m** and another with **weight 24 kg** are **individually** normal; however **a sample with height 1.82 m and weight 24 kg is an outlier**

Data Cleaning: Outliers

- + How can we **handle outliers**?
- Well! It **depends**

Conventional approaches to **handle outliers** are to

- ① **remove them from training dataset**
 - ↳ It might be a **good idea** for small NNs with **low model capacity**
 - ↳ It can **hinder generalization** of our model if we have detected outliers based on **poor statistics**, e.g., **not enough samples** to understand data distribution
- ② **use loss functions that are robust against outliers**
 - ↳ They give **higher weights** to **typical** data-points
 - ↳ They give **less weights** to **outliers**
 - ↳ An example is to use **Minkowski error** instead of **squared error** for regression

Take a look at **Python** library Pandas if you need to do any **data cleaning**