

Applied Deep Learning

Chapter 4: Convolutional Neural Networks

Ali Bereyhi

ali.bereyhi@utoronto.ca

Department of Electrical and Computer Engineering
University of Toronto

Fall 2025

Computer Vision

Computer vision has been a fundamental problem in machine learning

The aim is to design a machine that can recognize patterns in visual contents

Long study on biological vision systems has been conducted

- Hubel and Wiesel studied visual cortex in late 1950s
 - ↳ They won Nobel Prize in 1981 for their discoveries
- Inspired by that study Fukushima introduced Neocognitron in 1979
 - ↳ It was a convolutional NN for unsupervised learning
- Yann LeCun proposed LeNet for supervised learning in 1989
 - ↳ This is pretty much the beginning of modern CNNs

Convolutional NNs: CNNs

Let's make an agreement: *though we all know CNN News Channel we also say*

Convolutional NN \equiv *CNN*


CNN

CNN is an *FNN* which in addition to standard fully-connected layers uses *convolutional* and *pooling* layers for *feature extraction*

- + But, what is a *convolutional* layer? What is a *pooling* layer? What do you mean by *feature extraction*?
- Well! We get there *soon*! But first let's see what the main *motivation* is

Motivational Example: Pattern Recognition

We intend to train an **NN** that classifies **MNIST** dataset with **one difference** to our earlier classification problem: *we assume that MNIST images are now included in a larger white background*



Motivational Example: Pattern Recognition

- + Can't we simply use standard **fully-connected** FNN?
- Sure! We can








Say the photos with background are **3 times larger** in height and width

- **Images** are now $84 \times 84 \equiv$ we have **7056 pixels**
- We should have **input layer with 7056 pixels** and train the NN with **MNIST**

To do the training, we should first **convert** our **MNIST images** into new **larger-size images**

Motivational Example: Pattern Recognition

We do this by **zero-padding**: all images after conversion lie on **top left corner**




Do you think after training, NN classifies "3" in **lower right corner** correctly? **No!**

Motivational Example: Pattern Recognition

- + Why does this happen?
- Our NN is trained to only look at top left corner, it will miss information anywhere else including in lower right corner
- + Can we do anything about it?
- Yes! We learned it in the last chapter: Data Augmentation

We shift MNIST images in the large background left and right, up and down and add all those shifts with same label to the dataset

Motivational Example: Recognition with Augmented Data



We should get **too many** of them!

Using a Trained FNN

- + But it sounds like **too much work** and **computation!**
- Yes! It is! and frankly speaking **it is not worth it!**

Many scientist noted that our **brain** doesn't work **like that**

*once we learn “3” we can **recognize** it **anywhere** in our vision!*


We may **initially** note that

- ① Our brain **doesn't** process the **visual field as the whole**
- ② It **searches** for patterns in **smaller fields** within our vision
 - ↳ It constructs a pattern for “3” through **training**
 - ↳ After **training**, it **scans** any visual field to see if it finds that pattern

Let's try realizing it with a NN!

Using a Trained FNN

Let's assume we have trained the following FNN on MNIST




- It gets a 784-pixel image as input
- It passes it through three fully-connected layers
- It returns the class of the image
 - ↳ If it doesn't find a class it returns \emptyset


Scanning via MNIST Trained FNN

We can mimic what *our brain* does

we *scan* the *larger image* with *background*



Scanning via MNIST Trained FNN



We go through **windows** for size 28×28

- ① At each window, we give it **784 pixels** to the FNN to classify
 - ↳ If we find a class: we **save the class** and **return True**
 - ↳ If we don't find a class: we **return False**
- ② We compute **OR** of **outputs** for all **windows**
 - ↳ If **True**: we return **the saved class**
 - ↳ If **False**: we return **\emptyset**

CNNs: Scanning via Shared Weights

The above example is a *simple CNN*

- This CNN extracts features from the image using a 28×28 filter
 - ↳ the **filter's weights** are those given by the first layer of trained FNN
 - ↳ the **features** are affine values calculated in **first layer of trained FNN**
- The scanning procedure has a specific name: convolution
 - ↳ it goes through the image by **sliding** over it via a **smaller window**
 - ↳ it determines an **affine transform** of smaller subsets of pixels
- We can look at it as a **giant FNN** with **shared** weights
 - ↳ each pixel is connected to the next layer via **affine transform**
 - ↳ this affine transform has the **same weights** for many pixels
 - ↳ **not every feature** depends on **every pixels**
 - ↳ the first layer is **not fully-connected**: it's **locally-connected**


Let's make our understanding deeper by making our first CNN!

Recognizing X

In Assignment 1, we **trained** a perceptron with 9 inputs


it gets a 3×3 image and says whether it is “X” or not

Assume we have **weights and bias**: we want to **recognize “X” in a larger image**



Recognizing X

We follow our scanning idea to recognize 3×3 "X" in a 10×10 image: we put the weights on a 3×3 filter and slide it over the image



We slide **the filter** with **stride 1** and save the outputs of perceptron on a **map**

Recognizing X


*It's enough to have **only a single 1** to recognize "X"*

0	0	0	0	0	0	0	0
0	1	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

We therefore return the **OR** of all entries in the above **map**

Recognizing X: Convolution

We really don't need to determine the **activation** after each scan: we could only save the **affine transforms in the map**



Convolution with Stride 1

image

w_1	w_2	w_3							
w_4	w_5	w_6							
w_7	w_8	w_9							

filter

w_1	w_2	w_3
w_4	w_5	w_6
w_7	w_8	w_9

feature map

$z_{1,1}$	$z_{2,1}$	$z_{3,1}$	$z_{4,1}$	$z_{5,1}$	$z_{6,1}$	$z_{7,1}$	$z_{8,1}$
$z_{1,2}$	$z_{2,2}$	$z_{3,2}$	$z_{4,2}$	$z_{5,2}$	$z_{6,2}$	$z_{7,2}$	$z_{8,2}$
$z_{1,3}$	$z_{2,3}$	$z_{3,3}$	$z_{4,3}$	$z_{5,3}$	$z_{6,3}$	$z_{7,3}$	$z_{8,3}$
$z_{1,4}$	$z_{2,4}$	$z_{3,4}$	$z_{4,4}$	$z_{5,4}$	$z_{6,4}$	$z_{7,4}$	$z_{8,4}$
$z_{1,5}$	$z_{2,5}$	$z_{3,5}$	$z_{4,5}$	$z_{5,5}$	$z_{6,5}$	$z_{7,5}$	$z_{8,5}$
$z_{1,6}$	$z_{2,6}$	$z_{3,6}$	$z_{4,6}$	$z_{5,6}$	$z_{6,6}$	$z_{7,6}$	$z_{8,6}$
$z_{1,7}$	$z_{2,7}$	$z_{3,7}$	$z_{4,7}$	$z_{5,7}$	$z_{6,7}$	$z_{7,7}$	$z_{8,7}$
$z_{1,8}$	$z_{2,8}$	$z_{3,8}$	$z_{4,8}$	$z_{5,8}$	$z_{6,8}$	$z_{7,8}$	$z_{8,8}$

The above operation is **convolution** and we show it as below

$$\text{feature map} = \text{Conv}(\text{image} | \text{filter}, \text{stride} = 1)$$

Convolution with Stride 2

image

w_1	w_2	w_3	w_2	w_3	w_2	w_3	w_2	w_3
w_4	w_5	w_6	w_5	w_6	w_5	w_6	w_5	w_6
w_7	w_8	w_9	w_8	w_9	w_8	w_9	w_8	w_9
w_4	w_5	w_6	w_5	w_6	w_5	w_6	w_5	w_6
w_7	w_8	w_9	w_8	w_9	w_8	w_9	w_8	w_9

filter

w_1	w_2	w_3
w_4	w_5	w_6
w_7	w_8	w_9

feature map

$z_{1,1}$	$z_{2,1}$	$z_{3,1}$	$z_{4,1}$
$z_{1,2}$	$z_{2,2}$	$z_{3,2}$	$z_{4,2}$
$z_{1,3}$	$z_{2,3}$	$z_{3,3}$	$z_{4,3}$
$z_{1,4}$	$z_{2,4}$	$z_{3,4}$	$z_{4,4}$

We could also play with the *stride* \equiv the step-size by which we move filter

$$\text{feature map} = \text{Conv}(\text{image} | \text{filter}, \text{stride} = 2)$$

Convolution with Stride S

Let's formulate the convolution for a general filter: assume $\mathbf{W} \in \mathbb{R}^{F \times F}$ be a **filter**, we also call it **kernel**. Let $\mathbf{X} \in \mathbb{R}^{N \times N}$ be **pixel matrix** of the **image**. We want to find the output **feature map**, i.e.,

$$\mathbf{Z} = \text{Conv}(\mathbf{X} | \mathbf{W}, \text{stride} = S)$$

It's enough to find the corresponding sub-matrix for each entry of \mathbf{Z}


$$\mathbf{Z}[i, j] = \text{sum}(\mathbf{W} \odot \mathbf{X}_{i,j})$$

where $\mathbf{X}_{i,j}$ is the corresponding $F \times F$ sub-matrix, i.e.,

$$\mathbf{X}_{i,j} = \mathbf{X}[1 + (i - 1)S : F + (i - 1)S, 1 + (j - 1)S : F + (j - 1)S]$$

Recognizing X: Pooling

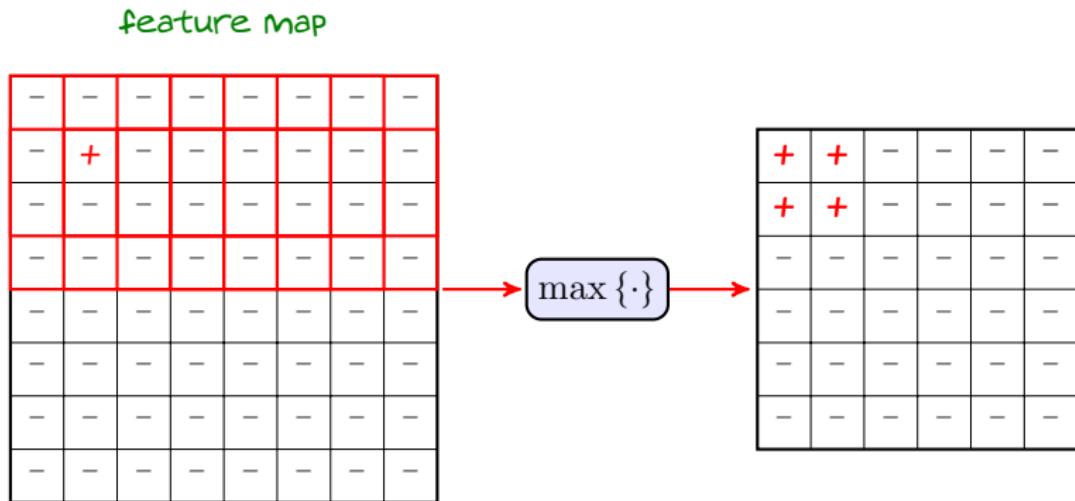
The next operation we did is called **pooling**



This is however **not conventional** to have a **pooling filter** of different size

Pooling: Max Pooling with Stride 1

The convention is to use *the same filter size as used in convolution layer*




We can now give the *feature map* after pooling to a *fully-connected FNN*: this is a *feature vector of reduced size!*

We can *repeat convolution and pooling over and over*

CNN: Simple Architecture

Our simple CNN looks like this



Convolution

Pooling

Flattening


Fully-Connected

This is a **general architecture** for CNNs, but of course we **go deeper!**

- We have more **convolutional** and **pooling layers**
- We do **high-dimensional convolutions** and **more advanced poolings**

CNN: Realistic Architectures

For instance the famous VGG-16 architecture looks like below



CNN: Connections to Biological Vision

Though it's **artificially** developed as a **computation model**: it is related to the initial model developed for description of **biological vision**

