

# Applied Deep Learning

## Chapter 4: Convolutional Neural Networks

Ali Bereyhi

[ali.bereyhi@utoronto.ca](mailto:ali.bereyhi@utoronto.ca)

Department of Electrical and Computer Engineering  
University of Toronto

Fall 2025

# Components of CNNs

CNNs consist of **three** major components

① **Convolutional layers**

- ↳ they are the **key** component
- ↳ they extract **features** from **input** by sweeping **filters** over **input**

② **Pooling layers**

- ↳ they **smooth** the extracted features

③ **Output FNN**

- ↳ they learn **label** from **final** extracted features

We now go through each of these components

# Convolution: 2D Arrays

We already know the definition of convolution for 2D arrays

## 2D Convolution

Let  $\mathbf{X} \in \mathbb{R}^{N \times M}$  be the *input matrix*: convolution of  $\mathbf{X}$  by *filter/kernel*  $\mathbf{W} \in \mathbb{R}^{F \times F}$  with *stride*  $S$  is denoted by

$$\mathbf{Z} = \text{Conv}(\mathbf{X} | \mathbf{W}, S)$$

The matrix  $\mathbf{Z}$  has  $\lfloor (N - F)/S \rfloor + 1$  rows and  $\lfloor (M - F)/S \rfloor + 1$  columns and its entry at row  $i$  and column  $j$  is computed as

$$\mathbf{Z}[i, j] = \text{sum}(\mathbf{W} \odot \mathbf{X}_{i,j})$$

with  $\mathbf{X}_{i,j}$  being the corresponding  $F \times F$  sub-matrix of  $\mathbf{X}$ , i.e.,

$$\mathbf{X}_{i,j} = \mathbf{X}[1 + (i - 1)S : F + (i - 1)S, 1 + (j - 1)S : F + (j - 1)S]$$

## Convolution: 2D Arrays

Let's see an example: assume  $\mathbf{X} \in \mathbb{R}^{4 \times 5}$  and  $\mathbf{W} \in \mathbb{R}^{2 \times 2}$

$$\mathbf{X} = \begin{bmatrix} X_{1,1} & X_{1,2} & X_{1,3} & X_{1,4} & X_{1,5} \\ X_{2,1} & X_{2,2} & X_{2,3} & X_{2,4} & X_{2,5} \\ X_{3,1} & X_{3,2} & X_{3,3} & X_{3,4} & X_{3,5} \\ X_{4,1} & X_{4,2} & X_{4,3} & X_{4,4} & X_{4,5} \end{bmatrix} \quad * \quad \mathbf{W} = \begin{bmatrix} W_{1,1} & W_{1,2} \\ W_{2,1} & W_{2,2} \end{bmatrix}$$

$\rightarrow \mathbf{X}_{1,1} \rightarrow \text{sum}(\mathbf{X}_{1,1} \odot \mathbf{W})$

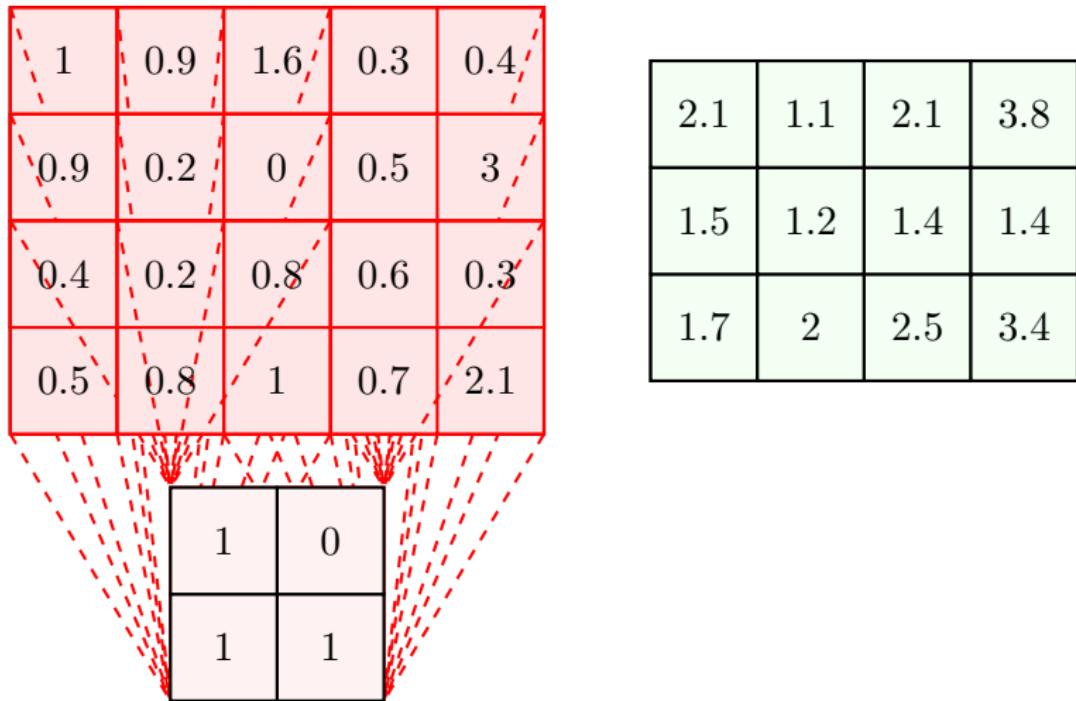
If we convolve with stride  $S = 1$

$$\mathbf{Z} = \begin{bmatrix} Z_{1,1} & Z_{1,2} & Z_{1,3} & Z_{1,4} \\ Z_{2,1} & Z_{2,2} & Z_{2,3} & Z_{2,4} \\ Z_{3,1} & Z_{3,2} & Z_{3,3} & Z_{3,4} \end{bmatrix}$$

Let's verify the dimensions of  $\mathbf{Z}$

$$\# \text{ rows} = \lfloor (4 - 2)/1 \rfloor + 1 = 3 \quad \# \text{ columns} = \lfloor (5 - 2)/1 \rfloor + 1 = 4$$

# Convolution: Numerical Example



## Convolution: 2D Arrays

Let's see an example: assume  $\mathbf{X} \in \mathbb{R}^{4 \times 5}$  and  $\mathbf{W} \in \mathbb{R}^{2 \times 2}$

$$\mathbf{X} = \begin{bmatrix} X_{1,1} & X_{1,2} \\ X_{2,1} & X_{2,2} \\ X_{3,1} & X_{3,2} \\ X_{4,1} & X_{4,2} \end{bmatrix} \quad \mathbf{W} = \begin{bmatrix} W_{1,1} & W_{1,2} \\ W_{2,1} & W_{2,2} \end{bmatrix}$$

\*

$\mathbf{X}_{1,1} \rightarrow \text{sum}(\mathbf{X}_{1,1} \odot \mathbf{W})$

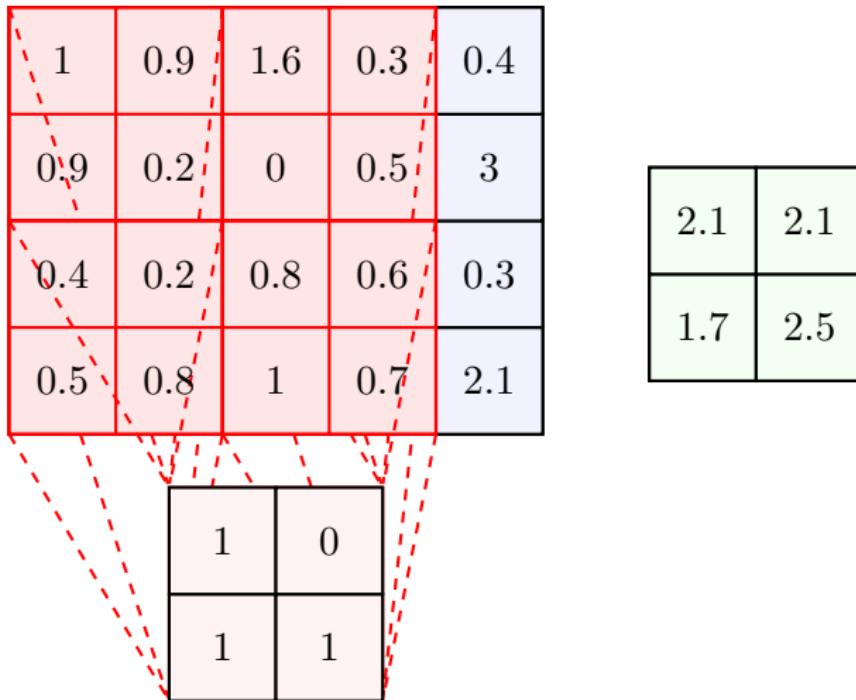
Now, we convolve with stride  $S = 2$

$$\mathbf{Z} = \begin{bmatrix} Z_{1,1} & Z_{1,2} \\ Z_{2,1} & Z_{2,2} \end{bmatrix}$$

Let's verify the dimensions of  $\mathbf{Z}$

$$\# \text{ rows} = \lfloor (4 - 2)/2 \rfloor + 1 = 2 \quad \# \text{ columns} = \lfloor (5 - 2)/2 \rfloor + 1 = 2$$

# Convolution: Numerical Example



## Convolution: Downsampling

Let's compare the result with strides 1 and 2

2.1		2.1	
1.7		2.5	

2.1	2.1
1.7	2.5

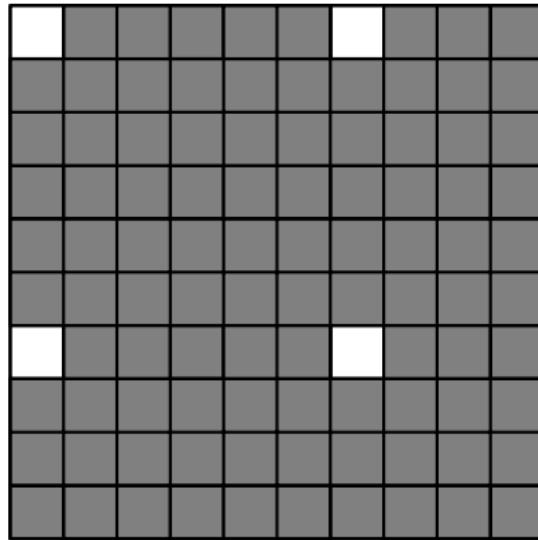
This is **downsampling** with factor 2!

### Downsampling

Downsampling with factor  $f$  drops the last  $f - 1$  of every  $f$  columns and rows

$$\text{dSample}(\mathbf{Z}|f)$$

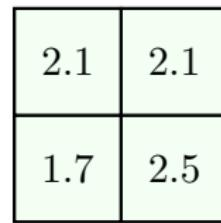
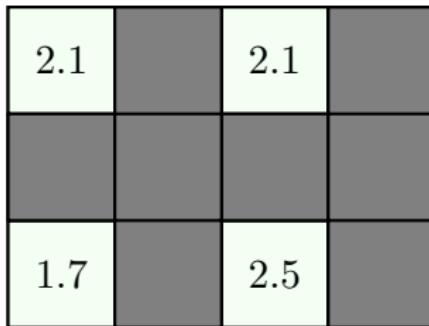
# Convolution: *Downsampling*



*Downsampling with factor 2* dSample ( $\mathbf{Z}|2$ ) *Downsampling with factor 3*  
dSample ( $\mathbf{Z}|3$ )

# Convolution: Downsampling

Let's compare the result with strides 1 and 2



## Stride as Downsampling

We can look at stride  $S$  as **downsampling** with factor  $S$ , i.e.,

$$\text{Conv}(\mathbf{X}|\mathbf{W}, S) = \text{dSample}(\text{Conv}(\mathbf{X}|\mathbf{W}, 1) | S)$$

## Convolution: *Downsampling*

So, we can make an agreement: *by default we consider unit stride, i.e.,  $S = 1$ , i.e., we drop  $S$  in convolution from now on*

$$\text{Conv}(\mathbf{X}|\mathbf{W}) = \text{Conv}(\mathbf{X}|\mathbf{W}, S=1)$$

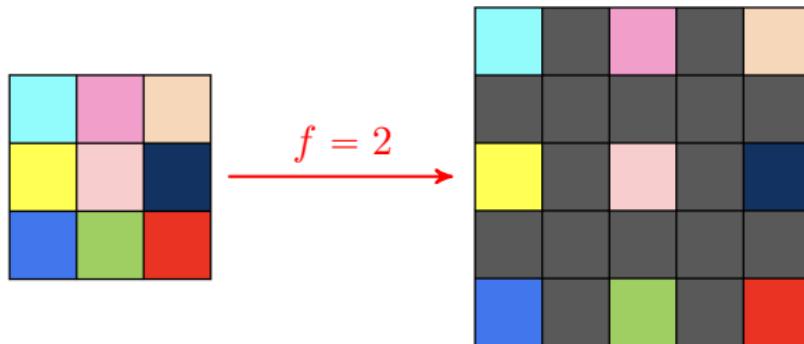
Whenever we need to convolve with  $\text{stride } S > 1$

We add an *downsampling* next to the convolution

- + Why do we do this?
- We'll see how easy things get when we want to *backpropagate*; however, it also helps to easily define having a *stride smaller than one!*

## Convolution: Upsampling

We can do the *sampling* in other way, i.e., add some zeros



### Upsampling

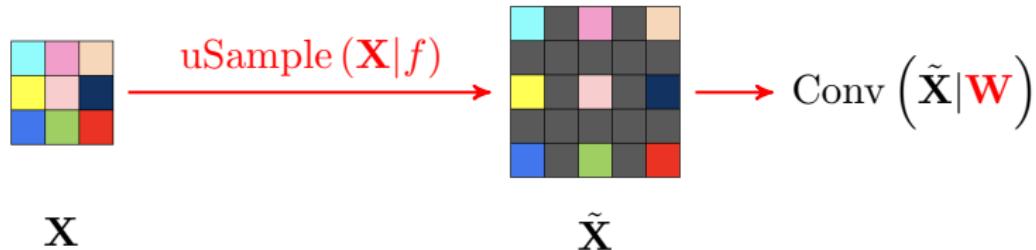
Upsampling with factor  $f$  adds  $f - 1$  rows and columns of zeros after every row and column

$$\text{uSample}(\mathbf{Z}|f)$$

## Convolution with Fractional Stride

- + Where do we use *upsampling*?
- If we want to **increase** the size of the feature map

To **increase** the size of the feature map, we can do following



This is called **fractionally-strided convolution**: for  $S < 1$  with  $1/S$  being integer

$$\text{Conv}(\mathbf{X}|\mathbf{W}, S) = \text{Conv}(\text{uSample}(\mathbf{X}|1/S)|\mathbf{W})$$

## Convolution: Resampling

We can extend **stride** to any fraction  $S = S_2/S_1$  with integer  $S_1$  and  $S_2$ : we first do **upsampling** with factor  $S_1$  and then **downsampling** with factor  $S_2$

### Moral of Story

Convolution with **stride** is equivalent with

*convolution with resampling*

Note that the **order of resampling** differs

- **Upsampling** is done always *before* convolution
- **Downsampling** is done always *after* convolution

The above interpretation of **stride** helps a lot in **backpropagation!**

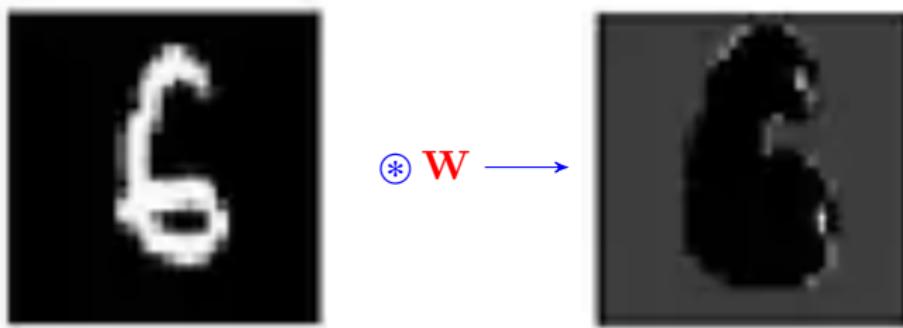
## Convolution: Padding

In our definition, feature map  $Z$  has *smaller* dimensions than  $X$  even with *stride  $S = 1$* . We can play with *dimensions of  $Z$*  via *zero-padding*

- + Why should we be interested in *changing dimensions of  $Z$* ?
- We'll see *multiple* reasons: a simple one is that we may like to have *a same-size feature map* to sketch it *as an image* and compare it to *input*

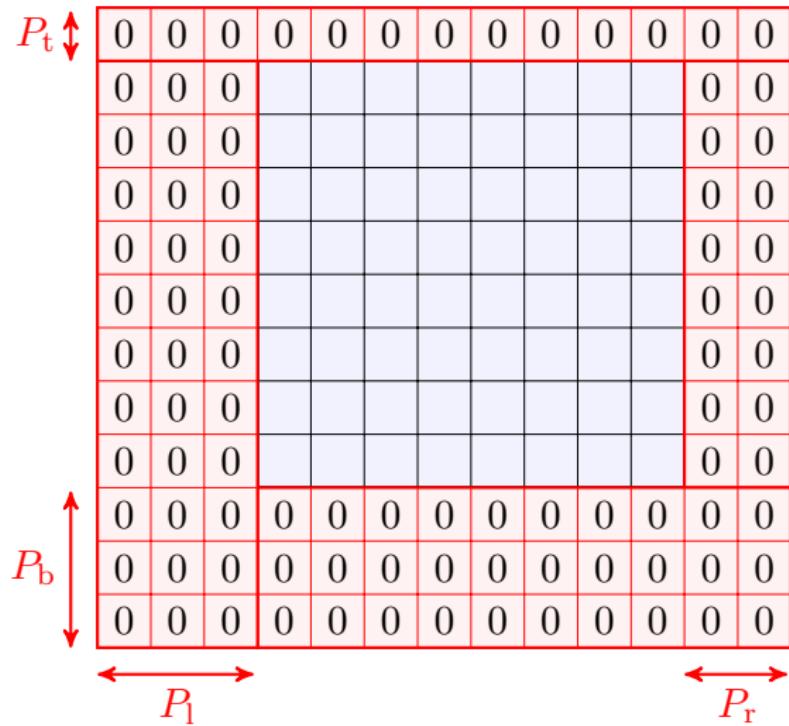
---

For instance in MNIST, we want to sketch the *feature map* as a  $28 \times 28$  image



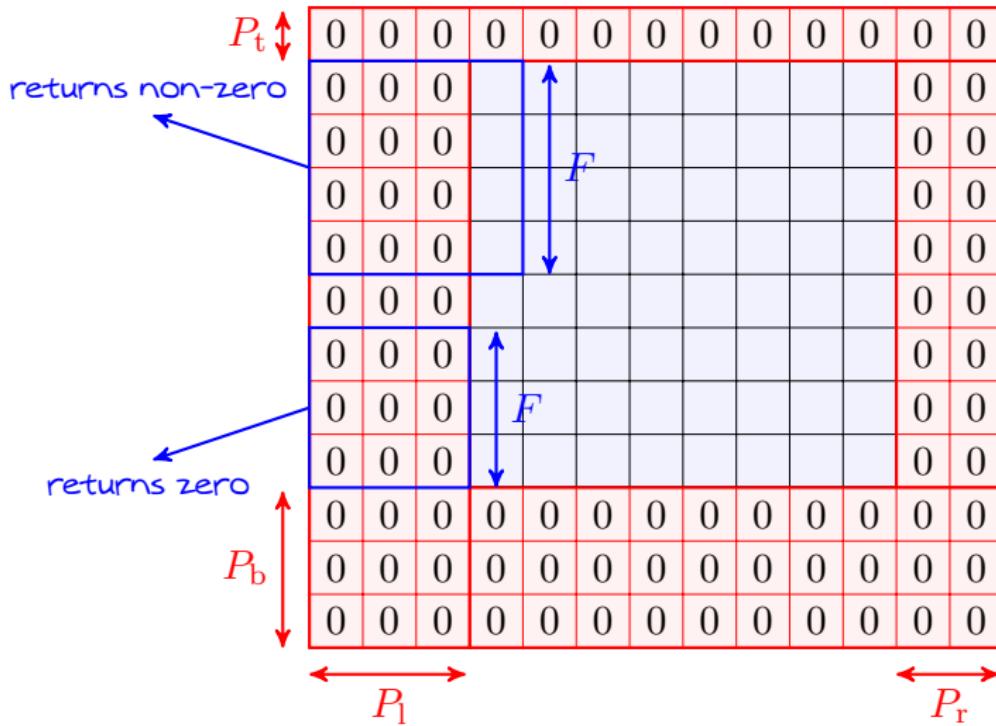
# Convolution: Zero-Padding

The trick to *resize feature map* is to *pad zeros at boundaries*



# Convolution: Zero-Padding

Typically we *pad* with widths smaller than *filter dimension F*



## Convolution: Zero-Padding

With **zero-padding**, the dimensions of **feature map** can be modified: say the **input** has  $N$  rows and  $M$  columns; then, at the feature map we have

- $\lfloor (N + P_t + P_b - F)/S \rfloor + 1$  rows
- $\lfloor (M + P_l + P_r - F)/S \rfloor + 1$  columns

In practice, we **pad symmetrically**, i.e.,  $P_b = P_t = P_r = P_l$

---

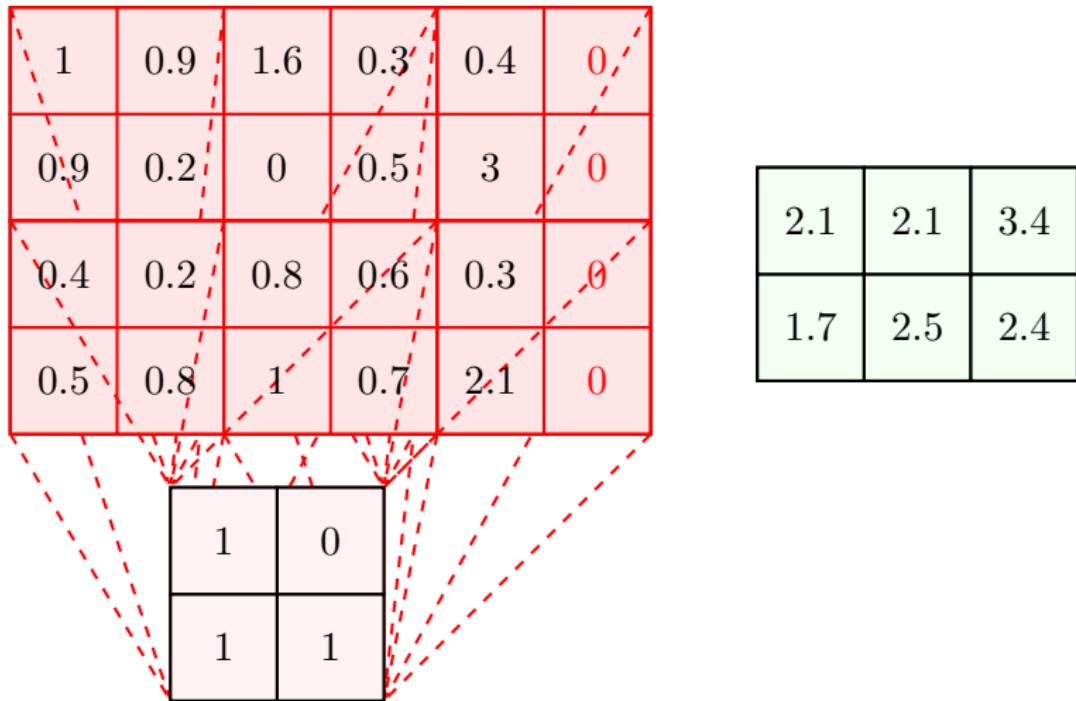
As we see the dimensions of **feature map** is a function of

**input dimensions, stride, padding length** and **filter size**

It is thus typical that we get some of these items **not specified**, e.g., we might get **input and output dimensions, stride and filter size** but not the **padding length**

we can find the **padding length** from other specifications

## Example: Stride = 2 and Padding



# Convolution: Activation

Convolution is a *spatial linear transform* on the input image

we should *activate* this *linear transform* if we go *deep*

A convolutional layer can hence be formally defined as below

## Convolutional Layer

Convolutional layer with *filter*  $\mathbf{W} \in \mathbb{R}^{F \times F}$ , *bias*  $b$  and *activation function*  $f(\cdot)$  transforms the *input map*  $\mathbf{X}$  to the *activated feature map*  $\mathbf{Y}$  as follows:

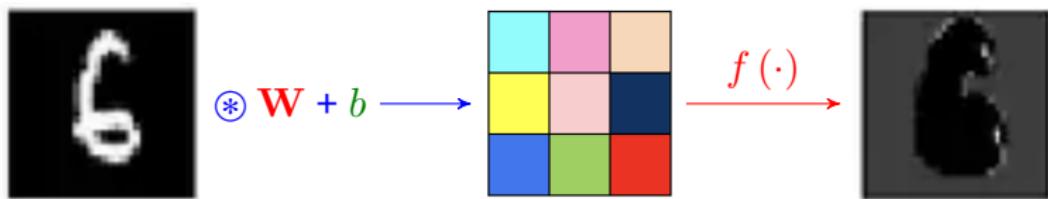
- ① it first applies *linear convolution* to find  $\mathbf{Z}$

$$\mathbf{Z} = \text{Conv}(\mathbf{X} | \mathbf{W}) + b \quad + \text{applied entry-wise}$$

- ② it then *activates*  $\mathbf{Z}$

$$\mathbf{Y} = f(\mathbf{Z}) \quad f(\cdot) \text{ applied entry-wise}$$

# Convolution: Activation

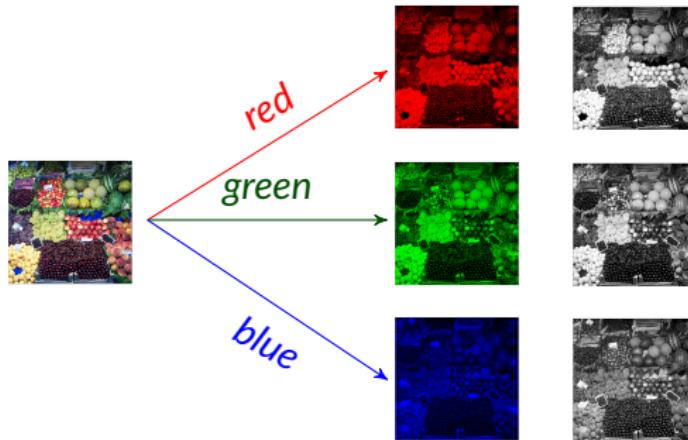


- + We did not discuss **bias** before!
- It's just a **scalar added to all entries**
- + What kinds of **activation** are used in CNNs?
- Similar to FNNs with **ReLU** being the **popular one**

## Convolution: Multi-Channel Input

What we discussed up to now holds for *single-channel images*: these are gray images that can be represented by *2D pixel arrays*, e.g., MNIST images

In practice, we have *multi-channel inputs*; for instance, *RGB images* have *three channels*: an  $N \times M$  color image is stored in the form of three  $N \times M$  matrices, one storing *red map*, another *green map*, and the other *blue map*



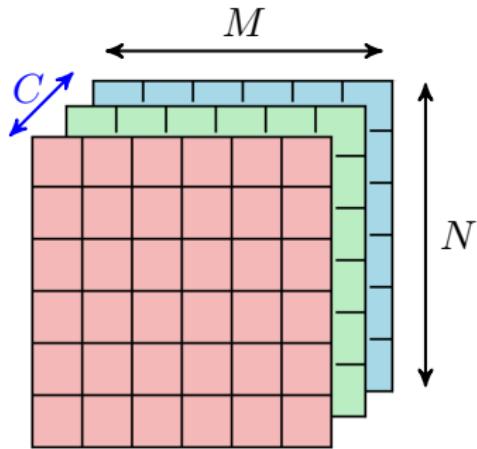
## Recap: 3D Tensors

We can think of multi-channel input as a **tensor of order 3**, i.e., a 3D array

### Reminder: Tensor of Order 3

Tensor  $\mathbf{X} \in \mathbb{R}^{C \times N \times M}$  is a collection of  **$C$  matrices** each of size  $N \times M$

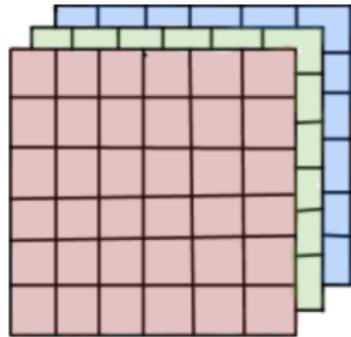
For instance, an  $N \times M$  pixel **RGB** image is tensor in  $\mathbb{R}^{3 \times N \times M}$



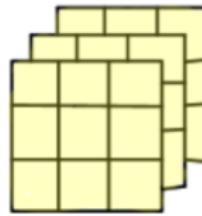
## Convolution: 3D Arrays

- + How can we extend convolution to these 3D tensors?
- We can look at 3D tensors as stack of 2D arrays

Let's try a visual example: we want to convolve RGB image with filter  $\mathbf{W}$



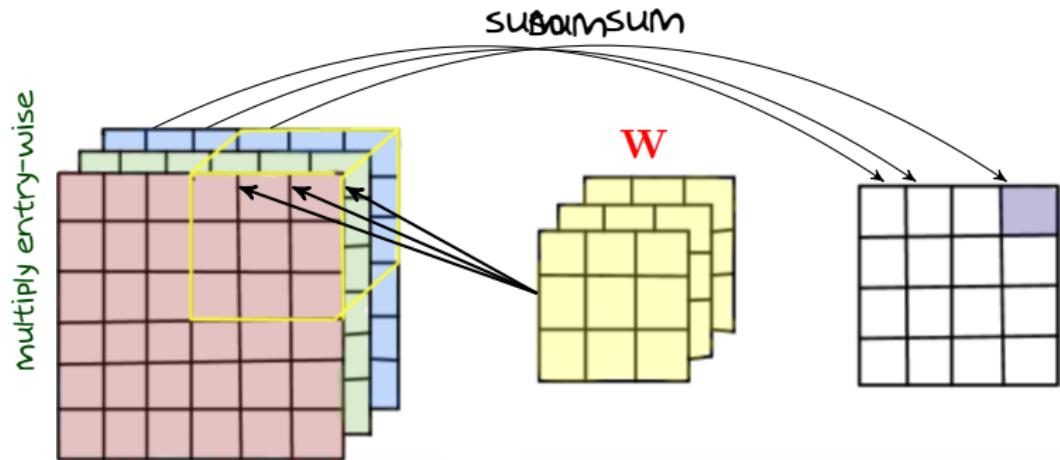
convolve with



Filter should have the same number of channels, i.e.,  $\mathbf{W} \in \mathbb{R}^{3 \times F \times F}$

## Convolution: General Form

We convolve every channel of **filter** with corresponding channel of **input**, and then *sum them up*



# Convolution: 3D Arrays

## 3D Convolution

Convolution of tensor  $\mathbf{X} \in \mathbb{R}^{C \times N \times M}$  by filter/kernel  $\mathbf{W} \in \mathbb{R}^{C \times F \times F}$  with stride  $S$  is denoted by

$$\mathbf{Z} = \text{Conv}(\mathbf{X} | \mathbf{W}, S)$$

The matrix  $\mathbf{Z}$  is a matrix with  $\lfloor (N - F)/S \rfloor + 1$  rows and  $\lfloor (M - F)/S \rfloor + 1$  columns and its entry at row  $i$  and column  $j$  is computed as

$$\mathbf{Z}[i, j] = \text{sum}(\mathbf{W} \odot \mathbf{X}_{i,j})$$

with  $\mathbf{X}_{i,j}$  being the corresponding  $C \times F \times F$  sub-tensor of  $\mathbf{X}$ , i.e.,

$$\mathbf{X}_{i,j} = \mathbf{X}[1:C, 1 + (i-1)S:F + (i-1)S, 1 + (j-1)S:F + (j-1)S]$$

## 3D Convolution: Summary

As for 2D arrays, we can *apply stride and by resampling*; thus, we write

$$\mathbf{Z} = \text{Conv}(\mathbf{X} | \mathbf{W})$$

from now on and keep in mind that

- $\mathbf{X}$  is a tensor with  $C$  channels
- $\mathbf{W}$  is a *tensor-like kernel* with  $C$  channels: some people say with depth  $C$
- $\mathbf{Z}$  is a matrix, i.e., it has a single channel

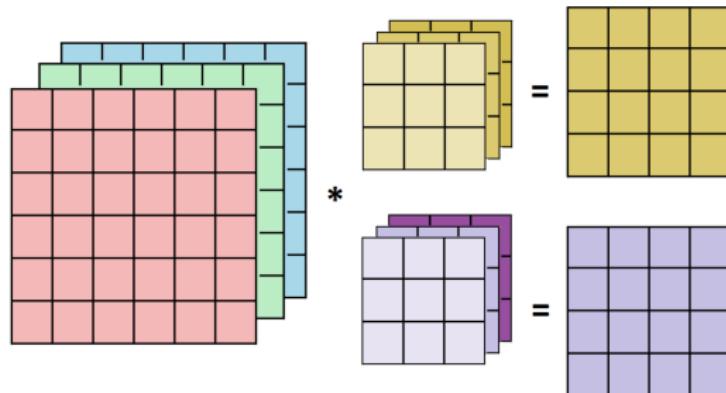
whenever needed we can

- apply a *fractional stride* by resampling
  - upsampling before convolution
  - downsampling after convolution
- adjust the size of  $\mathbf{Z}$  by zero-padding

## Convolution Layer: General Form

- + But, does it make sense to map a **large tensor** to a **small feature map**?
- This is a great point! This is why we compute **multiple feature maps**

A general convolutional layer has **multiple kernels**: each **kernel** computes a **separate feature map**



# Convolution Layer: General Form

## Multi-Channel Convolutional Layer

Convolutional layer with  $C$ -channel input and  $K$ -channel output consists of  $K$  filters  $\mathbf{W}_1, \dots, \mathbf{W}_K \in \mathbb{R}^{C \times F \times F}$ ,  $K$  biases  $b_1, \dots, b_K$  and an activation function  $f(\cdot)$ . It transforms the  $C$ -channel input  $\mathbf{X}$  to the activated  $K$ -channel feature tensor  $\mathbf{Y}$  as follows:

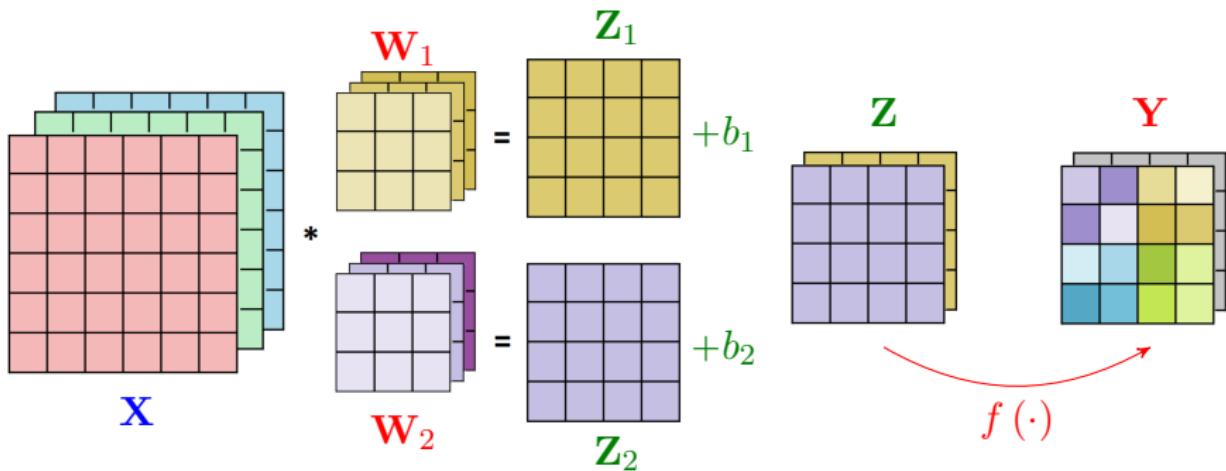
- ① it finds the feature tensor  $\mathbf{Z}$

$$\mathbf{Z} = [\text{Conv}(\mathbf{X} | \mathbf{W}_1) + b_1, \dots, \text{Conv}(\mathbf{X} | \mathbf{W}_K) + b_K]$$

- ② it then activates  $\mathbf{Z}$

$$\mathbf{Y} = f(\mathbf{Z}) \quad f(\cdot) \text{ applied entry-wise}$$

# Multi-Channel Convolution Layer: Visualization

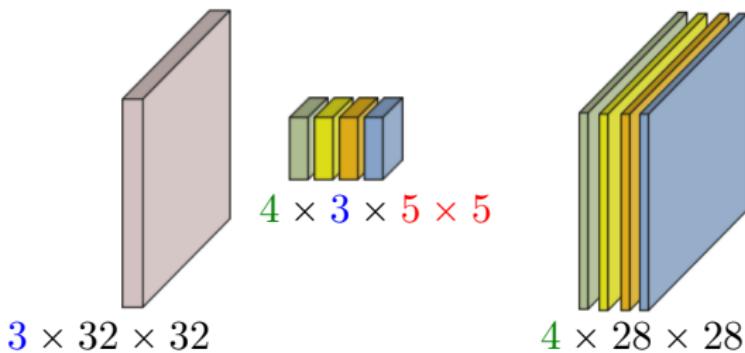


There are few points that we should keep in mind

- *Kernels have the same number of channels (also called depth) as input*
- *Number of kernels equals the number of channels in feature tensor*
- *If  $X$  is output of a convolutional layer it may have lots of channels!*
  - ↳ We **should not** think that “*input has at most 3 channels!*”

# How to Read Dimensions of Convolutional Layers

Many details are dropped as they are readily inferred from architecture



In this diagram, we see that  $C = 3$  which is the same in **kernels** and **input**

- We have  $K = 4$  **kernels**  $\leadsto$  **feature tensor** has 4 channels
- **Kernels** have width  $F = 5 \leadsto$  we see that  $32 - 5 + 1 = 28$ 
  - ↳ **stride** is one, i.e.,  $S = 1$
  - ↳ **no zero-padding** is applied  $P = 0$

## Idea of Pooling: Smoothing Filters

The output of convolution can be **jittering** which can come from *spatial correlation*: *Pooling acts as a filter by sliding over the extracted features and pooling out a function of each subpart*

- *Pooling can reduce the jittering behavior*
- *It can mix extracted features and potentially improve shift-invariance*

### Shift-Invariance

**shift-invariance** refers to robustness against simple geometric transform of input

- + How do we do the pooling?
- There are several pooling techniques but popular ones are **max-** and **mean-pooling**

# Max-Pooling

For max-pooling, we use a filter of size  $L \times L$  and slide over the feature map

$$\mathbf{Y} = \begin{bmatrix} Y_{1,1} & Y_{1,2} & \cdots & Y_{1,M} \\ Y_{2,1} & Y_{2,2} & \cdots & Y_{2,M} \\ \vdots & \vdots & \vdots & \vdots \\ Y_{N,1} & Y_{N,2} & \cdots & Y_{N,M} \end{bmatrix}$$

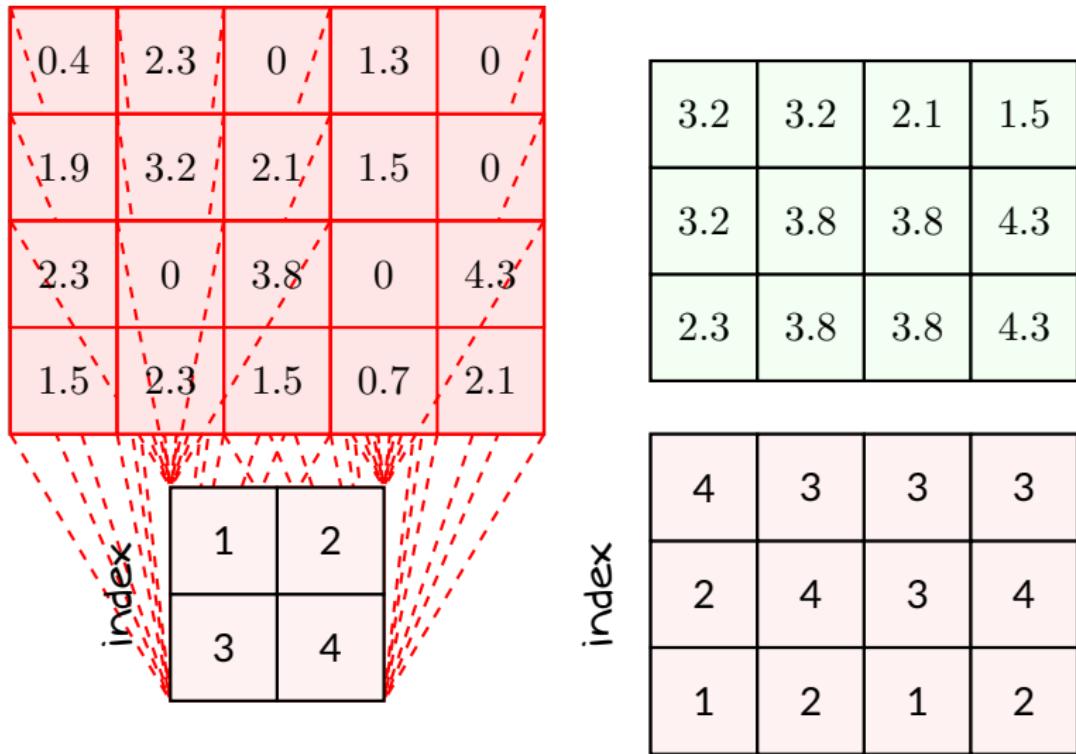
In each window, we pool the **maximum**

$$\hat{\mathbf{Y}} = \begin{bmatrix} \hat{Y}_{1,1} & \hat{Y}_{1,2} & \cdots & \cdots \\ \cdots & \cdots & \cdots & \cdots \\ \cdots & \cdots & \cdots & \cdots \end{bmatrix}$$

## Attention

In **max-pooling**, we also track **index of maximizer**: we need it in **backpropagation**

## Max Pooling: Numerical Example



# Mean-Pooling

In mean-pooling, we again use an  $L \times L$  filter and slide over the feature map

$$\mathbf{Y} = \begin{bmatrix} Y_{1,1} & Y_{1,2} & \dots & Y_{1,M} \\ Y_{2,1} & Y_{2,2} & \dots & Y_{2,M} \\ \vdots & \vdots & \vdots & \vdots \\ Y_{N,1} & Y_{N,2} & \dots & Y_{N,M} \end{bmatrix}$$

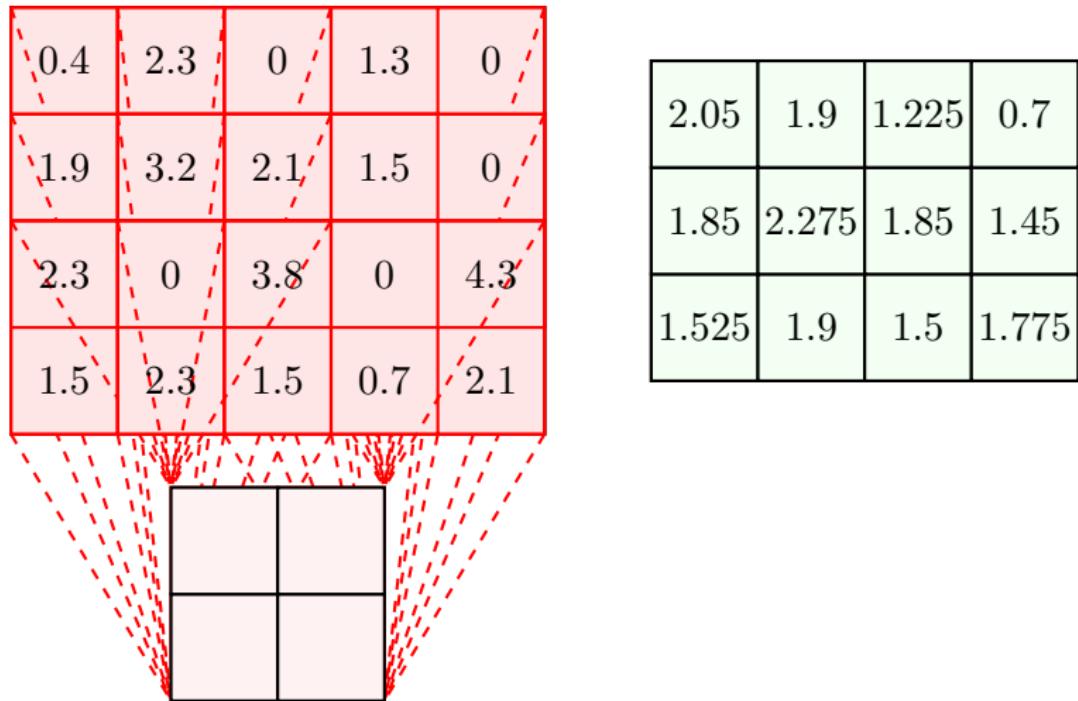
mean  $\{Y_{1,1} Y_{1,2} Y_{1,3} Y_{1,4}\}$

In each window, we pool the *average*

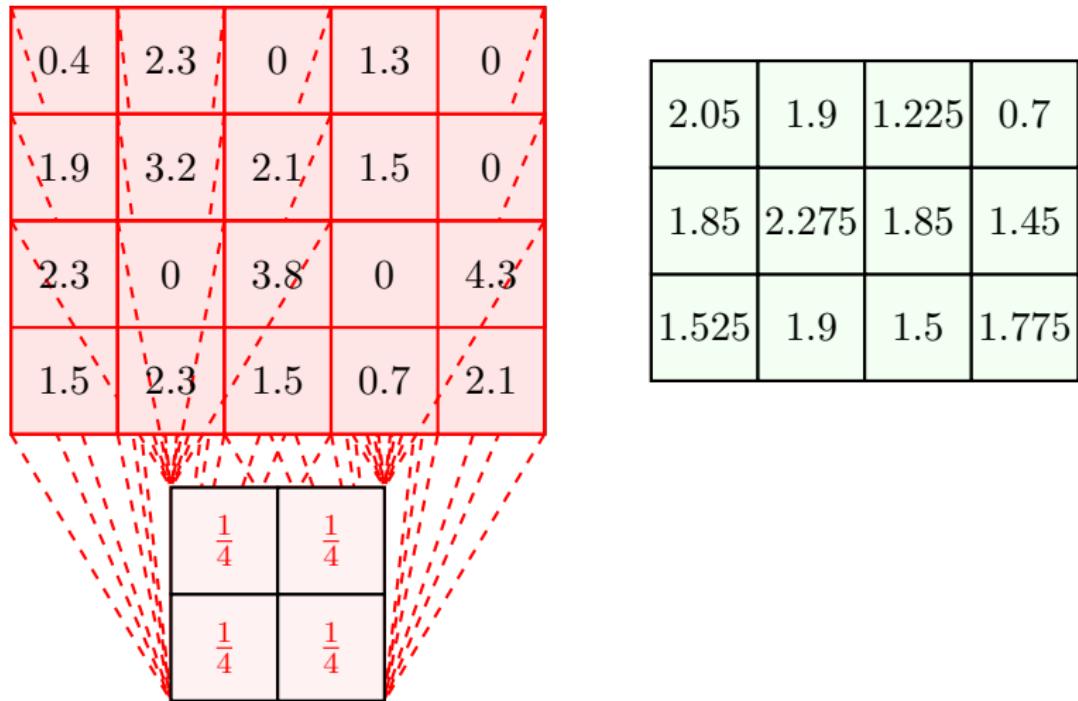
$$\hat{\mathbf{Y}} = \begin{bmatrix} \hat{Y}_{1,1} & \hat{Y}_{1,2} & \dots & \dots \\ \dots & \dots & \dots & \dots \\ \dots & \dots & \dots & \dots \end{bmatrix}$$

We can look at mean-pooling as a *convolution* with *uniform kernel*

# Mean-Pooling: Numerical Example



# Mean-Pooling: Numerical Example



mean-pooling  $\equiv$  convolution with uniform kernel

## Advanced Pooling: Use a General Function

We could in general replace **max or mean operator** with a general function

$$\mathbf{Y} = \begin{bmatrix} Y_{1,1} & Y_{1,2} & \dots & \dots & Y_{1,M} \\ \mathbf{Y}_{1,1} & Y_{2,2} & \dots & \dots & Y_{2,M} \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ Y_{N,1} & Y_{N,2} & \dots & \dots & Y_{N,M} \end{bmatrix} \xrightarrow{\Pi(\cdot) : \mathbb{R}^{L \times L} \mapsto \mathbb{R}}$$

Typical functions are **norms**, e.g.,  $\ell_2$ -norm

$$\hat{\mathbf{Y}} = \begin{bmatrix} \hat{Y}_{1,1} & \hat{Y}_{1,2} & \dots & \dots \\ \dots & \dots & \dots & \dots \\ \dots & \dots & \dots & \dots \end{bmatrix}$$

We can even replace it with a **small NN!**

# Pooling with Stride

Similar to convolution, we can apply *pooling* with *stride*

- We usually do not use *fractional stride* with *pooling*
  - ↳ Upsampling the input, only adds zeros
  - ↳ Extra zeros usually do not make any gain: think of max-pooling for instance
- Integer *stride* can be seen as downsampling
  - ↳ We apply pooling with *unit stride*
  - ↳ We down-sample output with *sampling factor* = *stride*
- In practice, we often leave integer strides for *pooling* layer
  - ↳ We apply convolution with *unit stride*
  - ↳ If we need to down-sample, we do it later in the *pooling* layer

## Pooling: Few Remarks

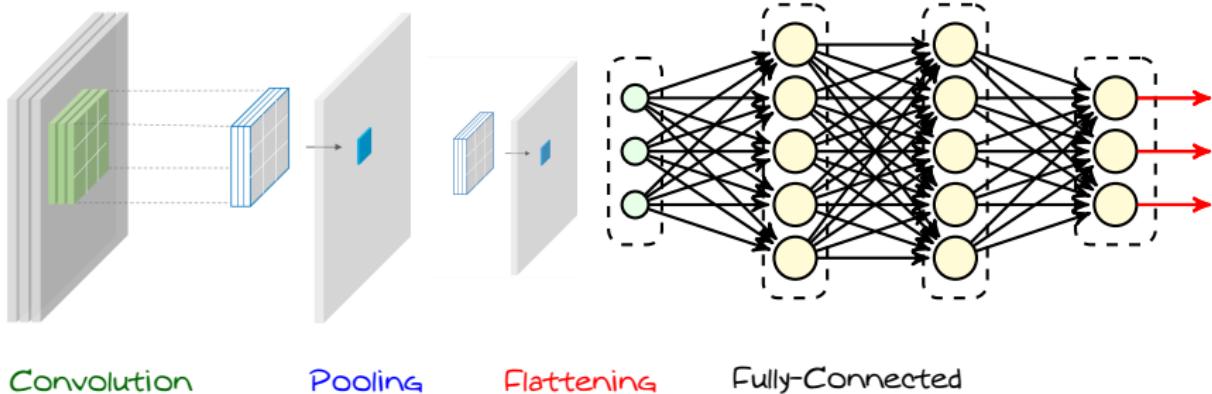
- + Do we always pool after convolution?
- Not really! In many architectures pooling is applied every couple of convolutional layers

It is worth noting that

- In many architectures **pooling** is applied **every couple of convolutional layers**
  - ↳ We apply **multiple** convolutional layers
  - ↳ We then apply a **pooling** layer **with stride**
- Most **poolings** have no **learnable parameters**; thus, **they are cheap**
  - ↳ **Max** and **mean-pooling** have **no weights**
- **Filter size for pooling** can be different from the **convolutional layers**
  - ↳ They are typically in **the same range**

# Output FNN: Flattening

Let's recall our simple CNN



Convolution

Pooling

Flattening

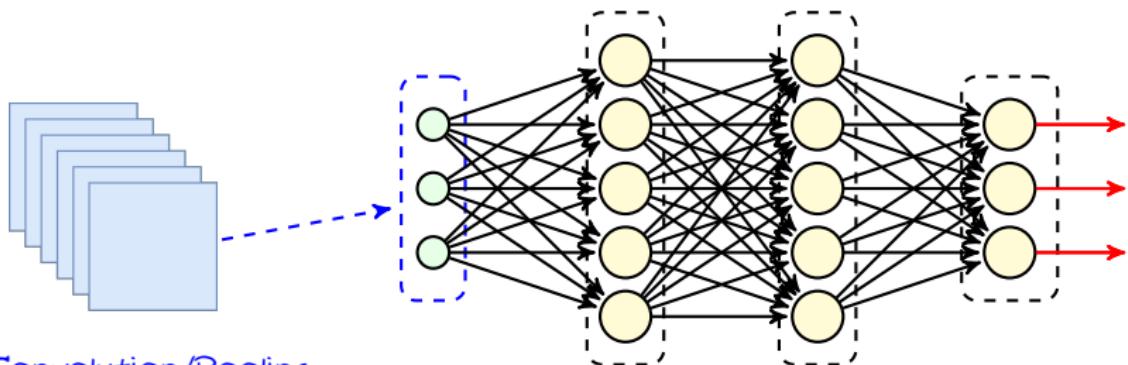
Fully-Connected

Once we are over with convolution and pooling we **flatten** final **feature tensor**

## Flattening

In **flattening**, we sort all entries of the **feature tensor** into a vector

# Flattening



Last Convolution/Pooling

Say we have a **feature tensor** with  $K$  channels

↳ each channel is an  $N \times M$  map after last **convolution** or **pooling**

We then have an input to the FNN with  $NMK$  entries

After **flattening** everything goes as before through the **fully-connected FNN**