

Applied Deep Learning

Chapter 4: Convolutional Neural Networks

Ali Bereyhi

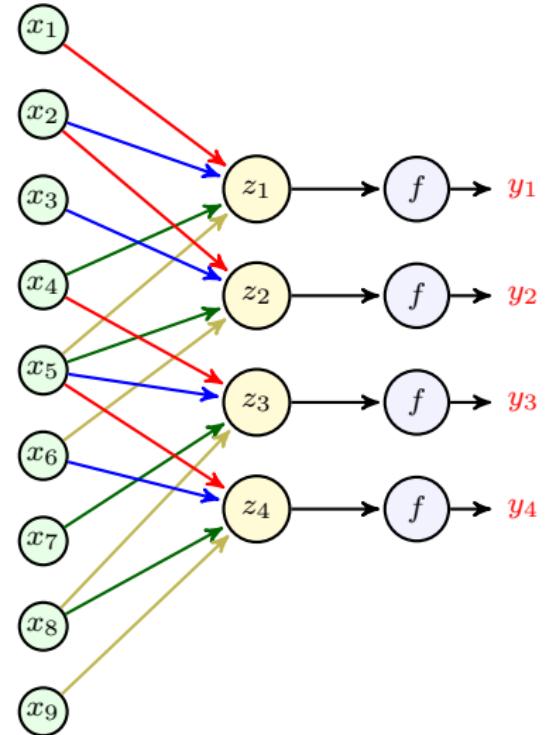
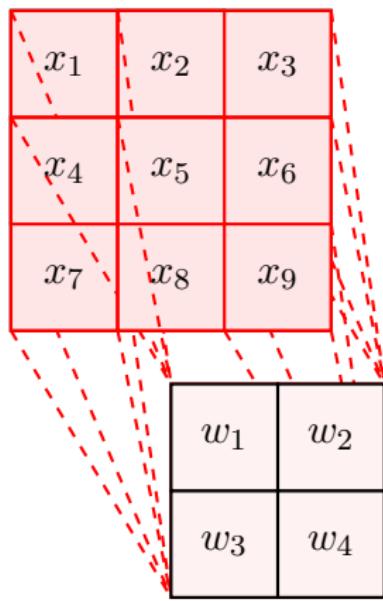
ali.bereyhi@utoronto.ca

Department of Electrical and Computer Engineering
University of Toronto

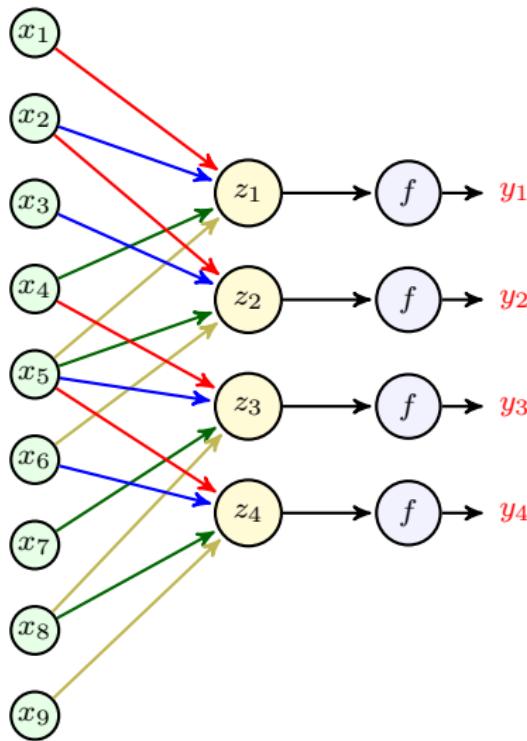
Fall 2025

Convolutional Layer as Set of Neurons

We can look at a *convolutional* layer as a *layer of neurons*



Convolutional Layer as Set of Neurons



In this viewpoint

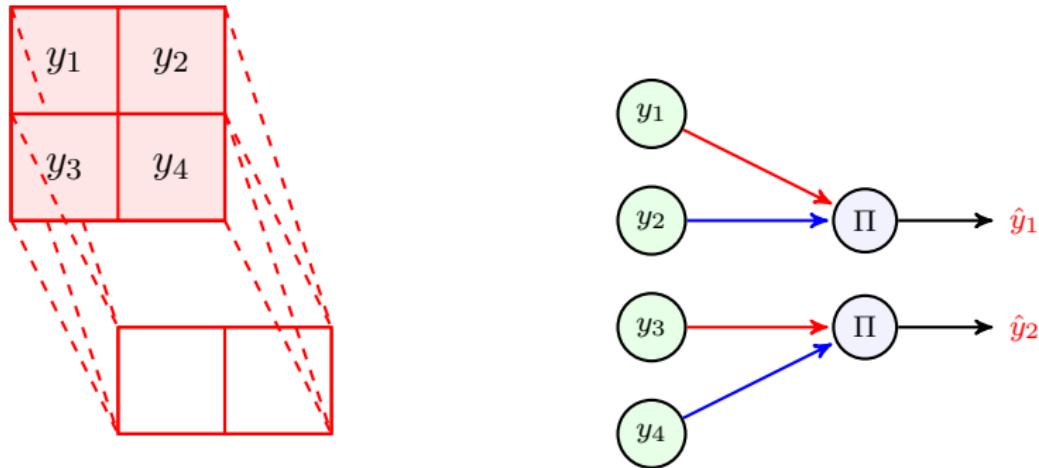
- ↳ each **feature** is output of a neuron
- ↳ neurons **share** same weights and bias
- ↳ neurons are **activated by $f(\cdot)$**

This is however not a **fully-connected** layer

- ↳ it is **locally-connected**
- ↳ neurons have **shared parameters**

Pooling Layer as Set of Neurons

We can extend our **viewpoint** to **pooling layers**



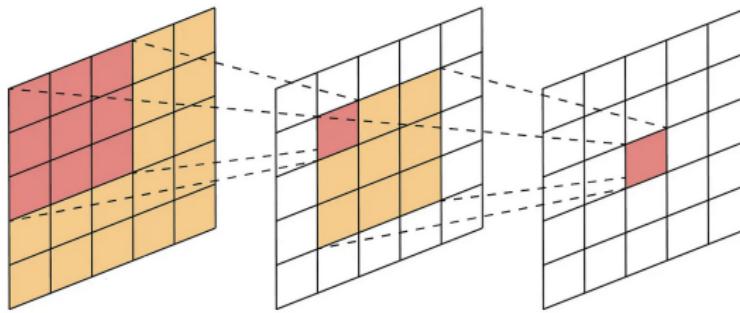
Similar to convolutional layer, pooling layer is a **feedforward layer** with
local connectivity and **shared** parameters

We usually have predefined parameters, i.e., **no weights to be learned**

Going Deep: Receptive Field

As we go **deep** in CNN, we mix **features**: *features in deeper layers depend on more pixels of input tensor. We often say*

*deeper layers have larger **receptive field***



Receptive Field

The input region that each neuron in a given convolutional layer responds to

Receptive Field

Few points we may note regarding the **receptive field**

- ↳ For first layer the receptive field is simply *where the filter screen*
- ↳ For *deeper* layers, actual receptive field is *not* immediately obvious and *must be calculated*

This is however obvious that

- ↳ *Receptive field increases* as we go *deeper*

Receptive field also depends on the filter sizes

- ↳ Large filters *increase receptive field* per layer
- ↳ Small filters *reduce receptive field* per layer

-
- + Why do we define **receptive field**? Does it have any particular meaning?!
 - It helps very much building intuition, especially as we go *deep*

Intuition on Features via Receptive Field

To design a **deep** CNN, we need to first *build a bit of more intuition*

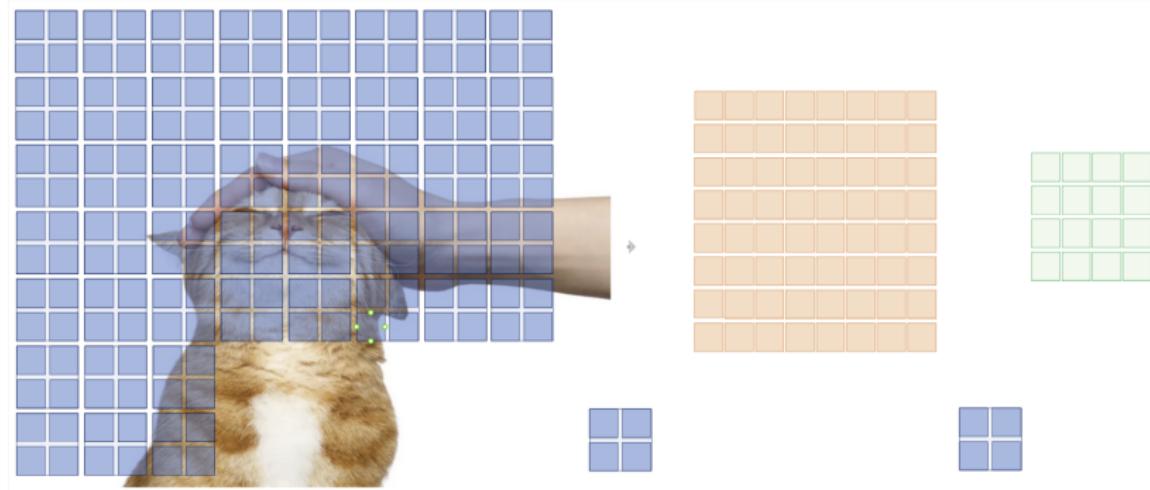
Let's make a closer look: *say we want to classify the following image*



We use two **convolutional** + **pooling** layers cascaded both with 2×2 filters

Intuition on Features via Receptive Field

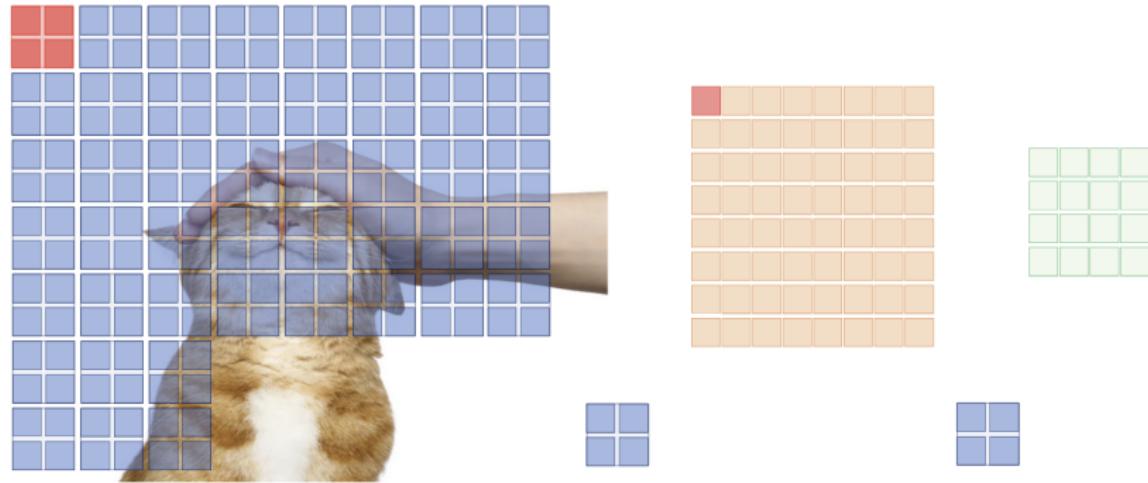
The first layer extracts features of small region of input image



*Say for instance, we look at the very first feature in the **orange feature map***

Intuition on Features via Receptive Field

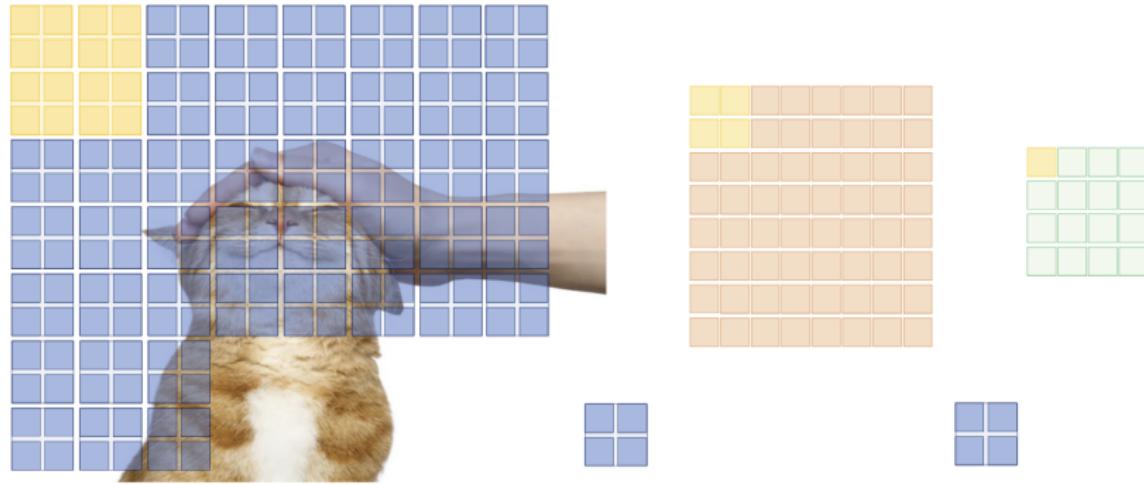
This **red feature** is calculated from input pixels at upper left corner



We could intuitively say that the **neuron** corresponding to **this feature** is looking **locally** at the upper left corner to extract the features of this region

Intuition on Features via Receptive Field

Now, let's look at the very first feature in the **green feature map**



This **neuron** is looking at a **larger region** of input image, but it extracts the features of this region through the **features** extracted via the **orange feature map**

General Perspective on Deep CNNs

The given example illustrates an *intuitive interpretation of deep* CNNs

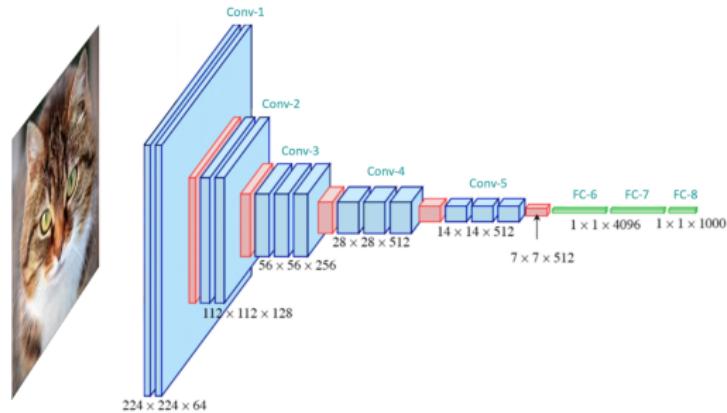
- ↳ Deep CNNs extract *gradually* the *features of input*
- ↳ In first layer they extract features of *smaller regions* of the input
- ↳ They *gradually expand* this region as they go *deeper*
 - ↳ *deeper* layers look at *high-level features*

This gives us a good idea on how a good *deep* CNN looks like

- First layer has *small* filters
 - ↳ *small* filters lead to *large feature maps*
 - ↳ the number of *output channels* is *small*
- As we go *deeper* in the CNN filters get a bit *larger*
 - ↳ *larger* filters return *smaller feature maps*
 - ↳ the number of *output channels* is *increased*

Deep CNN: Example

Recall the example of **VGG-16 architecture** we had a look on

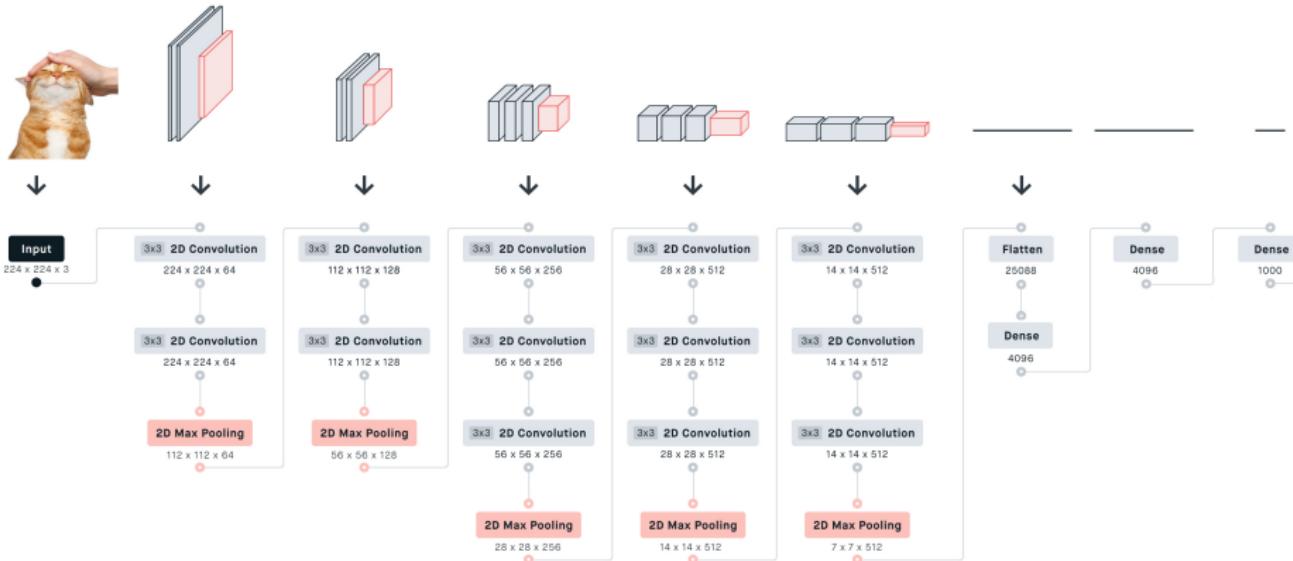


In the original proposal of this architecture, it is said

- convolutional filters are 3×3 and similar padding is used in all layers
- convolution performed with unit stride and max-pooling with stride 2

Deep CNN: Example

We can now break it down, as we know all the components



Deep CNN: Example

Let's count the number of layers with *learnable parameters*

$$\begin{aligned}\# \text{ layers with weights} &= 2 && \text{first two convolutions} \\ &+ 2 && \text{second two convolutions} \\ &+ 3 && \text{third round of convolutions} \\ &+ 3 && \text{fourth round of convolutions} \\ &+ 3 && \text{fifth round of convolutions} \\ &+ 3 && \text{fully-connected network} \\ &= 16\end{aligned}$$

This is actually why it's called VGG-16