

Applied Deep Learning

Chapter 4: Convolutional Neural Networks

Ali Bereyhi

ali.bereyhi@utoronto.ca

Department of Electrical and Computer Engineering
University of Toronto

Fall 2025

Big Picture: What to Train

In CNNs we need to train *all learnable parameters*, i.e.,

- *weights and biases* of output FNN
- *weights* in the *filters* of convolutional layers
- if we use *advanced pooling* with *weights*; then, we should find them as well

To train, we get dataset $\mathbb{D} = \{(\mathbf{X}_b, \mathbf{v}_b) : b = 1, \dots, B\}$ and train the CNN as


$$\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmin}} \hat{R}(\mathbf{w}) = \underset{\mathbf{w}}{\operatorname{argmin}} \frac{1}{B} \sum_{b=1}^B \mathcal{L}(\boldsymbol{\theta}_b, \mathbf{v}_b) \quad (\text{Training})$$

where $\boldsymbol{\theta}_b = \text{CNN}(\mathbf{X}_b | \mathbf{w})$ for the tensor-type data-point \mathbf{X}_b with label \mathbf{v}_b

- ↳ We solve this optimization via a *gradient-based method*
- ↳ This means we should be able to pass *forward* and *backward*

Computation Graph

Computation graph for single data-point \mathbf{X} and its true label \mathbf{v} is as follows



For **simplicity**, let's assume that we are doing **basic SGD**

- ↳ We need to pass **forward** \mathbf{X} over this graph to get **loss**
- ↳ We should then pass **backward** to compute **gradient**

Let's try both directions

Forward Pass over CNN


The forward pass is what we learned in the last section

- We pass \mathbf{X} through first convolution to get \mathbf{Z}_1
- We activate \mathbf{Z}_1 to get \mathbf{Y}_1
- We pool entries of \mathbf{Y}_1 and get $\hat{\mathbf{Y}}_1$
- \vdots
- We flatten and pass **forward** through the **output FNN**

Once we are over, we have

- ↳ *all convolution, activated and pooled values in convolutional layers*
- ↳ *all affine and activated values in the FNN*

Backward Pass



We now want to **backpropagate**

- ↳ We first compute $\nabla_{\theta} \hat{R}$
- ↳ We then **backpropagate** over the FNN till we get to its input


At this point, we sort $\nabla_{\hat{Y}_L} \hat{R}$ in tensor by **reversing the flattening**

we now have $\nabla_{\hat{Y}_L} \hat{R}$

We only need to learn how to **backpropagate** through **pooling** and **convolutional layers** and then we can **complete** the backward pass!

Backpropagate through Pooling: Max-Pooling

Let's start with a simple example: in the following **pooling layer**, we have partial derivatives with respect to **pooled variables** and want to compute the partial derivatives with respect to **input of pooling layer**



Backpropagation: Max-Pooling

We can use **chain rule**

$$\frac{\partial \hat{R}}{\partial y_1} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial y_1} = 0$$

$$\frac{\partial \hat{R}}{\partial y_7} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial y_7} = \frac{\partial \hat{R}}{\partial \hat{y}_3}$$

$$\frac{\partial \hat{R}}{\partial y_2} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial y_2} = \frac{\partial \hat{R}}{\partial \hat{y}_1} + \frac{\partial \hat{R}}{\partial \hat{y}_2}$$

$$\frac{\partial \hat{R}}{\partial y_8} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial y_8} = 0$$

$$\frac{\partial \hat{R}}{\partial y_3} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial y_3} = 0$$


$$\frac{\partial \hat{R}}{\partial y_9} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial y_9} = \frac{\partial \hat{R}}{\partial \hat{y}_4}$$

⋮

Let's see its visualization


Backpropagation: Max-Pooling

Each derivative in *green map* gets multiplied with its filter and added to corresponding entries on *blue map*



Backpropagate through Pooling: Mean-Pooling

Lets now consider **mean-pooling**: we have partial derivatives with respect to **pooled variables** and want to compute the partial derivatives with respect to **input of pooling layer**



Backward Pass: Mean-Pooling

We again use **chain rule**

$$\frac{\partial \hat{R}}{\partial y_1} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial y_1} = \frac{1}{4} \frac{\partial \hat{R}}{\partial \hat{y}_1}$$

$$\frac{\partial \hat{R}}{\partial y_2} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial y_2} = \frac{1}{4} \frac{\partial \hat{R}}{\partial \hat{y}_1} + \frac{1}{4} \frac{\partial \hat{R}}{\partial \hat{y}_2}$$

⋮


$$\frac{\partial \hat{R}}{\partial y_5} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial y_5} = \frac{1}{4} \frac{\partial \hat{R}}{\partial \hat{y}_1} + \frac{1}{4} \frac{\partial \hat{R}}{\partial \hat{y}_2} + \frac{1}{4} \frac{\partial \hat{R}}{\partial \hat{y}_3} + \frac{1}{4} \frac{\partial \hat{R}}{\partial \hat{y}_4}$$

⋮

$$\frac{\partial \hat{R}}{\partial y_9} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial y_9} = \frac{1}{4} \frac{\partial \hat{R}}{\partial \hat{y}_4}$$

Backward Pass: Mean-Pooling


*It has same visualization: the **filter** is however this time fixed*



But we can visualize it even better!

Backward Pass: Mean-Pooling

In fact, this is simply a convolution



Backpropagation through Pooling: Summary

Depending on **pooling** function, backpropagation can be different

↳ *It's however usually described by a convolution*

Moral of Story

Backpropagation through pooling can be done by a convolution-type operation

Backpropagation through Convolution: Activation

Convolutional layer has two operations

- ↳ *linear convolution*
- ↳ *entry-wise activation*

The *entry-wise activation* is readily backpropagate: say we have

$$\mathbf{Y} = f(\mathbf{Z})$$


Then, we can *backpropagate* as in the *fully-connected* FNNs

$$\nabla_{\mathbf{Z}} \hat{R} = \nabla_{\mathbf{Y}} \hat{R} \odot \dot{f}(\mathbf{Z})$$

Now, let's look at *linear convolution!*

Backpropagation: 2D Convolution

Let's find it out through a simple example: in the following convolutional layer, we have partial derivatives with respect to convolved variables and want to compute the partial derivatives with respect to input of convolutional layer



Backpropagation: 2D Convolution

Let's start with **chain rule**

$$\frac{\partial \hat{R}}{\partial x_1} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_1} = w_1 \frac{\partial \hat{R}}{\partial z_1}$$

$$\frac{\partial \hat{R}}{\partial x_2} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_2} = w_2 \frac{\partial \hat{R}}{\partial z_1} + w_1 \frac{\partial \hat{R}}{\partial z_2}$$

$$\frac{\partial \hat{R}}{\partial x_3} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_3} = w_2 \frac{\partial \hat{R}}{\partial z_2}$$

$$\frac{\partial \hat{R}}{\partial x_4} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_4} = w_3 \frac{\partial \hat{R}}{\partial z_1} + w_1 \frac{\partial \hat{R}}{\partial z_3}$$

$$\frac{\partial \hat{R}}{\partial x_5} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_5} = w_4 \frac{\partial \hat{R}}{\partial z_1} + w_3 \frac{\partial \hat{R}}{\partial z_2} + w_2 \frac{\partial \hat{R}}{\partial z_3} + w_1 \frac{\partial \hat{R}}{\partial z_4}$$

Backpropagation: 2D Convolution

We can keep on till finish with **chain rule**

$$\frac{\partial \hat{R}}{\partial x_6} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_6} = w_4 \frac{\partial \hat{R}}{\partial z_2} + w_2 \frac{\partial \hat{R}}{\partial z_4}$$

$$\frac{\partial \hat{R}}{\partial x_7} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_7} = w_3 \frac{\partial \hat{R}}{\partial z_3}$$


$$\frac{\partial \hat{R}}{\partial x_8} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_8} = w_4 \frac{\partial \hat{R}}{\partial z_3} + w_3 \frac{\partial \hat{R}}{\partial z_4}$$

$$\frac{\partial \hat{R}}{\partial x_9} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_9} = w_4 \frac{\partial \hat{R}}{\partial z_4}$$

Let's look at the visualization

Backpropagation: 2D Convolution


We can use the same visualization as in pooling



Or even a **better** visualization!

Backpropagation: 2D Convolution

It's again simply a convolution



with **filter** being **reversed up-to-down** and **right-to-left**

Backpropagation: 2D Convolution

To backpropagate a convolutional layer, we convolve the output gradient with the **filter** being **reversed up-to-down** and **right-to-left**. To match the dimension, we simply apply **zero-padding**

Backpropagation through 2D Convolution

Let $\mathbf{Z} = \text{Conv}(\mathbf{X}|\mathbf{W})$, where \mathbf{X} is the **input map**, i.e., a matrix, \mathbf{W} is **filter** and \mathbf{Z} is the **output map**. Assume that we know the gradient of loss \hat{R} with respect to the **output map**, i.e., $\nabla_{\mathbf{Z}} \hat{R}$. Then, the gradient of loss \hat{R} with respect to the **input map** is

$$\nabla_{\mathbf{X}} \hat{R} = \text{Conv}\left(\nabla_{\mathbf{Z}} \hat{R} | \check{\mathbf{W}}\right)$$


where $\check{\mathbf{W}}$ is the **up-to-down** and **left-to-right** reverse of **filter** \mathbf{W}

Backpropagation: 2D Convolution

- + What if the dimensions do **not** match?
- The dimensions can **only** not match if
 - ① we do the forward convolution with a **stride different from one**
 - ↳ We said that this is simply **resampling**
 - ↳ We are going to **discuss it later**
 - ↳ For now assume that we do the convolution with **stride one**
 - ② we did **no zero-padding** in the forward convolution
 - ↳ We simply **do zero-padding** in backward convolution to match the dimensions
- + What about the **gradient** with respect to **filter itself**? Don't we need it?
- Let's check it out

Backpropagation: Convolution Filter

We try it for **our example**: we have partial derivatives with respect to **convolved variables** and want to compute the partial derivatives with respect to **weights in the filter**



Backpropagation: Convolution Filter

As always, we use *chain rule*

$$\frac{\partial \hat{R}}{\partial w_1} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial w_1} = x_1 \frac{\partial \hat{R}}{\partial z_1} + x_2 \frac{\partial \hat{R}}{\partial z_2} + x_4 \frac{\partial \hat{R}}{\partial z_3} + x_5 \frac{\partial \hat{R}}{\partial z_4}$$


$$\frac{\partial \hat{R}}{\partial w_2} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial w_2} = x_2 \frac{\partial \hat{R}}{\partial z_1} + x_3 \frac{\partial \hat{R}}{\partial z_2} + x_5 \frac{\partial \hat{R}}{\partial z_3} + x_6 \frac{\partial \hat{R}}{\partial z_4}$$

$$\frac{\partial \hat{R}}{\partial w_3} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial w_3} = x_4 \frac{\partial \hat{R}}{\partial z_1} + x_5 \frac{\partial \hat{R}}{\partial z_2} + x_7 \frac{\partial \hat{R}}{\partial z_3} + x_8 \frac{\partial \hat{R}}{\partial z_4}$$

$$\frac{\partial \hat{R}}{\partial w_4} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial w_4} = x_5 \frac{\partial \hat{R}}{\partial z_1} + x_6 \frac{\partial \hat{R}}{\partial z_2} + x_8 \frac{\partial \hat{R}}{\partial z_3} + x_9 \frac{\partial \hat{R}}{\partial z_4}$$

Backpropagation: Convolution Filter

It's another convolution with $\nabla_{\mathbf{z}} \hat{R}$ being the filter!



Backpropagation: Convolution Filter

Once we backpropagate to the **output** of a **convolutional layer**, then we can find the gradient with respect to convolution filter by **convolving output gradient with the input map**

Gradient w.r.t. Filters

Let $\mathbf{Z} = \text{Conv}(\mathbf{X}|\mathbf{W})$, where \mathbf{X} is the **input map**, i.e., a matrix, \mathbf{W} is **filter** and \mathbf{Z} is the **output map**. Assume that we know gradient of loss \hat{R} with respect to the **output map**, i.e., $\nabla_{\mathbf{Z}} \hat{R}$. Then, gradient of loss \hat{R} with respect to the **filter** is

$$\nabla_{\mathbf{W}} \hat{R} = \text{Conv}\left(\mathbf{X} | \nabla_{\mathbf{Z}} \hat{R}\right)$$

- + But we checked only **2D case!** What about **multi-channel case??**
- Well, we can simply do it by **chain rule**

Backpropagation: Multi-Channel Convolution


Lets start with a simple case: we have a *C-channel input*, i.e., a *tensor* input, \mathbf{X} and we compute a *single feature map*, i.e., a *matrix*, \mathbf{Z}

↳ *Filter \mathbf{W} has also C channels*

Let's denote *channel c of input* and *filter* by $\mathbf{X}[c]$ and $\mathbf{W}[c]$; then, we can write


$$\mathbf{Z} = \sum_{c=1}^C \text{Conv}(\mathbf{X}[c] | \mathbf{W}[c])$$

Backpropagation: Multi-Channel Convolution



Backpropagation: Multi-Channel Convolution

Channel c of *input* is connected to the *output map* by a 2D convolution



So, we could say

$$\nabla_{\mathbf{X}[c]} \hat{R} = \text{Conv} \left(\nabla_{\mathbf{Z}} \hat{R} | \check{\mathbf{W}}[c] \right)$$

Backpropagation: Multi-Channel Convolution

To backpropagate a convolutional layer with **tensor input** and **matrix output**, we convolve the output gradient with the **filter of each channel** being **reversed up-to-down and right-to-left**


Backpropagation through 2D Convolution

Let $\mathbf{Z} = \text{Conv}(\mathbf{X}|\mathbf{W})$, where \mathbf{X} is a C -channel **input tensor**, \mathbf{W} is C -channel **filter** and \mathbf{Z} is a single-channel **output map**, i.e., a **matrix**. Assume that we know the gradient of loss \hat{R} with respect to the **output map**, i.e., $\nabla_{\mathbf{Z}} \hat{R}$. Then, the gradient of loss \hat{R} with respect to **input tensor** is

$$\nabla_{\mathbf{X}} \hat{R} = [\text{Conv}(\nabla_{\mathbf{Z}} \hat{R} | \check{\mathbf{W}}[1]), \dots, \text{Conv}(\nabla_{\mathbf{Z}} \hat{R} | \check{\mathbf{W}}[C])]$$

Note that $\nabla_{\mathbf{X}} \hat{R}$ is also a C -channel tensor

Backpropagation: Multi-Channel Convolution



Backpropagation: Multi-channel Convolution


We now go for the general case: we have a C -channel input tensor \mathbf{X} and we compute K -channel feature tensor \mathbf{Z}

- ↳ We have K filters $\mathbf{W}_1, \dots, \mathbf{W}_K$
- ↳ Each filter has C channels

Let's denote channel k of the output by \mathbf{Z} ; then, we can write

$$\mathbf{Z}[k] = \text{Conv}(\mathbf{X} | \mathbf{W}_k)$$

Backpropagation: Multi-Channel Convolution



Backpropagation: Multi-channel Convolution

We write the **chain rule** with a bit **cheating**: let's denote the **derivative** of an object A with respect to object B as $\nabla_B A$

- ↳ if A and B are both scalars then its simple **derivative**
- ↳ if A is a scalar and B is a vector it's **gradient**
- ↳ if A and B are both vectors then its **Jacobian**
- ↳ ...

When $\mathbf{Z} = f(\mathbf{X})$ and $\hat{R} = g(\mathbf{Z})$: **chain rule** says


$$\nabla_{\mathbf{X}} \hat{R} = \nabla_{\mathbf{Z}} \hat{R} \circ \nabla_{\mathbf{X}} \mathbf{Z}$$

where \circ is a kind of product

We now write the **chain rule** with this simplified notation

Backpropagation: Multi-Channel Convolution

We do know $\nabla_{\mathbf{Z}} \hat{R} = [\nabla_{\mathbf{Z}[1]} \hat{R}, \dots, \nabla_{\mathbf{Z}[K]} \hat{R}]$



Backpropagation: Multi-Channel Convolution

The output tensor can be seen as K functions of the input tensor

$$\mathbf{Z}[1] = f_1(\mathbf{X}) \quad \dots \quad \mathbf{Z}[K] = f_K(\mathbf{X})$$

So, the **chain rule** can be written as

$$\begin{aligned} \nabla_{\mathbf{X}} \hat{R} &= \sum_{k=1}^K \underbrace{\nabla_{\mathbf{Z}[k]} \hat{R} \circ \nabla_{\mathbf{X}} \mathbf{Z}[k]}_{\text{what we calculated for single output map}} \\ &= \sum_{k=1}^K \left[\text{Conv} \left(\nabla_{\mathbf{Z}[k]} \hat{R} | \check{\mathbf{W}}_k[1] \right), \dots, \text{Conv} \left(\nabla_{\mathbf{Z}[k]} \hat{R} | \check{\mathbf{W}}_k[C] \right) \right] \\ &= \left[\sum_{k=1}^K \text{Conv} \left(\nabla_{\mathbf{Z}[k]} \hat{R} | \check{\mathbf{W}}_k[1] \right), \dots, \sum_{k=1}^K \text{Conv} \left(\nabla_{\mathbf{Z}[k]} \hat{R} | \check{\mathbf{W}}_k[C] \right) \right] \end{aligned}$$


Backpropagation: Multi-Channel Convolution

The backward pass is therefore

$$\nabla_{\mathbf{X}} \hat{R} = \left[\sum_{k=1}^K \text{Conv} \left(\nabla_{\mathbf{Z}[k]} \hat{R} | \check{\mathbf{W}}_k[1] \right) \dots \sum_{k=1}^K \text{Conv} \left(\nabla_{\mathbf{Z}[k]} \hat{R} | \check{\mathbf{W}}_k[C] \right) \right]$$

Let's look at a particular input channel c

Backpropagation: Multi-Channel Convolution



Backpropagation: Multi-Channel Convolution

This is a new multi-channel convolution: let's define *K-channel filter* \mathbf{W}_c^\dagger for $c = 1, \dots, C$ as follows

$$\mathbf{W}_c^\dagger = [\check{\mathbf{W}}_1[c] \dots \check{\mathbf{W}}_K[c]]$$

Then, channel c of $\nabla_{\mathbf{X}} \hat{R}$ is the convolution of $\nabla_{\mathbf{Z}} \hat{R}$ with \mathbf{W}_c^\dagger


$$\nabla_{\mathbf{X}[c]} \hat{R} = \text{Conv}(\nabla_{\mathbf{Z}} \hat{R} | \mathbf{W}_c^\dagger)$$

Or shortly, we can write

$$\nabla_{\mathbf{X}} \hat{R} = \text{Conv}(\nabla_{\mathbf{Z}} \hat{R} | \mathbf{W}_1^\dagger, \dots, \mathbf{W}_C^\dagger)$$

Backpropagation: Multi-Channel Convolution

This means that we can look at channel c of $\nabla_{\mathbf{X}} \hat{R}$ as



Backpropagation: Multi-Channel Convolution

To backpropagate to channel c of **tensor input**, we convolve the output gradient **tensor** with the K -channel **filter tensor** that is constructed by collecting the channel c of all K forward filters, each being **reversed up-to-down and right-to-left**

Backpropagation through Multi-Channel Convolution

Let $\mathbf{Z} = \text{Conv}(\mathbf{X} | \mathbf{W}_1, \dots, \mathbf{W}_K)$, where \mathbf{X} is a C -channel **input tensor**, \mathbf{W}_k is C -channel **filter** and \mathbf{Z} is a K -channel **output tensor**. Assume that we know the gradient of loss \hat{R} with respect to the **output tensor**, i.e., $\nabla_{\mathbf{Z}} \hat{R}$. Then, the gradient of loss \hat{R} with respect to **input tensor** is

$$\nabla_{\mathbf{X}} \hat{R} = \text{Conv}\left(\nabla_{\mathbf{Z}} \hat{R} | \mathbf{W}_1^\dagger, \dots, \mathbf{W}_C^\dagger\right)$$

where $\mathbf{W}_c^\dagger = [\check{\mathbf{W}}_1[c] \dots \check{\mathbf{W}}_K[c]]$ is a K -channel filter

Backpropagation through Convolution: Summary

Moral of Story

We can always backpropagate through a **convolutional layer** with C input channels and K output channels by

- ① multiply output gradient entry-wise with derivative of **activation function**
- ② convolve **output gradient** with C different **K -channel filters** computed by **rearranging of the K forward filters**

To compute the gradient with respect to **weights of channel c in filter k**

- ↳ **convolve** channel c of input tensor with channel k of output gradient

- + Sounds great! But what about cases with $\text{stride} \neq 1$?
- Well! we said we can decompose them as convolution/pooling with $\text{stride } 1$ + a **resampling unit**. We just need to learn how to backpropagate through a **resampling unit**

Backpropagation: *Downsampling*

Let's again try an **example**: we have *partial derivatives with respect to down-sampled variables* and want to compute the *partial derivatives with respect to input variables*

x_1	x_2	x_3
x_4	x_5	x_6
x_7	x_8	x_9

$$\mathbf{Z} = \text{dSample}(\mathbf{X}|2)$$



$z_1=x_1$	$z_2=x_3$
$z_3=x_7$	$z_4=x_9$

Backpropagation: *Downsampling*

Let's write with **chain rule**

$$\frac{\partial \hat{R}}{\partial x_1} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_1} = \frac{\partial \hat{R}}{\partial z_1}$$

$$\frac{\partial \hat{R}}{\partial x_2} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_2} = 0$$

$$\frac{\partial \hat{R}}{\partial x_3} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_3} = \frac{\partial \hat{R}}{\partial z_2}$$

$$\frac{\partial \hat{R}}{\partial x_4} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_4} = 0$$

⋮

$$\frac{\partial \hat{R}}{\partial x_7} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_7} = \frac{\partial \hat{R}}{\partial z_3}$$

$$\frac{\partial \hat{R}}{\partial x_8} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_8} = 0$$

$$\frac{\partial \hat{R}}{\partial x_9} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_9} = \frac{\partial \hat{R}}{\partial z_4}$$

Backpropagation: *Downsampling*

This is simply an *upsampling* with the same factor

$\frac{\partial \hat{R}}{\partial z_1}$	0	$\frac{\partial \hat{R}}{\partial z_2}$
0	0	0
$\frac{\partial \hat{R}}{\partial z_3}$	0	$\frac{\partial \hat{R}}{\partial z_4}$

uSample $(\nabla_{\mathbf{z}} \hat{R} | 2)$

$\frac{\partial \hat{R}}{\partial z_1}$	$\frac{\partial \hat{R}}{\partial z_2}$
$\frac{\partial \hat{R}}{\partial z_3}$	$\frac{\partial \hat{R}}{\partial z_4}$

Backpropagation: *Upsampling*

Now, let's look into **upsampling**: we have partial derivatives with respect to **up-sampled variables** and want to compute the partial derivatives with respect to **input variables**

$z_1=x_1$	$z_2=0$	$z_3=x_2$
$z_4=0$	$z_5=0$	$z_6=0$
$z_7=x_3$	$z_8=0$	$z_9=x_4$

$$\mathbf{Z} = \text{uSample}(\mathbf{X}|2)$$



x_1	x_2
x_3	x_4

Backpropagation: *Downsampling*

Let's write with **chain rule**

$$\frac{\partial \hat{R}}{\partial x_1} = \sum_{i=1}^9 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_1} = \frac{\partial \hat{R}}{\partial z_1}$$

$$\frac{\partial \hat{R}}{\partial x_2} = \sum_{i=1}^9 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_2} = \frac{\partial \hat{R}}{\partial z_3}$$

$$\frac{\partial \hat{R}}{\partial x_3} = \sum_{i=1}^9 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_3} = \frac{\partial \hat{R}}{\partial z_7}$$

$$\frac{\partial \hat{R}}{\partial x_4} = \sum_{i=1}^9 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_4} = \frac{\partial \hat{R}}{\partial z_9}$$

Backpropagation: *Upsampling*

This is *downsampling* with the same factor

$\frac{\partial \hat{R}}{\partial z_1}$	$\frac{\partial \hat{R}}{\partial z_2}$	$\frac{\partial \hat{R}}{\partial z_3}$
$\frac{\partial \hat{R}}{\partial z_4}$	$\frac{\partial \hat{R}}{\partial z_5}$	$\frac{\partial \hat{R}}{\partial z_6}$
$\frac{\partial \hat{R}}{\partial z_7}$	$\frac{\partial \hat{R}}{\partial z_8}$	$\frac{\partial \hat{R}}{\partial z_9}$

dSample $\left(\nabla_{\mathbf{z}} \hat{R} | 2 \right)$

$\frac{\partial \hat{R}}{\partial z_1}$	$\frac{\partial \hat{R}}{\partial z_3}$
$\frac{\partial \hat{R}}{\partial z_7}$	$\frac{\partial \hat{R}}{\partial z_9}$

Backpropagation: Resampling

Backpropagation through Downsampling

To backpropagate through a **downsampling** unit with factor (stride) S we **up-sample** the output gradient with factor S

Backpropagation through Upsampling

To backpropagate through an **upsampling** unit with factor (stride) S we **down-sample** the output gradient with factor S

Forward and Backpropagation in CNNs: Summary

To each **forward** action, there is a **backward** counterpart

In **forward** pass we do

- **forward FNN**
- **pooling**
- **convolution**
- **upsampling**
- **downsampling**

In **backward** pass we do

- **backward FNN**
- **backward pooling** ~ convolution
- convolution with **reversed filters**
- **downsampling**
- **upsampling**

Once we over with **backward** pass, we compute **all required gradients** from **forward** and **backward** variables

We Use Advanced Techniques

For FNN, we looked into advanced techniques such as

- *Dropout and Regularization*
 - ↳ to reduce the impact of *overfitting*
- *Input Normalization*
 - ↳ to improve training with *unbalanced* datasets
- *Batch Normalization*
 - ↳ to make the training more stable against *feature variations*
- *Data Preprocessing*
 - ↳ to *clean* our training dataset

Same goes with CNNs

*we should use all **these methods** for same purposes in CNNs*

Other Forms of Convolution

The convolution operation we considered in this chapter is often called


2D Convolution

because it screens *input* in a 2D fashion, i.e., *left-to-right* and *up-to-down*

2D convolution is the *most popular form* of convolution; however,



- we could have *1D convolutions*
 - ↳ The filter *only slides in one direction*, i.e., *left-to-right* or *up-to-down*
 - ↳ This is useful with *audio data* where we slide the filter over *time*
- we could also have *3D convolutions*
 - ↳ The filter slides in *all three directions*, i.e., *left-to-right*, *up-to-down* and *front-to-back*
 - ↳ This is useful with *3D images*, e.g., *3D medical image of brain*

Other Forms of Convolution



1D Convolution

2D Convolution



3D Convolution

There are also **other forms**, e.g., depth-wise convolution

at the end of day, they all **slide** over **input** with **some filter**