

Applied Deep Learning

Chapter 5: Skip Connection and Residual Networks

Ali Bereyhi


ali.bereyhi@utoronto.ca

Department of Electrical and Computer Engineering
University of Toronto

Fall 2025

A Bit of History: LeNet

CNNs were first trained via gradient-based algorithms by Yan LeCun and his team: they could develop *backpropagation* through CNNs and hence they were able to *efficiently implement it*¹



32×32

Convolution

Pooling

Convolution

Pooling

FNN

¹Check out their paper at [this link!](#) The diagram is taken from the [the paper](#)

ILSVRC: ImageNet Large Scale Visual Recognition Challenge

The project ImageNet started a competition in 2010: it introduced a training dataset of 1.2 million images from 1000 different labels. The proposed trained architectures are then validated by a test dataset. The winner is the model with minimal classification error


The winners in 2010 and 2011 used shallow NNs!

- 2010: Team NEC-UIUC
- 2011: Team XRCE

In 2012, Supervision from U of T trained a deep CNN and won ILSVRC

First Deep Winner: AlexNet

Alex Krizhevsky, Ilya Sutskever and Geoffrey E. Hinton proposed **AlexNet** which could greatly reduce the classification error²



²Check it out in [their paper!](#)

AlexNet to ZFNet

AlexNet was a deep CNN with 8 learnable layers

- 5 convolutional layers
- 3 fully-connected layers

Zeiler and Fergus won ILSVRC in 2013 by using the same architecture but doing accurate hyperparameter tuning³

At this point, going deeper considered as the key to success

³You may find details in their paper

VGG Architectures

Visual Geometry Group at Oxford University developed deeper CNNs: *a class of architectures*⁴

- VGG-11
- VGG-13
- VGG-16
- VGG-19

They could win ILSVRC localization task and took second place in classification: their results also showed an interesting finding

As we go **deeper**, the **accuracy** gets **higher notably** up to **VGG-16**; however, **VGG-19** can only give **marginal improvements**!

⁴Check VGG architectures out in [their paper](#)


VGG Architectures



GoogLeNet: 2014 Classification Winner

In 2014, Google introduced GoogLeNet: *a deep CNN which uses fully-connected layer **only** at the **output** and is purely based on CNN*⁵

- They introduced a new module called “**inception module**”
 - ↳ It’s a collection of parallel **convolutions** and **poolings**



- Since there is no fully-connected layer it has much less model parameters
 - ↳ It hence requires **less memory** and is **trained faster**

GoogLeNet won ILSVRC classification task



⁵Check GoogLeNet in [their paper](#). The diagram is taken from [original paper](#)

A Hurdle in Going Deeper

Though deep CNNs were doing good job, as people kept going deeper they realized that the performance is getting saturated. Later studies showed that much deeper CNNs start to perform worse!



- Initial guess for this behavior is overfitting

↳ This was ruled out by Microsoft Research Lab by a study⁶



⁶Check it in [this paper](#) from which the figure is taken


A Hurdle in Going Deeper



- + How do we see this conclusion in this figure?
- Well! If it's coming from **overfitting** we should see **better training** and **worst test** risk. This is however **not** the case here!


People started to blame the **vanishing gradient** behavior of **deep NNs**

Depth vs Accuracy: ILSVRC Winners till 2014



Problem of Depth


- + But, what is the problem of *vanishing gradient*?
- This is a general behavior in deep NNs: as we go deeper, the *gradients determined by backpropagation at initial layers get smaller and smaller*, such that at some point they *stop* getting *updated* anymore, even *though they should*
- + How does it come?
- Let's see it! But we follow Albert Einstein advice!



“Everything should be made as *simple* as possible, but not *simpler*!”

Vanishing Gradients: Simple Example


Consider the following **dummy FNN**: *an FNN has a single scalar input, L hidden layers and a single scalar output. All neurons are activated by function $f(\cdot)$ and have no bias*



Let's write **forward pass**

- $y_1 = f(w_1 x)$
- $y_2 = f(w_2 y_1)$
- \dots
- $y_L = f(w_L y_{L-1})$
- $\theta = f(w_{L+1} y_L)$

Vanishing Gradients: Simple Example



Now, we go for **backward** pass: we start with $\overleftarrow{\theta} = \frac{d\hat{R}}{d\theta}$ and go backward


- $\theta = f(w_{L+1}y_L)$

$$\overleftarrow{y}_L = \frac{d\hat{R}}{dy_L} = \frac{d\hat{R}}{d\theta} \frac{d\theta}{dy_L} = \overleftarrow{\theta} w_{L+1} \dot{f}(w_{L+1}y_L)$$

- $y_L = f(w_L y_{L-1})$

$$\begin{aligned}\overleftarrow{y}_{L-1} &= \frac{d\hat{R}}{dy_{L-1}} = \frac{d\hat{R}}{d\theta} \frac{d\theta}{dy_L} \frac{d\theta}{dy_{L-1}} = \overleftarrow{y}_L w_L \dot{f}(w_L y_{L-1}) \\ &= \overleftarrow{\theta} w_{L+1} w_L \dot{f}(w_L y_{L-1}) \dot{f}(w_{L+1} y_L)\end{aligned}$$

Vanishing Gradients: Simple Example




As we keep on going **backward**, the multiplication terms expand

- $y_2 = f(w_2 y_1)$

$$\begin{aligned}\overleftarrow{y}_1 &= \frac{d\hat{R}}{dy_1} = \frac{d\hat{R}}{dy_2} \frac{dy_2}{dy_1} = \overleftarrow{y}_2 w_2 \dot{f}(w_2 y_1) \\ &= \overleftarrow{\theta} \prod_{\ell=2}^{L+1} w_\ell \dot{f}(w_\ell y_{\ell-1})\end{aligned}$$

Vanishing Gradients: Simple Example



Now, let's compute derivative of loss with respect to the **first weight w_1**

- $y_1 = f(w_1 x)$

$$\begin{aligned}
 \frac{d\hat{R}}{dw_1} &= \frac{d\hat{R}}{dy_1} \frac{dy_1}{dw_1} = \overleftarrow{y_1} x \dot{f}(w_1 x) \\
 &= \overleftarrow{\theta} x \dot{f}(w_1 x) \prod_{\ell=2}^{L+1} \textcolor{green}{w_\ell} \dot{f}(w_\ell \textcolor{red}{y_{\ell-1}})
 \end{aligned}$$

Vanishing Gradients: Simple Example

$$\frac{d\hat{R}}{dw_1} = \overleftarrow{\theta} x \dot{f}(w_1 x) \underbrace{\prod_{\ell=2}^{L+1} w_\ell \dot{f}(w_\ell y_{\ell-1})}_{\text{accumulated in Backpropagation}}$$

Now, consider the following cases

- **Case I:** We have a *sigmoid* activation and *all weight are smaller than 1*
 - ↳ Note that for sigmoid $\dot{f}(x) < 1$ for any x
 - ↳ There is one number $a < 1$ that all weights and derivatives are smaller than, e.g., $a = 1 - 10^{-10}$

$$\frac{d\hat{R}}{dw_1} = \overleftarrow{\theta} x \dot{f}(w_1 x) \prod_{\ell=2}^{L+1} w_\ell \dot{f}(w_\ell y_{\ell-1}) < \overleftarrow{\theta} x a^{2L+1}$$

Exploding Gradients: Simple Example

$$\frac{d\hat{R}}{dw_1} = \overleftarrow{\theta} x \dot{f}(w_1 x) \underbrace{\prod_{\ell=2}^{L+1} w_\ell \dot{f}(w_\ell y_{\ell-1})}_{\text{accumulated in Backpropagation}}$$

Now, consider the following cases

- **Case I:** We have a *sigmoid* activation and *all weight are smaller than 1*

$$\lim_{L \uparrow \infty} \frac{d\hat{R}}{dw_1} = 0$$

↳ Backpropagation accumulation concentrates at *zero!*

Gradient with respect to *first layer vanishes as the network gets too deep*

Exploding Gradients: Simple Example

$$\frac{d\hat{R}}{dw_1} = \overleftarrow{\theta} x \dot{f}(w_1 x) \underbrace{\prod_{\ell=2}^{L+1} w_\ell \dot{f}(w_\ell y_{\ell-1})}_{\text{accumulated in Backpropagation}}$$

Now, consider the following cases

- **Case II:** We have a ReLU activation and all weight are larger than 1
 - ↳ Assume $x > 0$; then, $\dot{f}(w_\ell y_{\ell-1}) = 1$ since all the sequence is positive
 - ↳ There is one number $a > 1$ that all weights are larger than, e.g.,
 $a = 1 + 10^{-10}$

$$\frac{d\hat{R}}{dw_1} = \overleftarrow{\theta} x \dot{f}(w_1 x) \prod_{\ell=2}^{L+1} w_\ell \dot{f}(w_\ell y_{\ell-1}) > \overleftarrow{\theta} x a^L$$

Vanishing Gradients: Simple Example

$$\frac{d\hat{R}}{dw_1} = \overleftarrow{\theta} x \dot{f}(w_1 x) \underbrace{\prod_{\ell=2}^{L+1} w_\ell \dot{f}(w_\ell y_{\ell-1})}_{\text{accumulated in Backpropagation}}$$

Now, consider the following cases

- **Case I:** We have a ReLU activation and *all weight are larger than 1*

$$\lim_{L \uparrow \infty} \frac{d\hat{R}}{dw_1} \rightarrow \infty$$

↳ Backpropagation accumulation *explodes!*

Gradient with respect to *first layer explodes* as the network *gets too deep*

Exploding-Vanishing Gradients: Summary

Moral of Story

As the network gets very **deep**, the gradients of initial layers can get **extremely small** or **large**

- The **vanishing** occurs more frequently
 - ↳ Weights are often adjusted by optimizer to get **small**
 - ↳ Once they all get **small**, the gradient starts to vanish
- **Weights and derivative of activation function** are key deciders
 - ↳ We need them to mostly around 1

The above observation also explains why we had some **specific preferences**

- We preferred **ReLU activation** in hidden layers
 - ↳ $\text{ReLU}(x) = 1$ when $x > 0$
- We preferred normalized features
 - ↳ *this can keep weights normalized*

Exploding-Vanishing Gradients: Solution

- + *But, is there any solution for that?*
- Yes! Actually, we already had some one them!

In practice, we can use different approaches to control this behavior

- We use **better** activations in **deep** NNs
 - ↳ *this is why in deep CNNs we use mostly ReLU*
- We apply batch-normalization
 - ↳ *we already have discussed it!*
- We use **skip connection**
 - ↳ *this helps us go even deeper!*

Let's understand what **skip connection** is!