

# Applied Deep Learning

## Chapter 6: Recurrent NNs

Ali Bereyhi

[ali.bereyhi@utoronto.ca](mailto:ali.bereyhi@utoronto.ca)

Department of Electrical and Computer Engineering  
University of Toronto

Fall 2025

# Learning from Sequence Data

In many applications, we have *sequence data*, e.g.,

- *speech data* which is usually a long *time series*
- *text data* which is *sequence of words* and letters
- *financial data* that is typically a *time-dependent sequence of values*

Learning from such data can be *inefficient* via FNNs, i.e., MLPs and CNNs

- On one hand, we have *long* sequence
  - ↳ This can easily make the NN size *infeasible*
- On another hand, we do *not* have *so much features*
  - ↳ Just think of a *long* text, where we need to *predict the next word in it*

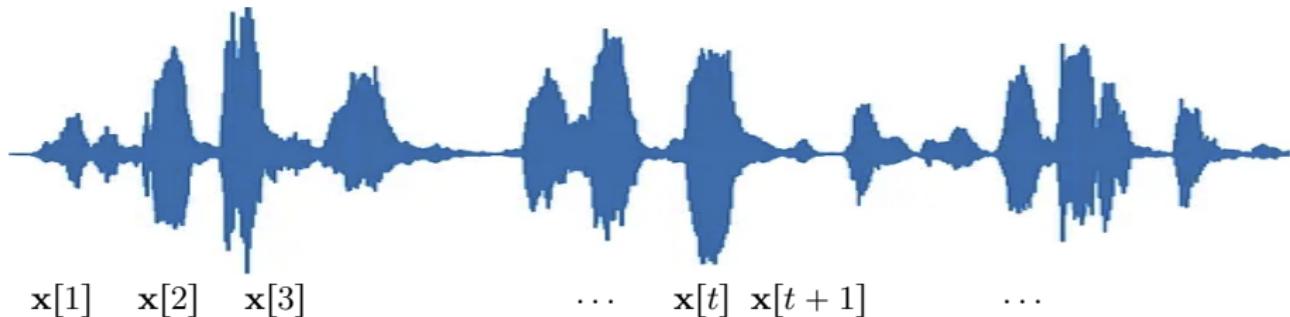
We need to develop some techniques to handle such data

*First, let's see some examples!*

## Learning from Sequence Data: Example 1

Assume we listen to a 15-minute **talk**: we want to find out whether it is about **sport** or **science**

- This is a classification problem
  - ↳ The data-point, i.e., **talk** is classified either as **sport** or **science**
- What about the data?
  - ↳ We **sample** the audio signal at rate 44.1 kHz and quantize the samples
  - ↳ We put every  $N$  successive samples in a **frames**  $\equiv$  **vector of samples**
  - ↳ We store the 15-minute talk as a sequence of time **frames**
  - ↳ We may also store frequency frames from the Fourier transform



## Learning from Sequence Data: Example 1

Assume we listen to a 15-minute **talk**: we want to find out whether it is about **sport** or **science**

- This is a **binary** classification problem
- What about the data? each data-point is a **sequence** of **vectors**

Say we make frames of 512 samples; then, we roughly have

$$77,587 \text{ frames} = 39,690,000 \text{ samples}$$

---

But do we need to pass them all together through an NN?

- It does **not** seem to be!
  - ↳ We can classify based of **simple words** and **expressions** in the talk
  - ↳ Processing all samples together seems to be an unnecessary hardness
- It does **not** generalize well
  - ↳ We want to classify **shorter and longer talks** as well

## Learning from Sequence Data: Example II

Now let's consider another example: we have a **long text** and want to learn what is the **next word** in the sentence

- This is a **prediction** task
  - ↳ Given **previous text** we **predict** the **next outcome**
- What about the data?
  - ↳ We **parse** the text into a sequence of words
  - ↳ We **represent** the letters of each word with their numeric, e.g., ASCII
  - ↳ We **save** each word into an  $N$ -dimensional frame

---

...therapy. UofT undergraduate students explore the use of AI to treat speech

$\mathbf{x}[1] \quad \mathbf{x}[2] \quad \mathbf{x}[3] \quad \dots \quad \mathbf{x}[t] \quad \mathbf{x}[t+1] \quad \dots \quad \mathbf{x}[T] \quad \mathbf{y} = ?$

## Learning from Sequence Data: Example III

Another example: we have a sequence of stock prices and are interested in the future price

- This is again a *prediction* task
  - ↳ Given *previous prices* we *predict* the *future price*
- What about the data?
  - ↳ We put daily prices in form of a *sequence*
  - ↳ We collect every couple of prices along with other indicators into a vector

---

In all these problems: we have a sequence of data and we intend to learn from them in a *generalizable way*

- ↳ We clearly need a *memory component* that can potentially be *infinite*
- ↳ We should keep track of this *infinitely large memory* via limited storage

# Predicting Next Word

Let's start with a simple example: we want to train a neural network that gets a sentence and complete the next word

$\mathbf{x}[t-6] \quad \mathbf{x}[t-5] \quad \mathbf{x}[t-4] \quad \mathbf{x}[t-3] \quad \mathbf{x}[t-2] \quad \mathbf{x}[t-1] \quad \mathbf{x}[t]$   $\mathbf{y} = \mathbf{x}[t+1]$

... Julia has been nominated to receive Alexander von Humboldt Prize for her

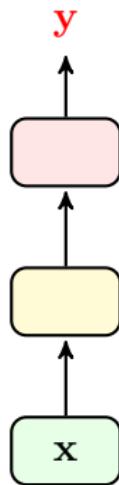
- + How does the training dataset look like?
- We are given with several long texts in the same context: at each entry of sequence in each of these texts the whole sequence is the data-point and the next word is label

# Predicting Next Word

Let's make some specification to clarify the problem

- Each entry is a **vector of dimension  $N$** , i.e.,  $\mathbf{x}[t] \in \mathbb{R}^N$
- We have an NN with **some hidden layers** to train

We show our NN **compactly** with the following diagram

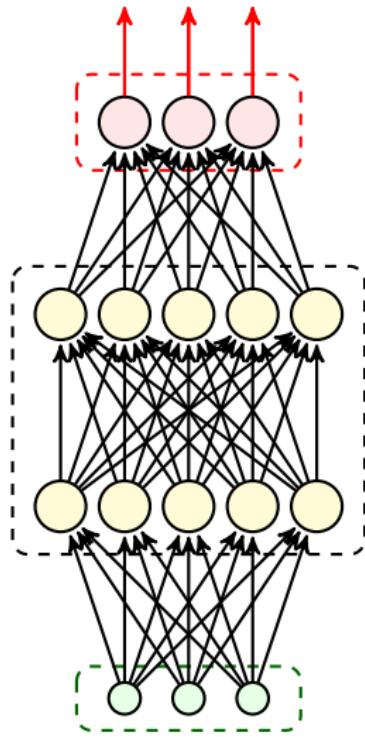
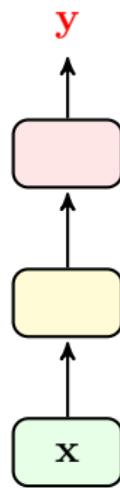


In this diagram

- **Green box** shows the input layer
- **Yellow box** includes hidden layers
  - ↳ It could be several layers
- **Red box** is the output layer
- Arrows refer to all links between the layers
  - ↳ They could be **learnable**

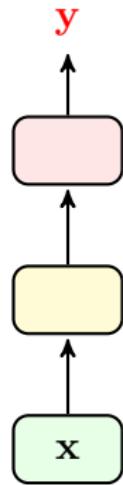
# Predicting Next Word

For instance, we could think of following equivalence



# Predicting Next Word: MLP

Let's try solving this problem with a *simple* MLP



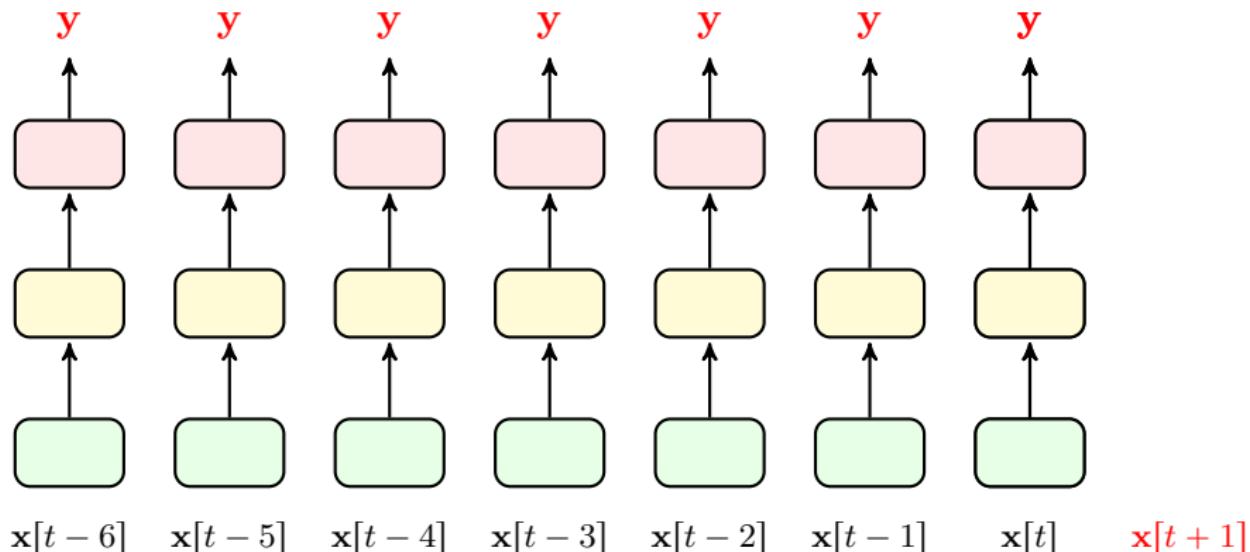
We have a fully-connected FNN that

- takes  $N$  inputs, i.e., *single entry*
- returns  $N$  outputs, i.e., *predicted next word*

We train this MLP

- we go over all text

# Predicting Next Word: MLP



... **Julia** has been nominated to receive Alexander von Humboldt Prize for her

Does FNN **predict** "her"? No! How can it **remember** we are talking about **Julia**?!

# Predicting Next Word: MLP

$x[t-6] \quad x[t-5] \quad x[t-4] \quad x[t-3] \quad x[t-2] \quad x[t-1] \quad x[t] \quad x[t+1]$

... **Julia** has been nominated to receive Alexander von Humboldt Prize for her

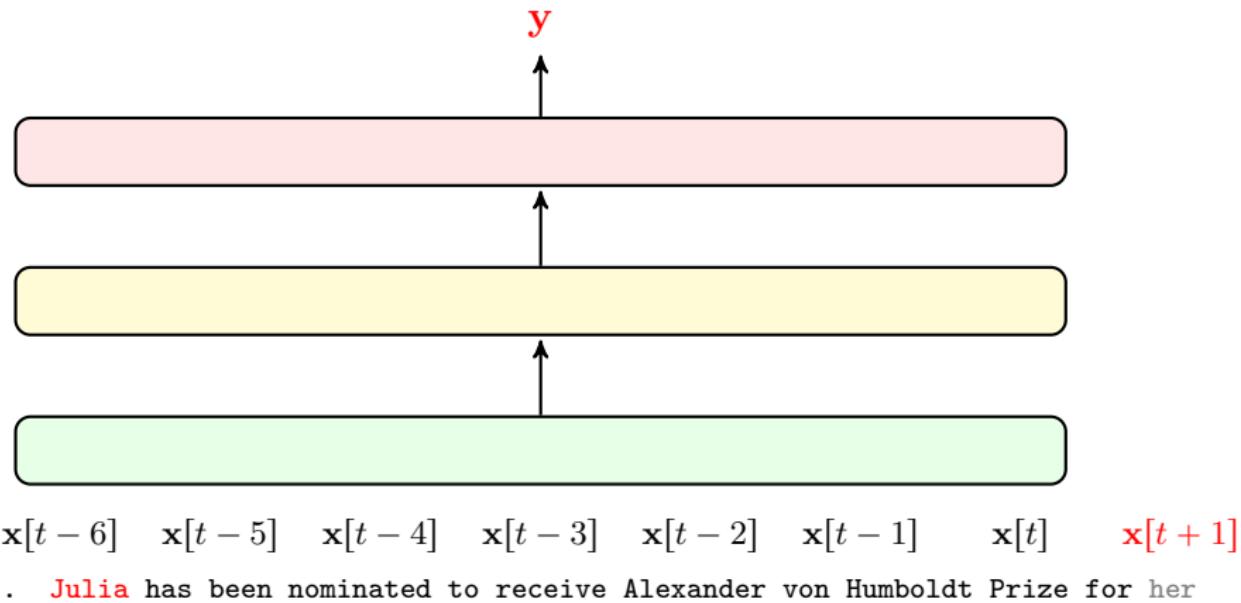
Each time the FNN gets trained with a new sample, it **forgets previous text**

- At the end, it has a set of weights that are average over all predictions
  - ↳ many of these predictions are irrelevant, e.g.,
    - ↳ “nominated to” is followed by “receive”: has nothing to say about “her”
- By the time we get to  $x[t]$ , the FNN gets no input that connects it to **Julia**

This indicates that we need to make a **memory component** for our NN

## Predicting Next Word: Large MLP

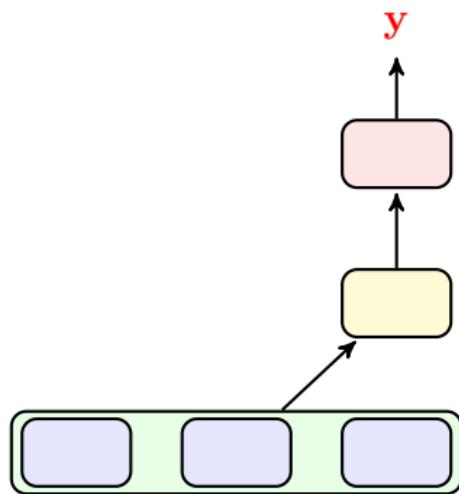
Maybe we could give **more inputs** to the FNN!



But, what if **Julia** has been mentioned 10 pages ago? Forget about **large** MLPs!

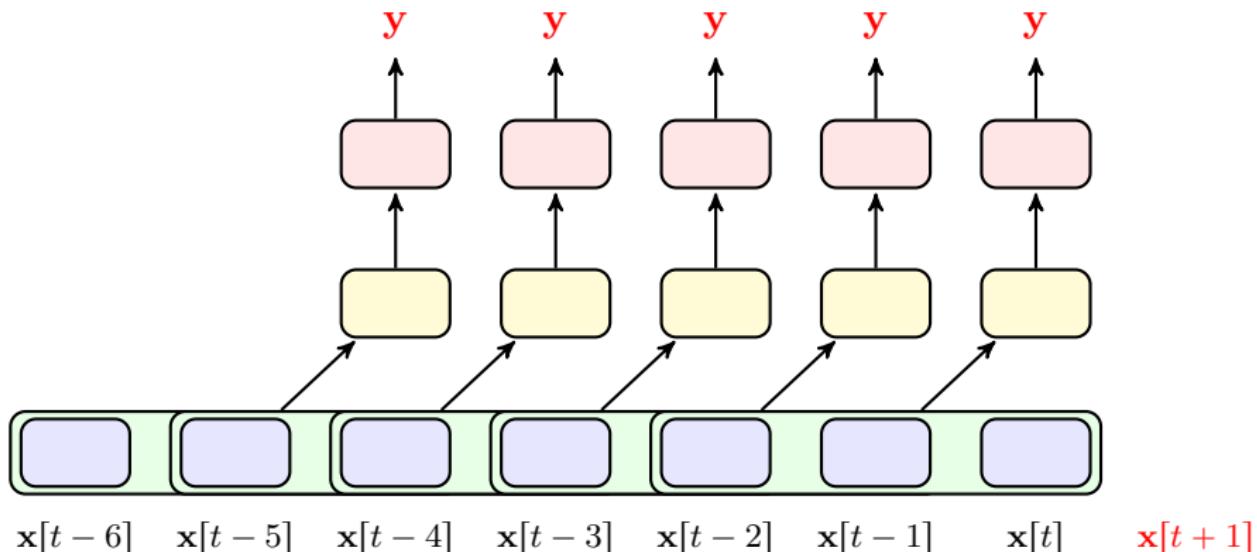
# Predicting Next Word: CNN

Let's now think about **CNNs**



- We have a fully-connected FNN that
- takes  $N$  inputs, i.e., *single entry*
  - returns  $N$  outputs, i.e., *predicted next word*
- We use convolution to
- look into a *larger input* by a *filter of size  $N$*
  - extract *features from a larger part of text* and pass it to hidden layers

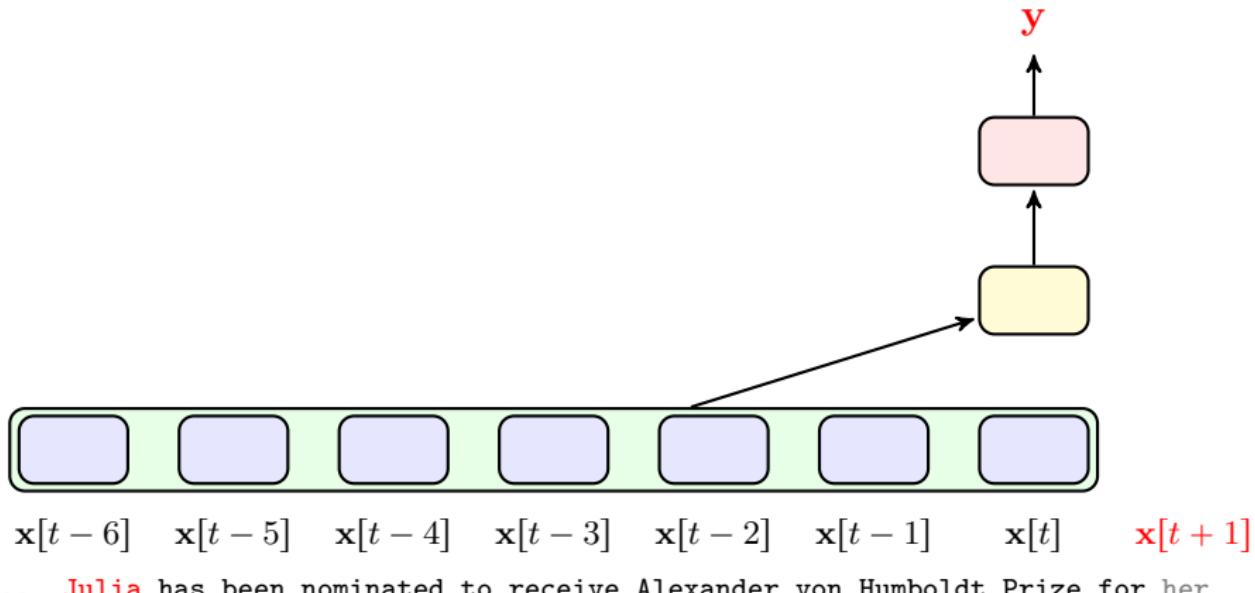
# Predicting Next Word: CNN



... **Julia** has been nominated to receive Alexander von Humboldt Prize for her

Yet, it doesn't seem to remember **Julia**! Unless we slide over the whole text!

# Predicting Next Word: Large CNN



Though better than MLP, it is still **infeasible** to track **the whole text**

# Finite Memory: Root Problem

$\mathbf{x}[t-6] \quad \mathbf{x}[t-5] \quad \mathbf{x}[t-4] \quad \mathbf{x}[t-3] \quad \mathbf{x}[t-2] \quad \mathbf{x}[t-1] \quad \mathbf{x}[t] \quad \mathbf{x}[t+1]$

... **Julia** has been nominated to receive Alexander von Humboldt Prize for her

The problem with *all* architectures we know is that they have **finite memory**

- They can only **remember** from **their input**
  - ↳ we always give them **independent inputs** with **similar features**
  - ↳ they gradually learn to connect any of **such inputs** to **their label**
- If we need to **remember** for **long time** we have to give them **huge inputs**
- But the **memory component** does **not** seem to be so **huge**
  - ↳ We may only remember that **Julia** is a “**single person**” and “**female**”
  - ↳ If text now switches to **Theodore** we should refresh our **memory** that we are talking about a “**single person**” and “**male**”

# Finite Memory Component with Infinite Response

## Component We Miss

We need to extract a **memory component** from our **data** that is **finite in size** but has been **influenced** (at least theoretically) **infinitely**

State-space model can help us building **such memory component**: it is widely used in control theory to describe **evolution** of a system **over time**

Assume  $\mathbf{x}[t]$  is an **input** to a system at time  $t$ : the system returns an **output**  $\mathbf{y}[t]$  to this input and a **state variable**  $\mathbf{s}[t]$ . The output in the next time, i.e.,  $t + 1$ , depends on the **new input** and **current state**, i.e.,

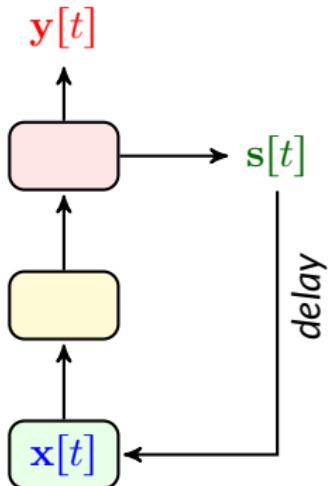
$$\mathbf{y}[t + 1], \mathbf{s}[t + 1] = f(\mathbf{x}[t + 1], \mathbf{s}[t])$$

In the above representation: the **state** is a **finite-size variable** that carries information for **infinitely long time**

# State-Space Model for NNs

We can look at a NN as a **state-dependent** system

*In this architecture, the NN*



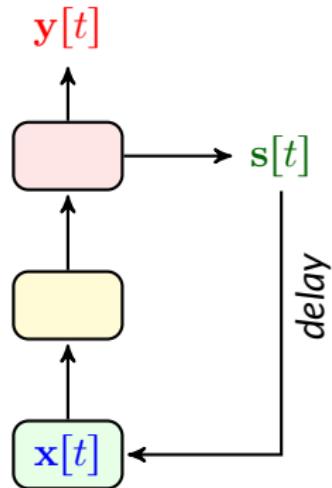
- takes  $x[t]$  and previous state  $s[t - 1]$  as inputs
- returns  $y[t]$  and new state  $s[t]$  as outputs

*This NN captures temporal behavior of data: starting from  $t = 1$ , the NN*

- initiates some  $s[0]$
  - computes  $y[1]$  and  $s[1]$  from time sample  $x[1]$  and  $s[0]$
  - computes  $y[2]$  and  $s[2]$  from  $x[2]$  and  $s[1]$
- ...

## State-Space Model for NNs

Theoretically, this NN has **infinite** time response



Say NN initiates with some  $s[0]$ . It takes only sample  $x[1]$  and no other time samples is given to it. At time  $t$ ,

- $y[t]$  depends on  $s[t - 1]$
- $s[t - 1]$  depends on  $s[t - 2]$
- ...
- $s[1]$  depends on  $x[1]$

This means that  $y[t]$  still remembers  $x[1]$

# State-Dependent NNs

It seems that for our purpose *state-space model* helps extracting a *good memory component*. The challenge is to design a good state-dependent NN

- + Why is it a *challenge*? We make an NN with input ( $\mathbf{x}, \mathbf{s}$ ) and output ( $\mathbf{y}, \mathbf{s}'$ ) and then train it!
- That sounds *easy*, but have some *challenges*

---

Design of state-dependent NNs has two main *challenges*

- ① *Defining the state variable*
  - ↳ how should we *specify* the *state* as we do not have a clear clue about it
- ② *Training the NN over time*: say we get a time sequence data and compute the *empirical risk* by *averaging* over time samples *up to time t*
  - ↳ *empirical risk* depends on *weights* in the NN
  - ↳ it also depends *previous memory components* which can be *learnable*
    - ↳ if we want to train the model *efficiently*, we need to *get back over time* and update all those *memory components*!

# State-Dependent NNs

Several attempts have been done: we look briefly into two of them

- *Jordan Network*
  - ↳ proposed by Michael Jordan<sup>1</sup> in 1986
  - ↳ it computes the *state variable* to be a *simple moving average*
- *Elman Network*
  - ↳ proposed by Jeffrey Elman<sup>2</sup> in 1990
  - ↳ it learns the *state variable* but *does not track back completely over time*

These models are simple forms of what we nowadays know as

Recurrent NNs  $\equiv$  RNNs

## RNN

RNN is an NN with *state-variable*: the term *recurrent* refers to the connection between *former state* and *new output*

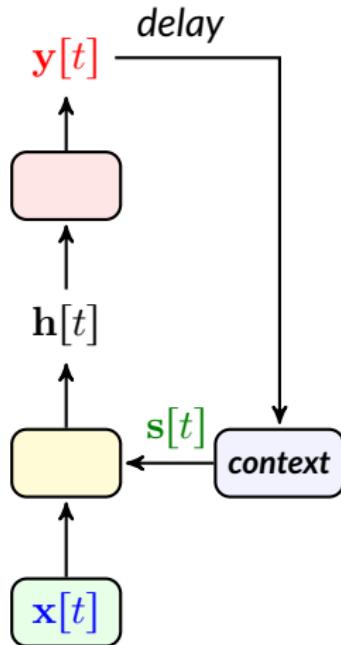
---

<sup>1</sup>Professor at CS Department of University of Berkeley; check his [page](#)

<sup>2</sup>Late Professor at UCSD ('48 - '18); check his [page](#)

# Jordan Model

*Jordan Network* uses a simple moving average of **output** as the state



The proposal was a shallow NN

- It starts with some **initial state  $s[1]$**
- In time  $t$ , it updates the **state  $s[t]$**  with fixed  $\mu$  as

$$s[t] = \mu s[t-1] + y[t-1]$$

- The hidden layer then computes

$$h[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m s[t])$$

- The output is

$$y[t] = f(\mathbf{W}_2 h[t])$$

## Jordan Network

Jordan Network has a **memory** component with **infinite** response: say we set **initial state to zero**, give  $\mathbf{x}[1]$ , and keep the input zero for the rest of time; then,

$$\mathbf{y}[1] = f(\mathbf{W}_2 f(\mathbf{W}_1 \mathbf{x}[1]))$$

$$\mathbf{y}[2] = f(\mathbf{W}_2 f(\mathbf{W}_m \mathbf{y}[1])) = f(\mathbf{W}_2 f(\mathbf{W}_m f(\mathbf{W}_2 f(\mathbf{W}_1 \mathbf{x}[1]))))$$

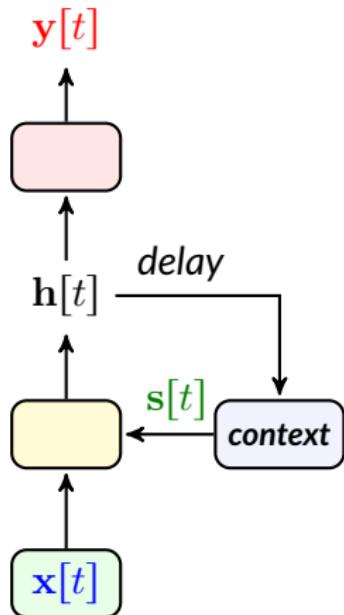
⋮

But, the network does **not** learn how to remember

- It takes a fixed **moving average** as **memory component**
- It only learns how to use this **pre-defined memory** for **prediction**
  - ↳ We could say that it **implicitly** learns to remember by learning  $\mathbf{W}_m$

# Elman Network

*Elman Network* uses *output of hidden layer* as state: also called *hidden state*



*The proposal was a shallow NN*

- It starts with some *initial hidden state  $\mathbf{h}[0]$*
- In time  $t$ , it updates the *state  $\mathbf{s}[t]$*  as

$$\mathbf{s}[t] = \mathbf{h}[t - 1]$$

- The hidden layer then computes

$$\mathbf{h}[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{s}[t])$$

- The output is

$$\mathbf{y}[t] = f(\mathbf{W}_2 \mathbf{h}[t])$$

# Elman Network

Similar to Jordan Network, Elman Network has a **memory** component with **infinite** response: with **zero initial hidden state**, we have

$$\mathbf{y}[1] = f(\mathbf{W}_2 f(\mathbf{W}_1 \mathbf{x}[1]))$$

$$\mathbf{y}[2] = f(\mathbf{W}_2 f(\mathbf{W}_m \mathbf{h}[1])) = f(\mathbf{W}_2 f(\mathbf{W}_m f(\mathbf{W}_1 \mathbf{x}[1])))$$

⋮

*Elman network also learns how to remember only implicitly*

# Challenge of Learning Through Time

Though Jordan and Elman Networks had memory, they did **not** get trained accurately over time, i.e., they simplified the solution to the second challenge

- + What is really this challenge?
- We are going to deal with it in next sections, but let's see it on these simple networks first

---

Let's assume a simple setting: we are to train our NN on single data sequence

- We have the sequence  $\mathbf{x}[1], \dots, \mathbf{x}[T]$  as the data-point
- For each entry of this sequence we have the true label
  - ↳ We have sequence  $\mathbf{v}[1], \dots, \mathbf{v}[T]$  with  $\mathbf{v}[t]$  being label of  $\mathbf{x}[t]$
- We are able to compute the loss between outputs and true labels as<sup>3</sup>

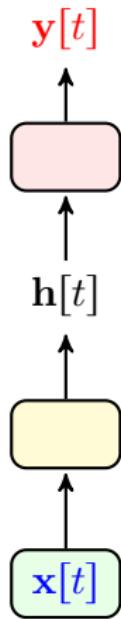
$$\hat{R} = \sum_{t=1}^T \mathcal{L}(\mathbf{y}[t], \mathbf{v}[t])$$

---

<sup>3</sup>We will see that this is not always this easy!

## Recap: Basic FNN

For sake of comparison, let's first train a basic FNN on this data sequence



We have a shallow FNN

- The hidden layer computes

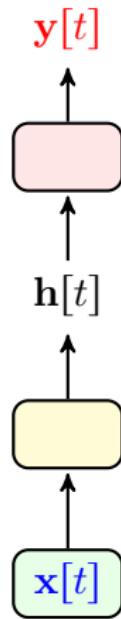
$$\mathbf{h}[t] = f(\mathbf{W}_1 \mathbf{x}[t])$$

- The output is

$$\mathbf{y}[t] = f(\mathbf{W}_2 \mathbf{h}[t])$$

## Recap: Basic FNN

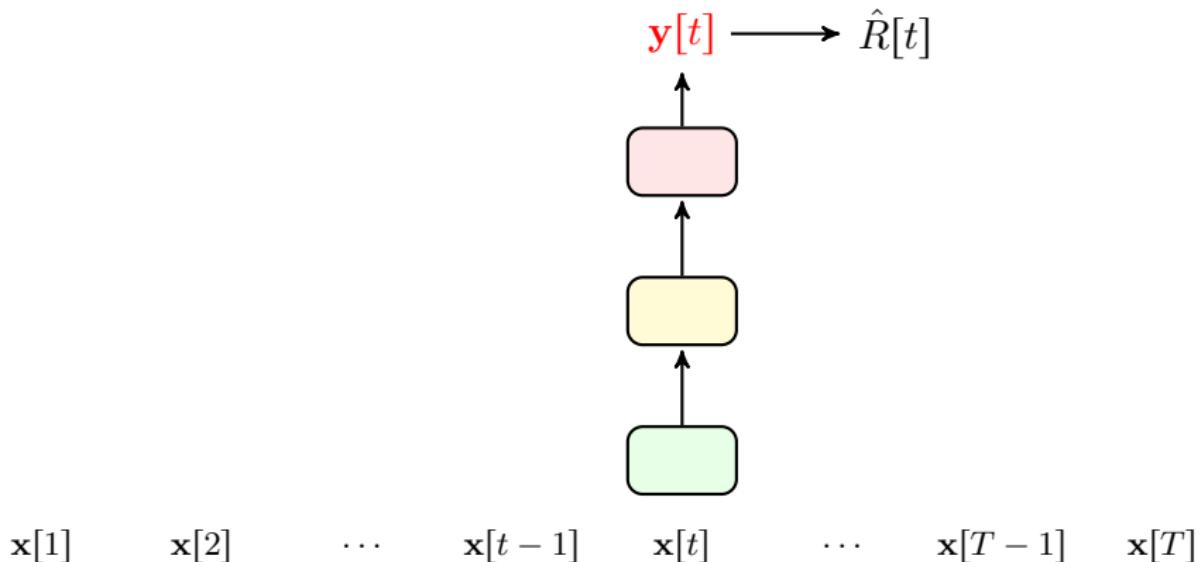
For sake of comparison, let's first train a basic FNN on this data sequence



*How do we train this FNN?*

- We compute the gradients  $\nabla_{\mathbf{W}_1} \hat{R}$  and  $\nabla_{\mathbf{W}_2} \hat{R}$ 
  - ↳ We do it via backpropagation
- We apply gradient descent

# Learning Through Time: FNNs



$$\hat{R} = \sum_{t=1}^T \underbrace{\mathcal{L}(\mathbf{y}[t], \mathbf{v}[t])}_{\hat{R}[t]} = \sum_{t=1}^T \hat{R}[t] \rightsquigarrow \nabla_{\mathbf{W}_i} \hat{R} = \sum_{t=1}^T \nabla_{\mathbf{W}_i} \hat{R}[t]$$

## Learning Through Time: FNNs

Let's learn  $\mathbf{W}_1$  and to ease computation we use our **cheating notation**, i.e., use  $\circ$  to show **any product**: to compute the gradient we start with the **output**

$$\begin{aligned}\nabla_{\mathbf{W}_1} \hat{R}[t] &= \nabla_{\mathbf{W}_1} \mathcal{L}(\mathbf{y}[t], \mathbf{v}[t]) \\ &= \nabla_{\mathbf{y}[t]} \hat{R}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{y}[t]\end{aligned}$$

We know that  $\mathbf{y}[t] = f(\mathbf{W}_2 \mathbf{h}[t])$ , so we can write

$$\begin{aligned}\nabla_{\mathbf{W}_1} \mathbf{y}[t] &= \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t] \\ &= \left( f'(\mathbf{W}_2 \mathbf{h}[t]) \circ \mathbf{W}_2 \right) \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t]\end{aligned}$$

What about  $\nabla_{\mathbf{W}_1} \mathbf{h}[t]$ ? We keep on backward!

# Learning Through Time: FNNs

Up to now, we have

$$\nabla_{\mathbf{W}_1} \hat{R}[t] = \nabla_{\mathbf{y}[t]} \hat{R}[t] \circ \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t]$$

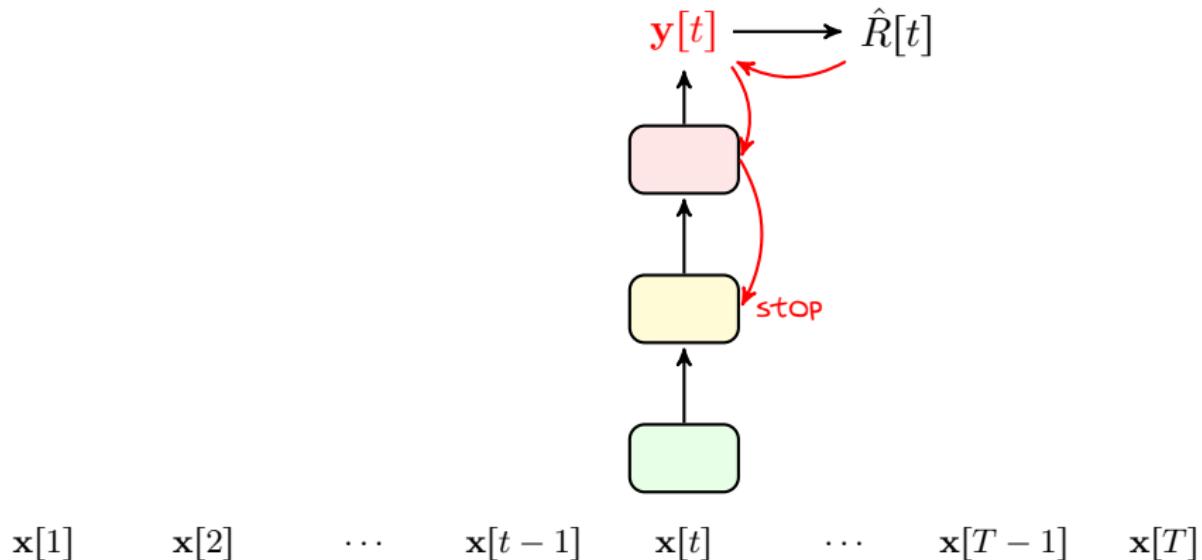
We use the fact that  $\mathbf{h}[t] = f(\mathbf{W}_1 \mathbf{x}[t])$

$$\begin{aligned} \nabla_{\mathbf{W}_1} \mathbf{h}[t] &= \nabla_{\mathbf{W}_1} f(\mathbf{W}_1 \mathbf{x}[t]) + \nabla_{\mathbf{x}[t]} \mathbf{h}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{x}[t] \\ &= \dot{f}(\mathbf{W}_1 \mathbf{x}[t]) \circ \mathbf{x}[t] + \left( \dot{f}(\mathbf{W}_1 \mathbf{x}[t]) \circ \mathbf{W}_1 \right) \circ \underbrace{\mathbf{0}}_{\mathbf{x}[t] \text{ is not a function of } \mathbf{W}_1} \end{aligned}$$

Therefore, we end chain rule here

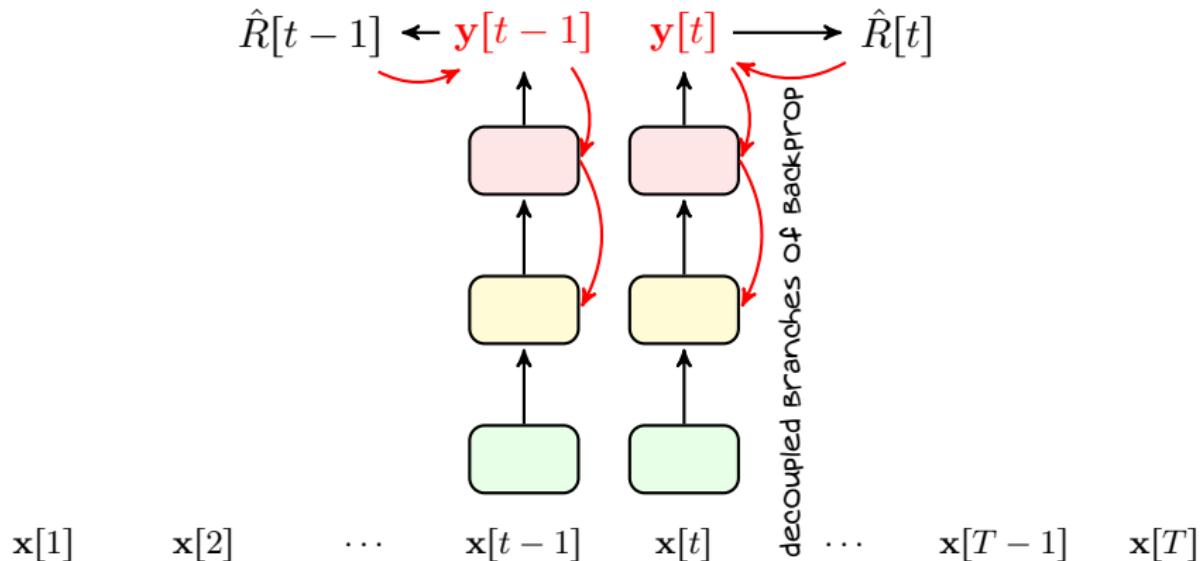
$$\nabla_{\mathbf{W}_1} \hat{R}[t] = \nabla_{\mathbf{y}[t]} \hat{R}[t] \circ \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t]$$

# Learning Through Time: FNNs



$$\nabla_{\mathbf{W}_1} \hat{R}[t] = \nabla_{\mathbf{y}[t]} \hat{R}[t] \circ \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t]$$

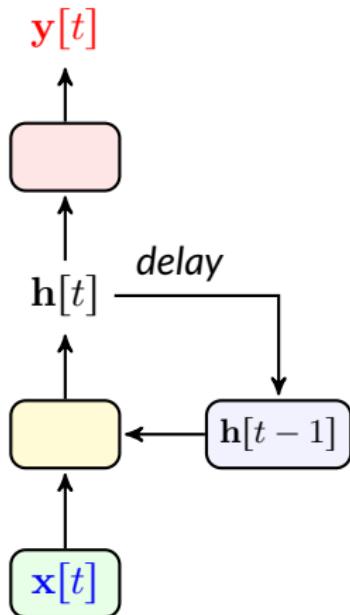
# Learning Through Time: FNNs



$$\nabla_{\mathbf{W}_2} \hat{R} = \sum_{t=1}^T \nabla_{\mathbf{W}_2} \hat{R}[t]$$

# Training a Basic RNN

Now, let's train Elman network on this sequence



*The proposal was a shallow NN*

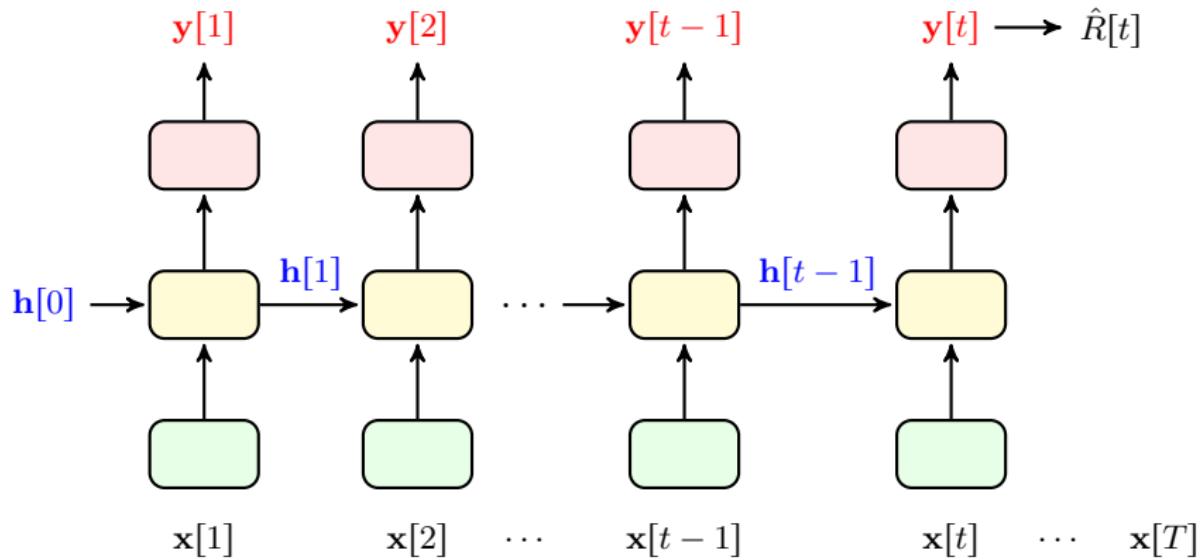
- It starts with some *initial hidden state  $h[0]$*
- The hidden layer then computes

$$h[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m h[t - 1])$$

- The output is

$$y[t] = f(\mathbf{W}_2 h[t])$$

# Inferring Through Time: Elman Network



$$\hat{R} = \sum_{t=1}^T \underbrace{\mathcal{L}(\mathbf{y}[t], \mathbf{v}[t])}_{\hat{R}[t]} = \sum_{t=1}^T \hat{R}[t]$$

# Learning Through Time: Elman Network

Let's again try to learn  $\mathbf{W}_1$ : we start with the *output*

$$\nabla_{\mathbf{W}_1} \hat{R}[t] = \nabla_{\mathbf{y}[t]} \hat{R}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{y}[t]$$

We know that  $\mathbf{y}[t] = f(\mathbf{W}_2 \mathbf{h}[t])$ , so we can write

$$\begin{aligned}\nabla_{\mathbf{W}_1} \mathbf{y}[t] &= \nabla_{\mathbf{h}[t]} \mathbf{y}[t] \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t] \\ &= \dot{f}(\mathbf{W}_2 \mathbf{h}[t]) \circ \mathbf{W}_2 \circ \nabla_{\mathbf{W}_1} \mathbf{h}[t]\end{aligned}$$

Next, we note that  $\mathbf{h}[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{h}[t-1])$

$$\begin{aligned}\nabla_{\mathbf{W}_1} \mathbf{h}[t] &= \nabla_{\mathbf{W}_1} f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{h}[t-1]) \\ &\quad + \nabla_{\mathbf{x}[t]} \mathbf{h}[t] \circ \underbrace{\nabla_{\mathbf{W}_1} \mathbf{x}[t]}_0 \\ &\quad + \nabla_{\mathbf{h}[t-1]} \mathbf{h}[t] \circ \underbrace{\nabla_{\mathbf{W}_1} \mathbf{h}[t-1]}_?\end{aligned}$$

## Learning Through Time: Elman Network

Well! We know that  $\mathbf{h}[t - 1] = f(\mathbf{W}_1 \mathbf{x}[t - 1] + \mathbf{W}_m \mathbf{h}[t - 2])$

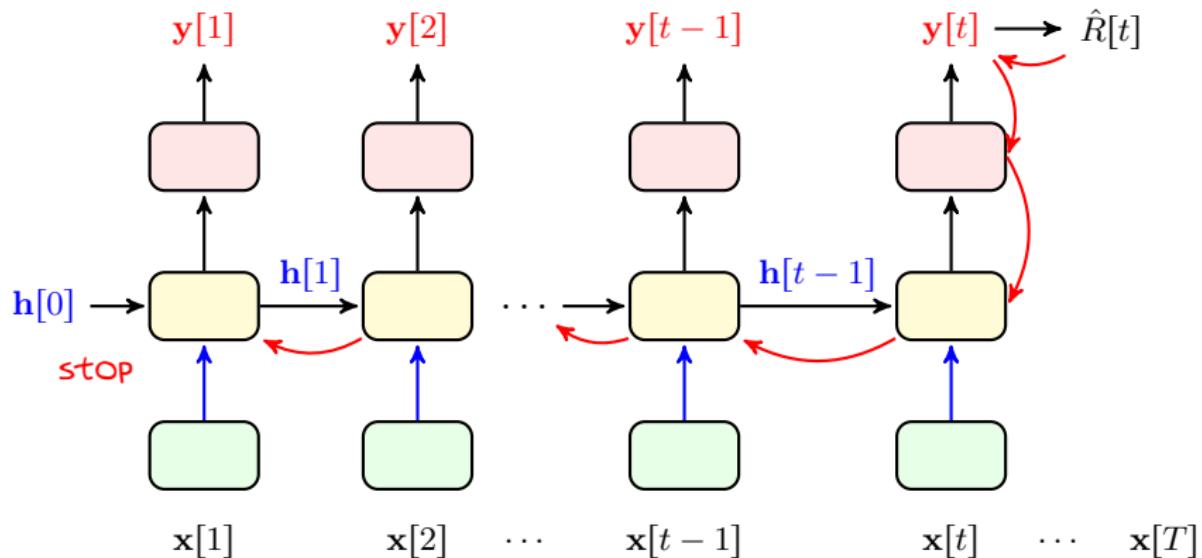
$$\begin{aligned}\nabla_{\mathbf{W}_1} \mathbf{h}[t - 1] &= \nabla_{\mathbf{W}_1} f(\mathbf{W}_1 \mathbf{x}[t - 1] + \mathbf{W}_m \mathbf{h}[t - 2]) \\ &\quad + \nabla_{\mathbf{x}[t-1]} \mathbf{h}[t - 1] \circ \underbrace{\nabla_{\mathbf{W}_1} \mathbf{x}[t - 1]}_0 \\ &\quad + \nabla_{\mathbf{h}[t-2]} \mathbf{h}[t - 1] \circ \underbrace{\nabla_{\mathbf{W}_1} \mathbf{h}[t - 2]}_?\end{aligned}$$

We are not still done! We know  $\mathbf{h}[t - 2] = f(\mathbf{W}_1 \mathbf{x}[t - 2] + \mathbf{W}_m \mathbf{h}[t - 3])$

$$\begin{aligned}\nabla_{\mathbf{W}_1} \mathbf{h}[t - 2] &= \nabla_{\mathbf{W}_1} f(\mathbf{W}_1 \mathbf{x}[t - 2] + \mathbf{W}_m \mathbf{h}[t - 3]) \\ &\quad + \nabla_{\mathbf{x}[t-2]} \mathbf{h}[t - 2] \circ \underbrace{\nabla_{\mathbf{W}_1} \mathbf{x}[t - 2]}_0 \\ &\quad + \nabla_{\mathbf{h}[t-3]} \mathbf{h}[t - 2] \circ \underbrace{\nabla_{\mathbf{W}_1} \mathbf{h}[t - 3]}_?\end{aligned}$$

# Learning Through Time: Elman Network

We should in fact pass all the way **back** to the initial time interval time!



Note that all blue edges are representing  $W_1$

# Learning Through Time

## Moral of Story

To learn how to **remember**, we need to **train** our RNN **through time**: at each time interval, we should move **all the way back to origin** to find out how **exactly** we should change the weights!

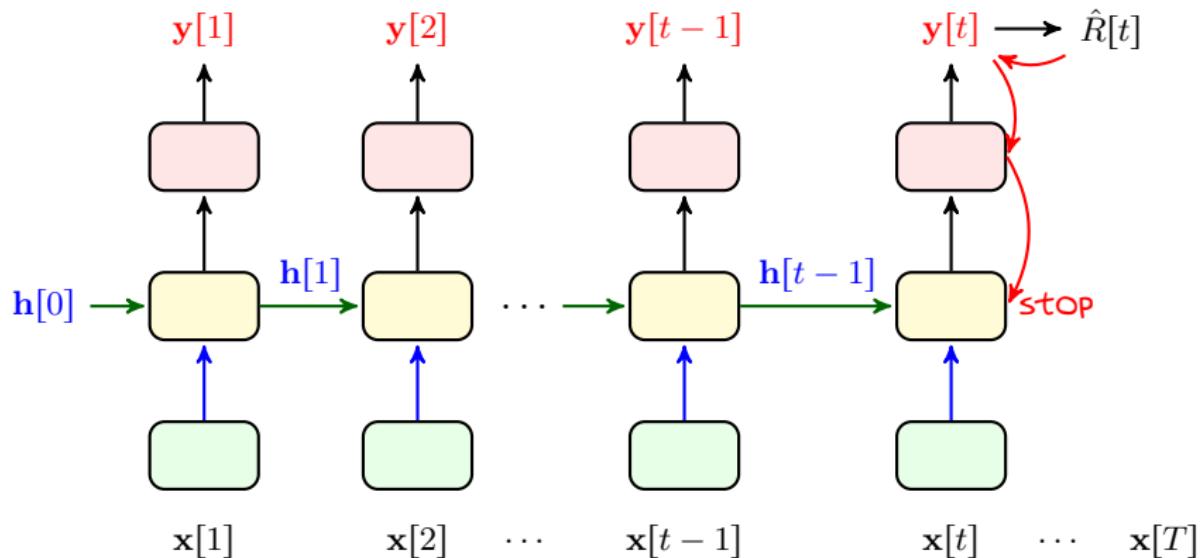
- + Did **Elman** did so?
- **Not** really!

For training **Elman** treated the hidden state  $\mathbf{h}[t - 1]$  as a fixed variable, i.e., he assumed  $\nabla_{\mathbf{W}_1} \mathbf{h}[t - 1] \approx \nabla_{\mathbf{W}_m} \mathbf{h}[t - 1] = 0$ ! So, he did not need to move backward in time!

This means that **Elman** did **not** really addressed **the second challenge**!

# Learning Through Time: Elman's Approximation

Elman treated it as a simple FNN with only one extra input!



Training  $\mathbf{W}_1$  is exactly as FNN. We just have one extra  $\mathbf{W}_m$  here!

## RNNs: Need to Learn Memory

Though appreciated, Elman and Jordan Networks did not do the job

- ① Their *memory component* is rather *simple*
  - ↳ We should use *deeper models* that enable advanced memory components
- ② They do *not* really *learn* how to *remember*
  - ↳ We should *train* the memory component *over time*

*These led us to development of RNNs!*

### RNN: Less Generic Definition

An RNN can be designed with *any known architecture* by letting NN also *learn from its past features* and *outputs*. This new enabling is called *recurrence*