

Applied Deep Learning

Chapter 6: Recurrent NNs

Ali Bereyhi

ali.bereyhi@utoronto.ca

Department of Electrical and Computer Engineering
University of Toronto

Fall 2025

RNN: Dataset and Learning Setting

We are now looking into a **supervised learning** problem where we are to learn **label** from a **sequence of data** that has generally a **temporal correlation**

Let's denote the **sequence** with $\mathbf{x}[1], \dots, \mathbf{x}[T]$

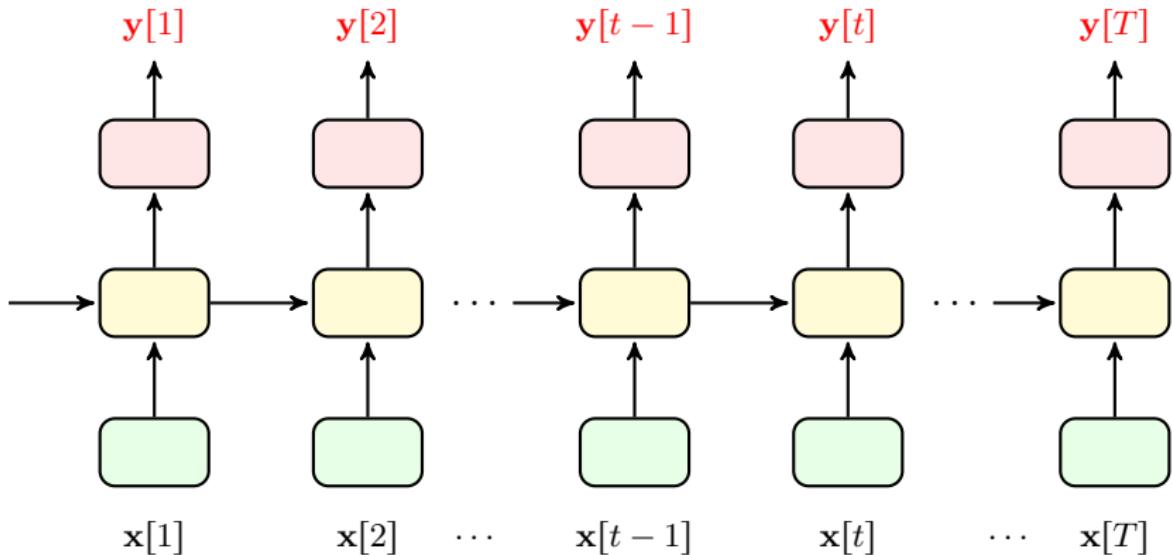
Temporal Correlation

By **temporal correlation** we mean that entries at **other time instances** carry information about one particular entry $\mathbf{x}[t]$

- + But how is a **label** assigned to this sequence?
- Well, that can be of various forms!

Types of Problems with Sequence Data

We considered a very simple case: **many-to-many type I**

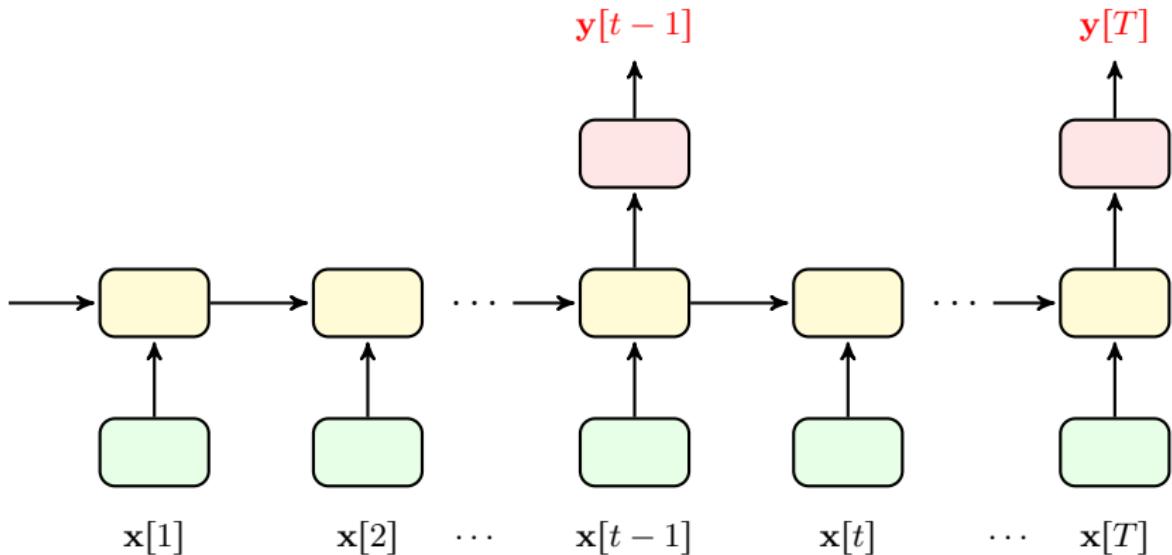


In this case, every entry has a **label**

↳ speech tagging: $x[t]$ is a part of speech and $y[t]$ is its tag

Types of Problems with Sequence Data

We considered a very simple case: **many-to-many type II**

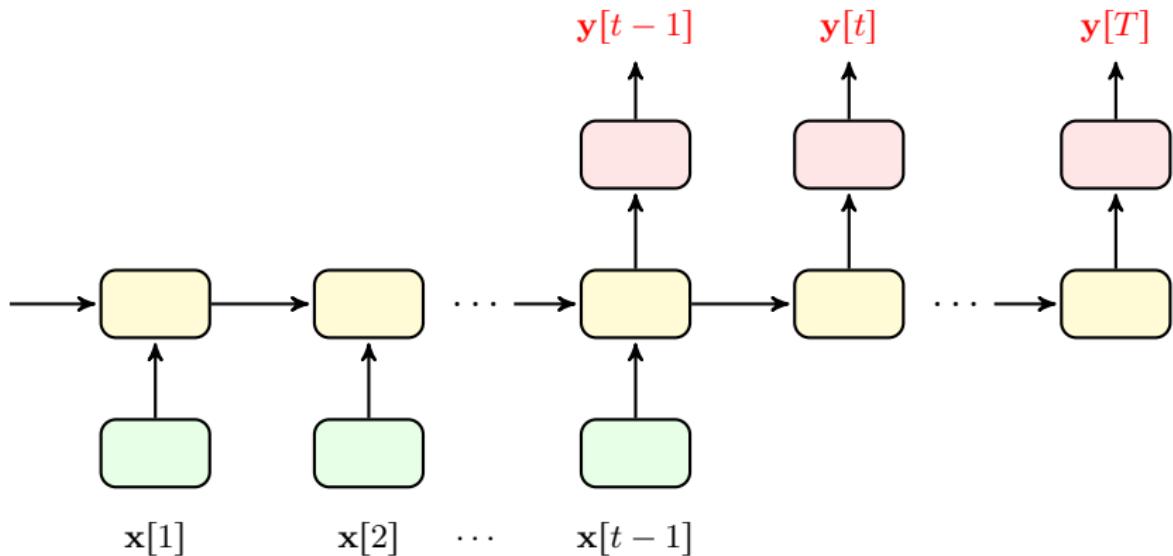


In this case, we get a **label** once after multiple entries

- ↳ speech recognition: $x[t]$ is a small part of speech and $y[t]$ says what is a every couple of minutes about

Types of Problems with Sequence Data

We considered a very simple case: **many-to-many type III**

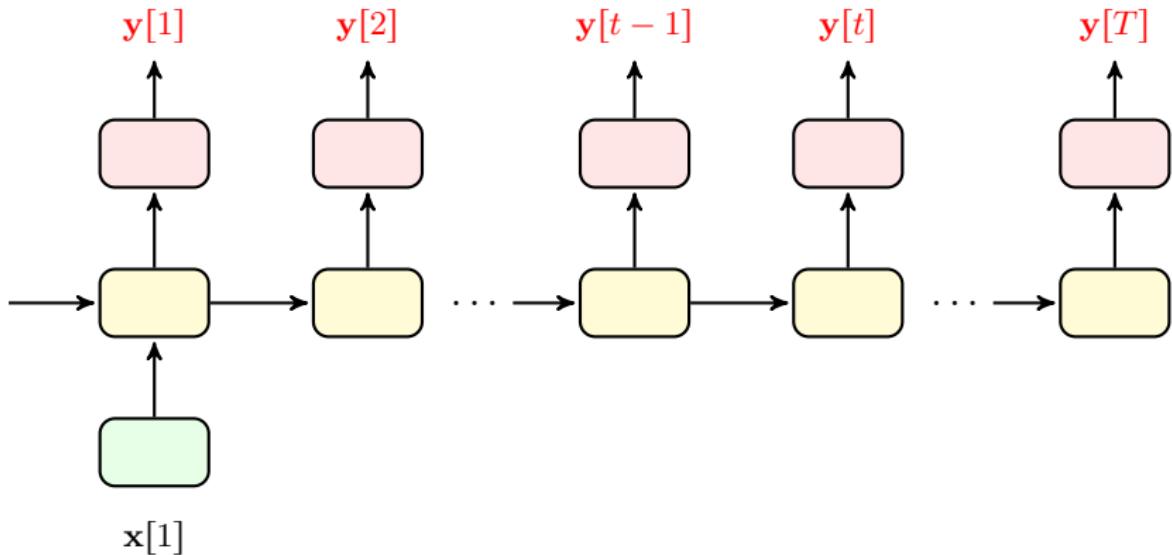


In this case, we start to get **labels** after some delay

- ↳ language translation: $x[1], \dots, x[t - 1]$ is a sentence in German and $y[t - 1], \dots, y[T]$ is its translation to English

Types of Problems with Sequence Data

We considered a very simple case: **one-to-many**

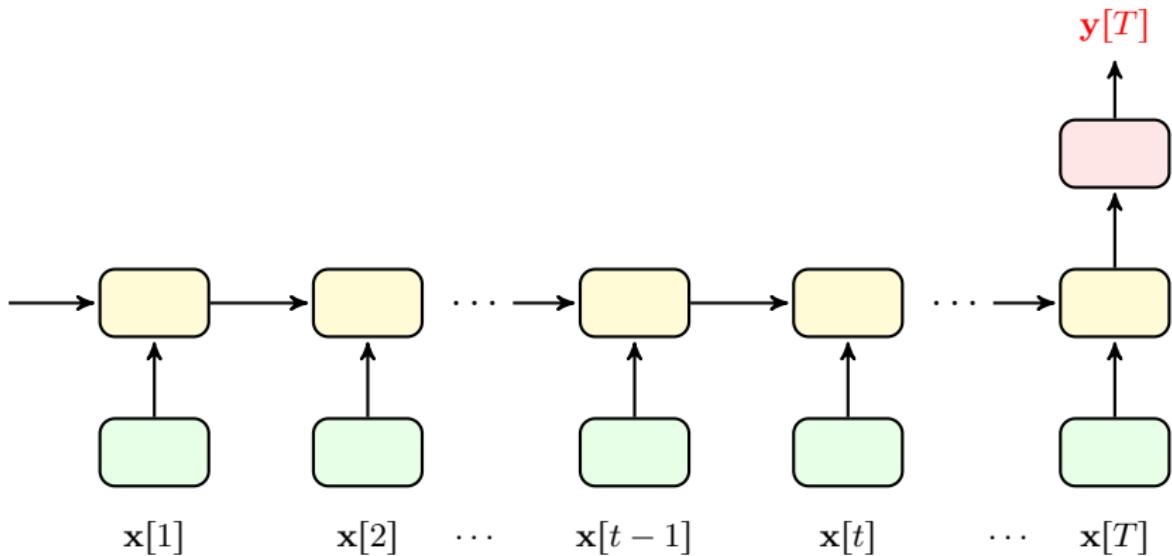


In this case, we get only one input data and have a sequence of **labels**

- ↳ *image captioning: $x[1]$ is an image and $y[1], \dots, y[T]$ is a caption describing what is inside the image*

Types of Problems with Sequence Data

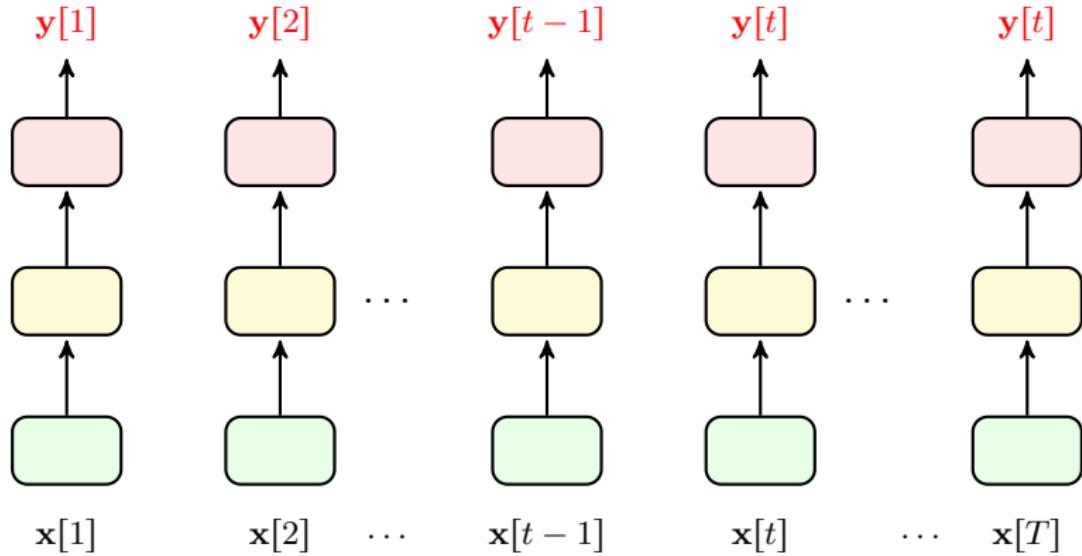
We considered a very simple case: **many-to-one**



In this case, we get only one **label** for a whole input sequence

- ↳ **sequence classification:** $x[1], \dots, x[T]$ is a speech and $y[T]$ says if this speech is constructive or destructive

Types of Problems with Sequence Data: FNNs



In fact, FNNs are **one-to-one** RNNs

- ↳ we can think of every data-point as *a sequence of length one*, or
- ↳ we may think of dataset as a long sequence with *no temporal correlation*

RNN: General Form

To construct an RNN, we can use any module that we have learned:

- we can use a **fully-connected layer**
- we can use a **convolutional layer**
- we can use a **residual unit**
- we may use an **inception unit** used in **GoogLeNet**

...

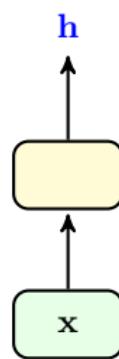
But, classical implementations use **fully-connected layers**

We can use **one** layer to make a **shallow** RNN or **multiple** to make a **deep** RNN

- + Don't we do any **change** to them?
- Not a serious change
 - ↳ We may change the **activation** to $\tanh(\cdot)$: we will see later why
 - ↳ We expand the input dimension: since we need to also give **memory** as input

Shallow RNN

Let's break it down a bit: say the **yellow box** is a **fully-connected layer**

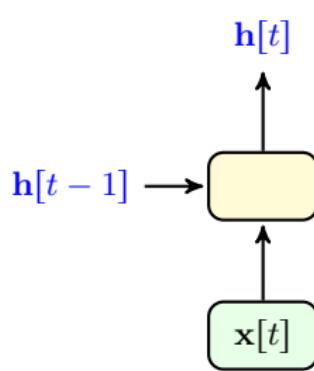


This layer gets input $\mathbf{x} \in \mathbb{R}^N$ and returns activated feature $\mathbf{h} \in \mathbb{R}^M$

- We replace its **activation** by $\tanh(\cdot)$
 - ↳ Not necessary, but usually suggested
- We modify it to get a new input in \mathbb{R}^{N+M}
 - ↳ N entries for **data inputs** \mathbf{x} in time t
 - ↳ M entries for **features** \mathbf{h} in time $t - 1$

Shallow RNN

Let's break it down a bit: say the **yellow box** is a **fully-connected layer**

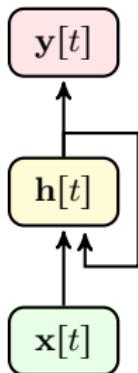


We show it this way now

- We call $h[t]$ usually the **hidden state**
- We pass the **hidden state** through an **output layer**
 - ↳ Output is **not necessarily** corresponding to label

Shallow RNN

We may show our shallow RNN compactly via the following diagram

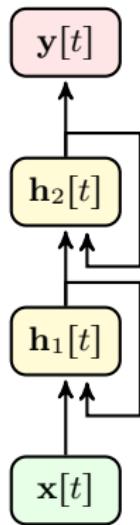


In this diagram

- *Each edge is a set of weights, e.g., a weight matrix*
- *The return edge also has a delay in time*
- + *But, isn't that simply Elman Network?*
- *If we use a fully-connected layer and sigmoid activation; then, Yes! But, Remember that*
 - *Elman did not train it over time*
 - *Elman in its model used sigmoid activation*

Deep RNN

We can add more layers to make a **deep RNN**

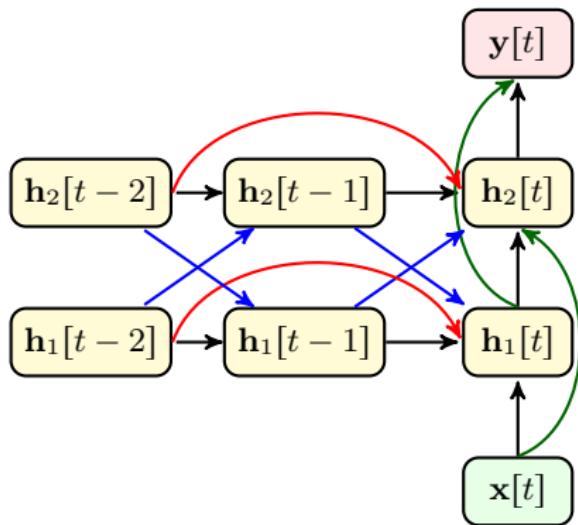


In this diagram

- *Each edge is a set of weights, e.g., a weight matrix*
- *The return edge also has a delay in time*
- *And, it's no more Elman Network*

Deep RNN

It might be easier to think of it as the following diagram



We can use any module that we like

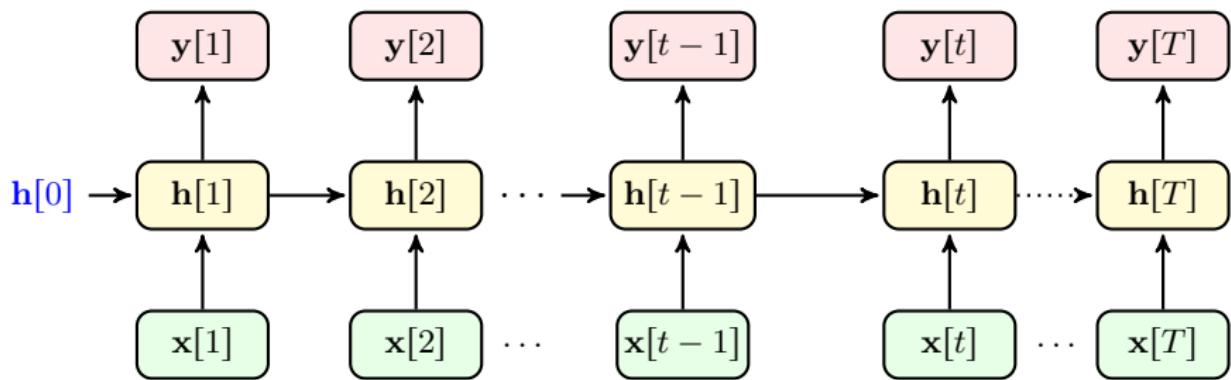
- We may add **skip connection**
- We could make **dense connections**
- We may **skip over time**
- We may replace hidden layers with **convolutions** or anything else

Shallow RNN: Elman-Like Network

Let's start with simple RNN: *a shallow RNN with **fully-connected** hidden layer*

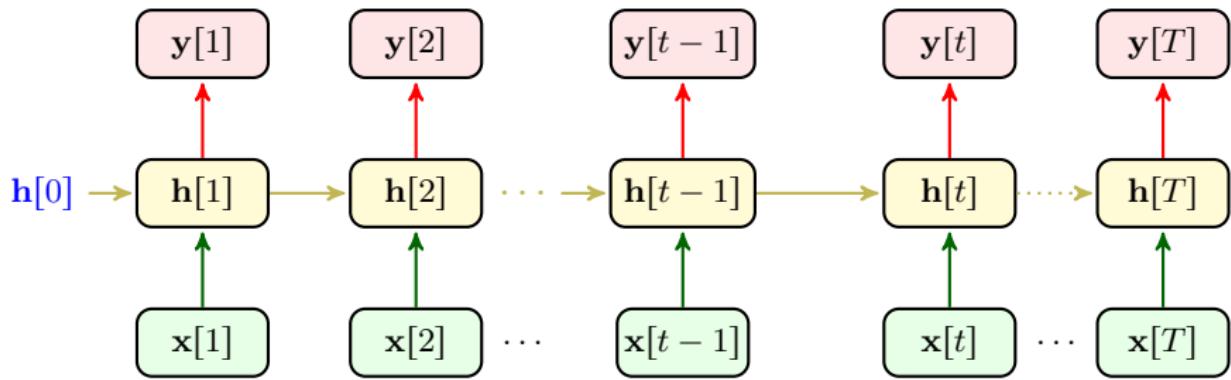
- + You mean *Elman Network*?
- Yes, but we now try to address the *challenges* that Elman did not address

Let's look at the flow of information once again



Shallow RNN: Forward Propagation

Let's specify the learnable parameters with some colors



Say we set the activation to $f(\cdot)$

- ① We have $y[t] = f(\mathbf{W}_2 \mathbf{h}[t])$: we *can learn* \mathbf{W}_2
- ② We have $\mathbf{h}[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{h}[t-1])$: we *can learn* \mathbf{W}_1 and \mathbf{W}_m
- ③ We can start with any *hidden state*: this means we *can learn* $\mathbf{h}[0]$ too