

# Applied Deep Learning

## Chapter 6: Recurrent NNs

Ali Bereyhi

[ali.bereyhi@utoronto.ca](mailto:ali.bereyhi@utoronto.ca)

Department of Electrical and Computer Engineering  
University of Toronto

Fall 2025

## Computing Loss: Challenge

We mentioned several times in this chapter that we *assume*  
*we can compute the loss between RNN's output sequence and label sequence*

However, it is in general a challenge!

---

- + Why is it a challenge? We did it *easily* in FNN and CNN chapters!
- Because the **problem** there was already *properly segmented!*
- + What do you mean by *segmented?*
- Let's break it down!

## Computing Loss: Motivating Example

Let's consider a simple example: we have an image that includes a sequence of handwritten digits, e.g.,

- The sequence includes five digits
- Each digit is either 1, 2, 3, or 4

Our task is to recognize this sequence, i.e., return the five digits in correct order

- This is a classification task
- How can we do it? We use NNs
  - ↳ We train an NN over lots of images: we have lots of sequence of digits
  - ↳ We then use it to recognize new images



23 241 ↘ 23241

Let's say we are going to use a CNN

# Computing Loss: Motivating Example


To use a CNN, we need to specify our input size

- We segment an input image into a sequence of **five images**
  - ↳ These images are all as large as CNN's input size



- We label each image with its **label**, e.g., 2 is **labeled as 2**
- We give these **five images** to our CNN and get **five outputs**
  - ↳ Assume we use **softmax** at the output layer
  - ↳ For each image, we get a **vector of size 4** as output
    - ↳ Each entry represents **probability** of **image** being one of digits 1, 2, 3, and 4
- To compute loss, we compare each **output** with its **corresponding label**

# Computing Loss: Motivating Example



↳  $y_j[t]$  is the **probability** of **digit** in time  $t$  being  $j$

# Computing Loss: Motivating Example

Here, we already have the data **segmented** into

a sequence that for each **time step** has a **label**

So, computing loss is easy as pie!

$$\begin{aligned}\hat{R} &= \mathcal{L}(\mathbf{y}[1], \dots, \mathbf{y}[5], \mathbf{v}[1], \dots, \mathbf{v}[5]) \\ &= \sum_{t=1}^5 \mathcal{L}(\mathbf{y}[t], \mathbf{v}[t]) = \sum_{t=1}^5 \hat{R}[t]\end{aligned}$$

When we compute gradients, we note that **only**  $\hat{R}[t]$  depends on  $\mathbf{y}[t]$ : so, for a given output at time  $t = i$  we can simply write

$$\nabla_{\mathbf{y}[i]} \hat{R} = \sum_{t=1}^5 \nabla_{\mathbf{y}[i]} \hat{R}[t] = \nabla_{\mathbf{y}[i]} \hat{R}[i] = \nabla_{\mathbf{y}[i]} \mathcal{L}(\mathbf{y}[i], \mathbf{v}[i])$$

# Computing Loss: One-to-One Correspondence

- + But is it practical to do segmentation **by hand**?
- **No!** This is why we built RNNs!

With RNNs, we address this learning task as below


- We look at the **complete image** as a sequence of data
  - ↳ We divide input into **multiple equal-size frames**
- We go over each frame separately
  - ↳ We give the frame as the input along with **previous state**
  - ↳ We compute a **new state** which can potentially give us the output

# Computing Loss: One-to-One Correspondence

If we are extremely lucky; then, our segmentation looks like this

  $\rightsquigarrow \mathbf{x}[1], \mathbf{x}[2], \mathbf{x}[3], \mathbf{x}[4], \mathbf{x}[5]$

and we have a label for each time step




# Computing Loss: One-to-One Correspondence

But, that's too good to happen! Usually we have

  $\rightsquigarrow \mathbf{x}[1], \mathbf{x}[2], \dots, \mathbf{x}[T]$


and we have a label in **some** time steps



# Computing Loss: One-to-One Correspondence

In this typical case, two questions seem non-trivial

- ① Where should we **put each label?**  $\equiv$  Where should we **read each label?**
- ② What should we do with **non-labeled outputs**, e.g.,  $y[1]$ ?



# Computing Loss: One-to-One Correspondence

The key challenge in computing the loss is that we do **not** have necessarily **one-to-one correspondence** with **sequence data**

## Correspondence Problem

With sequence data, we could have a data-sequence of length  $T$  that is labeled by a sequence of size  $K < T$  where

**no time index** is specified for any label in the  $K$ -long label sequence

Correspondence problem exists pretty much in **all practical** sequence data

- In speech recognition, **multiple time frames** correspond to a **single word**
- In text recognition, **multiple image frames** correspond to a **single letter**
- ...

## Correspondence Problem: Formulation

Let's formulate the problem clearly: Say we have

A sequence of data

$$\mathbf{x}[1 : T] = \mathbf{x}[1], \dots, \mathbf{x}[T]$$

that is labeled with the sequence of  $K$  true labels


$$\mathbf{v}[1 : K] = \mathbf{v}[1], \dots, \mathbf{v}[K]$$

where  $K$  and  $T$  can be different

For this setting, we want to train an RNN with this data sequence: starting with an initial state, this RNN returns an output sequence

$$\mathbf{y}[1 : T] = \mathbf{y}[1], \dots, \mathbf{y}[T]$$

# Correspondence Problem: Formulation




To be able to train this RNN, we need to

- ➊ define a loss function that computes  $\hat{R} = \mathcal{L}(\mathbf{y}[1 : T], \mathbf{v}[1 : K])$
- ↳ We need this loss function to be differentiable with respect to all outputs

$$\nabla_{\mathbf{y}[1]} \hat{R}, \dots, \nabla_{\mathbf{y}[T]} \hat{R}$$

# Correspondence Problem: Formulation



To use this RNN after training, i.e., for inferring, we need to

- ② know how to map **outputs** to **predicted labels**

↳ We need to extract  $K$  labels from  $y[1 : T]$ , i.e.,


$$y[1], \dots, y[T] \mapsto \hat{v}[1], \dots, \hat{v}[K]$$

Let's look into different settings

# Setting I: Perfectly Segmented

In some problems, we have our data perfectly segmented


- There is a separate label for each time step, i.e.,  $K = T$ 
  - ↳ many-to-many type I



# Setting I: Perfectly Segmented

In some problems, we have our data perfectly segmented

- There is a separate label for each time step, i.e.,  $K = T$ 
  - ↳ many-to-many type I and one-to-many



## Attention

We can always treat a non-existing input entry as an empty

- ↳ We are good as long as we have a label at each time  $t$

## Setting I: Defining Loss

In such settings, we define the loss to be *aggregated loss* over time

$$\hat{R} = \sum_{t=1}^T \mathcal{L}(\mathbf{y}[t], \mathbf{v}[t])$$

for some loss function  $\mathcal{L}(\cdot, \cdot)$

The gradients are then trivially computed

Gradient with respect to particular output  $\mathbf{y}[t]$  is

$$\nabla_{\mathbf{y}[t]} \hat{R} = \nabla_{\mathbf{y}[t]} \hat{R}[t] = \nabla_{\mathbf{y}[t]} \mathcal{L}(\mathbf{y}[t], \mathbf{v}[t])$$

## Setting I: Inference

Inference in such setting is performed by one-to-one mapping: at time  $t$ , we predict based on  $\mathbf{y}[t]$

$$\mathbf{y}[1] \mapsto \hat{\mathbf{v}}[1], \dots, \mathbf{y}[T] \mapsto \hat{\mathbf{v}}[T]$$

For instance, assume  $\mathbf{y}[t]$  is output of a softmax activation; then, we set

$$\hat{\mathbf{v}}[t] = \operatorname{argmax} \mathbf{y}[t]$$

where  $\operatorname{argmax}$  returns the index of the largest entry, e.g.,

$$\operatorname{argmax} \begin{bmatrix} 0.1 \\ 0.7 \\ 0.2 \\ 0 \end{bmatrix} = 2$$

## Setting II: Known Segments


In some problems, we have only one label for the whole sequence, i.e.,  $K = 1$

↳ It corresponds to **many-to-one** type of problems

↳ This can be that we have really only one label, e.g., content classification

↳ It can be that we know the time index  $t$  at which each label is assigned

↳ We split data to sub-sequences with each sub-sequence having only one label



## Setting II: Defining Loss for Dumb NN

A **naive** approach to define loss is to set it be the loss between **last output** and **label**, i.e.,

$$\hat{R} = \mathcal{L}(\mathbf{y}[T], \mathbf{v}[1])$$


With this loss, gradient with respect to particular output  $\mathbf{y}[t]$  is

$$\nabla_{\mathbf{y}[t]} \hat{R} = \begin{cases} \nabla_{\mathbf{y}[T]} \mathcal{L}(\mathbf{y}[T], \mathbf{v}[1]) & t = T \\ 0 & t \neq T \end{cases}$$

- + But does it make sense to ignore all other outputs?
- Not at all! We are training a **dumb** NN that can respond only when it's over with the whole sequence!

## Setting II: Loss for Smarter Training

An **extremely smart** NN is the one who knows the label **before** the input speaks!



For this NN, the loss is

$$\hat{R} = \sum_{t=1}^T \mathcal{L}(\mathbf{y}[t], \mathbf{v}[1])$$

But, we should be **careful!** We should not expect NN to know everything from potentially **irrelevant** input!

## Setting II: Defining Proper Loss

A **realistic** approach is to *define the loss via a weighted sum*, i.e.,

$$\hat{R} = \sum_{t=1}^T w_t \mathcal{L}(\mathbf{y}[t], \mathbf{v}[1])$$

where  $w_t$  is the weight at time  $t$

- initially  $w_t$  is **small**
  - ↳ we do not expect the NN to know everything from very beginning
- it gradually **increases** up to its maximum  $w_T$ 
  - ↳ by time  $T$  the NN should know the label

With this loss, *gradient with respect to particular output  $\mathbf{y}[t]$  is*

$$\nabla_{\mathbf{y}[t]} \hat{R} = w_t \nabla_{\mathbf{y}[t]} \hat{R}[t] = w_t \nabla_{\mathbf{y}[t]} \mathcal{L}(\mathbf{y}[t], \mathbf{v}[1])$$

## Setting II: Inference

Inference in such setting is performed by *many-to-one mapping*: we only predict based on  $\mathbf{y}[1 : T]$

$$\mathbf{y}[1 : T] \mapsto \hat{\mathbf{v}}[1]$$

For instance, assume  $\mathbf{y}[t]$  is output of a softmax activation; then, we set


$$\tilde{\mathbf{v}}[t] = \operatorname{argmax} \mathbf{y}[t]$$

and then take a (potentially weighted) majority vote:  $\hat{\mathbf{v}}[1]$  is the class that most often estimated with occurrence at each time being weighted by some weight

## Setting III: Unknown Segments

Most common case is that we have a label sequence **shorter** than our **data**

- ↳ Each label in this sequence is corresponding to a **segment of input**
  - ↳ We **do not know** where this **segment** begins and where it ends
    - ↳ There might be even no clear answer to that!




Note that we are dealing with a **sequence to sequence** model: we want to learn relation between **sequence  $x[1 : T]$**  and **sequence  $v[1 : K]$** !

## Setting II: Example

Assume we have image 121 that is divided into a sequence of five pixel vectors


- Since it is a training data, it is labeled as 121
  - ↳ We do not know after which output we should expect RNN to know first, second or third digit!



- + Sounds impossible!
- Only impossible is impossible! Let's carry on and see what we can do!

## Setting II: Genie-Defined Loss

Assume a genie has told us end of each segment




We can fill the *empty labels* with *repetition*, and then define the loss as

$$\hat{R} = \sum_{k=1}^K \sum_{t=i_{k-1}+1}^{i_k} w_t \mathcal{L}(\mathbf{y}[t], \mathbf{v}[k])$$

where  $i_k$  is where label  $\mathbf{v}_k$  ends, e.g., in above diagram  $i_1 = 2$

## Setting II: Defining Loss



We don't have the *genie*: we could assume that  $i_k$  is something to learn!

$$\hat{R}(\mathbf{i}) = \sum_{k=1}^K \sum_{t=i_{k-1}+1}^{i_k} w_t \mathcal{L}(\mathbf{y}[t], \mathbf{v}[k])$$

where  $\mathbf{i} = [0, i_1 \dots, i_K]$  is something we need to learn

## Setting II: Optimal Segmentation

- + How could we learn  $\mathbf{i}$ ? Should we compute also  $\nabla_{\mathbf{i}} \hat{R}$ ?
- Well! You may try! But, obviously  $i_k$  is an **integer!**

### Optimal Segmentation

Optimal approach for finding  $\mathbf{i}$  is to **train** the NN for **all possible choice for  $\mathbf{i}$**  and then find the final training loss  $\hat{R}(\mathbf{i})$ . The **optimal segmentation** is then given by

$$\mathbf{i}^* = \operatorname{argmin}_{\mathbf{i}} \hat{R}(\mathbf{i})$$

- + Is it **computationally feasible**?
- No! The number of **possible choice for  $\mathbf{i}$**  grows **exponentially with  $T$** ! We need to go for sub-optimal approaches

## Setting II: Number of Possible Segmentations

- + How is it exponentially large?
- Let's look at our example

In our example, we should assign **label sequence 121** to a **sequence of length 5**: each entry of **output sequence** in this case can be labeled by **1 (the first one)**, **2** or **1 (the last one)**. This means that we have **3 choices** of label for each time interval; thus, the total number of possible segmentations is around  $3^5$ .

In general number of segmentations grows exponentially with  $T$

- + But wait a moment! We have also counted the case of **labeling all outputs with 1!** This cannot be the case!
- This is right! It is in general **much less than  $3^5$**  but it's still **exponential**


Let's see the exact possible segmentations!

## Setting II: Number of Possible Segmentations


We intend to compare each of  $y[1], \dots, y[5]$  with a label

- We know that the *label sequence* is 121

- ↳ First output is definitely in the *first segment*: its *label* is definitely 1
- ↳ Second output could be *still in the first segment* or in the *second segment*
- ↳ Third output could be in the *first, second, or third segment*
- ↳ Our labels should finish by the end of output sequence: fourth output *cannot* be in *first segment*
- ↳ Last output could be only in the *third segment*




## Setting II: Showing Segmentations on Graph



Though it's exponentially large: we see that each segmentation corresponds to one path on this graph

Blue path corresponds to  $i_1 = 3$ ,  $i_2 = 4$ , and  $i_3 = 5$ , i.e.,  $\mathbf{i} = [0, 3, 4, 5]$

## Setting II: Loss on Segmentation Graph




We can **compute the loss for each segmentation directly on this graph**: let's say that we have  **$L$  different paths** on the graph. For **each path**, we can write an **expanded** label sequence, e.g.,

**Expanded** label sequence of **blue path** is { 1, 1, 1, 2, 1 }

This sequence is of length  $T$  and we show it for **path  $\ell$**  with  $\tilde{v}_\ell[t]$

## Setting II: Loss on Segmentation Graph



For each path  $\ell = 1, \dots, L$ , the loss is computed by aggregating the losses between outputs and extended labels

$$\hat{R}_\ell = \sum_{t=1}^T w_t \mathcal{L}(\mathbf{y}[t], \tilde{\mathbf{v}}_\ell[t]) = \sum_{t=1}^T \hat{R}_\ell[t]$$

It again decomposes into sum of  $T$  terms with only one being function of  $\mathbf{y}[t]$

## Setting II: Optimal Segmentation on Graph

We can represent the optimal segmentation on the graph as below

OptimalSegmentTraining():

- 1: Initiate with  $\hat{R} = +\infty$  and some random  $\ell^* = \emptyset$
- 2: **for**  $\ell = 1, \dots, L$  **do**
- 3:   Let the loss be  $\hat{R}_\ell$
- 4:   Train for sufficient epochs
- 5:   **if** After training  $\hat{R}_\ell < \hat{R}$  **then**
- 6:      $\hat{R} \leftarrow \hat{R}_\ell$  and  $\ell^* \leftarrow \ell$
- 7:   **end if**
- 8: **end for**
- 9: Return learnable parameters and  $\ell^*$


- + Say we could be over with this **infeasible** training! How do we use the trained RNN for inference?
- In this case, we have  $\ell^*$  which gives us optimal segmentation: we infer label of each segment based on its corresponding outputs

## Setting II: Maximum-Likelihood Segmentation

Since optimal segmentation is **infeasible**, people uses **maximum-likelihood approach** that is well-known *in detection and coding theory*

### Maximum-Likelihood Segmentation

Start with an **initial guess** for **optimal path** on segmentation graph and do one step of training; then, **improve the guess** based on the outputs of next forward pass and go for **next step of training**



Let's look at its pseudo-code

## Setting II: Maximum-Likelihood Segmentation

MaxLikelihoodTraining():

- 1: **for** Iteration  $i = 1, \dots, I$  **do**
- 2:   Pass forward through time: Compute **output sequence**  $\mathbf{y}[1 : T]$
- 3:   Compute  $p(\tilde{\mathbf{v}}_\ell[1 : T] | \ell)$  for each **path**  $\ell$  on segmentation graph
- 4:   Update  $\ell^* = \operatorname{argmax}_\ell p(\tilde{\mathbf{v}}_\ell[1 : T] | \ell)$
- 5:   Set loss to  $\hat{R}_{\ell^*}$  and backpropagate over RNN
- 6:   Update **learnable parameters**
- 7: **end for**
- 8: Return **learnable parameters** and  $\ell^*$


⊗

⊗

- + Why we call it **maximum likelihood**?
- Because we guess **path** by **maximizing** the **likelihood**  $p(\tilde{\mathbf{v}}_\ell[1 : T] | \ell)$
- + But how can find **likelihood** of a **path**?
- We can use **output sequence**  $\mathbf{y}[1 : T]$

## Setting II: Finding Likelihood on Segmentation Graph

$$\begin{bmatrix} y_1[1] \\ y_2[1] \\ y_3[1] \\ y_4[1] \end{bmatrix} \quad \begin{bmatrix} y_1[2] \\ y_2[2] \\ y_3[2] \\ y_4[2] \end{bmatrix} \quad \begin{bmatrix} y_1[3] \\ y_2[3] \\ y_3[3] \\ y_4[3] \end{bmatrix} \quad \begin{bmatrix} y_1[4] \\ y_2[4] \\ y_3[4] \\ y_4[4] \end{bmatrix} \quad \begin{bmatrix} y_1[5] \\ y_2[5] \\ y_3[5] \\ y_4[5] \end{bmatrix}$$




Assume that each label could be 1, 2, 3, or 4: *at each time  $t$  the RNN returns a 4-dimensional vector whose entries are probability of each class*

*we can multiply the probabilities of classes on the path*

## Setting II: Finding Likelihood on Segmentation Graph


$$\begin{bmatrix} y_1[1] \\ y_2[1] \\ y_3[1] \\ y_4[1] \end{bmatrix} \quad \begin{bmatrix} y_1[2] \\ y_2[2] \\ y_3[2] \\ y_4[2] \end{bmatrix} \quad \begin{bmatrix} y_1[3] \\ y_2[3] \\ y_3[3] \\ y_4[3] \end{bmatrix} \quad \begin{bmatrix} y_1[4] \\ y_2[4] \\ y_3[4] \\ y_4[4] \end{bmatrix} \quad \begin{bmatrix} y_1[5] \\ y_2[5] \\ y_3[5] \\ y_4[5] \end{bmatrix}$$



For instance, the **yellow path** has a likelihood

$$p(\tilde{\mathbf{v}}_{\ell}[1:T]|\ell) = \prod_{t=1}^T p(\tilde{\mathbf{v}}_{\ell}[t]|\ell) = y_1[1]y_1[2]y_1[3]y_2[4]y_1[5]$$

## Setting II: Finding Likelihood on Segmentation Graph



Or better to say: we just put *output entries* in graph and move on the *path*

$$p(\tilde{\mathbf{v}}_\ell[1 : T] | \ell) = \prod_{t=1}^T y_{\tilde{v}_\ell[t]}[t] = y_1[1]y_1[2]y_1[3]y_2[4]y_1[5]$$

## Setting II: Maximum-Likelihood Segmentation

- + OK! We can find the likelihood, but how can we maximize it? It's again an exponentially large search!

$$\ell^* = \operatorname{argmax}_{\ell} p(\tilde{\mathbf{v}}_{\ell}[1:T] | \ell)$$

- Well! If we only need the maximum, it turns not to be exponential

---

We can readily show that finding maximum likelihood on the graph is a dynamic programming problem and can be solved by the Viterbi algorithm

Maximum likelihood training can be implemented efficiently

## Setting II: Maximum-Likelihood Inference

MaxLikelihoodTraining():

- 1: **for** Iteration  $i = 1, \dots, I$  **do**
- 2:   Pass forward through time: Compute **output sequence**  $\mathbf{y}[1 : T]$
- 3:   Compute  $p(\tilde{\mathbf{v}}_\ell[1 : T] | \ell)$  for each **path**  $\ell$  on segmentation graph
- 4:   Update  $\ell^* = \operatorname{argmax}_\ell p(\tilde{\mathbf{v}}_\ell[1 : T] | \ell)$
- 5:   Set loss to  $\hat{R}_{\ell^*}$  and backpropagate over RNN
- 6:   Update **learnable parameters**
- 7: **end for**
- 8: Return **learnable parameters** and  $\ell^*$



- + How can we use our RNN for inference after training via **maximum likelihood** segmentation?
- We have access to  $\ell^*$ : we **predict the label** of each segment based on its **corresponding outputs**

## Setting II: Connectionist Temporal Classification

It turns out that *maximum-likelihood* could stick to a bad local minimum, i.e.,  
it quickly converges to a *path*  $\ell^*$  that is *much different* from  $\ell^*$

- + Is there any solution to this?
- Yes! We can use connectionist temporal classification (CTC) loss

### CTC Loss

Instead of searching for a best segmentation and then minimizing its loss, we learn directly from *unsegmented data* by minimizing the average loss over *all possible segmentations*, i.e., we define loss to be

$$\hat{R} = \mathbb{E}_{\ell} \left\{ \hat{R}_{\ell} \right\} = \sum_{\ell=1}^{L} p(\ell | \tilde{\mathbf{v}}_{\ell}[1 : T]) \hat{R}_{\ell}$$

and train the RNN by finding *learnable parameters* that minimize this loss

## Setting II: CTC Loss

- + But, why should it be a **better choice** of loss?
- Because we are **sure** that **optimal segmentation** is **contributing** to our loss

$$\begin{aligned}\hat{R} &= \mathbb{E}_{\ell} \left\{ \hat{R}_{\ell} \right\} = \sum_{\ell=1}^L p(\ell | \tilde{\mathbf{v}}_{\ell}[1 : T]) \hat{R}_{\ell} \\ &= p(\ell^* | \tilde{\mathbf{v}}_{\ell^*}[1 : T]) \hat{R}_{\ell^*} + \sum_{\ell \neq \ell^*} p(\ell | \tilde{\mathbf{v}}_{\ell}[1 : T]) \hat{R}_{\ell}\end{aligned}$$

- + Agreed! Now, how should we determine  $p(\ell | \tilde{\mathbf{v}}_{\ell}[1 : T])$  ?
- Just use the Bayes rule!
- + What about the expectation? It is at the end sum of **exponentially** large number of terms!
- We can again go on the **graph** and determine it via **dynamic programming**

## Setting II: CTC Loss

The CTC loss can be written as

$$\begin{aligned}\hat{R} &= \sum_{\ell=1}^L p(\ell | \tilde{\mathbf{v}}_\ell[1 : T]) \hat{R}_\ell = \sum_{\ell=1}^L p(\ell | \tilde{\mathbf{v}}_\ell[1 : T]) \sum_{t=1}^T w_t \mathcal{L}(\mathbf{y}[t], \tilde{\mathbf{v}}_\ell[t]) \\ &= \sum_{t=1}^T w_t \underbrace{\sum_{\ell=1}^L p(\ell | \tilde{\mathbf{v}}_\ell[1 : T]) \mathcal{L}(\mathbf{y}[t], \tilde{\mathbf{v}}_\ell[t])}_{\check{R}[t]} = \sum_{t=1}^T w_t \check{R}[t]\end{aligned}$$

This has been shown that  $\check{R}[t]$  can be *recursively computed*<sup>1</sup>:

by some *approximation* we are able to readily compute  $\nabla_{\mathbf{y}[t]} \check{R}[t]$

and we set  $\nabla_{\mathbf{y}[t']} \check{R}[t] \approx \mathbf{0}$  for  $t' \neq t$

---

<sup>1</sup>Check out the [original paper](#)

## Setting II: Training with CTC Loss

CTC\_Training() :

- 1: **for** iteration  $i = 1, \dots, I$  **do**
- 2:   Pass forward through time: Compute **output sequence**  $\mathbf{y}[1 : T]$
- 3:   Compute CTC loss  $\hat{R}$  and  $\nabla_{\mathbf{y}[t]} \hat{R}$  by **recursion**
- 4:   Backpropagate through time and update **learnable parameters**
- 5: **end for**
- 6: Return **learnable parameters**

This looks like **standard training loop** now

the **loss** is only replaced with **CTC loss**


- + What about inference?
- Well! We should figure it out, since the training loop **does not** compute any **segmentation path!**

## Setting II: Inference with CTC-Trained RNN

Let's get back to our simple example: assume that after training with **CTC loss** we give an image of handwritten  $121$  to the **RNN**

- RNN divides it into **5 frames** and is able to track optimal **segmentation**
  - ↳ The first three frames belong to the first segment
  - ↳ The remaining frames belong to the second and third segments
- RNN infers from output sequence  $\hat{v}[1 : 5]$  but does **not** return optimal path

$$\hat{v}[1] = 1 \quad \hat{v}[2] = 1 \quad \hat{v}[3] = 1 \quad \hat{v}[4] = 2 \quad \hat{v}[5] = 1$$



## Setting II: Inference with CTC-Trained RNN

We can conclude from  $\hat{v}[1 : 5]$  that the sequence is  $\{1, 2, 1\}$  if we are sure the sequence has no repetition

### Label Encoding and Decoding

CTC uses this fact and constructs following encoding and decoding method: it introduces a new label called “blank:-” which does not belong to set of classes

- While training, it adds blank between any two repetitions
  - ↳ For instance, we encode  $112 \mapsto 1\text{-}12$ , or  $111 \mapsto 1\text{-}1\text{-}1$
- For inference, it removes any repetition in inferred sequence  $\hat{v}[1 : T]$  and then drops blanks
  - ↳ For instance, we decode  $1\text{-}11\text{-}312 \mapsto 11312$ , or  $3333\text{-}3121 \mapsto 33121$

## Setting II: Training and Inference with CTC

CTC\_Training() :

- 1: **for** iteration  $i = 1, \dots, I$  **do**
- 2:   Add **blanks** to the label sequences with repetition
- 3:   Pass forward through time: Compute **output sequence**  $\mathbf{y}[1 : T]$
- 4:   Compute CTC loss  $\hat{R}$  and  $\nabla_{\mathbf{y}[t]} \hat{R}$  by **recursion**
- 5:   Backpropagate through time and update **learnable parameters**
- 6: **end for**
- 7: Return **learnable parameters**

CTC\_Inference() :

- 1: Pass forward through time the input and compute **output**  $\mathbf{y}[1 : T]$
- 2: Infer encoded sequence  $\hat{\mathbf{v}}[1 : T]$  from  $\mathbf{y}[1 : T]$
- 3: Remove repetitions from  $\hat{\mathbf{v}}[1 : T] \mapsto \hat{\mathbf{v}}[1 : T']$
- 4: Remove **blanks** from  $\hat{\mathbf{v}}[1 : T'] \mapsto \hat{\mathbf{v}}[1 : K]$
- 5: Return  $\hat{\mathbf{v}}[1 : K]$

# In PyTorch: CTC Loss

We can access CTC loss in `torch.nn` module as

```
torcn.nn.CTCLoss()
```

Few notes about CTC loss implementation

- We need to specify the index of **blank** label
  - ↳ It should be **out of our set of classes**
  - ↳ By default, it is set to `blank = 0`
- When we define our model, we should always take **blank** label into account
  - ↳ If we do classification with  $C$  **classes**, model should return  $C + 1$  **classes** with blank being one of them
- PyTorch considers cross-entropy loss function, i.e.,  $\mathcal{L}(\mathbf{y}, \tilde{\mathbf{v}}) = \text{CE}(\mathbf{y}, \tilde{\mathbf{v}})$
- As input to CTC loss: **y** should be **logarithm of probabilities**
  - ↳ We can activate the output layer with **logarithmic softmax**