

Applied Deep Learning

Chapter 7: Sequence-to-Sequence Models

Ali Bereyhi

ali.bereyhi@utoronto.ca

Department of Electrical and Computer Engineering
University of Toronto

Fall 2025

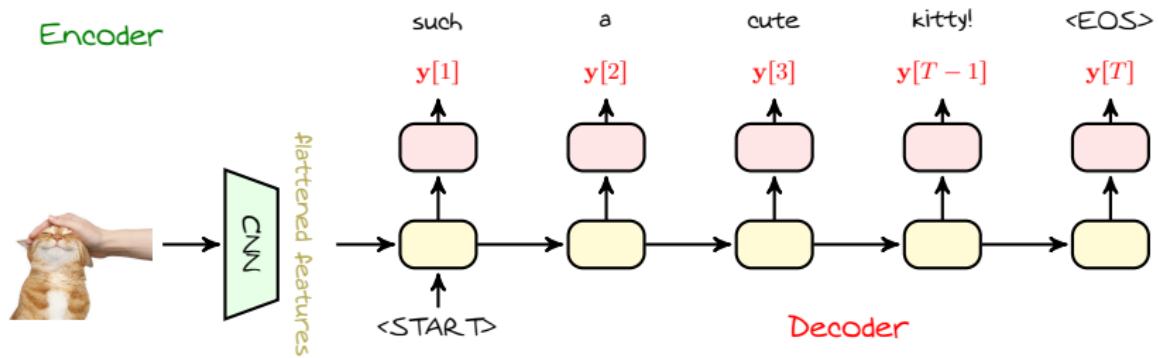
Encoder-Decoder Architecture

Encoder-Decoder

Encoder-decoder architecture comprises of two separate NNs

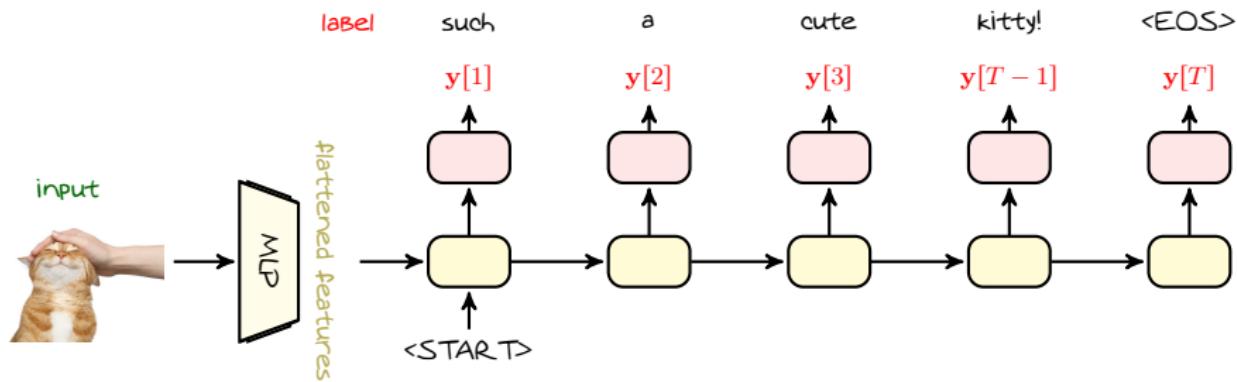
- ① **Encoder** takes the **input sequence** and encodes it into vector of features
- ② **Decoder** takes vector of features and decodes it into **output sequence**

- + What kind of NNs should we use?
- Pretty much everything is allowed!



Back to Caption Generation

We used a CNN for *encoding* and an RNN for *decoding*

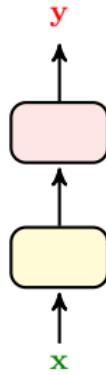


We can replace CNN with any other architecture the extracts features

RNN as Conditional Distribution

- + But, why should those features make RNN speak about cat?
- It makes RNN to generate random words *conditional* to input image

When we are dealing with classification: NN can be seen as a machine that computes *distribution* and based on its *input*, it generates *random outputs*

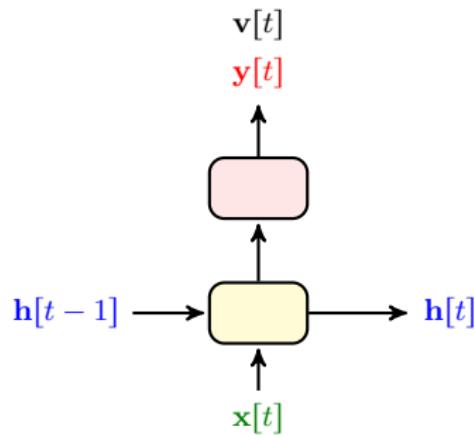


In classification \mathbf{y} is a vector of probability

- Its length equals to the number of classes
 - Its entry k represents the probability of class k
- ↳ We can say that

$$\mathbf{y} = \begin{bmatrix} y_1 \\ \dots \\ y_K \end{bmatrix} \iff y_k \propto \Pr \{\text{label} = k | \mathbf{x}\}$$

RNN as Conditional Distribution



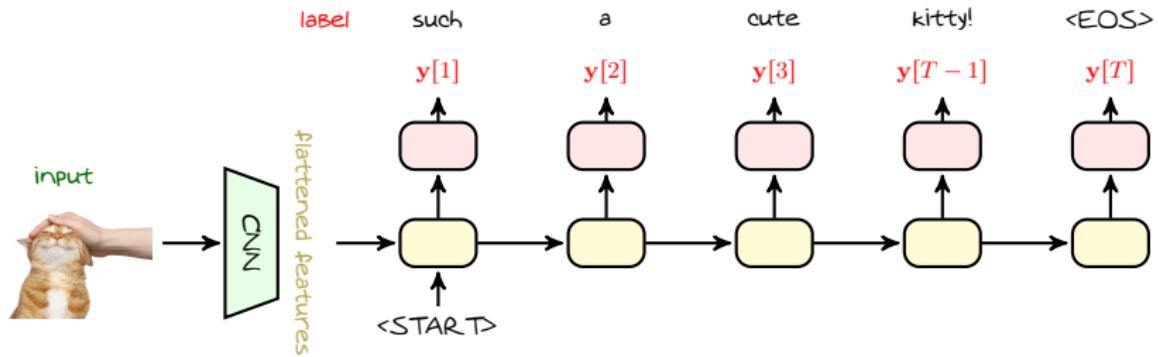
Similarly the output of RNN in each time can be seen as

$$y_k[t] \propto \Pr \{\text{label} = k | \mathbf{h}[t-1], \mathbf{x}[t]\} = p(\mathbf{v}[t] | \mathbf{h}[t-1], \mathbf{x}[t])$$

Since $\mathbf{h}[t-1]$ already contains memory about $\mathbf{x}[1 : t-1]$, we could say

$$y_k[t] \propto p(\mathbf{v}[t] | \mathbf{x}[1 : t])$$

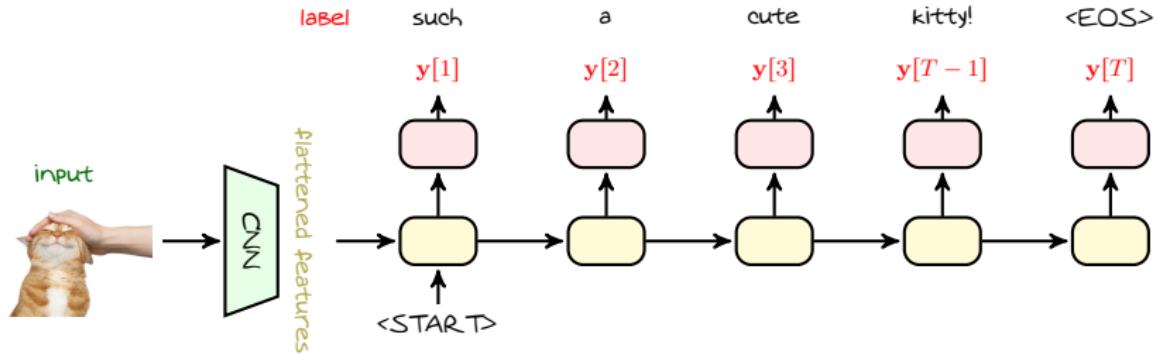
Caption Generation: Dataset



To train this architecture, we collect a dataset

- It contains several **images**
 - ↳ They could potentially be of different classes
- For each image, we have a **sample caption**
 - ↳ These sentences are again of different lengths

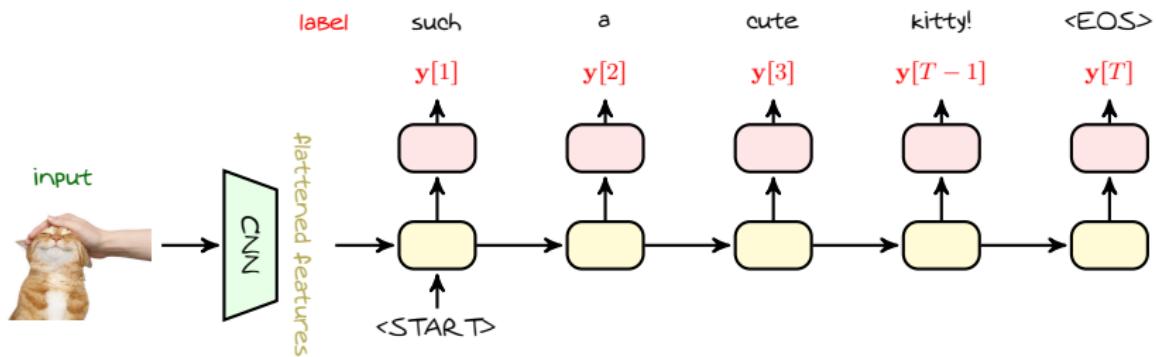
Caption Generation: Training



Say we want to train it on one sample: first we pass forward

- We first **tokenize the words** in captions to take them as one-hot labels
- We pass the image forward through CNN and get the **feature vector**
- We initiate the RNN with the feature vector and give input **<START>**
- We pass forward through time till we see have the **output sequence**

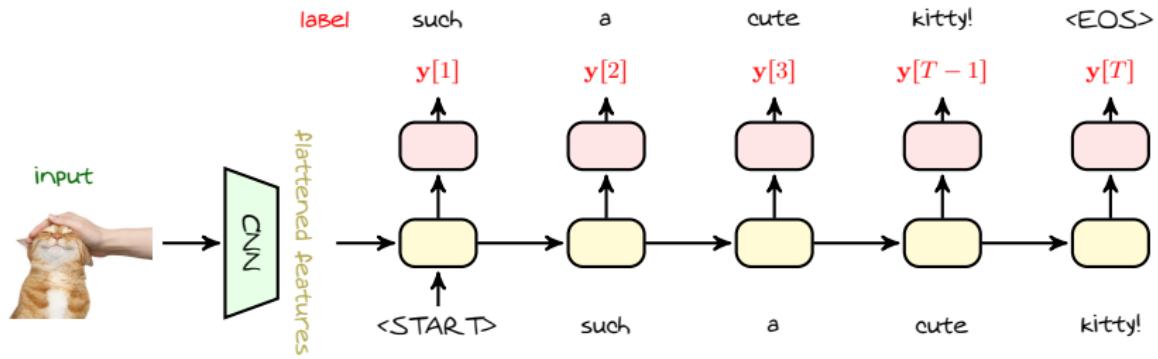
Caption Generation: Training



Say we want to train it on one sample: then we pass backward

- We compute the **loss** between the output sequence and one-hot labels
 - ↳ We can simply use the cross-entropy function
- We **backpropagate** through time till we arrive at the **beginning of decoder**
- We have $\nabla_{\text{features}} \hat{R}$
 - ↳ So we **backpropagate** through the CNN
- We update all weights and go for the next round

Caption Generation: Inference



Say we finished with training: we want to caption a *new image*

- We send it over network and read the *output sequence*
- We could also set output of each time step as *input* for next time
 - ↳ We could also do it while training
 - ↳ Intuitively, it could help the RNN writing *more coherent sentence*

S_{eq}2S_{eq} Model: Basic Translator

Now, let's take a step further:

we want to build a model that translates German sentences to English

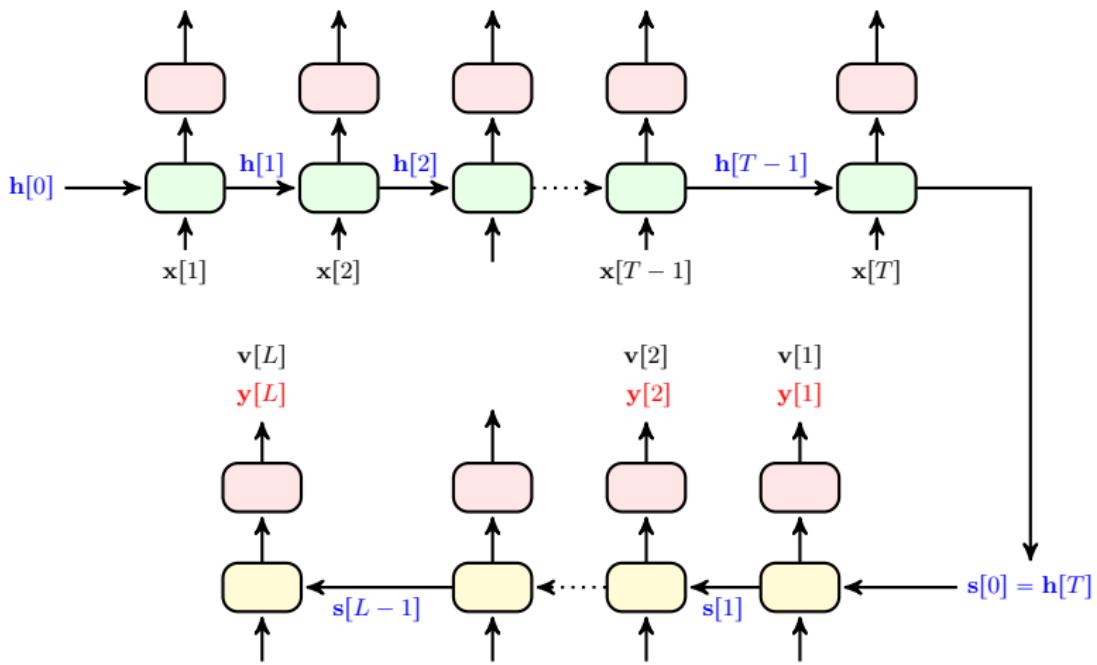
- We need a Seq2Seq model
 - ↳ We have a sequence of input German words
 - ↳ We need to return a sequence of English words
 - ↳ These sequences could be of different lengths

We know encoder-decoder model: we use it to build our translator

- We need an encoder that takes a German sentence
 - ↳ RNN is a good choice, since we have input sequence
- We need a decoder that returns the English translation
 - ↳ RNN is again the choice, since we have another sequence

Basic Translator: Encoder-Decoder Model

So, the model for our translator looks like this

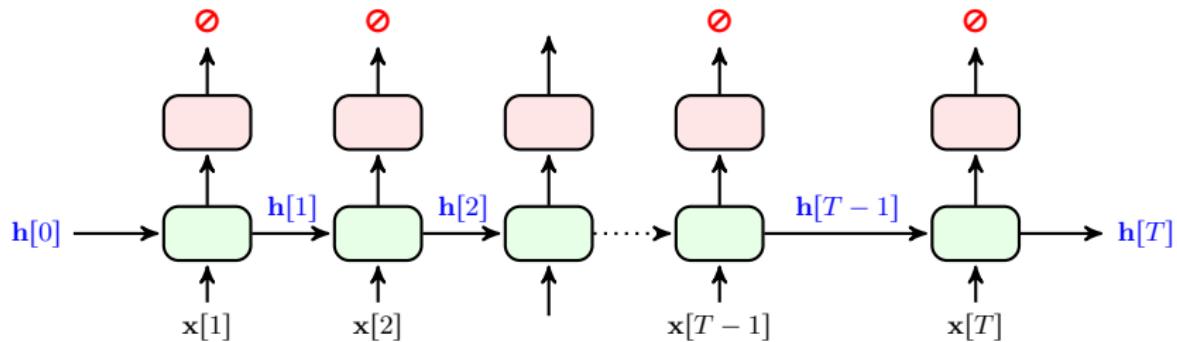


Basic Translator: Dataset

To train this model, we collect some dataset: in this dataset

- We have German sentences

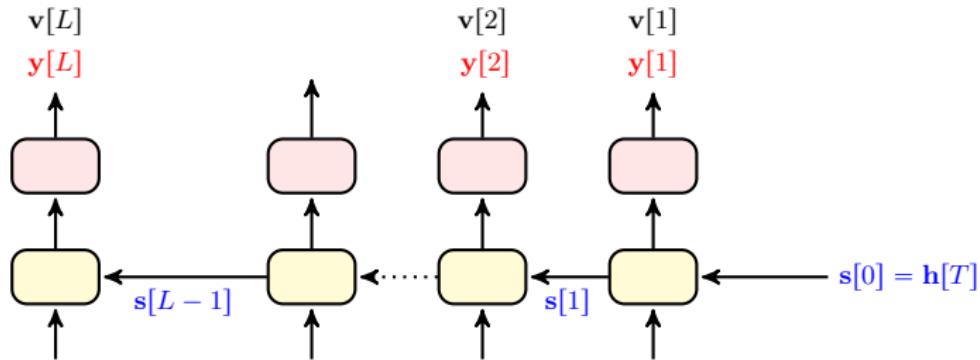
↳ We tokenize each sentence and represent it with a sequence $x[1 : T]$



Basic Translator: Dataset

To train this model, we collect some dataset: in this dataset

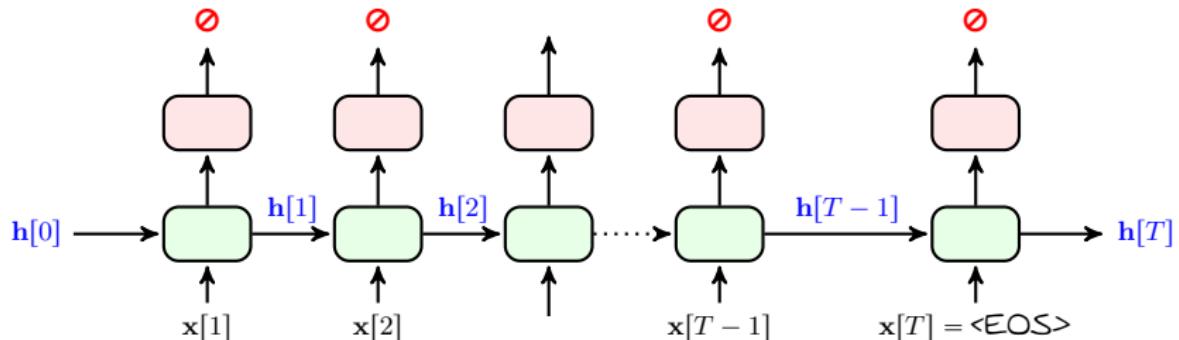
- Corresponding to each German sentence, we have the English translation
 - ↳ We tokenize it as well and represent it with a sequence $\mathbf{v}[1 : L]$
 - ↳ L and T are **not** of the same length



Basic Translator: Training

Say we want to train it for sample sentence: we start with forward pass

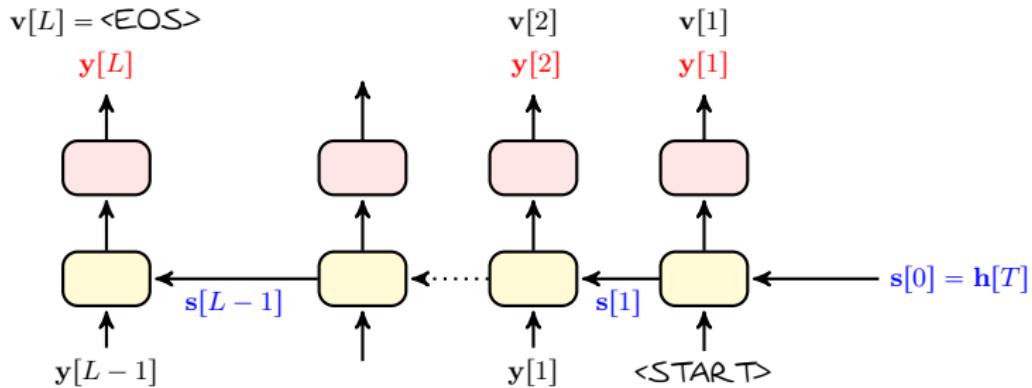
- We pass the tokens through time forward till we arrive at <EOS>
 - ↳ At this point we have $\mathbf{h}[T]$ at the output of encoder
- We have already computed all variables inside this encoder
 - ↳ We need them in them **backward pass**



Basic Translator: Training

Say we want to train it for sample sentence: we start with forward pass

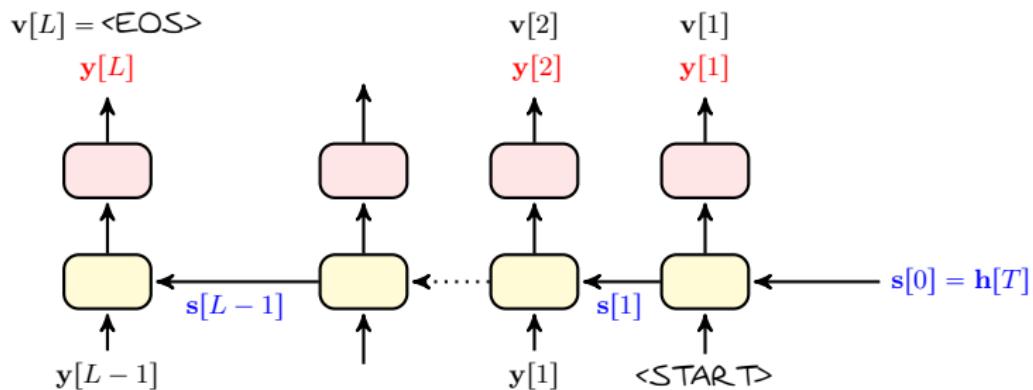
- We initiate the decoder with $s[0] = h[T]$
 - ↳ We could also give token of $\langle \text{START} \rangle$ as first input
 - ↳ We can give $y[\ell - 1]$ as the input at time ℓ
- We continue till we get to label $\langle \text{EOS} \rangle$



Basic Translator: Training

Say we want to train it for sample sentence: now we pass *backward*

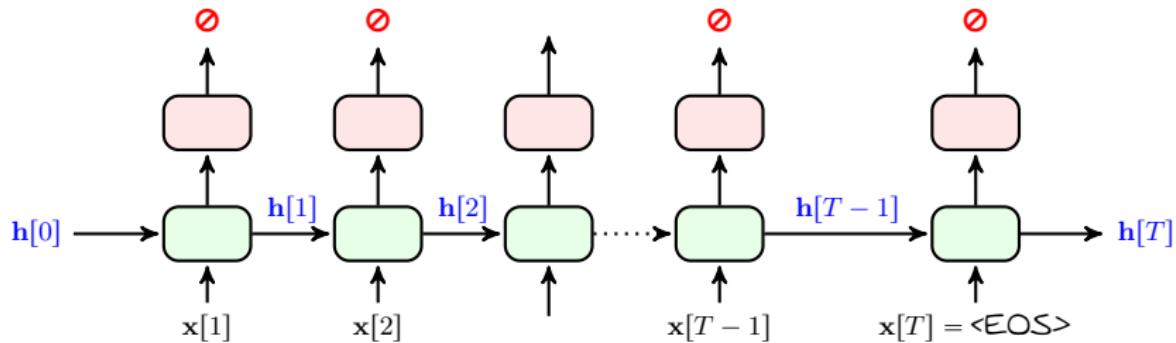
- We compute *loss* between the labels and outputs
 - ↳ We aggregate cross-entropy losses over time
- We *backpropagate* through time
 - ↳ We get to $\nabla_{\mathbf{s}[0]} \hat{R} = \nabla_{\mathbf{h}[T]} \hat{R}$



Basic Translator: Training

Say we want to train it for sample sentence: now we pass *backward*

- We have $\nabla_{\mathbf{h}[T]} \hat{R}$
 - ↳ We *backpropagate* again over time
- We update all the weights and start a new round

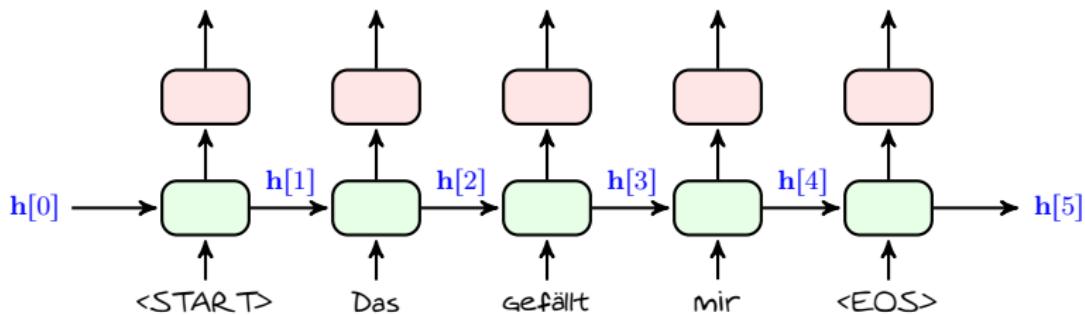


Basic Translator: Inference

Say we have trained this model and we want to use it to translate

“Das gefällt mir” \rightsquigarrow I like it

Let's start with encoding

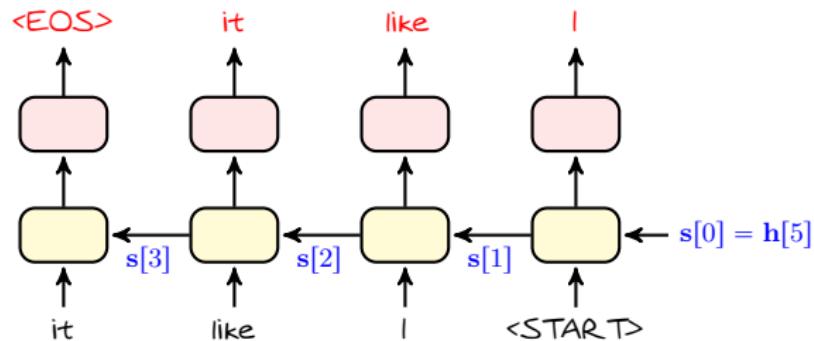


Basic Translator: Inference

Say we have trained this model and we want to use it to translate

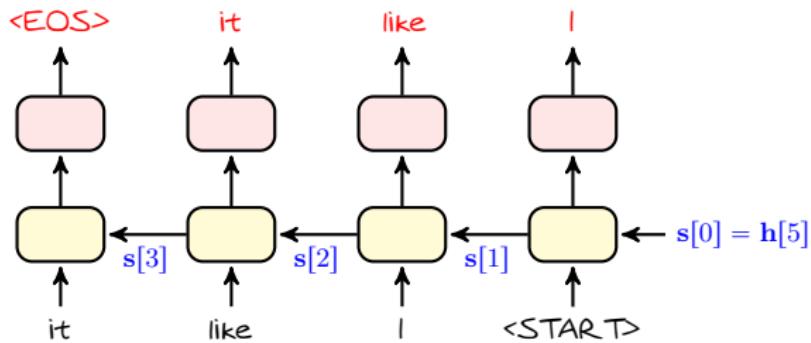
“Das gefällt mir” \rightsquigarrow I like it

Now, for **decoding** we start with $h[5]$



Basic Translator: Inference

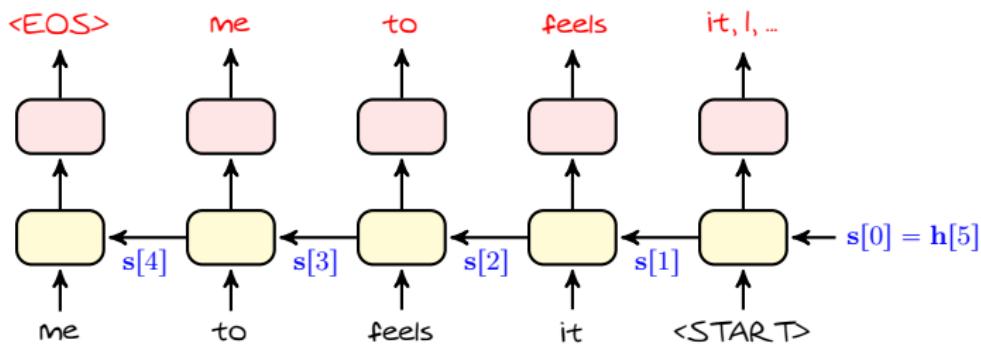
If we are lucky, the token of word "**I**" has highest probability in $\ell = 1$



and probably the RNN keeps generating a correct sentence

Basic Translator: Inference

If we are unlucky, another token might be slightly higher in probability at $\ell = 1$

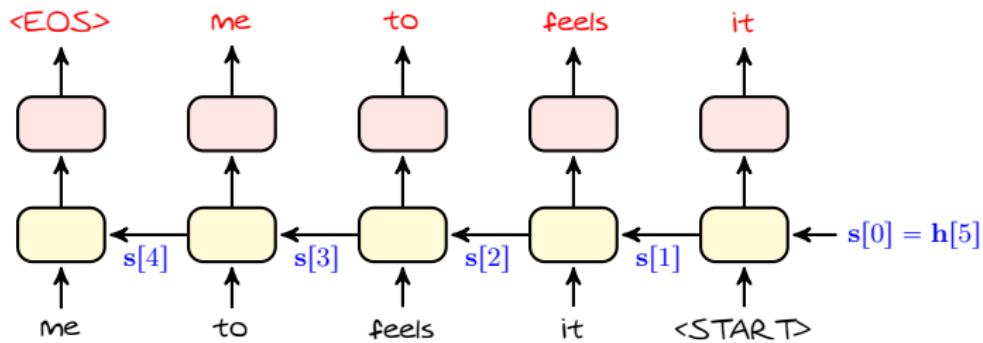


In this case, a small mistake can lead to a sequence of mistakes, e.g., say

- word "it" has slightly higher chance at the beginning
 - ↳ e.g., "it": 0.121 and "I": 0.119, thus, we choose it as the first word
- since we chose "it", RNN needs to generate the next word accordingly
 - ↳ with "it" as input: "feels" has higher chances than "like"

Basic Translator: Inference via Decoding

If we are unlucky, another token might be slightly higher in probability at $\ell = 1$



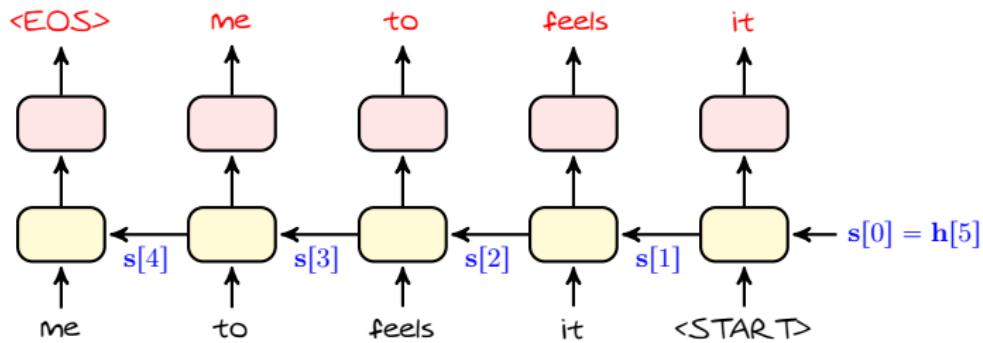
Decoding

To avoid error propagation, a better idea is to find the **sequence with highest probability**, i.e., sequence $\mathbf{v}^*[1 : L]$ that has highest conditional probability

$$\mathbf{v}^*[1 : L] = \operatorname{argmax}_{\mathbf{v}[1:L]} p(\mathbf{v}[1 : L] | \mathbf{x}[1 : T])$$

Basic Translator: Inference via Decoding

If we are unlucky, another token might be slightly higher in probability at $\ell = 1$



Intuitively, this means: we do not classify in each time

- We wait till the sentence is over
- We consider all possible combinations
- We find the one which has highest conditional probability
 - ↳ We need to compute this conditional probability

Basic Translator: Inference via Decoding

Let's try finding this probability first

$$\begin{aligned}
 p(\mathbf{v}[1:L]|\mathbf{x}[1:T]) &= p(\mathbf{v}[1:L]|\mathbf{h}[T]) \\
 &= \prod_{\ell=1}^L p(\mathbf{v}[\ell]|\mathbf{h}[T], \mathbf{v}[1:\ell-1]) \rightsquigarrow \text{Bayes chain rule} \\
 &\propto \prod_{\ell=1}^L p(\mathbf{v}[\ell]|\mathbf{s}[\ell-1], \mathbf{y}[\ell-1]) \propto \prod_{\ell=1}^L y_{\mathbf{v}[\ell]}[\ell]
 \end{aligned}$$

$y_{\mathbf{v}[\ell]}[\ell]$ means the entry of $\mathbf{y}[\ell]$ that corresponds to the class of $\mathbf{v}[\ell]$, e.g.,

$$\mathbf{v}[\ell] = \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix} \quad \text{and} \quad \mathbf{y}[\ell] = \begin{bmatrix} 0.1 \\ 0.3 \\ 0.4 \\ 0.2 \end{bmatrix} \rightsquigarrow y_{\mathbf{v}[\ell]}[\ell] = 0.3$$

Basic Translator: Inference via Decoding

Since we are more comfortable with sum, we can use log-likelihood

$$\begin{aligned}\mathbf{v}^*[1:L] &= \operatorname{argmax}_{\mathbf{v}[1:L]} \log p(\mathbf{v}[1:L] | \mathbf{x}[1:T]) \\ &= \operatorname{argmax}_{\mathbf{v}[1:L]} \sum_{\ell=1}^L \log y_{\mathbf{v}[\ell]}[\ell]\end{aligned}$$

- + But, is it feasible to find the sequence with highest sum?
- No! Since we deal here with an exponentially large case

For a vocabulary with D words in it, the number of possible sequences is

$$(D + 2)^L$$

Basic Translator: Inference via Decoding

In practice, however, we know that many of those combinations are **invalid**, e.g.,

“I are potato” is not an **invalid** English **sentence**!

So we can use sub-optimal search methods to find a good sequence

- We do **not** necessarily hit the best sequence
 - ↳ But, for fair length, we can be close
 - ↳ We will be definitely better than **naive instant decoding**
- Most famous approach is **beam search** via **top- k**
 - ↳ At time step ℓ we find k sequences with highest sum **log-likelihoods** till ℓ
 - ↳ We update **top- k sequences** by searching among children on sequence tree
 - ↳ We finally select the sequence with highest probability among those **top- k**
- Since decoded sequences can be of different length, we usually maximize **normalized log-likelihood**

$$\mathbf{v}^*[1:L] = \operatorname{argmax}_{\mathbf{v}[1:L]} \frac{1}{L} \sum_{\ell=1}^L \log y_{v[\ell]}[\ell]$$