

# Applied Deep Learning

## Chapter 7: Sequence-to-Sequence Models

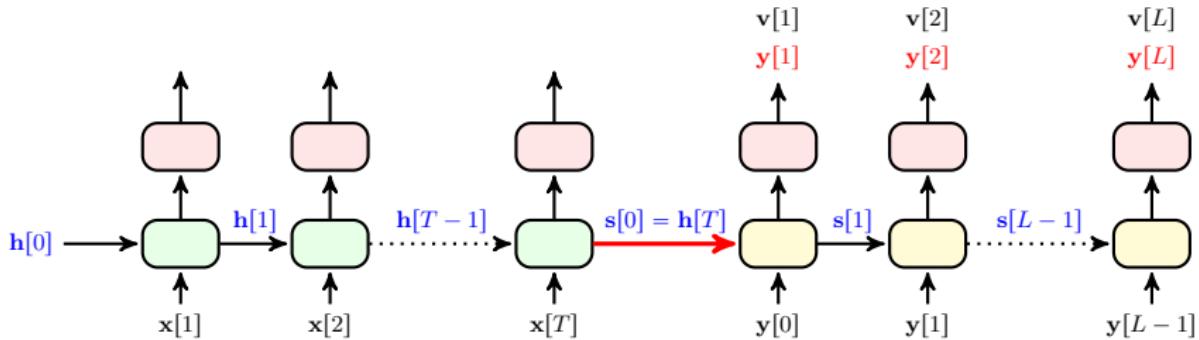
Ali Bereyhi

[ali.bereyhi@utoronto.ca](mailto:ali.bereyhi@utoronto.ca)

Department of Electrical and Computer Engineering  
University of Toronto

Fall 2025

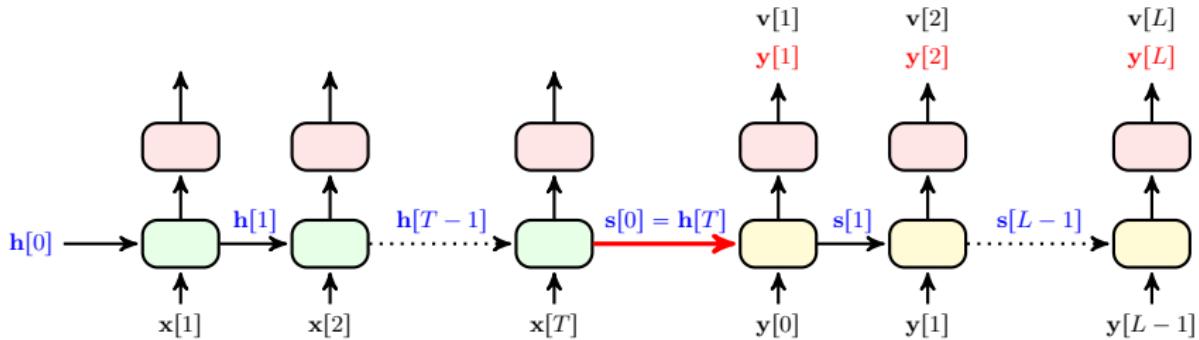
# Translating Long Texts



With standard RNN encoder and decoder: *model works up to some length*

- *The decoder could get lost at some point*
  - ↳ *It could miss the case of the word or its order*
  - ↳ *Imagine it wants to translate:*  
“Ich habe **den** Apfel genommen, über **den** wir beim letzten Mal gesprochen haben” ↵ “I took **the** apple, about **what** we talked about last time”

# Information Bottleneck



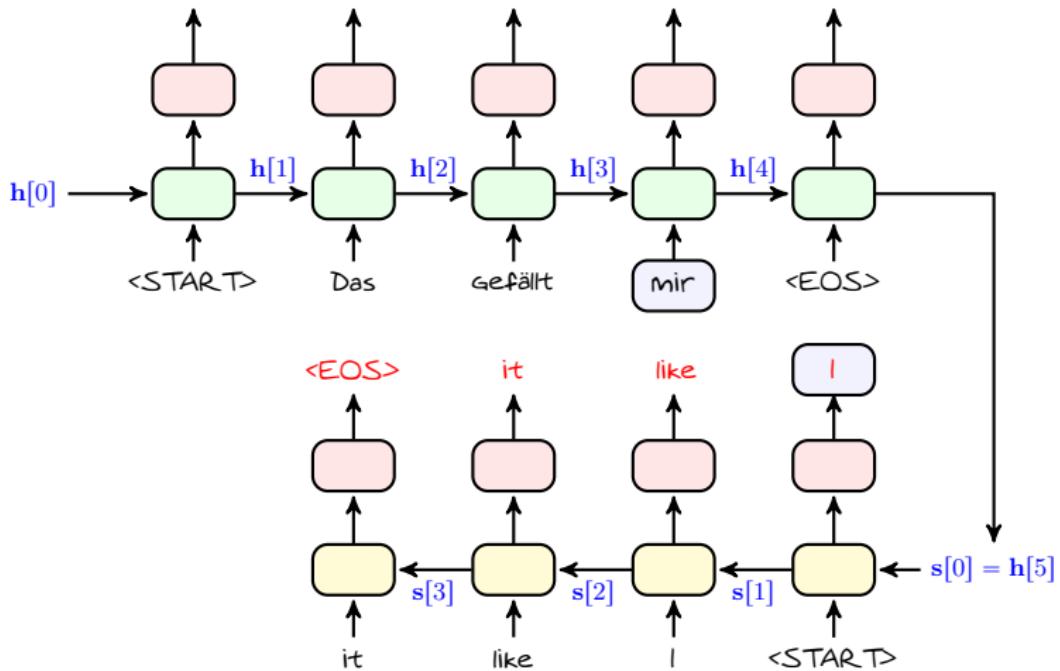
The source of this problem is that

*decoder gets all its information through a single bottleneck*

*This problem is known as information bottleneck problem*

# Attention: Finding Relevant Input

Let's look back at our translator: "**I**" is given in translation because of "**mir**"



## Attention: Finding Relevant Input

What if we could tell the **decoder**:

use **hidden state of time  $t = 4$**  to generate its **output word at time  $t = 1$**

We could intuitively say that in this case

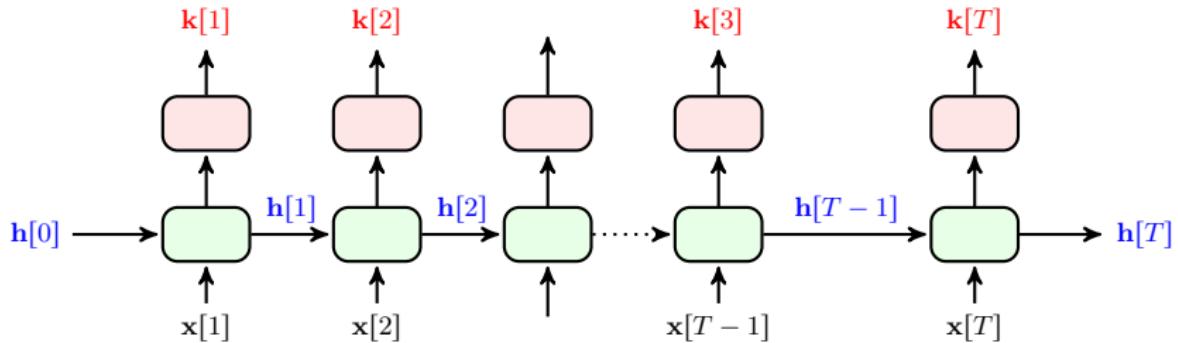
$$\mathbf{y}[1] = f(\mathbf{s}[0], \langle \text{START} \rangle, \text{mir})$$

which is more likely to be "**I**" as compared to the case in which

$$\mathbf{y}[1] = f(\mathbf{s}[0], \langle \text{START} \rangle)$$

**Attention mechanism** formulates mathematically this intuitive idea

# Attention: Generating Keys



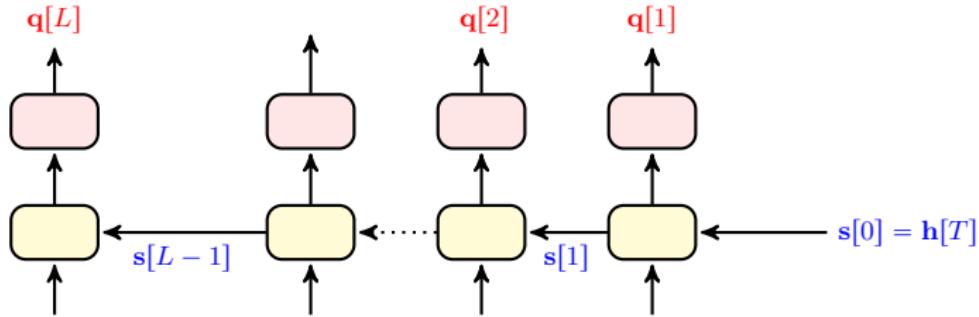
With **attention**, we generate a **key** for each time step at **encoding**

- Keys are learned by an arbitrary layer that is **fixed over time**, e.g.,

$$\mathbf{k}[t] = \sigma(\mathbf{W}_k \mathbf{h}[t])$$

- We can even use multiple layers
  - ↳ Not really needed in most applications

# Attention: Generating Queries



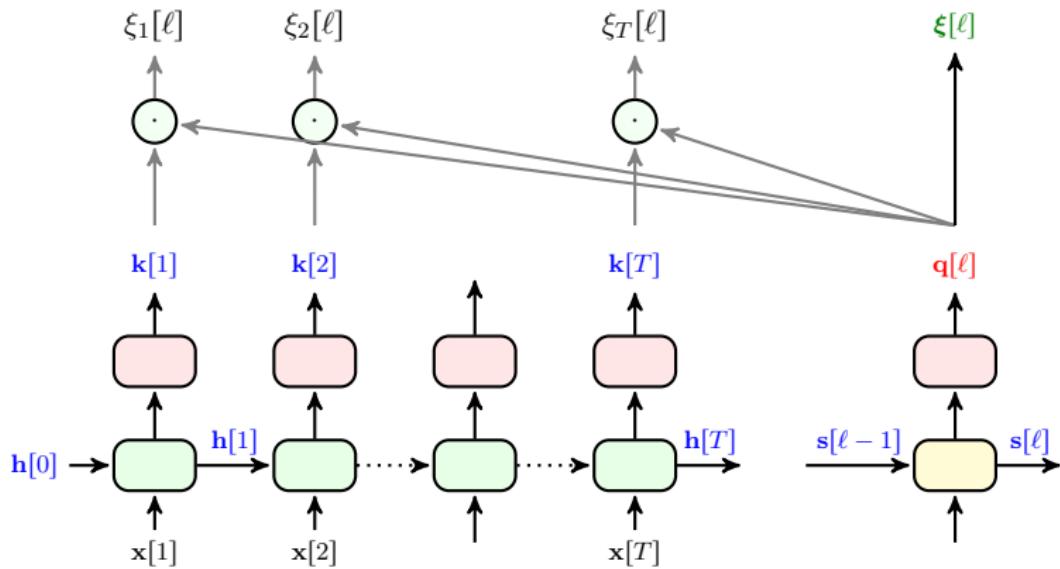
We next generate a **query** for each time step at **decoding**

- **Queries** are again learned by an **arbitrary layer**, e.g.,
- $$\mathbf{q}[t] = \sigma(\mathbf{W}_q \mathbf{s}[t])$$
- It's in general a **new learnable layer** different from the **key generator**

# Attention: Score at Time $\ell$

At *decoder*, we find **score** of **query** at time  $\ell$  by comparing it to *all keys*

$$\xi_t[\ell] = \mathbf{k}^T[t] \mathbf{q}[\ell] \rightsquigarrow \boldsymbol{\xi}[\ell] = [\xi_1[\ell], \dots, \xi_T[\ell]]$$



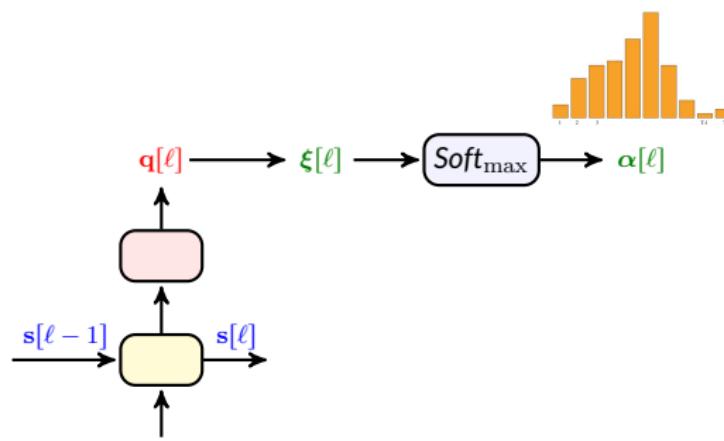
## Attention: Attention Weight at Time $\ell$

We can pass the scores through **softmax** to find

chance of  $v[\ell]$  being related to input entry  $x[t] \equiv \alpha_t[\ell]$

**Softmax** gives us

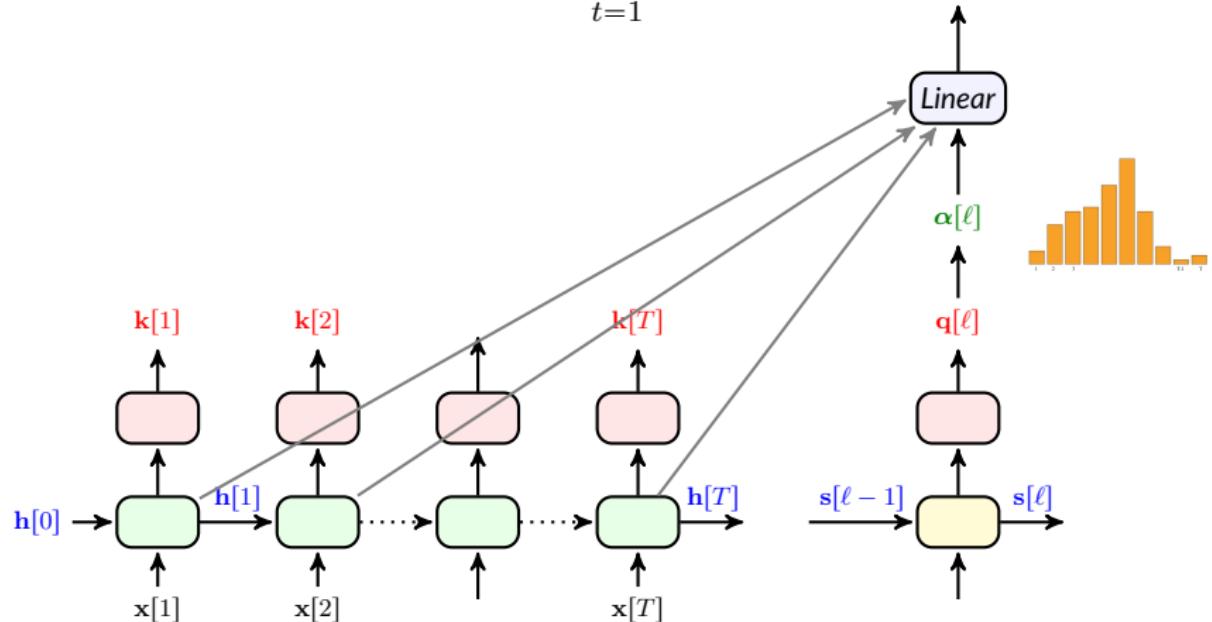
$$\alpha[\ell] = \text{Soft}_{\max}(\xi[\ell]) = [\alpha_1[\ell], \dots, \alpha_T[\ell]]$$



# Attention: Attention Feature

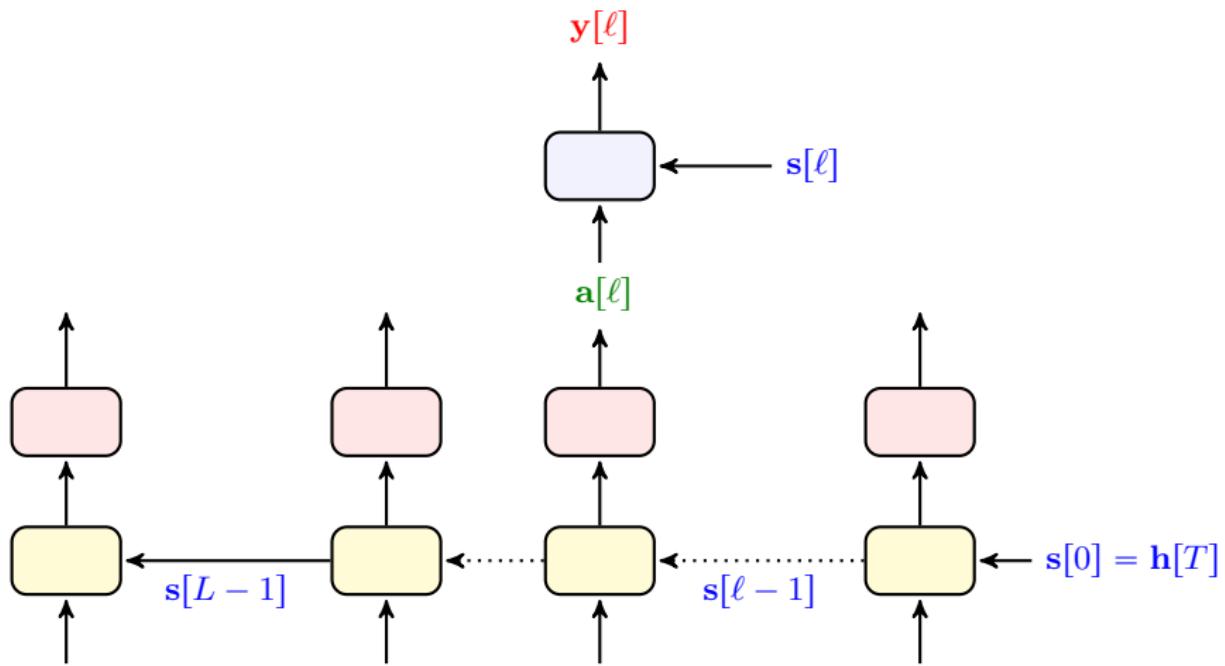
We now use these **weights** to make a vector of attention features

$$\mathbf{a}[\ell] = \sum_{t=1}^T \alpha_t[\ell] \mathbf{h}[t]$$



# Attention: Computing Output

We now use attention features to build the **output sequence**



# Attention: Computing Output

We use attention features to build the **output sequence**

- This can be any layer, as in standard RNN, e.g.,

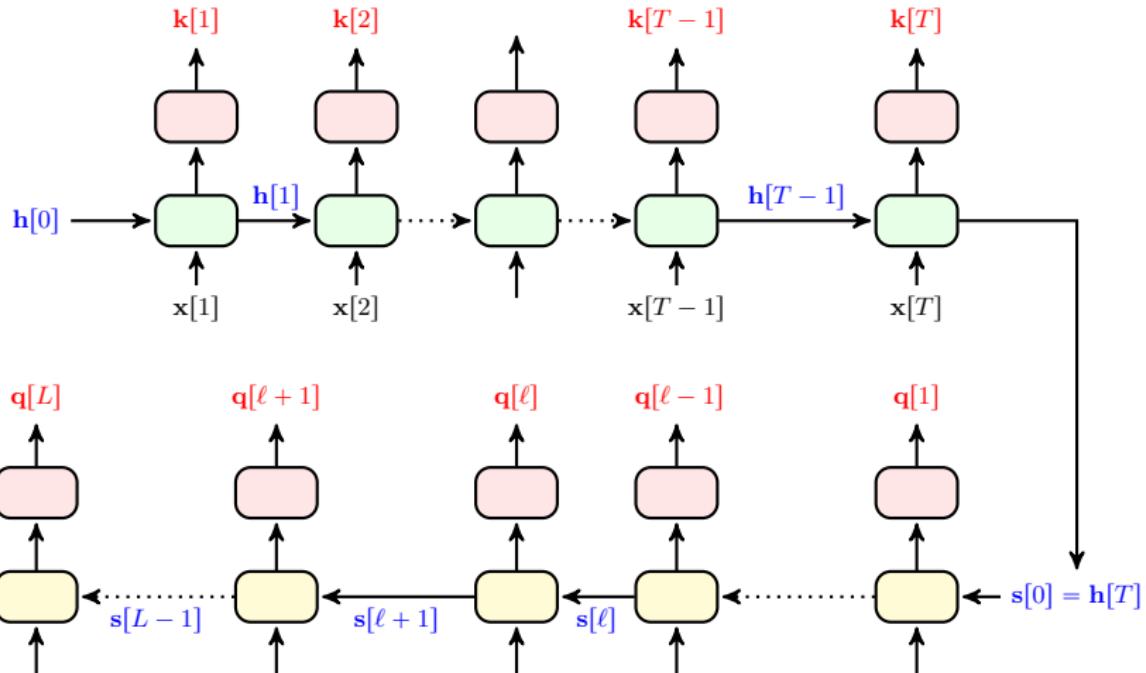
$$\mathbf{y}[\ell] = \text{Soft}_{\max} (\mathbf{W}_{\text{out}} \mathbf{s}[\ell] + \mathbf{W}_{\text{att}} \mathbf{a}[\ell])$$

- We could also use  $\mathbf{a}[t]$  as a new state

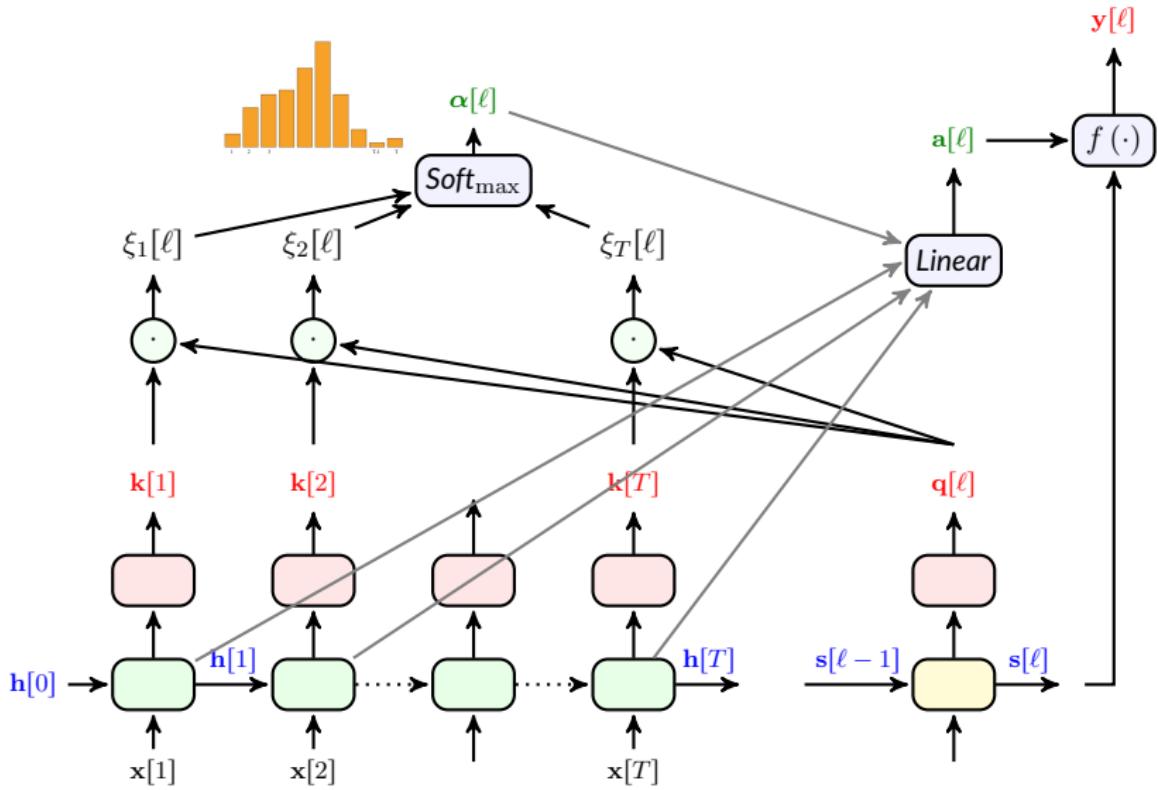
- ↳ We could combine it with the current state
  - ↳ We can pass it to higher layer

It turns out that attention can **hugely** help in practice!

# Attention: End to End Architecture



# Attention: End to End Architecture



## Attention: Training

Let's look at training: assume we want to train it over a single pair of sequences

- We have a sequence of inputs  $\mathbf{x}[t]$ 
  - ↳ For instance a German sentence
- We have a sequence of labels  $\mathbf{v}[\ell]$ 
  - ↳ For instance the English translation
  - ↳ We can compare each output  $\mathbf{y}[\ell]$  with its label

We start with forward pass

- Pass forward through the encoder
  - ↳ Also generate the keys
- Pass the encoder's state and its keys to the decoder
- Pass forward through the decoder
  - ↳ Generate queries and compare them to the keys
  - ↳ Compute attention and the outputs
- Compute loss by aggregating  $\mathcal{L}(\mathbf{y}[\ell], \mathbf{v}[\ell])$

# Attention: Training

Now we should pass backward

- Pass backward through the output layer
  - ↳ Compute  $\nabla_{\mathbf{a}[\ell]} \hat{R}[\ell]$  and  $\nabla_{\mathbf{s}[\ell]} \hat{R}[\ell]$
- Pass backward through the attention layer
  - ↳ Compute  $\nabla_{\boldsymbol{\alpha}[\ell]} \hat{R}[\ell]$ ,  $\nabla_{\boldsymbol{\xi}[\ell]} \hat{R}[\ell]$ ,  $\nabla_{\mathbf{k}[\ell]} \hat{R}[\ell]$  and  $\nabla_{\mathbf{q}[\ell]} \hat{R}[\ell]$
- Pass backward through time at the decoder

$$\nabla_{\mathbf{s}[\ell-1]} \hat{R}[\ell] = \nabla_{\mathbf{s}[\ell]} \hat{R}[\ell] \circ \nabla_{\mathbf{s}[\ell-1]} \mathbf{s}[\ell] + \nabla_{\mathbf{q}[\ell]} \hat{R}[\ell] \circ \nabla_{\mathbf{s}[\ell-1]} \mathbf{q}[\ell]$$

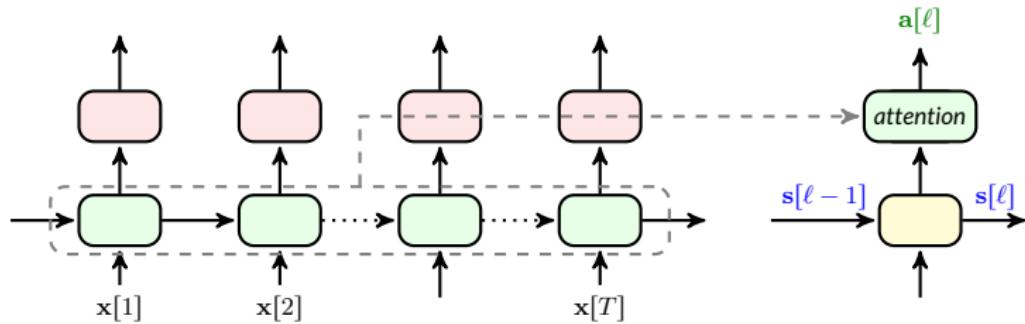
- Pass backward through time at the encoder

$$\begin{aligned} \nabla_{\mathbf{h}[t-1]} \hat{R}[\ell] &= \nabla_{\mathbf{h}[t]} \hat{R}[\ell] \circ \nabla_{\mathbf{h}[t-1]} \mathbf{h}[t] + \nabla_{\mathbf{k}[t]} \hat{R}[\ell] \circ \nabla_{\mathbf{h}[t-1]} \mathbf{q}[t] \\ &\quad + \nabla_{\mathbf{a}[\ell]} \hat{R}[\ell] \circ \nabla_{\mathbf{h}[t-1]} \mathbf{a}[\ell] \end{aligned}$$

- Aggregate over  $\ell$ , update all weights and go for the next round

# Attention as a Layer

We can look at the whole *attention mechanism* as a *layer*



This comes in handy once we want to look at *self-attention*