

Applied Deep Learning

Chapter 7: Sequence-to-Sequence Models

Ali Bereyhi

ali.bereyhi@utoronto.ca


Department of Electrical and Computer Engineering
University of Toronto

Fall 2025

Attention is All You Need!

This was the title of the paper introduced [transformers¹](#)

The attention layer already computes a sort of memory




- + Can't we then drop time connections?
- This is the idea of [self-attention](#)

¹Check the paper at [this link](#)

Self-Attention: Basic Architecture

In **self-attention**, we

- ① start with the **complete sequence**
- ② capture all **temporal correlation** via **attention mechanism**
- ③ return a sequence which better represents **time (positional) dependencies**



Typically, we have **same time length** for input and output

Self-Attention: Basic Architecture


- + How can we do this via attention?
- Input entries apply **attention mechanism** on the input sequence **itself**
- + But, how can we do this?
- Let's do it step by step

First, let's generate **hidden features**, **keys** and **queries**

Self-Attention() :

- 1: **for** time $t = 1, \dots, T$ **do**
- 2: Generate a **hidden feature (value)** $\mathbf{h}[t] = f(\mathbf{W}_h \mathbf{x}[t])$
- 3: Generate a **key** $\mathbf{k}[t] = f(\mathbf{W}_k \mathbf{x}[t])$
- 4: Generate a **query** $\mathbf{q}[t] = f(\mathbf{W}_q \mathbf{x}[t])$
- 5: **end for**

Self-Attention: Basic Architecture




Self-Attention: Basic Architecture

Next we generate **score** for **every time step**

Self-Attention():

- 1: **for** time $t = 1, \dots, T$ **do**
- 2: Generate a **hidden feature (value)** $\mathbf{h}[t] = f(\mathbf{W}_h \mathbf{x}[t])$
- 3: Generate a **key** $\mathbf{k}[t] = f(\mathbf{W}_k \mathbf{x}[t])$
- 4: Generate a **query** $\mathbf{q}[t] = f(\mathbf{W}_q \mathbf{x}[t])$
- 5: **for** time $j = 1, \dots, T$ **do**
- 6: Compare **query** of time t with **key** j by computing $\xi_j[t] = \mathbf{q}^T[t] \mathbf{k}[j]$
- 7: **end for**
- 8: **end for**

Self-Attention: Basic Architecture



Self-Attention: Basic Architecture

Finally, we find the **weights** and use them to compute **attention features**


Self-Attention():

- 1: **for** time $t = 1, \dots, T$ **do**
- 2: Generate a **hidden feature (value)** $\mathbf{h}[t] = f(\mathbf{W}_h \mathbf{x}[t])$
- 3: Generate a **key** $\mathbf{k}[t] = f(\mathbf{W}_k \mathbf{x}[t])$
- 4: Generate a **query** $\mathbf{q}[t] = f(\mathbf{W}_q \mathbf{x}[t])$
- 5: **for** time $j = 1, \dots, T$ **do**
- 6: Compare **query** of time t with **key** j by computing $\xi_j[t] = \mathbf{q}^T[t] \mathbf{k}[j]$
- 7: **end for**
- 8: Pass score $\xi[t]$ through **Softmax** to find **attention weights** $\alpha[t]$
- 9: Compute **attention features** from **values** and **attention weights**


$$\mathbf{a}[t] = \sum_{j=1}^T \alpha_j[t] \mathbf{h}[j]$$

- 10: **end for**

Self-Attention: Basic Architecture




Self-Attention: Complete Architecture



Self-Attention as a Layer


We can look at the whole process as a **single layer**

- ① it takes a **sequence of time as input**
- ② it processes it via **attention mechanism**
- ③ it returns an **output time sequence**



Self-Attention as a Layer

Note that *all the weights are fixed through time: think of it as in RNNs*



However, we should pay attention that

Self-attention returns the *complete output sequence only after going through the complete input sequence: No sequential order!*


Processing Sequences Only with Attention

At *this point*, people started to think: *what if*

we process the time sequence via multiple layers of self-attention

- Start with *input $x[t]$* and compute *an attention sequence $a_1[t]$*
 - ↳ $a_1[t]$ has captured temporal features of $x[t]$
- From *attention $a_1[t]$* , compute *a new attention sequence $a_2[t]$*
 - ↳ $a_2[t]$ has captured *better* temporal features of $x[t]$
- $\dots \times N$
- From *last attention $a_{N-1}[t]$* , we compute *an output sequence $y[t]$*
 - ↳ $y[t]$ has captured *almost all* temporal features of $x[t]$

Processing Sequences Only with Attention




Processing Sequences Only with Attention

Intuitive Observation

Self-attention does the same thing as convolution


- It has a fixed set of weights to compute **values**, **keys**, and **queries**
 - ↳ Convolutional layers have fixed **filters** (**kernels**)
 - It goes through the entries sequence with these weights using **attention**
 - ↳ Convolutional layers apply **convolution** with the **same filter**
 - It returns a **output sequence** that captures **temporal correlation**
 - ↳ Convolutional layers return **maps** that capture **spatial correlation**
-
-
- + But, we compute **multiple maps** via **multiple filters** in convolutional layers!
 - Well, we could do **the same** here!

Self-Attention: Multi-headed Architecture



- We can compute **multiple** self-attention **layers** in parallel
 - ↳ A **layer** in this context is called a **head**
- Each **head** computes a separate attention feature
 - ↳ Head h computes the attention feature $\mathbf{a}^{(h)}[t] \in \mathbb{R}^{D_0}$
- The **concatenation** of these features is the final **attention feature**
 - ↳ Attention feature at time t is $\mathbf{a}[t] = [\mathbf{a}^{(1)}[t], \dots, \mathbf{a}^{(H)}[t]] \in \mathbb{R}^{HD_0}$

Self-Attention: Multi-headed Architecture



Typical Practice

In practice, we often choose D_0 and H such that dimensionality is preserved: say every $\mathbf{x}[t] \in \mathbb{R}^N$

- We set $D_0 = N/H$ for an H -head self-attention layer
- The concatenated attention feature is then of length $HD_0 = N$


The key reason is that we often use this layer with skip connection

Self-Attention: Adding Nonlinearity

It turns out that although self-attention captures temporal correlation


it is not complex enough to capture content features accurately

- + *Can't we do it by having more complex layers!*
- *Sure! We often process the attention features with a nonlinear layer*



This layer further captures content features

Self-Attention: A Unit



Self-Attention Unit: Generic Form

A **self-attention unit** consists of a **multi-head attention layer** followed by an **activated layer**. These layers are all **fixed** through time


A **deep self-attention** architecture is made by cascading **self-attention units**

Transformer: Encoding and Decoding via Self-Attention

Experiments showed that


self-attention can better capture **temporal** and **content** features

- + Can't we replace them with RNNs in an **encoder-decoder** architecture?
- Sure! This is indeed a basic **transformer**!



Transformer: Encoding and Decoding via Self-Attention

- + Does this idea work that easily?!
- Well! Pretty much Yes, after some basic modifications




There are two challenges in using self-attention for encoding and decoding

- For encoding: changing time order at input does not impact the output
 - ↳ Self-attention captures temporal correlation, but not sequential order!
- For decoding: we should generate the entire output at once
 - ↳ We often want to generate $y[t]$ depending on previous outputs $y[1 : t - 1]$

Encoding via Self-Attention: Positional Encoding

A simple remedy is to **include time index** in the input!




We replace the **input** with **(input, time)**, i.e.,

$$\mathbf{x}[t] \leftarrow (\mathbf{x}[t], t)$$

Now, **time permutation changes output** \rightsquigarrow **output depends on sequential order**


Positional Encoding



Positional Encoding

To use *self-attention* at *encoder* we concatenate *input* with a *time-dependent (positional) feature* $p[t]$ that can be potentially *learned* from data itself

Decoding via Self-Attention: Masked Decoding




Recall that attention at time t is computed as

$$\mathbf{a}[t] = \sum_{j=1}^T \alpha_j[t] \mathbf{h}[j]$$

which depends on the **future** data through $\mathbf{h}[j]$ for $j = t + 1, \dots, T$

We could simply ignore them by **masking** $\alpha_j[t]$ for $j = t + 1, \dots, T$

Decoding via Self-Attention: Masked Decoding




We replace $\alpha_j[t] \leftarrow 0$ for $j = t + 1, \dots, T$ and compute

$$\mathbf{a}[t] = \sum_{j=1}^T \alpha_j[t] \mathbf{h}[j] = \sum_{j=1}^t \alpha_j[t] \mathbf{h}[j]$$

A Classic Transformer

A classic transformer looks like this²



²The diagram is from the paper [Attention is All You Need!](#)

Why Transformers?

- + Transformers seem to be **more challenging** to implement! Why should we use them?
- That's right! But they come with some benefits!

In general transformers

- process sequences **in parallel** \rightsquigarrow room for **parallel computing**
- preserve much better **long-term connections**
- can be **deepened** much more **straightforward**

But, it's true that

- they pose **more complexity** \rightsquigarrow we can handle it nowadays

Final Notes

Transformer is a core element of many modern learning models

- GPT stands for Generative Pre-trained Transformer
- They are the better choice for complicated Seq2Seq problems
- Pre-implemented transformer is available in PyTorch

```
torch.nn.Transformer(d_model, nhead, num_encoder_layers,  
                      num_decoder_layers, ...)
```

If you are interested to know more about transformers

- ↳ Consider taking NLP course next fall semester
- ↳ We also revisit it in the Generative Models course this summer
- ↳ Take a look at this nice online resource